

Multi-core RISC Processor Design and Implementation

ELEC5881M - Final Report

Ben David Lancaster

Student ID: 201280376

Submitted in accordance with the requirements for the degree of
Master of Science (MSc)
in Embedded Systems Engineering

Supervisor: Dr. David Cowell

Assessor: Mr David Moore

University of Leeds

School of Electrical and Electronic Engineering

August 29, 2019

Word count: 8963

Abstract

This interim report details the 4-month progress on a project to design, implement, and verify, a multi-core FPGA RISC processor. The project has been split into two stages: firstly to build a functional single-core RISC processor, and then secondly to add multiprocessor principles and functionality to it.

Current multiprocessor and network-on-chip communication methods have been discussed and how they could be included in this multi-core RISC design. To-date, a 16-bit instruction set architecture has been designed featuring common load/store instructions, comparison, and bitwise operations. A single-core processor has been implemented in Verilog and verified using simulations/test benches running various simple software programs.

Future tasks have been planned and will focus on the second stage of the project. Work will start on designing a loosely coupled multiprocessor communication interface and bringing them to the single-core processor.

Declaration of Academic Integrity

The candidate confirms that the work submitted is his/her own, except where work which has formed part of jointly-authored publications has been included. The contribution of the candidate and the other authors to this work has been explicitly indicated in the report. The candidate confirms that appropriate credit has been given within the report where reference has been made to the work of others.

This copy has been supplied on the understanding that no quotation from the report may be published without proper acknowledgement. The candidate, however, confirms his/her consent to the University of Leeds copying and distributing all or part of this work in any forms and using third parties, who might be outside the University, to monitor breaches of regulations, to verify whether this work contains plagiarised material, and for quality assurance purposes.

The candidate confirms that the details of any mitigating circumstances have been submitted to the Student Support Office at the School of Electronic and Electrical Engineering, at the University of Leeds.

Name: Ben David Lancaster

Date: August 29, 2019

Acknowledgements

I would like to thank my supervisor, David Cowell, and assessor, David Moore, for giving me the opportunity to explore a project of my own choosing, and for their support and guidance throughout this project.

Table of Contents

1	Introduction	10
1.1	Why Multi-core?	10
1.2	Why RISC?	11
1.3	Why FPGA?	11
2	Background	12
2.1	Amdahl's Law and Parallelism	12
2.2	Loosely and Tightly Coupled Processors	12
2.3	Network-on-chip Architectures	13
3	Project Overview	15
3.1	Project Deliverables	15
3.1.1	Core Deliverables (CD)	15
3.1.2	Extended Deliverables (ED)	16
3.2	Project Timeline	17
3.2.1	Project Stages	17
3.2.2	Project Stage Detail	17
3.2.3	Timeline	19
3.3	Resources	19
3.3.1	Hardware Resources	19
3.3.2	Software Resources	20
3.4	Legal and Ethical Considerations	21
4	Single-core Design	22
4.1	Introduction	22
4.2	Design and Implementation	22
4.2.1	Instruction Set Architecture	23
4.2.2	Memory Management Unit	24
4.2.3	Instruction and Data Memory	24
4.2.4	ALU Design	24
4.2.5	Decoder Design	26
4.2.6	Pipelining	26
4.2.7	Design Optimisations	27
4.3	Interrupts	27

4.3.1	Overview	28
4.3.2	Hardware Implementation	28
4.3.3	Software Interface	29
4.3.4	Design Improvements	30
4.4	Verification	31
5	Interconnect	32
5.1	Introduction	32
5.1.1	Comparison of On-chip Buses	32
5.2	Overview	33
5.2.1	Design Considerations	34
5.3	Interfaces	35
5.3.1	Multi-master Support	36
5.4	Further Work	37
6	Memory Mapping	39
6.1	Introduction	39
6.2	Address Decoding	39
6.2.1	Decoder Optimisations	40
6.3	Memory Map	42
7	Multi-core Communication	43
7.1	Introduction	43
7.1.1	Design Goals	43
7.1.2	Context Identification	43
7.1.3	Thread Synchronisation	44
8	Analysis & Results	47
8.1	Analysis	47
8.1.1	Design Area/Size Requirements	47
8.1.2	Maximum Frequency	49
8.2	Scenario Performance	49
8.2.1	Scenario Overview	50
8.2.2	Performance Measurements	50
8.2.3	Performance Results	51
8.2.4	Shared Instruction Memory Impact	52
8.3	Analysis Review	53
9	Conclusion	54
9.1	Overview	54
9.2	Review Against Project Deliverables	54
9.2.1	Core Deliverables	54
9.2.2	Extended Deliverables	55
9.3	Future Work	56

References	59
Appendices	60
A Peripheral Information	60
A.1 Special Registers	60
A.2 Watchdog Timer	61
A.3 GPIO Interface	61
A.4 Timer with Interrupt	62
A.5 UART Interface	62
B Additional Figures	63
B.1 Register Set Multiplex	63
B.2 Instruction Set Architecture	64
C Configuration Options	65
C.1 System-on-Chip Configuration Options	65
C.2 Core Options	66
C.3 Peripheral Options	67
D Viva Demonstration Examples	68
D.1 2-core Timer Interrupt and ISR	68
D.2 1-160 Core Parallel Summation	70
E Code Listing	72
E.1 SoC Code Listing	72
E.1.1 vmicro16_soc.config.v	72
E.1.2 top.ms.v	74
E.1.3 vmicro16_soc.v	75
E.1.4 vmicro16.v	82
E.2 Peripheral Code Listing	95
E.3 Assembly Compiler Listing	101
E.4 Text Compiler Listing	106

List of Figures

2.1	A loosely coupled multiprocessor system. Each node features it's own memory and IO modules and uses a Message Transfer System to perform inter-node communication. Image source: [1].	13
2.2	A tightly coupled multiprocessor system. Nodes are directly connected to memory and IO modules. Image source: [1].	13
2.3	A multiprocessor network-on-chip architecture with 16 processing nodes. Nodes are connected in a grid formation with routers and links. Image source: [2]. . .	14
3.1	Project stages in a Gantt chart.	19
3.2	Terasic DE1-SoC development board featuring the Altera Cyclone V FPGA and many peripherals. Image source: [3].	20
3.3	Minispartan-6+ development board featuring the Xilinx Spartan 6 XC6SLX9. Note that the XC6SLX9 and XC6SLX25 FPGAs share the same board. Image source: [4].	21
4.1	Vmicro16 RISC 5-stage RTL diagram showing: instruction pipelining (data passed forward through clocked register banks at each stage); branch address calculation; ALU operand calculation (rd2 or imm); and program counter incrementing.	23
4.2	Vmicro16 ALU diagram showing clocked inputs from the previous IDEX stage being	25
4.3	Time diagram showing the TIMR0 peripheral emitting a 1us periodic interrupt signal (out) to the processor. The processor acknowledges the interrupt (int_pending_ack) and enters the interrupt mode (regs_use_int) for a period of time. When the interrupt handler reaches the Interrupt Return instruction (indicated by w_intr) the processor returns to normal mode and restores the normal state.	29
4.4	The interrupt vector (0x0100 - 0x0107) consists of eight 16-bit values that point to memory addresses of the instruction memory to jump to.	29
4.5	Interrupt Mask register (0x0108). Each bit corresponds to an interrupt source. 1 signifies the interrupt is enabled for/visible to the core. Bits [7:2] are left to the designer to assign. Bit 0 is assigned to TIMR0's interval timer. Bit 1 is assigned to the UART0's receiver (unassigned if DEF_USE_REPROG is enabled).	30
5.1	Waveform showing an APB read transaction.	33

5.2	Block diagram of the Vmicro16 system-on-chip.	34
5.3	Source: [5]	35
5.4	Vmicro16 master/slave interface using AMBA APB	35
5.5	Multi-master schematic for the Vmicro16 system-on-chip.	37
6.1	Schematic showing the address decoder (addr_dec) accepting the active PADDR signal and outputting PSEL chip enable signals to each peripheral.	40
6.2	Example 4-bit binary comparator which compares the bits (a, b, c, d) to the constant value 1010. The 0s of the constant are inverted and then all are passed to a wide-AND.	40
6.3	Bits [7:3] of an 8-bit PADDR signal are used as inputs to 5-bit LUTs to generate a PSEL signal. In addition, a default error case is shown allowing the address decoder to detect incorrect PADDR values (e.g. if no PSEL signals are generated).	41
6.4	Partial address decoding used by the Vmicro16 SoC design. Each peripheral shown only needs to decode a signal bit to determine if it is enabled.	41
6.5	Memory map showing addresses of various memory sections.	42
7.1	Block diagram showing the main multi-processing components: the CPU array and a peripheral interconnect used for core synchronisation.	44
7.2	Vmicro16 Special Registers layout (0x0080 - 0x008F).	44
7.3	Assembly code for locking a mutex. r1 is the address to lock. r3 is zero. r4 is the branch address.	45
8.1	Per-core instruction memory schematic and performance.	48
8.2	Shared instruction memory schematic and performance.	49
8.3	Cyclone V maximum design frequency for various core count configurations.	49
8.4	Chart showing how the communication times (Tbus) and serial times (Tsum) changes with core count.	51
8.5	Similar to Figure 8.4 but using shared instruction memory to reduce block memory requirements per core.	52
A.1	Vmicro16 Special Registers layout (0x0080 - 0x008F).	61
B.1	Normal mode (bottom) and interrupt mode (top) register sets are multiplexed to switch between contexts.	63
B.2	Vmicro16 instruction set architecture.	64

List of Tables

3.1	Project stages throughout the life cycle of the project.	18
5.1	Decoding of the LE and SE parameters to determine the global shared memory operation.	36
C.1	SoC Configuration Options	65
C.2	Core Options	66
C.3	Peripheral Options	67

List of Listings

1	ALU branch detection using flags: zero (Z), overflow (V), and negative (N). . . .	25
2	Vmicro16's ALU implementation named vmicro16_alu. vmicro16.v	26
3	Vmicro16's ALU implementation named vmicro16_alu. vmicro16.v	26
4	Vmicro16's decoder module code showing nested bit switches to determine the intended opcode. vmicro16.v	26
5	RAM and lock memories instantiated by the shared memory peripheral.	45
6	Assembly code for a memory barrier. Threads will wait in the barrier_wait function until all other threads have reached that code point.	46
7	Variable size inputs and outputs to the interconnect.	63

Chapter 1

Introduction

1.1 Why Multi-core?	10
1.2 Why RISC?	11
1.3 Why FPGA?	11

This project will detail the design, implementation, and verification, of a new multi-core RISC processor aimed at FPGA devices. This project was chosen due to my interest in processor design, in which I have only previously designed single-core RISC processors, and wish to extend this knowledge to gain a basic understanding of multi-core communication, design considerations, and the challenges of software and hardware parallelism first hand.

I will use this opportunity to further develop my knowledge of FPGA and processor design by implementing, designing, and verifying, a multi-core RISC processor from scratch, including the design of a communication interface between multiple cores.

1.1 Why Multi-core?

Moore's Law states that the number of transistors in a chip will double every 2 years. CPU designers would utilize the additional transistors to add more pipeline stages in the processor to reduce the propagation delay which would allow for higher clock frequencies.

The size of transistors have been decreasing and today can be manufactured in sub-10 nanometer range. However, the extremely small transistor size increases electrical leakage and other negative effects resulting in unreliability and potential damage to the transistor. The high transistor count produces large amounts of heat and requires increasing power to supply the chip. These trade-offs are currently managed by reducing the input voltage, utilising complex cooling techniques, and reducing clock frequency. These factors limit the performance of the chip significantly. These are contributing factors to Moore's Law *slowing* down. The capacity limit of the current-generation planar transistors is approaching and so in order for performance increases to continue, other approaches such as alternate transistor technologies like Multigate transistors [6], software and hardware optimisations, and multi-processor architectures are employed.

This report will focus on the latter: to produce a small multi-core processor that can utilise software-based parallelism to gain performance benefits, compared to a larger single-core

design.

1.2 Why RISC?

RISC architectures feature simpler and fewer instructions compared to CISC, which emphasises instructions that perform larger tasks. A single CISC instruction might be performed with multiple RISC instructions. Because of the fewer and simpler instructions, RISC machines rely heavily on software optimisations for performance. RISC instruction sets are based on load/store architectures, where most instructions are either register-to-register or memory reading and writing [7]. This constraint greatly reduces complexity.

RISC architectures are easier to design implement, especially for beginners, due to their simpler instructions that share the same pipeline, compared to CISC where there may be different pipeline for each instruction, which would greatly consume FPGA resources.

1.3 Why FPGA?

Field programmable gate arrays (FPGA) are a great choice for prototyping digital logic designs due to their programmable nature and quick development times.

My previous experience with FPGAs in previous projects will reduce risk and learning times and allow for more time to be spent on adding and extending features (discusses further in section 3).

FPGAs, however, may not be suitable for prototyping all register-transistor logic (RTL) projects. Larger RTL projects, such as large commercial processors, may greatly exceed the logic cell resources available in today's high-end FPGA devices and may only be prototyped through silicon fabrication, which can be expensive. This resource limitation will not be problem as the project aims to produce a small and minimal design specifically for learning about multi-core architectures.

Chapter 2

Background

2.1 Amdahl's Law and Parallelism	12
2.2 Loosely and Tightly Coupled Processors	12
2.3 Network-on-chip Architectures	13

2.1 Amdahl's Law and Parallelism

In many applications, not restricted to software, there may exist many opportunities for processes or algorithms to be performed in parallel. These algorithms can be split into two parts: a serial part that cannot be parallelised, and a part that can be parallelised. Amdahl's Law defines a formula for calculating the maximum *speedup* of a process with potential parallelism opportunities when ran in parallel with n many processors. Speedup is a term used to describe the potential performance improvements of an algorithm using an enhanced resource (in this case, adding parallel processors) compared to the original algorithm. Amdahl's Law is defined below, where the potential speedup S_p is dependant on the portion of program that can be parallelised p and the number of processing cores n :

$$S_p = \frac{1}{(1 - p) + \frac{p}{n}} \quad (2.1)$$

This formula will be used throughout the project to gauge the performance of the multi-core design running various software algorithms.

2.2 Loosely and Tightly Coupled Processors

Multiprocessor systems can be generalised into two architectures: loosely and tightly coupled, and each architecture has advantages and disadvantages. In loosely coupled systems, each processing node is self-contained – each node has its own dedicated memory and IO modules. Communication between nodes is performed over a *Message Transfer System (MTS)* [1] in a master-slave control architecture.

Scalability in loosely coupled systems is generally easier to implement as each node can simply be appended to the shared MTS interface without large modifications to the rest of

the system. Scalability is an important concern in this project as I wish to test the developed solution with a range of processing nodes.

As loosely coupled system's nodes feature their own memory and IO modules, they generally perform better in cases where interaction between nodes is not prominent – each node can store a separate part of the software program in its memory module allowing simultaneous executing of the program.

In scenarios where inter-node communication is prominent however, access to the MTS interface must be scheduled to avoid access conflicts which introduces delays and idle times in the software programs execution, resulting in lower throughput. Figure 2.1 shows a general layout of a loosely coupled multiprocessor system.

Tightly coupled systems feature processing nodes that do not have their own dedicated memory or IO modules – each node is directly connected to a shared memory module using a dedicated port. In scenarios where inter-node communication is prominent, tightly coupled systems are generally better suited as nodes are directly connected to a shared memory and do not need to wait to use a shared bus.

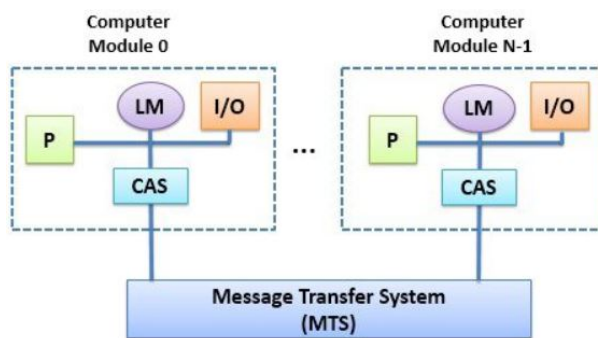


Figure 2.1: A loosely coupled multiprocessor system. Each node features its own memory and IO modules and uses a Message Transfer System to perform inter-node communication. Image source: [1].

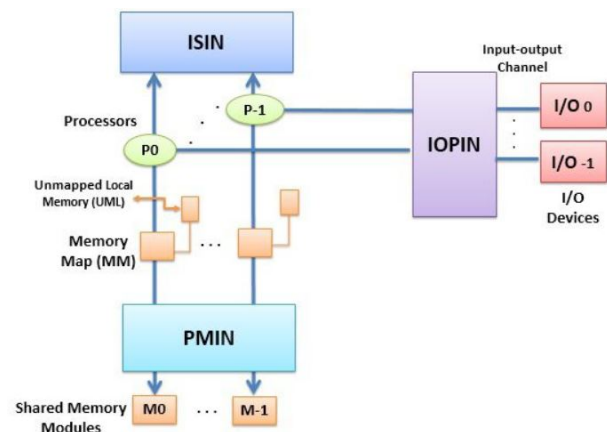


Figure 2.2: A tightly coupled multiprocessor system. Nodes are directly connected to memory and IO modules. Image source: [1].

This project will utilise a loosely coupled architecture due to its easier scalability implementation and my previous experience with the design of single-core processors. Although it will require a scheduler to access the MTS, the experience and knowledge gained from this task will be greatly beneficial for future projects.

2.3 Network-on-chip Architectures

Network-on-chip (NoC) architectures implement on-chip communication mechanisms that are based on network communication principles, such as routing, switching, and massive scalability [8]. NoC's can generally support hundreds to millions of processing cores. Figure 2.3 shows an example 16-core network-on-chip architecture. NoC's can scale to very large sizes while not sacrificing performance because each processor core is able to drive the network rather than needing to wait for a shared bus to become free before doing so.

The greater the number of cores in a network-on-chip design, the greater quality of service

(QoS) problems arise. As such, network-on-chip architectures suffer the same problems as networks, such as fairness and throughput [9].

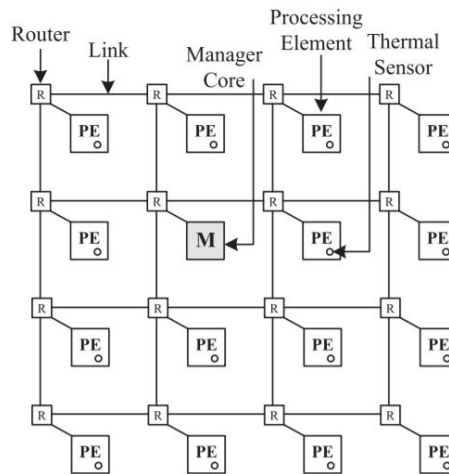


Figure 2.3: A multiprocessor network-on-chip architecture with 16 processing nodes. Nodes are connected in a grid formation with routers and links. Image source: [2].

Chapter 3

Project Overview

3.1	Project Deliverables	15
3.1.1	Core Deliverables (CD)	15
3.1.2	Extended Deliverables (ED)	16
3.2	Project Timeline	17
3.2.1	Project Stages	17
3.2.2	Project Stage Detail	17
3.2.3	Timeline	19
3.3	Resources	19
3.3.1	Hardware Resources	19
3.3.2	Software Resources	20
3.4	Legal and Ethical Considerations	21

This chapter discusses the the project’s requirements, goals, and structure.

3.1 Project Deliverables

The project’s deliverables are split into two sections: core deliverables (CD) – each deliverable must be satisfied for the project to be a minimum viable product (MVP), and extended deliverables (ED) – deliverables that are not required for a MVP – features that only improve upon an existing feature.

3.1.1 Core Deliverables (CD)

The project’s core deliverables are described below.

CD1 Design a compact 16-bit RISC instruction set architecture.

The instruction set will be the primary interface to control the processor from software. An instruction set will be required to implement the custom multi-core communication interface.

It was decided to design a new instruction set rather than to extend an existing architecture as this will increase my knowledge of the constraints to consider when designing instruction sets and processors.

CD2 Design and implement a Verilog RISC core that implements the ISA in CD1.

The Verilog RISC core will be able to run software program written for the instruction set architecture.

CD3 Design and implement an on-chip interconnect for multi-core processing (2 to 32 cores) using the RISC core from CD2.

The interconnect will be a chief requirement to enable multi-core communication. The interconnect should support up to 32 cores, however FPGA implementation constraints may limit this due to limited resources.

The interconnect will control communication between the cores to enable software parallelism.

CD4 Analyse performance of serial and parallel software algorithms, such as parallel DFT, on the processor.

To evaluate the effectiveness of the developed solution, a serial and parallel implementation of a simple computing algorithm (parallel reduction, sorting) will be ran on the processor and it's performance analysed. Effectiveness will be rated on total algorithm run-time and the speed-up gained by adding more cores.

CD5 Allow the RISC core to be easily compiled to multiple FPGA vendors (Xilinx, Altera).

The developed solution should be generic and portable to allow it to be used across a wide-range of FPGA vendors and devices.

Verilog is a generic implementation-independent hardware-description language and so designing implementation specific modules is recommended.

A key consideration for this requirement is to consider the varying hard IP provided by the FPGA vendors (such as BRAM, ethernet, and PCIe [10, 11]). To overcome this problem, the developed Verilog code will conditionally compile where vendor specific requirements are present.

3.1.2 Extended Deliverables (ED)

The project's extended deliverables are described below.

ED1 Design a RISC core with an instructions-per-clock (IPC) rating of at least 1.0 (a single-cycle CPU).

ED2 Design a RISC core with a pipe-lined data path to increase the design's clock speed.

ED3 Design a scalable multi-core interconnect supporting arbitrary (more than 32) RISC core instances (manycore) using Network-on-Chip (NoC) architecture.

ED4 Design a compiler-backend for the PRCO304 [12] compiler to support the ISA from CD1. This will make it easier to build complex multi-core software for the processor.

ED5 The RISC core can communicate to peripherals via a memory-mapped addresses using the Wishbone bus.

ED6 Implement various memory-mapped peripherals such as UART, GPIO, LCD, to aid visual representation of the processor during the demonstration viva.

ED7 Store instruction memory in SPI flash.

ED8 Reprogram instruction memory at runtime from host computer.

ED9 Processor external debugger using host-processor link.

3.2 Project Timeline

3.2.1 Project Stages

The project is split up into many stages to aid planning and management of the project. There are 8 unique stage areas: 1. Initial project conception; 2 Basic RISC core development; 3. Extended RISC core development; 4. Multi-core development; 5. Processor quality-of-life (QoL) improvements; 6. Compiler development; 7. Demo preparation, and 8. Final report.

The project stages are shown in Table [3.1](#).

3.2.2 Project Stage Detail

Stages 1.0 through 1.2 – Research and Project Conception

These stages cover initial research of existing problems and solutions in the multiprocessor area. The instruction set architecture is also proposed that later stages will implement.

Stages 2.1 through 2.3 – Processor module Design, Implementation, and Integration

These stages cover the design, implementation, and integration of key processor core modules such as the instruction decoder, register sets and local memory. Integration of all the modules is a challenging task because some modules have both asynchronous and synchronous signals that need to be timed correctly in order for other modules to receive valid data. An example of this is the register set which has asynchronous read ports that are later clocked in the instruction decode stage.

Stages 3.1 through 3.4 – Advanced Processor Implementation

These stages add advanced features to the processor to provide a more functional product. Although these stages are classified as extended, their technical requirement to design and implement is not great and so are have time allocations in the project schedule. The extended features that these stages introduce are: pipelined processor stages – to drastically increase processor performance; provide a memory-mapped peripheral interface through the MMU; provide a Wishbone master interface to the MMU – allowing external peripherals such as GPIO and LCD displays to be utilised in a modular fashion; and to implement a cache memory for each processor core.

Stage	Title	Start Date	Core	Status
1.0	Research	Feb 04	x	Completed
1.1	Requirement gathering/review	Feb 11	x	Completed
1.1	Processor specification, architecture, ISA	Feb 18	x	Completed
1.2	Stage/Time Allocation Planning	Feb 25	x	Completed
2.1	Decoder, Register Set, impl & integration	Feb 25	x	Completed
2.2	Register set impl & integration	Mar 04	x	Completed
2.3	Local memory impl & integration	Mar 11	x	Completed
3.1	Memory mapped register layout & impl	Apr 01		On-going
3.2	Wishbone peripheral bus connected to MMU	Apr 08		On-going
3.3	Pipeline implementation and verification	Apr 15		On-going
3.4	Cache memory design & impl	Apr 22		Cancelled
4.1	Multi-core communication interface	Jun 05	x	Planned
4.2	Shared-memory controller	Jun 05	x	Planned
4.3	Scalable multi-core interface (10s of cores)	Jul 01	x	Planned
4.4	Multi-core example program (reduction)	Jul 10	x	Planned
5.1	SPI-FPGA interface for OTG programming	TBD		Cancelled
5.2	FPGA-PC interfacing	TBD		Cancelled
5.3	FPGA-PC debugging (instruction breakpoints)	TBD		Cancelled
6.1	Compiler backend for vmicro16	TBD		Unknown
6.2	Compiler support for multi-core codegen	TBD		Unknown
7.1	Wishbone peripherals for demo	Aug 01	x	Planned
8.1	Final Report	Jun 05	x	Planned

Table 3.1: Project stages throughout the life cycle of the project.

Stages 4.1 through 4.4 – Multiprocessor Functionality

These stages are dedicated to adding multiprocessor functionality using a loosely coupled architecture to the processor.

Stages 5.1 through 5.3 – Debugging Features

These stages cover debugging features and are classified as extended due to the large development time required to implement them as well as not being related to multiprocessor systems.

Stages 6.1 through 6.2 – Compiler Backends

These stages cover the implementation of a compiler backend to ease software writing and programming of the processor.

Stage 7.1 – Wishbone Peripherals

Additional Wishbone peripherals, such as SPI and timers will be added to produce a more useful multiprocessor system.

Stage 8.1 – Final Report

This stage is dedicated to the final report write-up. It is expected to be an iterative task that is active throughout the lifespan of the project.

3.2.3 Timeline

The project stages from Table 3.1 are displayed below in a Gantt chart.

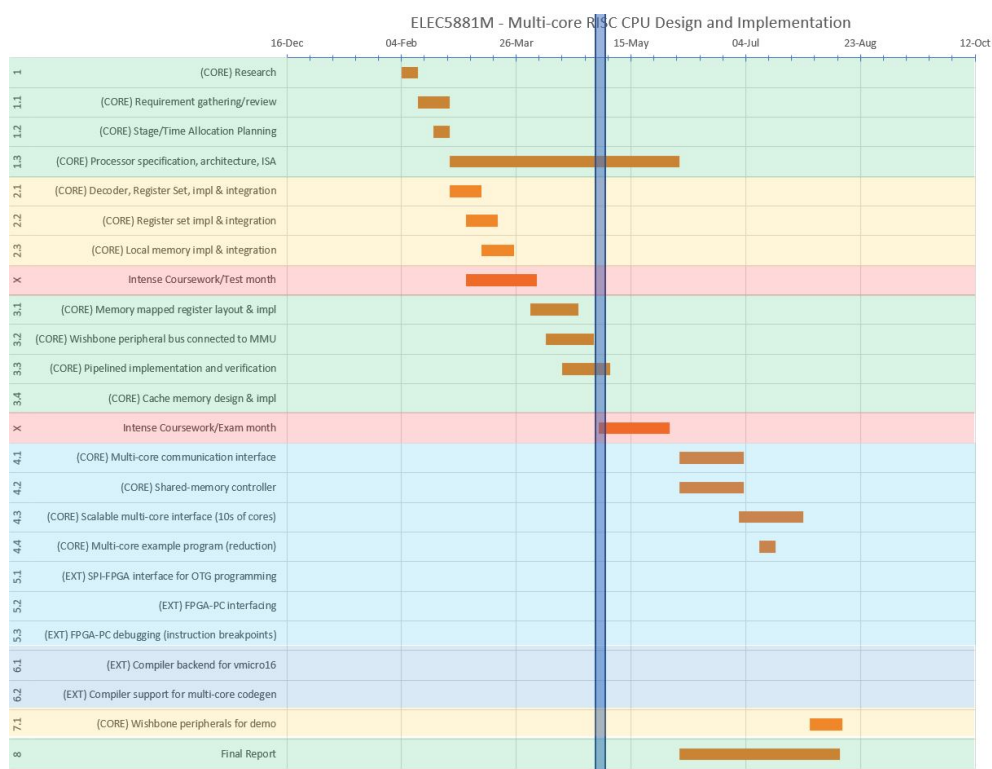


Figure 3.1: Project stages in a Gantt chart.

3.3 Resources

This section describes the hardware and software resources required to fulfil the project.

3.3.1 Hardware Resources

Core deliverable **CD5** requires the designed RISC core to be implemented and demonstrated on multiple FPGA devices. Although my design should synthesise for physical IC implementation, due to high costs and lengthy production times, it is not a primary development target. Due to having past experience with Xilinx FPGAs from my placement work and experience with

Altera from university modules it was decided to target the Xilinx Spartan 6 XC6SLX9 and the Altera Cyclone V.

Terasic DE1-SoC Development Board

The Terasic DE1-SoC development board features a large Cyclone V FPGA and many peripherals, such as seven-segment displays, 64 MB SDRAM, ADCs, and buttons and switches, which will aid demonstration of the project. The development board is available through the university so the cost is negligible. Figure 3.2 shows the peripherals (green) available to the FPGA.




Figure 3.2: Terasic DE1-SoC development board featuring the Altera Cyclone V FPGA and many peripherals. Image source: [3].

Minispartan 6+ FPGA Development Board

The Minispartan 6+ is a hobbyist FGPA development board with fewer peripherals than the DE1-SoC. The board features a Xilinx Spartan 6 XC6LX9 which has far fewer resources than the DE1-SoC's Cyclone V however it's simplicity and my familiarity with Xilinx's software suite will speed up development. The development board is shown in Figure 3.3.

3.3.2 Software Resources

Intel Quartus

Intel Quartus Prime is a paid-for SoC, CPLD, and FPGA software suite targeting Intel's Stratix, Arria, and Cyclone based FPGAs. The university provides student licences which will be used via VPN.

Xilinx ISE Webpack

Xilinx ISE Webpack is Xilinx's free software suite for FPGA development for Spartan 6 based FPGAs. Due to ISE's intuitive and fast work flow, most of the initial simulation and verification

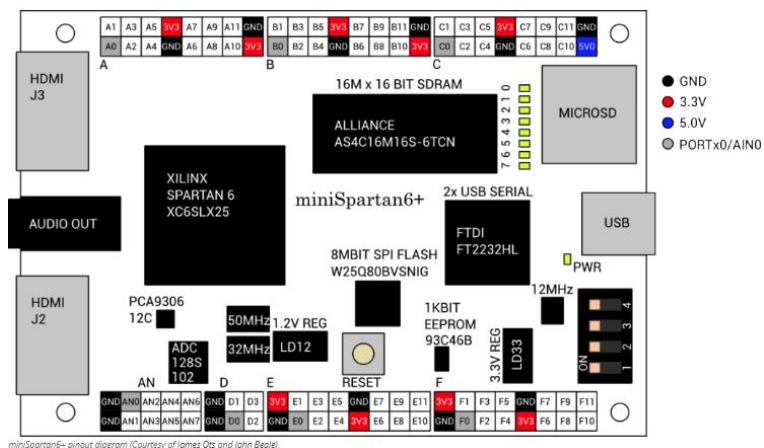


Figure 3.3: Minispartan-6+ development board featuring the Xilinx Spartan 6 XC6SLX9. Note that the XC6SLX9 and XC6SLX25 FPGAs share the same board. Image source: [4].

processes will be performed using ISE. This will greatly improve development times.

Verilator

Verilator is an open-source Verilog to C++ transpiler which provides a C++ interface to simulate Verilog modules and read/write values similar to a test bench. Verilator will be used for specific modules within the RISC core such as the ALU and decoder as Verilator is useful when performing exhaustive verification.

3.4 Legal and Ethical Considerations

The RISC core is designed to be used as an academic research and educational tool to aid learning and understanding of RISC and multi-core machines. It should not be use for roles where mission critical or safety is a factor.

The processor does not provide any memory protection features and any software running on the processor has full access to all memory.

The processor does not store/track/predict software instructions. The processor uses pipelining techniques to improve performance which results in future instructions entering the pipeline even if the software's logical sequence does not include these instructions. This could result in security vulnerabilities similar to Intel's Spectre vulnerability [13].

Chapter 4

Single-core Design

4.1	Introduction	22
4.2	Design and Implementation	22
4.2.1	Instruction Set Architecture	23
4.2.2	Memory Management Unit	24
4.2.3	Instruction and Data Memory	24
4.2.4	ALU Design	24
4.2.5	Decoder Design	26
4.2.6	Pipelining	26
4.2.7	Design Optimisations	27
4.3	Interrupts	27
4.3.1	Overview	28
4.3.2	Hardware Implementation	28
4.3.3	Software Interface	29
4.3.4	Design Improvements	30
4.4	Verification	31

4.1 Introduction

While the majority of this report will focus on the multi-processing functionality of this project, it is important to understand the design decisions of the single core to understand the features and limitations of the multi-core system-on-chip as a whole.

4.2 Design and Implementation

The single-core design is a traditional 5-stage RISC processor (fetch, decode, execute, memory, write-back). The core uses separate instruction and data memories in the style of a Harvard architecture.

To satisfy [CD5](#), the Verilog code will be self-contained in a single file. This reduces the hierarchical complexity and eases cross-vendor project set-up as only a single file is required to be included.

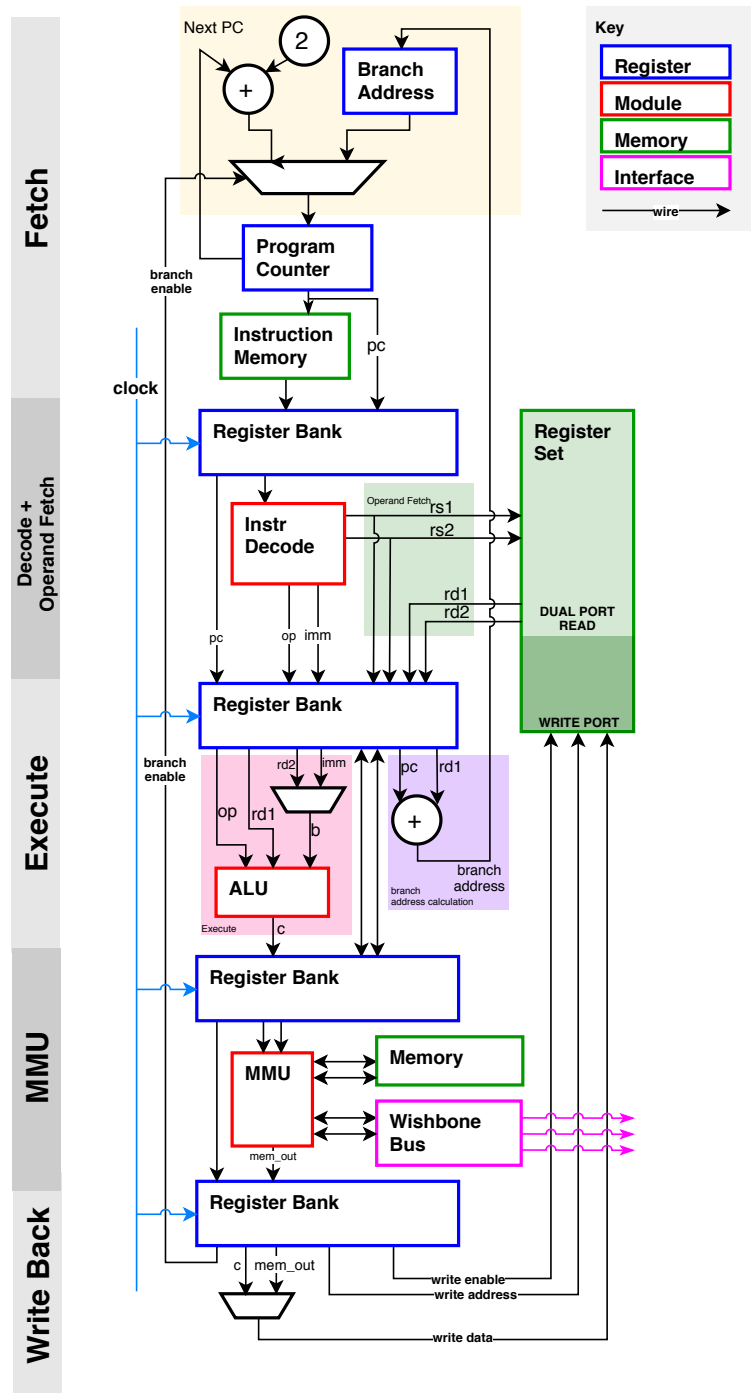


Figure 4.1: Vmicro16 RISC 5-stage RTL diagram showing: instruction pipelining (data passed forward through clocked register banks at each stage); branch address calculation; ALU operand calculation (rd2 or imm); and program counter incrementing.

As this is a multi-core system, a small reduction in size within the single-core will result in substantial size reductions for designs with many cores.

4.2.1 Instruction Set Architecture

Core deliverable [CD1](#) details the background for the requirement of a custom instruction set architecture. The 16-bit instruction set listing is shown in [Figure B.2](#).

In this proposed architecture, most instructions are *destructive* meaning that source operands

also act as the destination, hence effectively *destroying* the original source data. This design decision reduces the complexity of the ISA as traditional three operand instructions, for example `add r0, r0, r1`, can be encoded using only two operands `add r0, r1`. However, this does increase the complexity of compilers as they may need to make temporary copies of registers as the instructions will *destroy* the original source data.

The instruction set is split into 7 categories (highlighted by colours in [Figure B.2](#)):

- Special instructions, such as halting and interrupt returns;
- Bitwise operations, such as XOR and AND;
- Signed arithmetic;
- Unsigned arithmetic;
- Conditional branches and compare instructions;
- and Load/store instructions, with their atomic equivalents.

4.2.2 Memory Management Unit

It was decided to use a memory management unit (MMU) to make it easier and extensible to communicate with external peripherals or additional registers. This method transparently uses the existing `LW[EX]` / `SW[EX]` to easily provide an arbitrary number of peripherals/special purpose addresses to the software running on the processor.

4.2.3 Instruction and Data Memory

The design uses separate instruction and data memories similar to a Harvard architecture computer. This architecture was chosen due because it is generally easier to implement, however later resulted in design challenges in large multi-core designs. This is discussed later in the report.

Each single-core has it's own *scratch* memory – a small RAM-like memory which can be used for stack-space and arrays too large to fit into the 8 registers. These memories are provided as is – meaning it's up to the software to implement and provide any stack-frame, function, and calling, functionality. Each core also features it's own read-only instruction memory that is programmed at compile time of the design, or via the UART0 receiver interface (discussed later). Both of these memories map onto synchronous, read-first, single-port, FPGA block RAMs to minimise LUT requirements.

Users can customise the size of these memories by tweaking the following parameters in the `vmicro16_soc_config.v` file: `DEF_MEM_INSTR_DEPTH` for the instruction memory, and `DEF_MEM_SCRATCH_DEPTH` for the scratch memory.

4.2.4 ALU Design

The Vmicro16's ALU is an asynchronous module that has 3 inputs: data a; data b; and opcode op; and outputs data c. The ALU is able to operate on both register data (`rd1` and `rd2`) and

immediate values. A switch is used to set the *b* input to either the *rd2* or *imm* value from the previous stage.

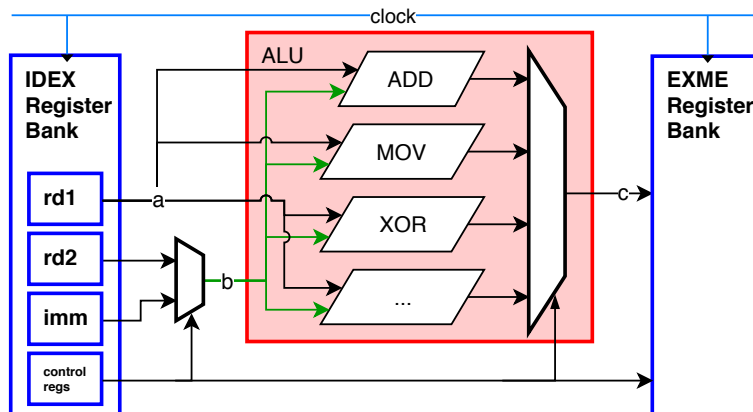


Figure 4.2: Vmicro16 ALU diagram showing clocked inputs from the previous IDEX stage being

The ALU also performs comparison (CMP) operations in which it returns flags similar to X86's overflow, signed, and zero, flags. The combination of these flags can be used to easily compute relationships between the two input operands. For example, if the zero flag is not equal to the signed flag, then the relationship between inputs *a* and *b* is that $a < b$.

```

1  module branch (
2      input [3:0] flags,
3      input [7:0] cond,
4      output reg en
5  );
6      always @(*)
7          case (cond)
8              `VMICRO16_OP_BR_U: en = 1;
9              `VMICRO16_OP_BR_E: en = (flags[`VMICRO16_SFLAG_Z] == 1);
10             `VMICRO16_OP_BR_NE: en = (flags[`VMICRO16_SFLAG_Z] == 0);
11             `VMICRO16_OP_BR_G: en = (flags[`VMICRO16_SFLAG_Z] == 0) &&
12                 (flags[`VMICRO16_SFLAG_N] == flags[`VMICRO16_SFLAG_V]);
13             `VMICRO16_OP_BR_L: en = (flags[`VMICRO16_SFLAG_Z] != flags[`VMICRO16_SFLAG_N]);
14             `VMICRO16_OP_BR_GE: en = (flags[`VMICRO16_SFLAG_Z] == flags[`VMICRO16_SFLAG_N]);
15             `VMICRO16_OP_BR_LE: en = (flags[`VMICRO16_SFLAG_Z] == 1) ||
16                 (flags[`VMICRO16_SFLAG_N] != flags[`VMICRO16_SFLAG_V]);
17             default: en = 0;
18         endcase
19     endmodule

```

Listing 1: ALU branch detection using flags: zero (Z), overflow (V), and negative (N).

The Verilog implementation of the ALU is shown in Listing 2. The ALU's asynchronous output is clocked with other registers, such as destination register *rs1* and other control signals, in the EXME register bank.

```

1  always @(*) case (op)
2      // branch/nop, output nothing
3      `VMICRO16_ALU_BR,
4      `VMICRO16_ALU_NOP:      c = {DATA_WIDTH{1'b0}};
5      // load/store addresses (use value in rd2)
6      `VMICRO16_ALU_LW,
7      `VMICRO16_ALU_SW:      c = b;
8      // bitwise operations
9      `VMICRO16_ALU_BIT_OR:      c = a | b;
10     `VMICRO16_ALU_BIT_XOR:      c = a ^ b;
11     `VMICRO16_ALU_BIT_AND:      c = a & b;
12     `VMICRO16_ALU_BIT_NOT:      c = ~(b);
13     `VMICRO16_ALU_BIT_LSHFT:    c = a << b;
14     `VMICRO16_ALU_BIT_RSHFT:    c = a >> b;

```

Listing 2: Vmicro16's ALU implementation named vmicro16_alu. vmicro16.v

4.2.5 Decoder Design

Instruction decoding occurs in the between the IFID and IDEX stages. The decoder extracts register selects and operands from the input instruction. The decoder outputs are asynchronous which allows the register selects to be passed to the register set and register data to be read asynchronously. The register selects and register read data is then clocked into the IDEX register bank.

```

1  always @(*) case (instr[15:11])
2      `VMICRO16_OP_BR:      alu_op = `VMICRO16_ALU_BR;
3      `VMICRO16_OP_MULT:     alu_op = `VMICRO16_ALU_MULT;
4
5      `VMICRO16_OP_CMP:      alu_op = `VMICRO16_ALU_CMP;
6      `VMICRO16_OP_SETC:     alu_op = `VMICRO16_ALU_SETC;
7
8      `VMICRO16_OP_BIT:      casez (instr[4:0])
9          `VMICRO16_OP_BIT_OR:      alu_op = `VMICRO16_ALU_BIT_OR;
10         `VMICRO16_OP_BIT_XOR:      alu_op = `VMICRO16_ALU_BIT_XOR;
11         `VMICRO16_OP_BIT_AND:      alu_op = `VMICRO16_ALU_BIT_AND;
12         `VMICRO16_OP_BIT_NOT:      alu_op = `VMICRO16_ALU_BIT_NOT;
13         `VMICRO16_OP_BIT_LSHFT:    alu_op = `VMICRO16_ALU_BIT_LSHFT;
14         `VMICRO16_OP_BIT_RSHFT:    alu_op = `VMICRO16_ALU_BIT_RSHFT;
15     default:                alu_op = `VMICRO16_ALU_BAD; endcase

```

Listing 4: Vmicro16's decoder module code showing nested bit switches to determine the intended opcode. vmicro16.v

In Listing 4, it can be seen that the first 4 opcode cases (BR, MULT, CMP, SETC) are represented using the same 15-11 (opcode) bits, however the BIT instructions share the same opcode and so require another bit range to be compared to determine the output function.

4.2.6 Pipelining

In the interim progress update, the processor design featured *instruction pipelining* to meet requirement ED1. Instruction pipelining allows instructions executions to be overlapped in the pipeline, resulting in higher throughput (up to one instruction per clock) at the expense of 5-6 clocks of latency and *significant* code complexity. As the development of the project shifted from single-core to multi-core, it became obvious that the complexity of the pipelined processor would inhibit the integration of the multi-core functionality. It was decided to remove the instruction pipelining functionality and use a simpler state-machine based pipeline that is much simpler to extend and would cause fewer challenges later in the project.

4.2.7 Design Optimisations

In a design that has many instantiations of the same component, a small resource saving improvement within the component can have a significant overall savings improvement if it is instantiated many times. Project requirement [CD5](#) requires the design to be compiled for a range of FPGA sizes, and so space saving optimisations are considered.

Register Set Size Improvements

A register set in a CPU is a fast, temporary, and small memory that software instructions directly manipulate to perform computation. In the Vmicro16 instruction set, eight registers named r0 to r7 are available to software. The instruction set allows up to two registers to be references in most instructions, for example the instruction `add r0, r1` tells the processor to perform the following actions:

- Clock 1.** Fetch r0 and r1 from the register set
- Clock 2.** Add the two values together in the ALU
- Clock 3.** Store the result back the register set in r0

For Clock 1, it was originally decided to use a dual port register set (meaning that two data reads can be performed in a single clock, in this case r0 and r1), however due to the asynchronous design of the register set (for speed) the RTL produced consumed a significant amount of FPGA resources, approximately 256 flip-flops ($16 \text{ (data width)} * 8 \text{ (registers)} * 2 \text{ (ports)}$). To reduce this, it was decided to split task 1 into two steps over two clock cycles using a single-port register set. This required the processor pipe-line to use another clock cycle resulting in slightly lower performance, however the size improvements will allow for more cores to be instantiated in the design. This optimisation is also applied to the interrupt register set, resulting in a saving of approximately 256 flip-flops per core (128 in the normal mode register set, and 128 in the interrupt register set). As shown, adding a single clock delay saves a significant amount of LUTs. This saving will be amplified in designs with many cores.

4.3 Interrupts

Interrupts are a technique used by processors to run software functions when an event occurs within the processor, such as exceptions, or signalled from an external source, such as a UART receiver signalling it has received new data. Today, it is common for micro-controllers, soft-processors, and desktop processors, to all feature interrupts. Modern implementations support an *interrupt vector* which is a memory array that contains addresses to different *interrupt handlers* (a software function called when a particular interrupt is received).

Although interrupts are not a requirement for a multi-core system, it was decided to implement this functionality to boost my understanding of such systems. In addition, example demos provided with this project are better visualised with a interrupt functionality.

4.3.1 Overview

The interrupt functionality in this project supports the following:

- Per-core 8 cell interrupt vector accessible to software.
Software programs running on the Vmicro16 processor can edit the interrupt vector to add their own interrupt handlers at runtime.
- Fast context switching.
A dedicated interrupt register set is multiplexed with the normal mode register set to provide faster context switching. It should be noted that only the registers are saved during a context switch. This means that the stack is not saved. A schematic of the register multiplex is shown in [Figure B.1](#).
- Parametrised interrupt sources and widths.
Users can configure the width of the interrupt in signals and the data width per interrupt source via the `vmicro16_soc_config.v`. By default, 8 interrupt sources are available and each can provide 8-bits of data.

4.3.2 Hardware Implementation

Context Switching

When acting upon an incoming interrupt the current state the processor must be saved so that changes from the interrupt handler, such as register writes and branches, do not affect the current state. After the interrupt handler function signals it has finished (by using the *Interrupt Return* INTR instruction) the saved state is restored. In the case of the Vmicro16 processor, the program counter `r_pc[15:0]` and register set `regs` instance are the only states that are saved. Going forth, the terms *normal mode* and *interrupt mode* are used to describe what registers the processor should use when executing instructions.

When saving the state, to avoid clocking 128 bits (8 registers of 16 bits) into another register (which would increase timing delays and logic elements), a dedicated register set for the interrupt mode (`regs_isr`) is multiplexed with the normal mode register set (`regs`). Then depending on the mode (identified by the register `regs_use_int`) the processor can easily switch between the two large states without significantly affecting timing.

The timing diagram in [Figure 4.3](#) shows the behavioural logic for the TIMR0 interrupt source.



Figure 4.3: Time diagram showing the TIMR0 peripheral emitting a 1us periodic interrupt signal (out) to the processor. The processor acknowledges the interrupt (int_pending_ack) and enters the interrupt mode (regs_use_int) for a period of time. When the interrupt handler reaches the Interrupt Return instruction (indicated by w_intr) the processor returns to normal mode and restores the normal state.

4.3.3 Software Interface

A memory-mapped software interface is provided through the MMU to allow easy software control of the interrupt behaviour. The interface is provided at the address range 0x0100 to 0x0108. This interface is per-core allowing each core to individually control what interrupts it receives and what functions to call upon an interrupt. This enables complex functionality, such as allowing each core to execute different functions upon the same interrupt.




Figure 4.4: The interrupt vector (0x0100 - 0x0107) consists of eight 16-bit values that point to memory addresses of the instruction memory to jump to.

Interrupt Vector (0x0100-0x0107)

The interrupt vector is a per-core register that is used to store the addresses of interrupt handlers. An interrupt handler is simply a software function residing in instruction memory that is branched to when a particular interrupt is received.

Interrupt Mask (0x0108)

The interrupt mask is a per-core register that is used to mask/listen specific interrupt sources. This enables processing cores to individually select which interrupts they respond to. This allows for multi-processor designs where each core can be used for a particular interrupt




Figure 4.5: Interrupt Mask register (0x0108). Each bit corresponds to an interrupt source. 1 signifies the interrupt is enabled for/visible to the core. Bits [7:2] are left to the designer to assign. Bit 0 is assigned to TIMR0's interval timer. Bit 1 is assigned to the UART0's receiver (unassigned if DEF.USE_REPROG is enabled).

source, improving the time response to the interrupt for time critical programs. The Interrupt Mask register is an 8-bit read/write register where each bit corresponds to a particular interrupt source and each bit corresponds with the interrupt handler in the interrupt vector. The interrupt mask register is shown in [Figure 4.5](#).

Software Example

To better understand the usage of the described interrupt registers, a simple software program is described below. The following software program produces a simple and power efficient routine to initialise the interrupt vector and interrupt mask.

```

1  setup_interruptions:
2      // Set interrupt vector at 0x100
3      // Move address of isr0 function to vector[0]
4      movi    r0, isr0
5      // create 0x100 value by left shifting 1 8 bits
6      movi    r1, #0x1
7      movi    r2, #0x8
8      lshft   r1, r2
9      // write isr0 address to vector[0]
10     sw      r0, r1
11
12  enable_interruptions:
13     // enable all interrupts by writing 0x0f to 0x108
14     movi    r0, #0x0f
15     sw      r0, r1 + #0x8 // (0x100 + 0x8 = 0x108)
16     halt
17     // enter low power idle state
18
19  isr0:
20     movi    r0, #0xff // arbitrary name
21     // do something
22     intr    // return from interrupt

```

A more complex example software program utilising interrupts and the TIMR0 interrupt is described in section [D.1](#).

4.3.4 Design Improvements

The hardware and software interrupt design have changed throughout the projects cycle. In initial versions of the interrupt implementation, the software program, while waiting for an interrupt, would be in a tight infinite loop (branching to the same instruction). This resulted in the processor using all pipeline stages during this time. The pipeline stages produce many logic transitions and memory fetches which raise power consumption and temperatures. This is quite noticeable especially when running on the Spartan-6 LX9 FPGA.

To improve this, it was decided to implement a new state within the processor's state machine that, when entered, did not produce high frequency logic transitions or memory

fetches. The HALT instruction was modified to enter this state and the only way to leave is from an interrupt or top-level reset. This removes the need for a software infinite loop that produces high frequency logic transitions (decoding, ALU, register reads, etc.) and memory fetches.

4.4 Verification

Various verification techniques are employed to ensure correct operation of the processor.

The first technique involves using static assertions to identify incorrect configuration parameters at compile time, such as having zero instruction memory and scratch memory depth. These assertions use the `static_assert` for top level checks and `static_assert_ng` for checks inside generate blocks.

The second verification technique is to use assertions in always blocks to identify incorrect behavioural states. This is done using the `rassert` (run-time assert) macro.

The third verification technique is to use automatic verifying test benches. These test benches drive components of the processor, such as the ALU and decoder, and check the output against the correct value. This uses the `rassert` macro.

The final method of verification is to verify the complete design via a behavioural test bench. The design is passed a compiled software program with a known expected output, and is ran until the `r_halt` signal is raised. The test bench then checks the value on the `debug0`, `debug1`, and `debug2` signals against the expected value. If this matches, then it is assumed that sub-components of the design also operate correctly. This technique does not monitor the states of sub-components and statistics (such as time taken to execute an instruction), there leaves the possibility that some components could have entered an illegal state.

Chapter 5

Interconnect

5.1	Introduction	32
5.1.1	Comparison of On-chip Buses	32
5.2	Overview	33
5.2.1	Design Considerations	34
5.3	Interfaces	35
5.3.1	Multi-master Support	36
5.4	Further Work	37

5.1 Introduction

The Vmicro16 processor needs to communicate with multiple peripheral modules (such as UART, timers, GPIO, and more) to provide useful functionality for the end user.

Previous peripheral interface designs of mine have been directly connected to a main driver with unique inputs and outputs that the peripheral required. For example, a timer peripheral would have dedicated wires for it's load and prescaler values, wires for enabling and resetting, and wires for reading. A memory peripheral would have wires for it's address, read and write data, and a write enable signal. This resulted in each peripheral having a unique interface and unique logic for driving the peripheral, which consumed significant amounts of limited FPGA resources.

It can be seen that many of the peripherals need similar inputs and outputs (for example read and write data signals, write enables, and addresses), and because of this, a standard interface can be used to interface with each peripheral. Using a standard interface can reduce logic requirements as each peripheral can be driven by a single driver.

5.1.1 Comparison of On-chip Buses

The choice of on-chip interconnect has changed multiple times over the life-cycle of this project, primary due to ease of implementation and documentation quality. Originally, it was planned to use the Wishbone bus [14] due to it's popularity within open-source FPGA modules, primarily available from opencores.org, and good quality documentation. This

choice would enable easier integration with common peripherals, such as SRAM, UARTs, and timers.

Late in the project, it was decided to use the AMBA APB protocol [15] as it is more commonly used in large commercial designs and understanding how the interface worked would better benefit myself. APB describes an intuitive and easy to implement 2-state interface aimed at communicating with low-throughput devices, such as UARTs, timers, and watchdogs.

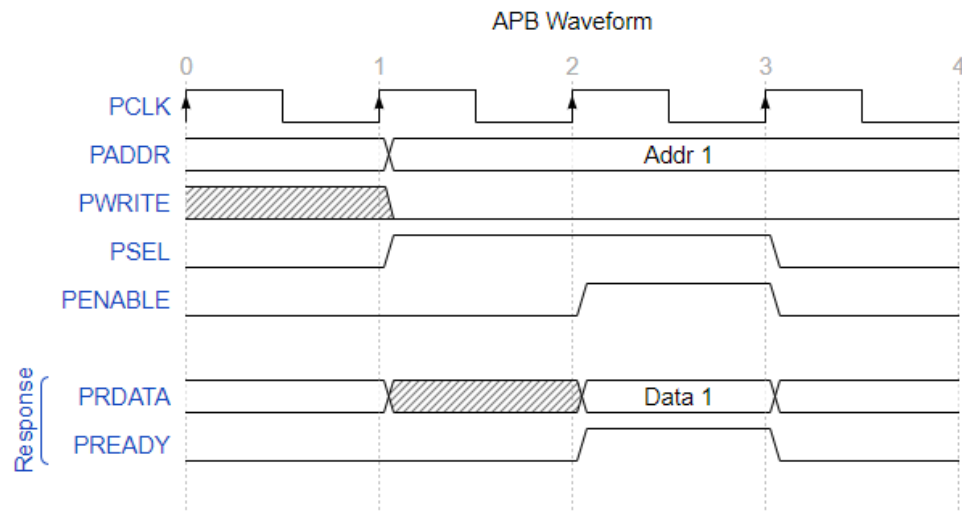


Figure 5.1: Waveform showing an APB read transaction.

Although AMBA AHB-Lite [5] might be better suited for processor-to-processor and processor-to-memory transactions due to its higher throughput and data bursts, the complexity of the interconnect is significantly greater and implementing it from scratch is out of scope for this project's timeline.

5.2 Overview

The system-on-chip design is split into 3 main parts: peripheral interconnect (red), CPU array (gray), and the instruction memory interconnect (green).

A block diagram of this project is shown in Figure 5.2

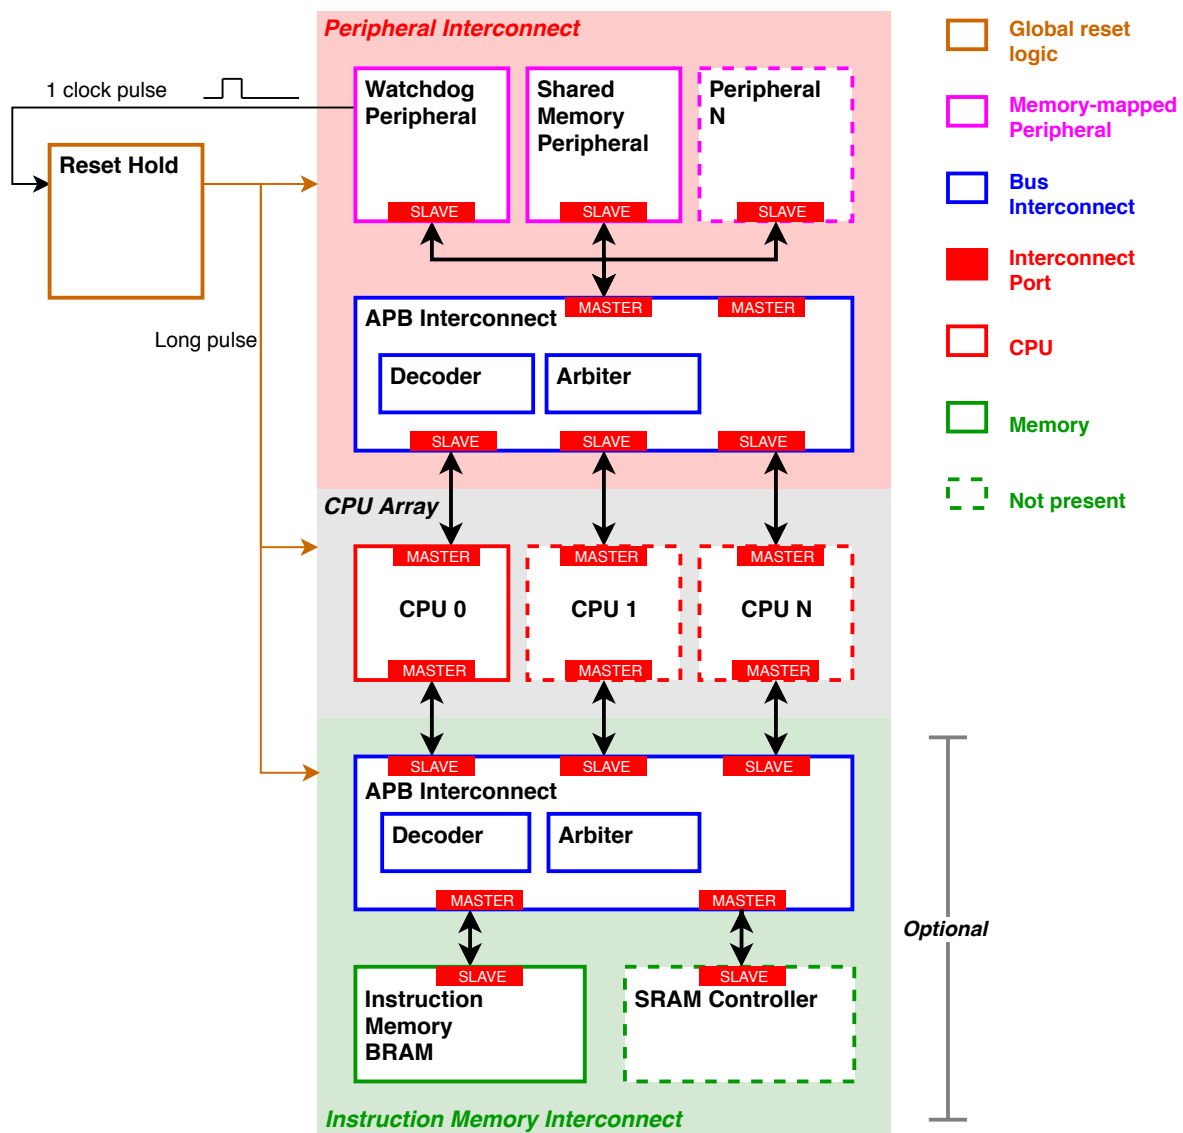


Figure 5.2: Block diagram of the Vmicro16 system-on-chip.

5.2.1 Design Considerations

There are several design issues to consider for this project. These are listed below:

- **Design size limitations**

The target devices for this project are small to medium sized FPGAs (featuring approximately 10,000 to 30,000 logic cells). Because of this, it is important to use a bus interconnect that has a small logic footprint yet is able to scale reasonably well.

- **Ease of implementation**

The interconnect and any peripherals should be easy to implement within the time allocations specified in Figure 3.1.

- **Scalable**

The interconnect should allow for easy scalability of master and slave interfaces with minimal code changes.

5.3 Interfaces

The master and slave interface ports shown in Figure 5.2 are each similar to the AMBA AHB schematic shown in Figure 5.3. This project's interconnect bus requires multi-master support which the figure does not show.

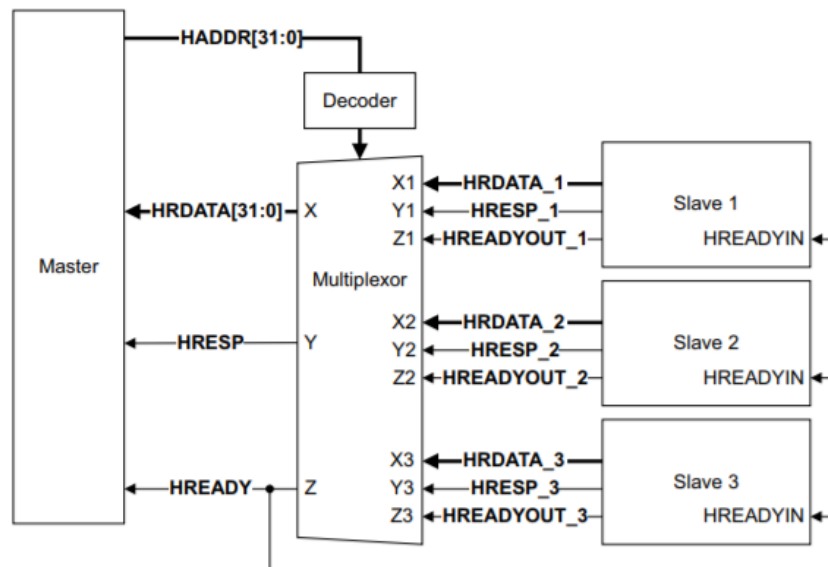


Figure 5.3: Source: [5]

In our multi-master design, the interface between masters and slaves use the signals shown in Figure 5.4. The interface extends the AMBA APB interface to provide more information regarding the owner of the transaction and any special operations.

20	19	18	17	16	15	0		
LE	SE	CORE_ID	Address				PADDR[20:0]	
			Write data				PWDATA[15:0]	
			Read Data				PRDATA[15:0]	
							WE	PWRITE[0:0]
								PSELx[0:0]
							EN	PENABLE[0:0]
								PREADY[0:0]

Figure 5.4: Vmicro16 master/slave interface using AMBA APB

This interface is designed to be backwards compatible with existing APB interfaces – the additional data parts are after the MSB of the data width. The CORE_ID parameter is the unsigned numerical representation of the core's ID that originated the transaction. SE, store exclusive, and LE, load exclusive, are two parameters used to distinguish between the SW/LW instructions and their exclusive counterparts, as they share the same bus. These extra parameters shown in Figure 5.4 (from [20:16]) are only used by the global shared memory peripheral to provide memory exclusivity. The decoding of the LE and SE parameters are shown in Table 5.1.

LE	SE	Operation
0	0	LW/SW depending on PWRITE
0	1	SWEX
1	0	LWEX
1	1	Not valid

Table 5.1: Decoding of the LE and SE parameters to determine the global shared memory operation.

5.3.1 Multi-master Support

In this design, each processor can act as an APB master to communicate with peripherals, for example to write a value to UART or to the shared memory peripheral. Because each core runs independently from other cores, it is likely, especially in many-core systems, that two or more processors will want to use the peripheral bus at the same time.

As the peripheral and instruction interconnects use a shared one-to-many (one master to many slaves) bus architecture, only one master can use the bus at any-time. To enable multiple masters to use the bus, a device called an *arbiter* must be used to control which master gets access to drive the shared interconnect.

Arbiters can vary in complexity, mostly relative to throughput requirements.

An ideal arbiter for this interconnect, which ideally features many, possibly tens of, high-throughput masters, would likely feature a priority-based and pipelined arbiter with various devices to improve performance such as cache-coherencies.

Overview

Due to this project's limited time, and my personal knowledge in this area, a simple rotating arbiter is used. This arbitration scheme is likely the simplest that can be thought of. A schematic of arbiter interconnect is shown in [Figure 5.5](#).

In this scheme, access to the bus is given incrementally to each master port, even if the master port has not requested to use the bus. The active master port can use the bus for as long as it requires, and signals it has finished by lowering the PSEL signal. When the PSEL signal is lowered, the arbiter grants access to the next master port. If this next master port has not raised its PSEL signal (i.e. it has not requested access to the bus) then the arbiter grants access to the next master port, and so on. In Verilog, this is simply an incremental counter which is used to index the master ports array. To support a variable number of master ports, the width of each APB signal is multiplied by the number of cores, as shown in [Listing 7](#).

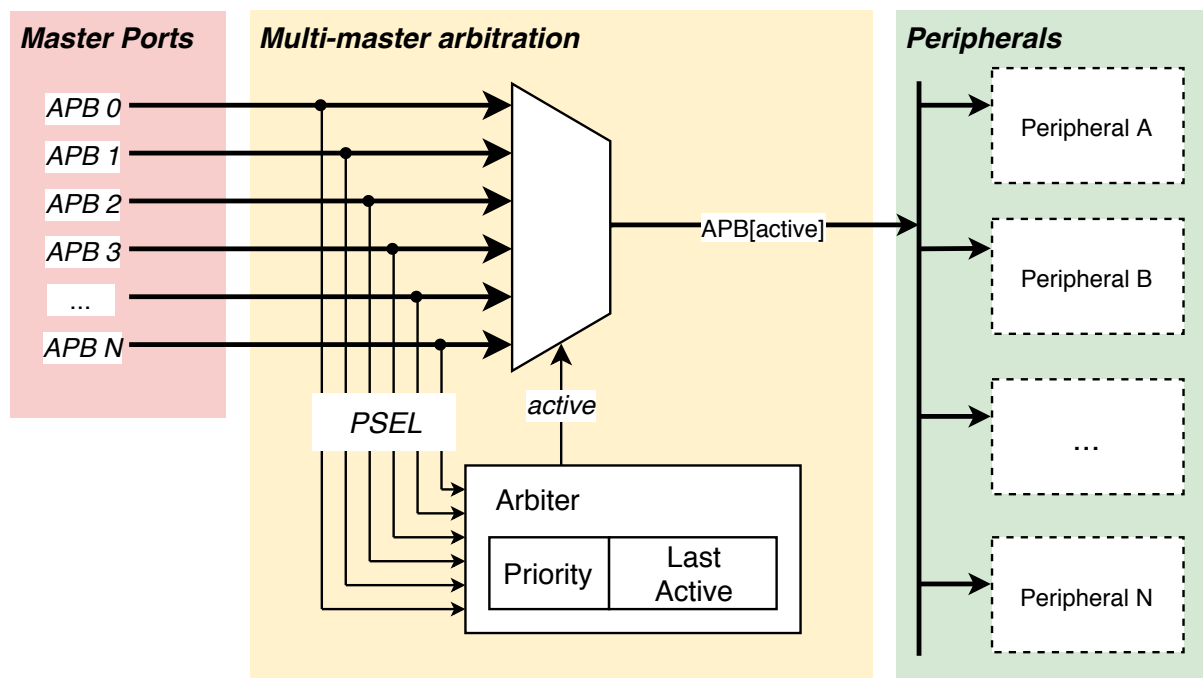


Figure 5.5: Multi-master schematic for the Vmicro16 system-on-chip.

5.4 Further Work

The submitted design is acceptable for a multi-core system as it fulfils the following requirements:

- Support an arbitrary number of peripherals.
- Supports memory-mapped address decoding.
- Supports multiple master interfaces.

Arbiter Performance Improvements

However, it fails in the performance aspect. A one clock penalty occurs if the next master port has not requested the bus. This may seem a small price to pay for such a simple arbiter design, however it can add up significantly in many-core designs. For example, if core #0 performs some action on the bus, but core #10 is the next master that wants to use the bus, then the arbiter will waste time incrementally granting access to cores #1 to #9 which do not need the bus. This is also made worse when one of the cores is blocking access to a peripheral resource, such as through a mutex or semaphore.

To overcome this penalty, a scheme could use an algorithm to find the next master port requesting access, and grant access directly to it when the current master has finished. Another scheme could be to use a priority encoder. Here, a hard-coded lookup table (LUT) could be used, where the inputs are each master port's PSEL signal (acting as a bus request line) and the output being which master to grant access to. As this is targeting FPGA devices, this implemented would require few LUT resources for the arbiter, due to the hard-coded LUT

approach. An example of this is given in M. Weber's *Arbiter: Design Ideas and Coding Styles* [16, p. 2].

APB Bus Errors and Recovery

This project's implementation of a multi-master APB interconnect does not provide a method of detecting errors and stalls. This is mainly due to time constraints.

An easy error that could be detected is PADDR addresses that do not fall into a memory-mapped address range. This can easily and cheaply be detected in the address decoding module. This will be discussed in detail in the next chapter.

As previously stated, the active bus master can take control of the bus for as long as it wants to. This is useful for high-throughput transactions, such as memory operations to global memory, but detecting a stalled or glitched operation is not immediately identifiable. If an active master stalls or glitches, it may not be able to lower the PSEL line which appears to the arbiter that the transaction is still happening normally. To overcome this, a timer could be used to detect stalled operations and reset the affected peripheral (essential a watchdog but for an interconnect).

Chapter 6

Memory Mapping

6.1	Introduction	39
6.2	Address Decoding	39
6.2.1	Decoder Optimisations	40
6.3	Memory Map	42

The Vmicro16 processor uses a memory-mapping scheme to communicate with peripherals and other cores. This chapter describes the design decisions and implementation of the memory-map used in this project.

6.1 Introduction

Memory mapping is a common technique used by CPUs, micro-controllers, and other system-on-chip devices, that enables peripherals and other devices to be accessed via a memory address on a common bus. In a processor use-case, this allows for the reuse of existing instructions (commonly memory load/store instructions) to communicate with external peripherals with little additional logic.

6.2 Address Decoding

An address decoder is used to determine the peripheral that the address is requesting. The address decoder module, `addr_dec` in `apb_intercon.v`, takes the 16-bit `PADDR` from the active APB interface and checks for set bits to determine which peripheral to select. The decoder outputs a chip enable signal `PSEL` for the selected peripheral. For example, if bit 12 is set in `PADDR` then the shared memory peripheral's `PSEL` is set high and others to low. A schematic for the decoder is shown in [Figure 6.1](#).

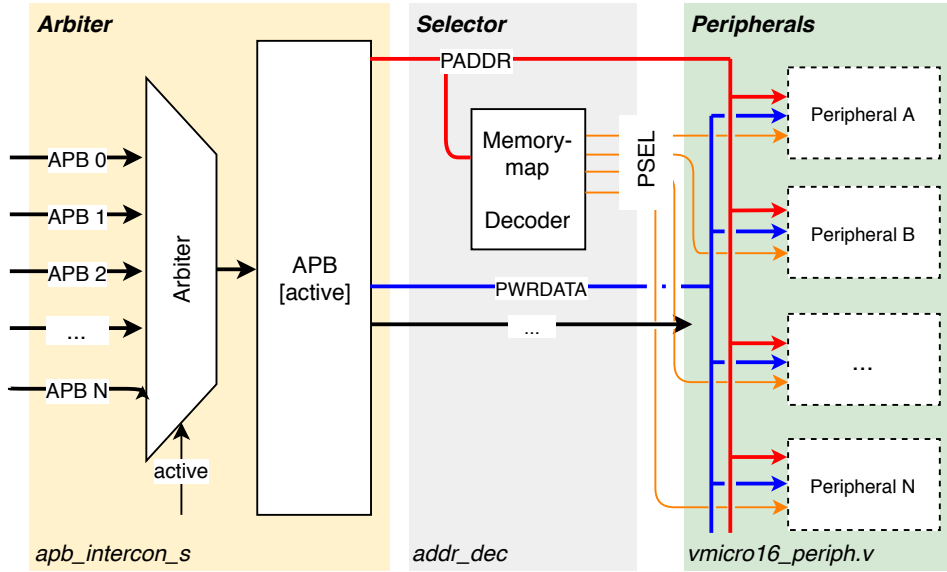


Figure 6.1: Schematic showing the address decoder (`addr_dec`) accepting the active `PADDR` signal and outputting `PSEL` chip enable signals to each peripheral.

6.2.1 Decoder Optimisations

Performing a 16-bit equality comparison of the `PADDR` signal against each peripheral memory address consumes a significant amount of logic. Depending on the synthesis tools and FPGA features, a 16-bit comparator might require a fixed 16-bit value input to compare against (where the 0s are inverted) and a wide-AND to reduce and compare [17, 18]. An example 4-bit comparator is shown below in Figure 6.2.




Figure 6.2: Example 4-bit binary comparator which compares the bits (a, b, c, d) to the constant value 1010. The 0s of the constant are inverted and then all are passed to a wide-AND.

As we are targeting FPGAs, which use LUTs to implement combinatorial logic, we can conveniently utilise Verilog's `==` operator on fairly large operands without worrying about consuming too many resources. The targeted FPGA devices in this project, the Cyclone V and Spartan 6, feature 6-input LUTs which allow 64 different configurations [19, 20]. Knowing this, we can design the address decoder to utilise the FPGA's LUTs more effectively and reduce its footprint significantly.

We can use part of the `PADDR` signal as a chip select and the other bits as sub-addresses to interface with the peripheral. The addressing bits are passed into the FPGA's 6-input LUTs which are programmed (via the bitstream) to output 1 or 0 depending on the address. Figure 6.3 below shows a LUT based approach to address decoding which will utilise approximately one ALM/CLB module per peripheral chip select (`PSEL`) and one for error detection. This method

of comparison (LUT based) is utilised in the `addr_dec` module in `apb_intercon.v`.




Figure 6.3: Bits [7:3] of an 8-bit PADDR signal are used as inputs to 5-bit LUTs to generate a PSEL signal. In addition, a default error case is shown allowing the address decoder to detect incorrect PADDR values (e.g. if no PSEL signals are generated).

The address decoding methods discussed above are examples of *full-address* decoding, where each bit (whether required or not) is compared. It is possible to further reduce the required logic by utilising *partial-address* decoding [21]. Partial-address decoding can reduce logic requirements by not using all bits. For example, if bits in address 0x0100 do not conflict with bits in other addresses (i.e. bit 8 is high in more than 1 address), then the address decoder needs only concern bit 8, not the other bits. This is visualised in Figure 6.4 below. This method is utilised in the MMU's address decoder (module `vmicro16_mmu` in `vmicro16.v`:181). As this is an optimisation per core, significant resources can be saved when a large number of cores are used.




Figure 6.4: Partial address decoding used by the Vmicro16 SoC design. Each peripheral shown only needs to decode a signal bit to determine if it is enabled.

6.3 Memory Map

The system-on-chip's memory map is shown below in [Figure 6.5](#). The addresses for each peripheral have been carefully chosen for both:

- Easy software access – creating addresses via software requires few instructions (normally one to four MOVI and LSHIFT instructions to address 0x0000 to 0xffff), which increases software performance.
- and Reducing address decoding logic – most addresses can be decoded using partial decoding techniques.

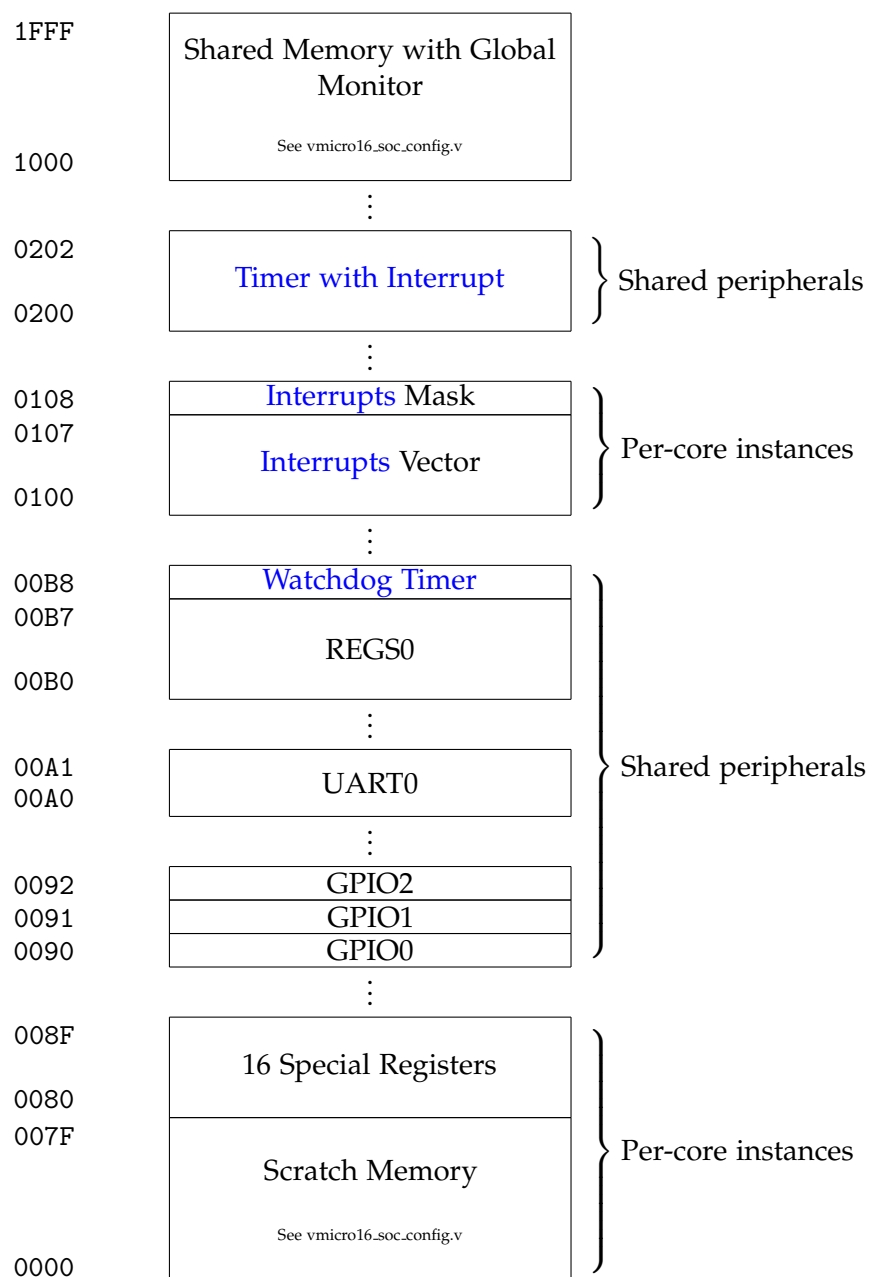


Figure 6.5: Memory map showing addresses of various memory sections.

Chapter 7

Multi-core Communication

7.1	Introduction	43
7.1.1	Design Goals	43
7.1.2	Context Identification	43
7.1.3	Thread Synchronisation	44

7.1 Introduction

So far we have discussed the features and design of the Vmicro16 system-on-chip. This section will discuss the multi-processing functionality that is critical to the goals of the project, and how to use it from software.

7.1.1 Design Goals

- **Support common synchronisation primitives.**

Software should be able to implement common synchronisation primitives, such as mutexes, semaphores, and memory barriers, to perform atomic operations and avoid race conditions, which are critical in parallel and concurrent software applications.

- **Context identification.**

The SoC should expose configuration information such as: the number of processing cores, amount of shared and scratch memory, and the `CORE_ID`, to each thread.

7.1.2 Context Identification

A goal of the multi-processing functionality of this project is allow software written for it to be run on any number of cores. This means that a software program will scale to use all cores in the SoC without needing to rewrite the software. To enable this functionality, the software must be able to read contextual information about the SoC, such as the number of cores, how much global and scratch memory is available, and what the `CORE_ID` of the current core is.

This information is provided through the Special Registers peripheral (0x0080 - 0x008F), shown in [Figure 7.1](#). This register set provides relevant information for writing software that can dynamically scale for various SoC configurations.

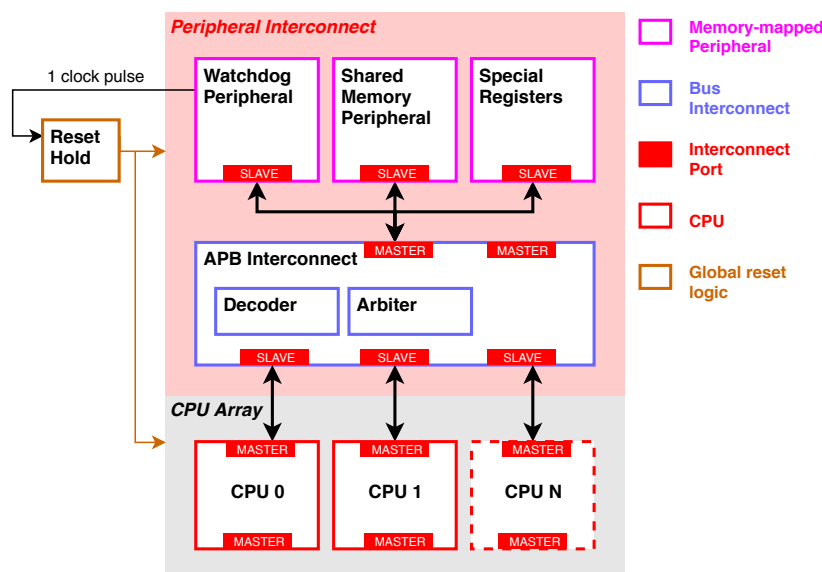


Figure 7.1: Block diagram showing the main multi-processing components: the CPU array and a peripheral interconnect used for core synchronisation.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
								CORE_ID								0080	R
								NUM_CORES								0081	R
SHARED_MEMORY cells (default 4096)																0082	R
								NUM_PERIPHERALS								0083	R
SCRATCH_MEMORY cells (default 64)																0084	RW
User defined																0085	RW
⋮																	
User defined																008F	RW

Figure 7.2: Vmicro16 Special Registers layout (0x0080 - 0x008F).

7.1.3 Thread Synchronisation

Mutexes

In software, a mutex is an object used to control access to a shared resource. The term *object* is used as its implementation is normally platform dependant, meaning that the processor may provide a hardware mechanism or is left for the operating system to provide.

In this project, mutexes are provided by the processor through the Shared Memory Peripheral (0x1000 to 0x1FFF) which provides a large RAM-style memory accessible by all cores through the peripheral interconnect bus. This large memory is explicitly defined to use the FPGA's BRAM blocks using Xilinx's Verilog `ram_style="block"` attribute to avoid wasting LUTs when using high core counts. The peripheral allows each memory cell to be *locked*, meaning that only the cell owner can modify its contents. This is implemented by using another large memory, *locks*, to store the `CORE_ID + 1` of the owner, as shown in [Listing 5](#). In this system, a lock containing the value 0 indicates an unlocked cell. As `CORE_ID`s are indexed

from zero, 1 is arithmetically added to each cell. For example, if core #2 wants to lock a memory cell, the value 3 is written to the lock.

```

1  reg [15:0]          ram    [0:8191]; // 16KB large RAM memory
2  reg [clog2(CORES):0] locks [0:8181]; // memory cell owner

```

Listing 5: RAM and lock memories instantiated by the shared memory peripheral.

To lock and unlock cells, the instructions LWEX and SWEX instructions are used. These instructions are similar to the LW/SW instructions but provide locking functionality. The *EX* in the instruction names indicate *exclusive access*. LWEX is used to read memory contents (like LW) and also lock the cell if not already locked. If a core attempts to lock an already locked cell, the lock does not change. Unlocking is done by the SWEX instruction, which conditionally writes to the memory cell if it is locked by the same core. Unlike SW, SWEX returns a zero for success and one for failure if it is locked by another core.

Figure 7.3 shows a simple assembly function to lock a memory cell.

```

1  lock_mutex:
2      // attempt lock
3      lwex r0, r1
4      // check success
5      swex r0, r1
6      cmp r0, r3
7      // if not equal (NE), retry
8      movi r4, lock_mutex
9      br r4, BR_NE
10 critical:
11     // core has the mutex

```

Figure 7.3: Assembly code for locking a mutex. r1 is the address to lock. r3 is zero. r4 is the branch address.

Barriers

Barriers are a useful software sequence used to block execution until all other threads (or a subset) have reached the same point. Barriers are often used for broadcast and gather actions (sending values to each core or receiving them). They are also used to synchronise program execution if some threads have more work to do than others.

The Vmicro16 processor provides barrier synchronisation through the Shared Memory

Peripheral. Like the mutex code, the barrier code uses the LWEX and SWEX instructions to lock a memory cell. Instead of immediately checking the lock as an abstract object, the barrier code treats the cell as a normal memory cell containing a numeric value. Listing 6 shows a software example of this. When the `barrier_reached` code is reached, the code will increment the shared memory value by 1, indicating that the number of threads that have reached this point has increased by one (r5). The `barrier_wait` function is then entered which waits until this numeric value (r5) is equal to the number of threads (r7) in the system. If this is true, then all threads have reached the `barrier_wait` function and can continue with normal program execution.

```
1  barrier_reached:
2      // load latest count
3      lwex    r0, r5
4      // try increment count
5      // increment by 1
6      addi    r0, r3 + #0x01
7      // attempt store
8      swex    r0, r5
9
10     // check success (== 0)
11     cmp      r0, r3
12     // branch if failed
13     movi     r4, barrier_reached
14     br       r4, BR_NE
15
16  barrier_wait:
17     // load the count
18     lw        r0, r5
19     // compare with number of threads
20     cmp       r0, r7
21     // jump back to barrier if not equal
22     movi     r4, barrier_wait
23     br       r4, BR_NE
```

Listing 6: Assembly code for a memory barrier. Threads will wait in the `barrier_wait` function until all other threads have reached that code point.

Chapter 8

Analysis & Results

8.1	Analysis	47
8.1.1	Design Area/Size Requirements	47
8.1.2	Maximum Frequency	49
8.2	Scenario Performance	49
8.2.1	Scenario Overview	50
8.2.2	Performance Measurements	50
8.2.3	Performance Results	51
8.2.4	Shared Instruction Memory Impact	52
8.3	Analysis Review	53

8.1 Analysis

So far the system's design, implementation, and example usage, has been presented and discussed. This chapter presents a critical discussion of the synthesised top-level implementation of the design and potential improvements.

8.1.1 Design Area/Size Requirements

On a minimal system-on-chip configuration, with one core and minimal peripherals and features (no reprogramming, no interrupts, no UART), the design requires as few as 700 LUTs with the processor core requiring approximately 300-400 LUTs.

Memory Constraints

As discussed in [Chapter 4 Single-core Design](#), each processor core features two memories: instruction and scratch memory, which can both map onto synchronous, single-port, FPGA BRAM blocks. While this will reduce LUT requirements in designs with few cores, it becomes a non-trivial problem as the core counts increase. FPGAs have a fixed number of hard-BRAM blocks available for inference by the HDL compiler, for example the low-end Xilinx Spartan-6 XC6SLX9 FPGA features 32 18 Kb BRAM blocks [22, p. 2], and the Cyclone V 5CSEMA5F31C6N (used in the DE1-SoC) has 397 10 Kb blocks [23, p. 22].




Figure 8.1: Per-core instruction memory schematic and performance.

As shown in Figure 8.1a, as the number of processor cores increases, they eventually outnumber the available BRAM blocks resulting in their memories being implemented in either distributed RAMs or ALMs, both of which can consume significant logic resources of the FPGA which reduces the maximum possible core count.

Figure 8.1b shows the FPGA resource requirements for the DE1-SoC board featuring the Cyclone V FPGA. Approximately 32 cores can be instantiated before the all the available registers and ALMs are consumed.

Reducing Memory Requirements

As shown in Figure 8.1a, each core has its own instruction read-only memory. These memories have identical contents which presents an opportunity for optimisation. In the proposed design in Figure 8.2a, this memory is removed from each core and is instead available through a dedicated shared bus. This approach can be configured to be used in the Vmicro16 SoC through the `DEF_CORE_HAS_INSTR_MEM` parameter in `vmicro16_soc_config.v`, which enables the *Instruction Memory Interconnect* shown in Figure 5.2.

As shown in Figure 8.2b, the resource requirements using this shared memory approach is significantly less than having an instruction memory per-core. On the DE1-SoC, 64 cores can now be instantiated with a few thousand regs and ALMs left for other logic.

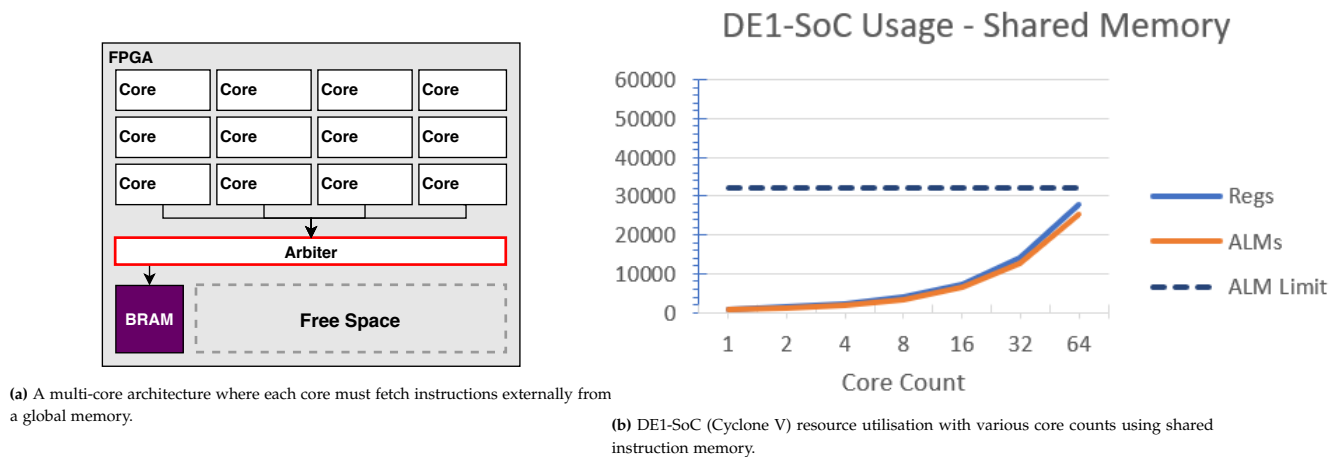


Figure 8.2: Shared instruction memory schematic and performance.

Whilst this is a significant resource saving opportunity, it does have significant drawbacks. In the shared instruction memory approach, each core must now fetch its instruction from the instruction memory interconnect which is subject to the arbiter and its scheduling algorithm. The arbiter uses the same algorithm as the peripheral interconnect arbiter meaning that cores receive access incrementally, and as discussed in [Section 5.3.1](#), this results in significant delays in many-core designs. This drawback is further explained in [Section 8.2.3](#).

8.1.2 Maximum Frequency

[Figure 8.3](#) below shows the maximum clock frequency for the design (F_{max}) on the Cyclone V on the DE1-SoC. As expected, having more cores results in more logic and thus a longer propagation delays to each core. System designers should consider the number of cores in their design as having fewer, faster, cores may outperform having more, slower, cores in some use-cases.




Figure 8.3: Cyclone V maximum design frequency for various core count configurations.

8.2 Scenario Performance

To evaluate the performance of the system-on-chip, scenarios encompassing computational problems that are reflective of real-world applications are compiled and ran on the design.

8.2.1 Scenario Overview

The scenario is a software program that runs a parallel implementation of the summation function, i.e. `sum [1..10]` which returns 55. While this may seem too simple at first to measure performance of a multi-core system-on-chip, the function is actually quite appropriate as it encompasses various parallel problems, such as: a fixed time/size serial part; broadcasting of the data set (in this case the range of the summation); thread synchronisation (to know when the data is ready and to schedule gathering of intermediary results); and is highly scalable.

The summation task flow is as follows:

1. Root (core #0) broadcasts the range of the summation (i.e. sum 1 to 10) to all cores via the global shared memory.
2. Non-root cores wait for this broadcast to finish (memory barrier), then calculate their own subset of the range to sum. For example, if Root broadcasts that there are 240 samples and 10 cores in the system, each core calculates the subset size:

$$240/10 = 24 \quad (8.1)$$

calculations starting from:

$$ID_{CORE} * 24 \quad (8.2)$$

For example, Core #5 will start its 24 sample subset summation from

$$5 * 24 = 120 \quad (8.3)$$

effectively performing `sum [120..123]`.

3. All cores perform an intermediary summation over their subset of the range (serial part).
4. All cores attempt to add their intermediary result to a global sum value in global shared memory (mutex).
5. All cores halt, signalling that their work has been committed to the global shared memory and have finished the program.

This program is written in assembly in the file `sw/demos/asm/sum64.s` and can be compiled using the assembly compiler (developed for deliverable [ED4](#)) using the command below. The assembly compiler outputs the file `asm.s.hex` containing hex instruction words for use in Verilog's `$readmemh` function. This data is used for each core's instruction memory. The assembly program is also shown in [Section D.2](#).

```
python sw/asm.py sw/demos/asm/sum64.s
```

8.2.2 Performance Measurements

Behavioural simulation is used to measure the following metrics to estimate general performance of the system-on-chip:

- Total program run-time.
This is the time from when the reset signal is de-asserted to when all cores have halted. Each core has an output halt signal which the SoC can use to determine if all cores have halted using `wire all_halted = &core_halts;`.
- Time spent on the serial part.
The serial part of this scenario consists of the intermediary summation of it's subset range. As each core is performing this task, the average will be used.
- Time spent on communication.
This includes time spent on thread synchronisation, i.e. waiting for the global memory to become available and waiting on the root to finish broadcast. Again, the average time will be used.
- Time spent fetching instructions.
Instruction fetches occur during stage STAGE_IF of the pipeline. The behavioural test bench will record the number of clock cycles each core spends in this state, then calculate the average time spent fetching instructions.

These measurements are recorded using non-synthesisable Verilog code in both the test-bench and module code (`vmicro16_soc.v`).

8.2.3 Performance Results

The scenario program was simulated on system configuration with 1 to 30 cores with a 50 MHz clock. Figure 8.4 shows the time breakdown of the multi-core system-on-chip running the scenario problem with various core counts. In these measurements all cores feature a small instruction memory which is accessed in constant time (one clock), and so this fetch time is not shown in the chart.




Figure 8.4: Chart showing how the communication times (T_{bus}) and serial times (T_{sum}) changes with core count.

The chart shows the expected shape for software parallel performance – as the number of cores increase, each core's problem space is reduced resulting in faster completion of the

summation part, and an increasing amount of time is spent on thread communication. This result matches my CUDA and MPI parallelism performance analysis results conducted in 2018 [24].

It can be seen that the total run-time is fastest near 8 cores and increases at this point when using more cores. This is likely due to the small summation range per core – with 30 cores just 8 samples per core are summed, which is extremely overkill and not representative of what a 30-core plus system should be used for.

If a much larger summation range was used (effectively representing a more appropriate scenario for systems with high core counts), the chart shape would stretch horizontally, resulting in the fastest time being on a higher core count than 8.

With high-core counts (16+) it can be seen that the communication time ($T_{bus.T}$) increases significantly. This is likely due to the rotating arbiter design. As discussed in [Section 5.3.1](#), the arbiter grants bus access to the next core incrementally after the previous core finishes. For large core systems, this can result in large time penalties. For example, if core #30 is blocking all other cores and the arbiter is lagging behind granting access to core #5, it will take a significant amount of time for the arbiter to reach core #30 to unblock the system.

8.2.4 Shared Instruction Memory Impact

As previous discussed in [Section 8.1.1](#), using a shared instruction memory approach reduces the size of the design but will result in increasingly long instruction fetches as the core count increases. [Figure 8.5](#) below shows the same scenario but using shared instruction memory. The shape of the chart remains similar, however the total run-time is close to double the time of the per-core instruction memory results, and increases significantly with more cores.




Figure 8.5: Similar to [Figure 8.4](#) but using shared instruction memory to reduce block memory requirements per core.

As with [Figure 8.4](#), using too many cores for a small data set can result in the scenario taking longer than a single core.

8.3 Analysis Review

There are several takeaways from these results:

1. Use an appropriate number of cores for the dataset size.
Too few cores result in longer work times and shorter communication times, and too many cores results in shorter work times and longer communication times.
2. Use an appropriate arbitration scheme to prevent blocking the system for too long.
In this design, and likely others, the blocking core is known by the global shared memory (via the locking cells) meaning that this information can be passed to the arbiter to give priority, while still avoiding deadlocks, to the blocking core.
3. Use an appropriate number of cores and still have space for other business logic.
4. More cores may result in lower clock frequencies.
From [Figure 8.3](#), the single-core design can be ran at ≈ 95 MHz while the 4-core design only ≈ 65 MHz (a $\approx 30\%$ decrease). The parallel speed improvements from having more cores may be less than a single fast core.

System designers should experiment with their algorithm and these takeaways to determine the approximate number of processor cores for their requirements, be that algorithm time, size, clock frequency, or compilation time.

Chapter 9

Conclusion

9.1 Overview	54
9.2 Review Against Project Deliverables	54
9.2.1 Core Deliverables	54
9.2.2 Extended Deliverables	55
9.3 Future Work	56

9.1 Overview

This project's primary goal was to explore and understand the design considerations of multi-core systems designing and implementing one from scratch. A 16-bit processor core was created and connected to a scalable multi-master bus interconnect to enable communication to peripherals. More processor cores were added and software was written to implement multi-threaded communication primitives, such as mutexes and memory barriers, to enable core-to-core communication.

9.2 Review Against Project Deliverables

This section will compare against the project deliverables initially outlined in [Chapter 3 Project Overview](#).

9.2.1 Core Deliverables

This project has achieved all core deliverables. A brief overview of each accomplishment is described below.

CD1 Design a compact 16-bit RISC instruction set architecture.

Achieved. A new 16-bit instruction set has been designed and implemented, nicknamed *Vmicro16* (Verilog microprocessor 16-bits).

CD2 Design and implement a Verilog RISC core that implements the ISA in [CD1](#).

Achieved. The implementation is found in `vmicro16.v`.

CD3 Design and implement an on-chip interconnect for multi-core processing (2 to 32 cores) using the RISC core from CD2.

Achieved. The AMBA APB bus was chosen and adapted for multi-master support. While the design is simple and lacks in performance, it does not consume a significant amount of area, leaving space for more cores and peripherals.

CD4 Analyse performance of serial and parallel software algorithms, such as parallel DFT, on the processor.

Achieved. A summation computation was chosen to demonstrate effectiveness in various parallel computational problems, such as subset calculation, thread communication, and broadcast/gather procedures.

CD5 Allow the RISC core to be easily compiled to multiple FPGA vendors (Xilinx, Altera).

Achieved. The design has been successfully compiled on Quartus Prime, Xilinx Vivado, Xilinx ISE, ModelSim, and Icarus Verilog. The design has been successfully implemented on the Cyclone V and Spartan-6 FPGA devices.

9.2.2 Extended Deliverables

ED1 Design a RISC core with an instructions-per-clock (IPC) rating of at least 1.0 (a single-cycle CPU).

Not achieved. Although present in the interim update, it was decided to remove this feature due to its complexity and possibility of it hindering multi-core and interrupt integration.

ED2 Design a RISC core with a pipe-lined data path to increase the design's clock speed.

Achieved. Data in the processor's pipeline is clocked through stages of flip-flops to reduce the critical path.

ED3 Design a scalable multi-core interconnect supporting arbitrary (more than 32) RISC core instances (manycore) using Network-on-Chip (NoC) architecture.

Not achieved. The implemented design uses traditional multi-core communication techniques (shared memory communication, mutexes, etc.) rather than a NoC architecture. While the NoC design would be more scalable, it would be of much higher complexity and therefore out-of-scope for this project's timeline.

ED4 Design a compiler-backend for the PRCO304 [12] compiler to support the ISA from CD1. This will make it easier to build complex multi-core software for the processor.

Somewhat achieved. While a compiler backend was produced that allowed high-level code to be compiled and run on the Vmicro16 processor, it lacked many multi-core and interrupt features. In addition, many unforeseen bugs were found in the original compiler, and it was decided instead to build and use a simpler text-assembly based compiler, `asm.py`.

ED5 The RISC core can communicate to peripherals via a memory-mapped addresses using the Wishbone bus.

Achieved. Although not Wishbone, AMBA APB was used instead as it's interface is more intuitive.

ED6 Implement various memory-mapped peripherals such as UART, GPIO, LCD, to aid visual representation of the processor during the demonstration viva.

Somewhat achieved. Modular UART, GPIO, watchdog, bus recovery, and memory peripherals were created to visually demonstrate features of the design. An LCD peripheral was not made due to time-constraints, however it would have been beneficial to do so to demonstrate conditional compilation for different development boards. A collection of software demos utilising these modules can be found in [Appendix D](#).

ED7 Store instruction memory in SPI flash.

Not achieved.

ED8 Reprogram instruction memory at runtime from host computer.

Achieved. The global instruction memory (if enabled, see [Appendix C](#)) supported run-time programming over the UART0 interface. An example host-programmer application is provided in `sw/prog.py`.

ED9 Processor external debugger using host-processor link.

Not achieved.

9.3 Future Work

There are two types of improvements that could be performed to improve the project. The first is to perform various refinements to existing project features to produce a higher quality, more efficient and performant design. The second is to add features for the purpose of gaining knowledge.

A few of these improvements are described below:

1. Use AMBA AHB for processor-to-processor and processor-to-memory communication. This will allow for higher throughput and larger data transfers between these devices. If instructions can be fetched in bulk from the instruction memory via AHB, it would enable more processor optimisations such as out-of-order execution and branch prediction. The AMBA APB interface would still be used for communicating with the peripherals.
2. Re-add instruction pipelining to each processor. While this was removed due to it's complexity and fears of hindering multi-core and interrupt integration, now that these features have been accomplished and understood, instruction pipelining should be added to improve instruction throughput.
3. As discussed in [Section 5.3.1](#), a more complex arbiter that features priority and inputs from the global shared memory would have resulted in significantly lower wait times and therefore performance.
4. As discussed in [Section 8.2.3 Performance Results](#), using a shared memory significantly reduces design size but results in lower performance. To overcome this, a design could

cluster cores into groups, where each group has a instruction memory that it's core use. This would reduce the demand on a single instruction memory from all the cores to only a few. This is a trade-off using the advantages and disadvantages of both designs.

References

- [1] Tech Differences, "Difference between loosely coupled and tightly coupled multiprocessor system (with comparison chart)," Jul 2017. [Online]. Available: <https://techdifferences.com/difference-between-loosely-coupled-and-tightly-coupled-multiprocessor-system.html> (Accessed 2019-04-20).
- [2] N. Chatterjee, S. Paul, and S. Chattopadhyay, "Fault-tolerant dynamic task mapping and scheduling for network-on-chip-based multicore platform," *ACM Transactions on Embedded Computing Systems*, vol. 16, pp. 1–24, 05 2017.
- [3] Terasic Technologies, "SoC Platform - Cyclone - DE1-SoC Board." [Online]. Available: <https://www.terasic.com.tw/cgi-bin/page/archive.pl?Language=English&No=836> (Accessed 2019-04-20).
- [4] *MiniSpartan6+*, Scarab Hardware, 2014. [Online]. Available: <https://www.scarabhardware.com/minispartan6/> (Accessed 2019-04-20).
- [5] ARM, *AMBA 3 AHB-Lite Protocol*, ARM.
- [6] V. Subramanian, "Multiple gate field-effect transistors for future CMOS technologies," *IETE Technical review*, vol. 27, no. 6, pp. 446–454, 2010.
- [7] M. J. Flynn, *Computer architecture: Pipelined and parallel processor design*. Jones & Bartlett Learning, 1995.
- [8] L. Benini and G. De Micheli, "Networks on Chips: A new SoC paradigm," *Computer*, vol. 35, pp. 70–78, 02 2002.
- [9] D. Zhu, L. Chen, S. Yue, T. M. Pinkston, and M. Pedram, "Balancing On-Chip Network Latency in Multi-application Mapping for Chip-Multiprocessors," in *2014 IEEE 28th International Parallel and Distributed Processing Symposium*, May 2014, pp. 872–881.
- [10] Xilinx, *Spartan-6 FPGA Block RAM Resources*, Xilinx.
- [11] Altera, *Recommended HDL Coding Styles - QII51007-9.0.0*, Altera.
- [12] B. Lancaster, "FPGA-based RISC Microprocessor and Compiler," vol. 3.14, pp. 37–50. [Online]. Available: <https://github.com/bendl/prco304> (Accessed March 2018).
- [13] P. Kocher, D. Genkin, D. Gruss, W. Haas, M. Hamburg, M. Lipp, S. Mangard, T. Prescher, M. Schwarz, and Y. Yarom, "Spectre attacks: Exploiting speculative execution," *arXiv preprint arXiv:1801.01203*, 2018.
- [14] OpenCores.org, *WISHBONE System-on-Chip (SoC) Interconnection Architecture for Portable IP Cores*, OpenCores.org. [Online]. Available: https://cdn.opencores.org/downloads/wbspec_b3.pdf
- [15] ARM, *AMBA 3 APB Protocol*, ARM.

- [16] M. Weber, "Arbiters: Design Ideas and Coding Styles," p. 2.
- [17] A. Palchoudhuri and R. S. Chakraborty, *High Performance Integer Arithmetic Circuit Design on FPGA: Architecture, Implementation and Design Automation*. Springer, 2015, vol. 51.
- [18] V. Salauyou and M. Gruszeński, "Designing of hierarchical structures for binary comparators on fpga/soc," in *IFIP International Conference on Computer Information Systems and Industrial Management*. Springer, 2015, pp. 386–396.
- [19] Xilinx, *Spartan-6 FPGA Configurable Logic Block - User Guide - UG384*, Xilinx.
- [20] Altera, *Cyclone V Device Handbook - Device Interfaces and Integration - CV-5V2*, Altera.
- [21] A. S. Tanenbaum, *Structured Computer Organization*. Pearson Education India, 2016.
- [22] Xilinx, *Spartan-6 Family Overview - DS160*, Xilinx.
- [23] Intel, *Cyclone V Device Overview - CV-51001*, Intel.
- [24] B. Lancaster, "Parallel Computing and Distributed Systems." [Online]. Available: <https://github.com/bendl/Parallel-DFT-CUDA-MPI/raw/master/report/SOFT354%20-%20Ben%20Lancaster%2010424877.pdf>

Appendix A

Peripheral Information

A.1 Special Registers	60
A.2 Watchdog Timer	61
A.3 GPIO Interface	61
A.4 Timer with Interrupt	62
A.5 UART Interface	62
B.1 Register Set Multiplex	63
B.2 Instruction Set Architecture	64

To provide user's with useful functionality, common system-on-chip peripherals were created. This section describes each peripheral and it's design decisions. The full memory-map is shown in [Figure 6.5](#).

A.1 Special Registers

From the software perspective, it is important for both the developer and software algorithms to know the target system's architecture to better utilise the resources available to them. Software written for one architecture with N cores must also run on an architecture with M cores. To enable such portability, the software must query the system for information such as: number of processor cores and the current core identifier. Without this information, the developer would be required to produce software for each individual architecture (e.g. an Intel i5 with 4 cores or an Intel i7 with 8 cores, or an NVIDIA GTX 970 with 1664 CUDA cores.

The special register peripheral is shown below.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
								CORE_ID								0080	R
								NUM_CORES								0081	R
SHARED_MEMORY cells (default 4096)																0082	R
								NUM_PERIPHERALS								0083	R
SCRATCH_MEMORY cells (default 64)																0084	RW
User defined																0085	RW
⋮																	
User defined																008F	RW

Figure A.1: Vmicro16 Special Registers layout (0x0080 - 0x008F).

A.2 Watchdog Timer

In any multi-threaded system there exists the possibility for a deadlock – a state where all threads are in a waiting state – and algorithm execution is forever blocked. This can occur either by poor software programming or incorrect thread arbitration by the processor. A common method of detecting a deadlock is to make each thread signal that it is not blocked by resetting a countdown timer. If the countdown timer is not reset, it will eventually reach zero and it is assumed that all threads are blocked as none have reset the countdown.

In this system-on-chip design, software can reset the watchdog timer by writing any 16-bit value to the address 0x00B8.

This peripheral is optional and can be enabled using the configuration parameters described in [Configuration Options](#).

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Reset Watchdog																00B8 W

A.3 GPIO Interface

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
								GPIO0 Output								0090 RW
GPIO1 Output																0091 RW
								GPIO2 Output								0092 RW
								GPIO3 Input								0093 R

On the DE1-SoC board, GPIO0 is assigned to the LEDs, and GPIO1 and GPIO2 to the 6 seven-segment displays.

A.4 Timer with Interrupt

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Load Value																0200 RW
													I	R	S	0201 W
Prescaler																0202 W

Clock Frequency Uses top level FPGA clock (normally 50 MHz).

Load Value Value to count down from each clock.

I Interrupt enable bit. Default 0.

R Reset Load Value and Prescaler values to their last written value.

S Start the timer countdown. 1 = start. 0 = stop.

Prescaler Number of clocks per FPGA clock to wait between each decrement.

A.5 UART Interface

15	8	7	1	0			
					Transmit Data	00A0 W	
					Receive Data	00A1 R	
					E	I	00A2 RW

E Enable the UART component.

I Enable an interrupt upon receiving new data. Default 1.

Note: If DEF_USE_REPROG is enabled in vmicro16_soc_config.v then the Receive Data register will be reserved for programming the instruction memory, resulting in reads and writes to addresses 0x00A1 and 0x00A2 to return 0.

Appendix B

Additional Figures

```
1  input      [MASTER_PORTS*BUS_WIDTH-1:0] S_PADDR,  
2  input      [MASTER_PORTS-1:0]          S_PWRITE,  
3  input      [MASTER_PORTS-1:0]          S_PSELx,  
4  input      [MASTER_PORTS-1:0]          S_PENABLE,  
5  input      [MASTER_PORTS*DATA_WIDTH-1:0] S_PWDATA,  
6  output reg [MASTER_PORTS*DATA_WIDTH-1:0] S_PRDATA,  
7  output reg [MASTER_PORTS-1:0]          S_PREADY,
```

Listing 7: Variable size inputs and outputs to the interconnect.

B.1 Register Set Multiplex

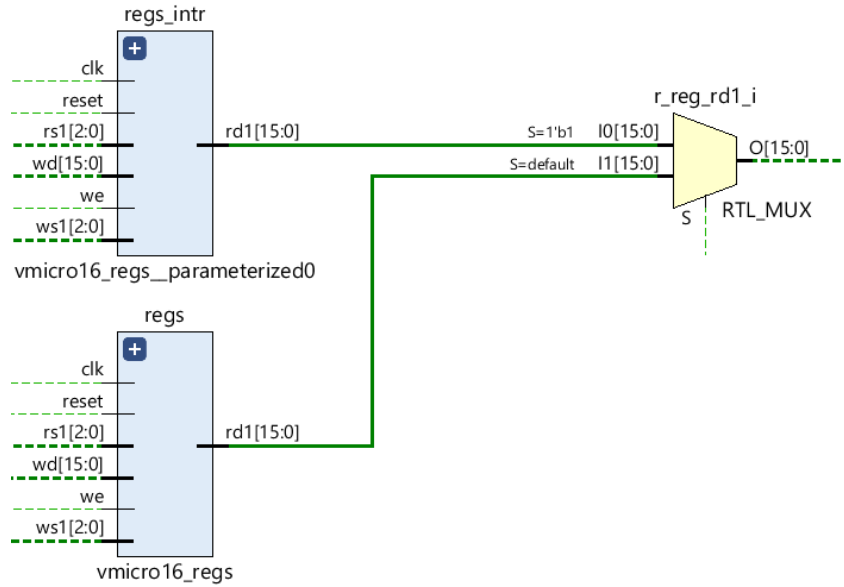


Figure B.1: Normal mode (bottom) and interrupt mode (top) register sets are multiplexed to switch between contexts.

B.2 Instruction Set Architecture

	15-11	10-8	7-5	4-0	rd ra simm5
	15-11	10-8	7-0		rd imm8
	15-11	10-0			nop
	15	14:12	11:0		extended immediate
SPCL	00000	11 bits			NOP
SPCL	00000	11h'000			NOP
SPCL	00000	11h'001			HALT
SPCL	00000	11h'002			Return from interrupt
LW	00001	Rd	Ra	s5	Rd <= RAM[Ra+s5]
SW	00010	Rd	Ra	s5	RAM[Ra+s5] <= Rd
BIT	00011	Rd	Ra	s5	bitwise operations
BIT_OR	00011	Rd	Ra	00000	Rd <= Rd Ra
BIT_XOR	00011	Rd	Ra	00001	Rd <= Rd ^ Ra
BIT_AND	00011	Rd	Ra	00010	Rd <= Rd & Ra
BIT_NOT	00011	Rd	Ra	00011	Rd <= ~Ra
BIT_LSHFT	00011	Rd	Ra	00100	Rd <= Rd << Ra
BIT_RSHFT	00011	Rd	Ra	00101	Rd <= Rd >> Ra
MOV	00100	Rd	Ra	X	Rd <= Ra
MOVI	00101	Rd		i8	Rd <= i8
ARITH_U	00110	Rd	Ra	s5	unsigned arithmetic
ARITH_UADD	00110	Rd	Ra	11111	Rd <= uRd + uRa
ARITH_USUB	00110	Rd	Ra	10000	Rd <= uRd - uRa
ARITH_UADDI	00110	Rd	Ra	0AAAA	Rd <= uRd + Ra + AAAA
ARITH_S	00111	Rd	Ra	s5	signed arithmetic
ARITH_SADD	00111	Rd	Ra	11111	Rd <= sRd + sRa
ARITH_SSUB	00111	Rd	Ra	10000	Rd <= sRd - sRa
ARITH_SSUBI	00111	Rd	Ra	0AAAA	Rd <= sRd - sRa + AAAA
BR	01000	Rd		i8	conditional branch
BR_U	01000	Rd		0000 0000	Any
BR_E	01000	Rd		0000 0001	Z=1
BR_NE	01000	Rd		0000 0010	Z=0
BR_G	01000	Rd		0000 0011	Z=0 and S=0
BR_GE	01000	Rd		0000 0100	S=0
BR_L	01000	Rd		0000 0101	S != 0
BR_LE	01000	Rd		0000 0110	Z=1 or (S != 0)
BR_S	01000	Rd		0000 0111	S=1
BR_NS	01000	Rd		0000 1000	S=0
CMP	01001	Rd	Ra	X	SZO <= CMP(Rd, Ra)
SETC	01010	Rd		Imm8	Rd <= (Imm8_f_SZO) ? 1 : 0
MULT	01011	Rd	Ra	X	Rd <= uRd * uRa
HALT	01100			X	
LWEX	01101	Rd	Ra	s5	Rd <= RAM[Ra+s5]
SWEX	01110	Rd	Ra	s5	RAM[Ra+s5] <= Rd Rd <= 0 1 if success

Figure B.2: Vmicro16 instruction set architecture.

Appendix C

Configuration Options

C.1 System-on-Chip Configuration Options	65
C.2 Core Options	66
C.3 Peripheral Options	67

The following configuration options are defined in `vmicro16_soc_config.v`.

Defaults with empty/blank values signifies that the preprocessor define is commented out/not defined/disabled by default/computed by other parameters.

C.1 System-on-Chip Configuration Options

Macro	Default	Purpose
CORES	4	Number of CPU cores in the SoC
SLAVES	9	Number of peripherals
DEF_USE_WATCHDOG	//	Enable watchdog module to recover from deadlocks and infinite loops. Requires DEF_GLOBAL_RESET.
DEF_GLOBAL_RESET	//	Enable synchronous reset logic. Will consume more LUT resources. Does not reset BRAM blocks.
DEF_USE_BUS_RESET	//	Detect bus stalls or errors to soft-reset the whole design. Requires DEF_USE_WATCHDOG.

Table C.1: SoC Configuration Options

C.2 Core Options

Macro	Default	Purpose
DATA_WIDTH	16	Width of CPU registers in bits
DEF_CORE_HAS_INSTR_MEM	//	Enable a per core instruction memory cache
DEF_MEM_INSTR_DEPTH	64	Instruction memory cache per core
DEF_MEM_SCRATCH_DEPTH	64	RW RAM per core
DEF_ALU_HW_MULT	1	Enable/disable HW multiply (1 clock)
FIX_T3	//	Enable a T3 state for the APB transaction
DEF_USE_REPROG	//	Programme instruction memory via UART0. Requires DEF_GLOBAL_RESET. Enabling this will reserve the UART0 RX port for exclusive use for programming the instruction memory. Software reads of UART0 RX will return 0.

Table C.2: Core Options

C.3 Peripheral Options

Macro	Default	Purpose
APB.WIDTH		AMBA APB PADDR signal width
APB.PSELX_GPIO0	0	GPIO0 index
APB.PSELX_UART0	1	UART0 index
APB.PSELX_REGS0	2	REGS0 index
APB.PSELX_BRAM0	3	BRAM0 index
APB.PSELX_GPIO1	4	GPIO1 index
APB.PSELX_GPIO2	5	GPIO2 index
APB.PSELX_TIMR0	6	TIMR0 index
APB.BRAM0_CELLS	4096	Shared memory words
DEF_MMU_TIM0_S	16'h0000	Per core scratch memory start/end address
DEF_MMU_TIM0_E	16'h007F	"
DEF_MMU_SREG_S	16'h0080	Per core special registers start/end address
DEF_MMU_SREG_E	16'h008F	"
DEF_MMU_GPIO0_S	16'h0090	Shared GPIO0 start/end address
DEF_MMU_GPIO0_E	16'h0090	"
DEF_MMU_GPIO1_S	16'h0091	"
DEF_MMU_GPIO1_E	16'h0091	"
DEF_MMU_GPIO2_S	16'h0092	"
DEF_MMU_GPIO2_E	16'h0092	"
DEF_MMU_UART0_S	16'h00A0	Shared UART start/end address
DEF_MMU_UART0_E	16'h00A1	"
DEF_MMU_REGS0_S	16'h00B0	Shared registers start/end address
DEF_MMU_REGS0_E	16'h00B7	"
DEF_MMU_BRAM0_S	16'h1000	Shared memory with global monitor start/end address
DEF_MMU_BRAM0_E	16'h1FFF	"
DEF_MMU_TIMR0_S	16'h0200	Shared timer peripheral start/end address
DEF_MMU_TIMR0_E	16'h0202	"

Table C.3: Peripheral Options

Appendix D

Viva Demonstration Examples

D.1 2-core Timer Interrupt and ISR	68
D.2 1-160 Core Parallel Summation	70

D.1 2-core Timer Interrupt and ISR

This example demo, shown during the viva, blinks an LED every 0.5 seconds via a timer interrupt. Core 0 sets up the interrupt vector (by writing the isr0 function address to the interrupt vector) and enables all interrupt sources. Core 1 sets up the timer interval peripheral to produce an interrupt every 0.5 seconds. Core 1 also performs the interrupt handler (isr0): toggle an LED, write the state to UART0, and resets the watchdog.

```
1 // interrupts.s
2 // Toggle LED in ISR
3
4 entry:
5 // get core idx 0x80 in r7
6 movi r7, #0x80
7 lw r7, r7
8
9 // core1 sets up the timer
10 // Core0 enables interrupts and performs the isr
11 cmp r7, r0
12 movi r0, timer
13 br r0, BR_NE
14
15 // Set interrupt vector (0)
16 movi r0, isr0
17 movi r1, #0x1
18 movi r2, #0x08
19 lshft r1, r2
20 sw r0, r1
21
22 // enable all interrupts
23 movi r0, #0x0f
24 sw r0, r1 + #0x8
25
26 // enter idle state
27 halt r0, r0
28
29 timer:
30 // set timr0 address 0x200 into r0
31 // shift left 8 places
32 movi r0, #0x01
33 movi r1, #0x09
34 lshft r0, r1
35
36 // Set load value
37 //movi r1, #0x31
38 //sw r1, r0
39 // test we the expected value back
40 //lw r2, r0
41
42 // set load = 0x3000
```

```

43     movi    r1, #0x3
44     movi    r2, #0x0C
45     //movi   r2, #0x04
46     lshft   r1, r2
47     sw      r1, r0
48
49     // Set prescale value to 0x1000
50     // 20ns * load * prescaler = nanosecond delay
51     // 20ns * 10000 * 5000     = 1.0s
52     // 20.0 * 0x3000 * 0x1000  = ~1.0s
53     movi    r1, #0x1
54     // 1.0 second
55     //movi   r2, #0x0C
56     // 0.5 second
57     movi    r2, #0x0B
58     // 0.25 second
59     //movi   r2, #0x0a
60     // 0.0625 second
61     //movi   r2, #0x04
62     lshft   r1, r2
63     sw      r1, r0 + #0x02
64
65     // Start the timer (write 0x0001 to 0x0101)
66     movi    r1, #0x01
67     sw      r1, r0 + #0x01
68
69     exit:
70     // enter idle state
71     halt    r0, r0
72
73     isr0:
74     movi    r0, #0x90
75     lw      r1, r0
76     // xor with 1
77     movi    r2, #0x1
78     xor     r1, r2
79     // write back
80     sw      r1, r0
81
82     // write ascii value to uart0
83     movi    r0, #0xa0
84     movi    r2, #0x30
85     add     r1, r2
86     sw      r1, r0
87
88     // reset watchdog
89     movi    r0, #0xb8
90     sw      r1, r0
91
92     // return from interrupt
93     intr    r0, r0

```

D.2 1-160 Core Parallel Summation

This example demo performs a parallel summation of numbers 1 to 320. The algorithm *assigns* each core a subset of the summation space. It does this using the core's ID and the number of cores in the system. The following formulas determine where the subset begins and ends for each core. Core 0 broadcasts the number to sum to then each core calculates its subset start and end positions. Each core then performs a summation over it's subset then adds the result to a global shared value. After pushes it's results, the global shared value will contain the final summation result.

$$N_{samples} = 320 \quad (D.1)$$

$$N_{threads} = 64 \quad (D.2)$$

$$subset = N_{samples} / N_{threads} \quad (D.3)$$

$$start = ID * subset \quad (D.4)$$

$$end = start + subset \quad (D.5)$$

```

1  // sum64.s
2  // Simple 1-160 core summation program
3
4  // Set up common values, such as: Core id (r6),
5  // number of threads (cores) (r7), shared memory addresses (r5)
6  entry:
7  // Core id in r6
8  movi    r0, #0x80
9  lw      r0, r0
10 // store in r6
11 mov     r6, r0
12
13 // get number of threads
14 movi    r0, #0x81
15 lw      r0, r0
16 // store in r7
17 mov     r7, r0
18
19 // BRAMO shared memory 0x1000
20 movi    r5, #0x01
21 movi    r2, #0x0C
22 lshft   r5, r2
23
24 jmp_to_barrier:
25 // NOT_ROOT
26 // wait at barrier
27 cmp     r6, r3
28 movi    r4, barrier_arrive
29 br      r4, BR_NE
30
31 // ROOT
32 // calculates nsamples_per_thread
33 // ns = 100
34 // nst = ns / (num_threads)
35 // nst = ns >> (num_threads - 1)
36 // r0 = (num_threads - 1) WRONG!!!
37
38 root_broadcast:
39 // The root (core idx 0) broadcasts the number of samples
40 // 16 cores
41 //movi    r4, #0x14
42 // 32 cores
43 //movi    r4, #0x0a
44 // 64 cores
45 movi    r4, #0x05
46 // 80 cores
47 //movi    r4, #0x04
48 // 160 cores
49 //movi    r4, #0x02
50
51 // ROOT
52 // Do the broadcast
53 // write nsamples_per_thread to shared bram (broadcast)
54 // 0x1001
55 sw      r4, r5 + #0x01
56

```



```

57 // Reach the barrier to tell everyone
58 // that we have arrived
59 barrier_arrive:
60 // load latest count
61 lwex    r0, r5
62 // try increment count
63 // increment by 1
64 addi    r0, r3 + #0x01
65 // attempt store
66 swex    r0, r5
67 // check success (== 0)
68 cmp     r0, r3
69 // branch if failed
70 movi    r4, barrier_arrive
71 br      r4, BR_NE
72
73 // Wait in an infinite loop
74 // for all cores to 'arrive'
75 barrier:
76 // load the count
77 lw      r0, r5
78 // compare with number of threads
79 cmp     r0, r7
80 // jump back to barrier if not equal
81 movi    r4, barrier
82 br      r4, BR_NE
83
84 // EACH CORE
85 // All cores have arrived and in sync
86 synced1:
87 // Retrieve load the nsamples_per_thread
88 lw      r4, r5 + #0x01
89 // Calculate nstart = idx * nsamples_per_thread
90 // in r2
91 mov     r2, r6
92 mult    r2, r4
93
94 // Loop limit in r4
95 // samples_per_thread -> samples_per_thread + nstart
96 add     r4, r2
97
98 // Perform the summation in a tight for loop
99 // Sum numbers from nstart to limit
100 sum_loop:
101 // sum += i
102 add     r1, r2
103 // increment i
104 addi    r2, r3 + #0x01
105 // check end
106 cmp     r2, r4
107 movi    r0, sum_loop
108 br      r0, BR_NE
109
110 // Summation of the subset finished, result is in r1
111 // Now use a mutex to add it to the global sum value in shared mem
112 sum_mutex:
113 // load latest count
114 lwex    r0, r5 + #0x2
115 // try increment count
116 // increment by 1
117 add     r0, r1
118 // make copy as swex has a return value
119 mov     r2, r0
120 // attempt store
121 swex    r0, r5 + #0x02
122 // check success (== 0)
123 cmp     r0, r3
124 // branch if failed
125 movi    r4, sum_mutex
126 br      r4, BR_NE
127
128 // Write the latest global sum value to gpio1
129 write_gpio:
130 movi    r3, #0x91
131 sw      r2, r3
132
133 // Write the latest global sum value to uart0 tx
134 write_uart_done:
135 movi    r3, #0xa0
136 movi    r2, #0x30
137 add     r2, r6
138 sw      r2, r3
139
140 // This core has finished
141 // Enter a low power state
142 exit:
143 halt    r0, r0

```

Appendix E

Code Listing

E.1	SoC Code Listing	72
E.1.1	vmicro16_soc_config.v	72
E.1.2	top_ms.v	74
E.1.3	vmicro16_soc.v	75
E.1.4	vmicro16.v	82
E.2	Peripheral Code Listing	95
E.3	Assembly Compiler Listing	101
E.4	Text Compiler Listing	106

E.1 SoC Code Listing

E.1.1 vmicro16_soc_config.v

Configuration file for configuring the vmicro16_soc.v and vmicro16.v features.

```
1 // Configuration defines for the vmicro16_soc and vmicro16 cpu.
2
3 `ifndef VMICRO16_SOC_CONFIG_H
4 `define VMICRO16_SOC_CONFIG_H
5
6 `include "clog2.v"
7
8 `define FORMAL
9
10 `define CORES 32
11 `define SLAVES 9
12
13 ///////////////////////////////////////////////////////////////////
14 // Core parameters
15 ///////////////////////////////////////////////////////////////////
16 // Per core instruction memory
17 // Set this to give each core its own instruction memory cache
18 //`define DEF_CORE_HAS_INSTR_MEM
19
20 // Top level data width for registers, memory cells, bus widths
21 `define DATA_WIDTH 16
22
23 // Set this to use a workaround for the MMU's APB T2 clock
24 //`define FIX_T3
25
26 // Instruction memory (read only)
27 // Must be large enough to support software program.
28 `ifdef DEF_CORE_HAS_INSTR_MEM
29 // 64 16-bit words per core
30 `define DEF_MEM_INSTR_DEPTH 64
31 `else
32 // 4096 16-bit words global
33 `define DEF_MEM_INSTR_DEPTH 4096
34 `endif
```

```

35
36 // Scratch memory (read/write) on each core.
37 // See `DEF_MMU_TIMO_*` defines for info.
38 `define DEF_MEM_SCRATCH_DEPTH 64
39
40 // Enables hardware multiplier and mult rr instruction
41 `define DEF_ALU_HW_MULT 1
42
43 // Enables global reset (requires more luts)
44 `define DEF_GLOBAL_RESET
45
46 // Enable a watch dog timer to reset the soc if threadlocked
47 `define DEF_USE_WATCHDOG
48
49 // Enable to detect bus communication stalls or errors.
50 // If detected, the whole SoC will be soft-reset, as if by a watchdog
51 `define DEF_USE_BUS_RESET
52
53 // Enables instruction memory programming via UART0
54 `define DEF_USE_REPROG
55
56 `ifndef DEF_USE_REPROG
57     `ifndef DEF_GLOBAL_RESET
58         `error_DEF_USE_REPROG_requires_DEF_GLOBAL_RESET
59     `endif
60 `endif
61
62 ///////////////////////////////////////////////////
63 // Memory mapping
64 ///////////////////////////////////////////////////
65 `define APB_WIDTH (2 + `clog2(`CORES) + `DATA_WIDTH)
66
67 `define APB_PSELX_GPIO0 0
68 `define APB_PSELX_UART0 1
69 `define APB_PSELX_REGSO 2
70 `define APB_PSELX_BRAMO 3
71 `define APB_PSELX_GPIO1 4
72 `define APB_PSELX_GPIO2 5
73 `define APB_PSELX_TIMRO 6
74 `define APB_PSELX_WDOGO 7
75 `define APB_PSELX_PERRO 8
76
77 `define APB_GPIO0_PINS 8
78 `define APB_GPIO1_PINS 16
79 `define APB_GPIO2_PINS 8
80
81 // Shared memory words
82 `define APB_BRAMO_CELLS 4096
83
84 ///////////////////////////////////////////////////
85 // Memory mapping
86 ///////////////////////////////////////////////////
87 // TIMO
88 // Number of scratch memory cells per core
89 `define DEF_MMU_TIMO_CELLS 64
90 `define DEF_MMU_TIMO_S 16'h0000
91 `define DEF_MMU_TIMO_E 16'h007F
92 // SREG
93 `define DEF_MMU_SREG_S 16'h0080
94 `define DEF_MMU_SREG_E 16'h008F
95 // GPIO0
96 `define DEF_MMU_GPIO0_S 16'h0090
97 `define DEF_MMU_GPIO0_E 16'h0090
98 // GPIO1
99 `define DEF_MMU_GPIO1_S 16'h0091
100 `define DEF_MMU_GPIO1_E 16'h0091
101 // GPIO2
102 `define DEF_MMU_GPIO2_S 16'h0092
103 `define DEF_MMU_GPIO2_E 16'h0092
104 // UART0
105 `define DEF_MMU_UART0_S 16'h00A0
106 `define DEF_MMU_UART0_E 16'h00A1
107 // REGSO
108 `define DEF_MMU_REGSO_S 16'h00B0
109 `define DEF_MMU_REGSO_E 16'h00B7
110 // WDOGO
111 `define DEF_MMU_WDOGO_S 16'h00B8
112 `define DEF_MMU_WDOGO_E 16'h00B8
113 // BRAMO
114 `define DEF_MMU_BRAMO_S 16'h1000
115 `define DEF_MMU_BRAMO_E 16'h1fff
116 // TIMRO
117 `define DEF_MMU_TIMRO_S 16'h0200
118 `define DEF_MMU_TIMRO_E 16'h0202
119
120 ///////////////////////////////////////////////////
121 // Interrupts
122 ///////////////////////////////////////////////////
123 // Enable/disable interrupts
124 // Disabling will free up resources for other features

```

```

125 //`define DEF_ENABLE_INT
126 // Number of interrupt in signals
127 `define DEF_NUM_INT 8
128 // Default interrupt bitmask (0 = hidden, 1 = enabled)
129 `define DEF_INT_MASK 0
130 // Bit position of the TIMRO interrupt signal
131 `define DEF_INT_TIMRO 0
132 // Interrupt vector memory location
133 `define DEF_MMU_INTSV_S 16'h0100
134 `define DEF_MMU_INTSV_E 16'h0107
135 // Interrupt vector memory location
136 `define DEF_MMU_INTSM_S 16'h0108
137 `define DEF_MMU_INTSM_E 16'h0108
138
139 `endif

```

E.1.2 top_ms.v

Top level module that connects the SoC design to hardware pins on the FPGA.

```

1  module seven_display # (
2      parameter INVERT = 1
3  ) (
4      input  [3:0] n,
5      output [6:0] segments
6  );
7      reg [6:0] bits;
8      assign segments = (INVERT ? ~bits : bits);
9
10     always @(n)
11     case (n)
12         4'h0: bits = 7'b0111111; // 0
13         4'h1: bits = 7'b0000110; // 1
14         4'h2: bits = 7'b1011011; // 2
15         4'h3: bits = 7'b1001111; // 3
16         4'h4: bits = 7'b1100110; // 4
17         4'h5: bits = 7'b1101101; // 5
18         4'h6: bits = 7'b1111101; // 6
19         4'h7: bits = 7'b0000111; // 7
20         4'h8: bits = 7'b1111111; // 8
21         4'h9: bits = 7'b1100111; // 9
22         4'hA: bits = 7'b1110111; // A
23         4'hB: bits = 7'b1111100; // B
24         4'hC: bits = 7'b0111001; // C
25         4'hD: bits = 7'b1011110; // D
26         4'hE: bits = 7'b1111001; // E
27         4'hF: bits = 7'b1110001; // F
28     endcase
29 endmodule
30
31
32 // minispartan6+ XC6SLX9
33 module top_ms # (
34     parameter GPIO_PINS = 8
35 ) (
36     input          CLK50,
37     input  [3:0]   SW,
38     // UART
39     input          RXD,
40     output         TXD,
41     // Peripherals
42     output [7:0]   LEDS,
43
44     // 3v3 input from the s6 on the deisoc
45     input          S6_3v3,
46
47     // SSDs
48     output [6:0]   ssd0,
49     output [6:0]   ssd1,
50     output [6:0]   ssd2,
51     output [6:0]   ssd3,
52     output [6:0]   ssd4,
53     output [6:0]   ssd5
54 );
55     //wire [15:0] M_PADDR;
56     //wire M_PWRITE;
57     //wire [5-1:0] M_PSEL; // not shared
58     //wire M_PENABLE;
59     //wire [15:0] M_PWDATA;
60     //wire [15:0] M_PRDATA; // input to intercon
61     //wire M_PREADY; // input to intercon
62
63     wire [7:0] gpio0;

```

```

64     wire [15:0] gpio1;
65     wire [7:0]  gpio2;
66
67     vmicro16_soc soc (
68         .clk      (CLK50),
69         .reset    (~SW[0]),
70
71         // .M_PADDR      (M_PADDR),
72         // .M_PWRITE     (M_PWRITE),
73         // .M_PSELx      (M_PSELx),
74         // .M_PENABLE    (M_PENABLE),
75         // .M_PWDATA     (M_PWDATA),
76         // .M_PRDATA     (M_PRDATA),
77         // .M_PREADY     (M_PREADY),
78
79         // UART
80         .uart_tx (TXD),
81         .uart_rx (RXD),
82
83         // GPIO
84         .gpio0  (LEDS[3:0]),
85         .gpio1  (gpio1),
86         .gpio2  (gpio2),
87
88         // DEBUG
89         .debug0 (LEDS[4])
90         // .debug1 (LEDS[7:4])
91     );
92
93     assign LEDS[7:5] = {TXD, RXD, S6_3v3};
94
95     // SSD displays (split across 2 gpio ports 1 and 2)
96     wire [3:0] ssd_chars [0:5];
97     assign ssd_chars[0] = gpio1[3:0];
98     assign ssd_chars[1] = gpio1[7:4];
99     assign ssd_chars[2] = gpio1[11:8];
100    assign ssd_chars[3] = gpio1[15:12];
101    assign ssd_chars[4] = gpio2[3:0];
102    assign ssd_chars[5] = gpio2[7:4];
103    seven_display ssd_0 (.n(ssd_chars[0]), .segments (ssd0));
104    seven_display ssd_1 (.n(ssd_chars[1]), .segments (ssd1));
105    seven_display ssd_2 (.n(ssd_chars[2]), .segments (ssd2));
106    seven_display ssd_3 (.n(ssd_chars[3]), .segments (ssd3));
107    seven_display ssd_4 (.n(ssd_chars[4]), .segments (ssd4));
108    seven_display ssd_5 (.n(ssd_chars[5]), .segments (ssd5));
109
110    endmodule

```

E.1.3 vmicro16_soc.v

```

1    //
2    //
3
4    `include "vmicro16_soc_config.v"
5    `include "clog2.v"
6    `include "formal.v"
7
8    module pow_reset # (
9        parameter INIT = 1,
10       parameter N    = 8
11    ) (
12       input      clk,
13       input      reset,
14       output reg resethold
15    );
16       initial resethold = INIT ? (N-1) : 0;
17
18       always @(*)
19           resethold = |hold;
20
21       reg [`clog2(N)-1:0] hold = (N-1);
22       always @(posedge clk)
23           if (reset)
24               hold <= N-1;
25           else
26               if (hold)
27                   hold <= hold - 1;
28    endmodule
29
30    // Vmicro16 multi-core SoC with various peripherals
31    // and interrupts
32    module vmicro16_soc (
33       input clk,
34       input reset,
35

```

```

36 // UART0
37 input          uart_rx,
38 output        uart_tx,
39 //
40 output [`APB_GPIO0_PINS-1:0] gpio0,
41 output [`APB_GPIO1_PINS-1:0] gpio1,
42 output [`APB_GPIO2_PINS-1:0] gpio2,
43 //
44 output        halt,
45 //
46 output      [`CORES-1:0]      dbug0,
47 output      [`CORES*8-1:0]    dbug1
48 );
49 wire [`CORES-1:0] w_halt;
50 assign halt = &w_halt;
51
52 assign dbug0 = w_halt;
53
54 // Watchdog reset pulse signal.
55 // Passed to pow_reset to generate a longer reset pulse
56 wire wdreset;
57 wire prog_prog;
58 // Set high if a bus stall or error occurs.
59 // This will reset the whole SoC!
60 wire bus_reset;
61
62 // soft register reset hold for brams and registers
63 wire soft_reset;
64 `ifndef DEF_GLOBAL_RESET
65     pow_reset # (
66         .INIT      (1),
67         .N          (8)
68     ) por_inst (
69         .clk        (clk),
70         `ifndef DEF_USE_WATCHDOG
71         .reset      (reset | wdreset | prog_prog
72         `ifndef DEF_USE_BUS_RESET
73             | bus_reset),
74         `else
75             ),
76         `endif
77         `else
78         .reset      (reset),
79         `endif
80         .resethold  (soft_reset)
81     );
82 `else
83     assign soft_reset = 0;
84 `endif
85
86 // Peripherals (master to slave)
87 wire [`APB_WIDTH-1:0] M_PADDR;
88 wire M_PWRITE;
89 wire [`SLAVES-1:0] M_PSELx; // not shared
90 wire M_PENABLE;
91 wire [`DATA_WIDTH-1:0] M_PWDATA;
92 wire [`SLAVES*DATA_WIDTH-1:0] M_PRDATA; // input to intercon
93 wire [`SLAVES-1:0] M_PREADY; // input
94
95 // Master apb interfaces
96 wire [`CORES*APB_WIDTH-1:0] w_PADDR;
97 wire [`CORES-1:0] w_PWRITE;
98 wire [`CORES-1:0] w_PSELx;
99 wire [`CORES-1:0] w_PENABLE;
100 wire [`CORES*DATA_WIDTH-1:0] w_PWDATA;
101 wire [`CORES*DATA_WIDTH-1:0] w_PRDATA;
102 wire [`CORES-1:0] w_PREADY;
103
104 // Interrupts
105 `ifndef DEF_ENABLE_INT
106 wire [`DEF_NUM_INT-1:0] ints;
107 wire [`DEF_NUM_INT*DATA_WIDTH-1:0] ints_data;
108 assign ints[7:1] = 0;
109 assign ints_data[`DEF_NUM_INT*DATA_WIDTH-1:DATA_WIDTH] =
110     {`DEF_NUM_INT*(DATA_WIDTH-1){1'b0}};
111 `endif
112
113 apb_intercon_s # (
114     .MASTER_PORTS  (`CORES),
115     .SLAVE_PORTS    (`SLAVES),
116     .BUS_WIDTH      (`APB_WIDTH),
117     .DATA_WIDTH     (`DATA_WIDTH),
118     .HAS_PSELX_ADDR (1)
119 ) apb (
120     .clk            (clk),
121     .reset          (soft_reset),
122     // APB master to slave
123     .S_PADDR        (w_PADDR),
124     .S_PWRITE        (w_PWRITE),

```

```

125     .S_PSELx      (w_PSELx),
126     .S_PENABLE   (w_PENABLE),
127     .S_PWDATA     (w_PWDATA),
128     .S_PRDATA     (w_PRDATA),
129     .S_PREADY     (w_PREADY),
130     // shared bus
131     .M_PADDR      (M_PADDR),
132     .M_PWRITE     (M_PWRITE),
133     .M_PSELx      (M_PSELx),
134     .M_PENABLE    (M_PENABLE),
135     .M_PWDATA     (M_PWDATA),
136     .M_PRDATA     (M_PRDATA),
137     .M_PREADY     (M_PREADY)
138 );
139
140 `ifdef DEF_USE_BUS_RESET
141     vmicro16_psel_err_apb error_apb (
142         .clk          (clk),
143         .reset        (),
144         // apb slave to master interface
145         .S_PADDR      (),
146         .S_PWRITE     (),
147         .S_PSELx      (M_PSELx[`APB_PSELX_PERRO]),
148         .S_PENABLE    (M_PENABLE),
149         .S_PWDATA     (),
150         .S_PRDATA     (),
151         .S_PREADY     (M_PREADY[`APB_PSELX_PERRO]),
152         // Error interrupt to reset the bus
153         .err_i        (bus_reset)
154     );
155 `endif
156
157 `ifdef DEF_USE_WATCHDOG
158     vmicro16_watchdog_apb # (
159         .BUS_WIDTH    (`APB_WIDTH),
160         .NAME          ("WDOGO")
161     ) wdog0_apb (
162         .clk          (clk),
163         .reset        (),
164         // apb slave to master interface
165         .S_PADDR      (),
166         .S_PWRITE     (M_PWRITE),
167         .S_PSELx      (M_PSELx[`APB_PSELX_WDOGO]),
168         .S_PENABLE    (M_PENABLE),
169         .S_PWDATA     (),
170         .S_PRDATA     (),
171         .S_PREADY     (M_PREADY[`APB_PSELX_WDOGO]),
172
173         .wdreset      (wdreset)
174     );
175 `endif
176
177     vmicro16_gpio_apb # (
178         .BUS_WIDTH    (`APB_WIDTH),
179         .DATA_WIDTH    (`DATA_WIDTH),
180         .PORTS        (`APB_GPIO0_PINS),
181         .NAME          ("GPIO0")
182     ) gpio0_apb (
183         .clk          (clk),
184         .reset        (soft_reset),
185         // apb slave to master interface
186         .S_PADDR      (M_PADDR),
187         .S_PWRITE     (M_PWRITE),
188         .S_PSELx      (M_PSELx[`APB_PSELX_GPIO0]),
189         .S_PENABLE    (M_PENABLE),
190         .S_PWDATA     (M_PWDATA),
191         .S_PRDATA     (M_PRDATA[`APB_PSELX_GPIO0*`DATA_WIDTH +: `DATA_WIDTH]),
192         .S_PREADY     (M_PREADY[`APB_PSELX_GPIO0]),
193         .gpio          (gpio0)
194     );
195
196     // GPIO1 for Seven segment displays (16 pin)
197     vmicro16_gpio_apb # (
198         .BUS_WIDTH    (`APB_WIDTH),
199         .DATA_WIDTH    (`DATA_WIDTH),
200         .PORTS        (`APB_GPIO1_PINS),
201         .NAME          ("GPIO1")
202     ) gpio1_apb (
203         .clk          (clk),
204         .reset        (soft_reset),
205         // apb slave to master interface
206         .S_PADDR      (M_PADDR),
207         .S_PWRITE     (M_PWRITE),
208         .S_PSELx      (M_PSELx[`APB_PSELX_GPIO1]),
209         .S_PENABLE    (M_PENABLE),
210         .S_PWDATA     (M_PWDATA),
211         .S_PRDATA     (M_PRDATA[`APB_PSELX_GPIO1*`DATA_WIDTH +: `DATA_WIDTH]),
212         .S_PREADY     (M_PREADY[`APB_PSELX_GPIO1]),
213         .gpio          (gpio1)

```

```

214 );
215
216 // GPIO2 for Seven segment displays (8 pin)
217 vmicro16_gpio_apb # (
218     .BUS_WIDTH    (`APB_WIDTH),
219     .DATA_WIDTH    (`DATA_WIDTH),
220     .PORTS         (`APB_GPIO2_PINS),
221     .NAME          ("GPIO2")
222 ) gpio2_apb (
223     .clk            (clk),
224     .reset          (soft_reset),
225     // apb slave to master interface
226     .S_PADDR        (M_PADDR),
227     .S_PWRITE        (M_PWRITE),
228     .S_PSELx         (M_PSELx[`APB_PSELX_GPIO2]),
229     .S_PENABLE        (M_PENABLE),
230     .S_PWDATA        (M_PWDATA),
231     .S_PRDATA        (M_PRDATA[`APB_PSELX_GPIO2*`DATA_WIDTH +: `DATA_WIDTH]),
232     .S_PREADY        (M_PREADY[`APB_PSELX_GPIO2]),
233     .gpio            (gpio2)
234 );
235
236 apb_uart_tx # (
237     .DATA_WIDTH    (8),
238     .ADDR_EXP      (4) //2~4 = 16 FIFO words
239 ) uart0_apb (
240     .clk            (clk),
241     .reset          (soft_reset),
242     // apb slave to master interface
243     .S_PADDR        (M_PADDR),
244     .S_PWRITE        (M_PWRITE),
245     .S_PSELx         (M_PSELx[`APB_PSELX_UART0]),
246     .S_PENABLE        (M_PENABLE),
247     .S_PWDATA        (M_PWDATA),
248     .S_PRDATA        (M_PRDATA[`APB_PSELX_UART0*`DATA_WIDTH +: `DATA_WIDTH]),
249     .S_PREADY        (M_PREADY[`APB_PSELX_UART0]),
250     // uart wires
251     .tx_wire         (uart_tx),
252     .rx_wire         ()
253 );
254
255 timer_apb timr0 (
256     .clk            (clk),
257     .reset          (soft_reset),
258     // apb slave to master interface
259     .S_PADDR        (M_PADDR),
260     .S_PWRITE        (M_PWRITE),
261     .S_PSELx         (M_PSELx[`APB_PSELX_TIMRO]),
262     .S_PENABLE        (M_PENABLE),
263     .S_PWDATA        (M_PWDATA),
264     .S_PRDATA        (M_PRDATA[`APB_PSELX_TIMRO*`DATA_WIDTH +: `DATA_WIDTH]),
265     .S_PREADY        (M_PREADY[`APB_PSELX_TIMRO])
266     //
267     `ifdef DEF_ENABLE_INT
268     ,.out             (ints [`DEF_INT_TIMRO]),
269     .int_data         (ints_data[`DEF_INT_TIMRO*`DATA_WIDTH +: `DATA_WIDTH])
270     `endif
271 );
272
273 // Shared register set for system-on-chip info
274 // RO = number of cores
275 vmicro16_regs_apb # (
276     .BUS_WIDTH        (`APB_WIDTH),
277     .DATA_WIDTH        (`DATA_WIDTH),
278     .CELL_DEPTH        (8),
279     .PARAM_DEFAULTS_RO (`CORES),
280     .PARAM_DEFAULTS_R1 (`SLAVES)
281 ) regs0_apb (
282     .clk            (clk),
283     .reset          (soft_reset),
284     // apb slave to master interface
285     .S_PADDR        (M_PADDR),
286     .S_PWRITE        (M_PWRITE),
287     .S_PSELx         (M_PSELx[`APB_PSELX_REGS0]),
288     .S_PENABLE        (M_PENABLE),
289     .S_PWDATA        (M_PWDATA),
290     .S_PRDATA        (M_PRDATA[`APB_PSELX_REGS0*`DATA_WIDTH +: `DATA_WIDTH]),
291     .S_PREADY        (M_PREADY[`APB_PSELX_REGS0])
292 );
293
294 vmicro16_bram_ex_apb # (
295     .BUS_WIDTH        (`APB_WIDTH),
296     .MEM_WIDTH         (`DATA_WIDTH),
297     .MEM_DEPTH         (`APB_BRAMO_CELLS),
298     .CORE_ID_BITS      (`clog2(`CORES))
299 ) bram_apb (
300     .clk            (clk),
301     .reset          (soft_reset),
302     // apb slave to master interface

```



```

303     .S_PADDR      (M_PADDR),
304     .S_PWRITE     (M_PWRITE),
305     .S_PSELx      (M_PSELx[`APB_PSELX_BRAMO]),
306     .S_PENABLE     (M_PENABLE),
307     .S_PWDATA      (M_PWDATA),
308     .S_PRDATA      (M_PRDATA[`APB_PSELX_BRAMO*`DATA_WIDTH +: `DATA_WIDTH]),
309     .S_PREADY      (M_PREADY[`APB_PSELX_BRAMO])
310 );
311
312 // There must be atleast 1 core
313 `static_assert(`CORES > 0)
314 `static_assert(`DEF_MEM_INSTR_DEPTH > 0)
315 `static_assert(`DEF_MMU_TIMO_CELLS > 0)
316
317
318 // Single instruction memory
319 `ifndef DEF_CORE_HAS_INSTR_MEM
320 // slave input/outputs from interconnect
321 wire [`APB_WIDTH-1:0]      instr_M_PADDR;
322 wire                      instr_M_PWRITE;
323 wire [1-1:0]              instr_M_PSELx; // not shared
324 wire                      instr_M_PENABLE;
325 wire [`DATA_WIDTH-1:0]    instr_M_PWDATA;
326 wire [1*`DATA_WIDTH-1:0]  instr_M_PRDATA; // slave response
327 wire [1-1:0]              instr_M_PREADY; // slave response
328
329 // Master apb interfaces
330 wire [`CORES*`APB_WIDTH-1:0] instr_w_PADDR;
331 wire [`CORES-1:0]          instr_w_PWRITE;
332 wire [`CORES-1:0]          instr_w_PSELx;
333 wire [`CORES-1:0]          instr_w_PENABLE;
334 wire [`CORES*`DATA_WIDTH-1:0] instr_w_PWDATA;
335 wire [`CORES*`DATA_WIDTH-1:0] instr_w_PRDATA;
336 wire [`CORES-1:0]          instr_w_PREADY;
337
338 `ifndef DEF_USE_REPROG
339 wire [`clog2(`DEF_MEM_INSTR_DEPTH)-1:0] prog_addr;
340 wire [`DATA_WIDTH-1:0] prog_data;
341 wire prog_we;
342 uart_prog rom_prog (
343     .clk      (clk),
344     .reset    (reset | wdreset),
345     // input stream
346     .uart_rx  (uart_rx),
347     // programmer
348     .addr     (prog_addr),
349     .data     (prog_data),
350     .we       (prog_we),
351     .prog     (prog_prog)
352 );
353 `endif
354
355 `ifndef DEF_USE_REPROG
356 vmicro16_bram_prog_apb
357 `else
358 vmicro16_bram_apb
359 `endif
360 # (
361     .BUS_WIDTH      (`APB_WIDTH),
362     .MEM_WIDTH      (`DATA_WIDTH),
363     .MEM_DEPTH      (`DEF_MEM_INSTR_DEPTH),
364     .USE_INITS      (1),
365     .NAME           ("INSTR_ROM_G")
366 ) instr_rom_apb (
367     .clk      (clk),
368     .reset    (reset),
369     .S_PADDR  (instr_M_PADDR),
370     .S_PWRITE (0),
371     .S_PSELx  (instr_M_PSELx),
372     .S_PENABLE (instr_M_PENABLE),
373     .S_PWDATA (0),
374     .S_PRDATA (instr_M_PRDATA),
375     .S_PREADY (instr_M_PREADY)
376
377     `ifndef DEF_USE_REPROG
378     ,
379     .addr     (prog_addr),
380     .data     (prog_data),
381     .we       (prog_we),
382     .prog     (prog_prog)
383     `endif
384 );
385
386 apb_intercon_s # (
387     .MASTER_PORTS  (`CORES),
388     .SLAVE_PORTS    (1),
389     .BUS_WIDTH      (`APB_WIDTH),
390     .DATA_WIDTH      (`DATA_WIDTH),
391     .HAS_PSELX_ADDR (0)

```

```

392 ) apb_instr_intercon (
393     .clk          (clk),
394     .reset        (soft_reset),
395     // APB master from cores
396     // master
397     .S_PADDR      (instr_w_PADDR),
398     .S_PWRITE     (instr_w_PWRITE),
399     .S_PSELx      (instr_w_PSELx),
400     .S_PENABLE    (instr_w_PENABLE),
401     .S_PWDATA     (instr_w_PWDATA),
402     .S_PRDATA     (instr_w_PRDATA),
403     .S_PREADY     (instr_w_PREADY),
404     // shared bus slaves
405     // slave outputs
406     .M_PADDR      (instr_M_PADDR),
407     .M_PWRITE     (instr_M_PWRITE),
408     .M_PSELx      (instr_M_PSELx),
409     .M_PENABLE    (instr_M_PENABLE),
410     .M_PWDATA     (instr_M_PWDATA),
411     .M_PRDATA     (instr_M_PRDATA),
412     .M_PREADY     (instr_M_PREADY)
413 );
414 `endif
415
416 genvar i;
417 generate for(i = 0; i < `CORES; i = i + 1) begin : cores
418
419     vmicro16_core # (
420         .CORE_ID          (i),
421         .DATA_WIDTH       (`DATA_WIDTH),
422
423         .MEM_INSTR_DEPTH  (`DEF_MEM_INSTR_DEPTH),
424         .MEM_SCRATCH_DEPTH (`DEF_MMU_TIMO_CELLS)
425     ) c1 (
426         .clk          (clk),
427         .reset        (soft_reset),
428
429         // debug
430         .halt         (w_halt[i]),
431
432         // interrupts
433         .ints         (ints),
434         .ints_data    (ints_data),
435
436         // Output master port 1
437         .w_PADDR      (w_PADDR  [`APB_WIDTH*i +: `APB_WIDTH] ),
438         .w_PWRITE     (w_PWRITE [i] ),
439         .w_PSELx      (w_PSELx  [i] ),
440         .w_PENABLE    (w_PENABLE [i] ),
441         .w_PWDATA     (w_PWDATA [`DATA_WIDTH*i +: `DATA_WIDTH]),
442         .w_PRDATA     (w_PRDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
443         .w_PREADY     (w_PREADY [i] )
444
445     `ifndef DEF_CORE_HAS_INSTR_MEM
446         // APB instruction rom
447         , // Output master port 2
448         .w2_PADDR     (instr_w_PADDR  [`APB_WIDTH*i +: `APB_WIDTH] ),
449         // .w2_PWRITE (instr_w_PWRITE [i] ),
450         .w2_PSELx     (instr_w_PSELx  [i] ),
451         .w2_PENABLE   (instr_w_PENABLE [i] ),
452         // .w2_PWDATA (instr_w_PWDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
453         .w2_PRDATA    (instr_w_PRDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
454         .w2_PREADY    (instr_w_PREADY [i] )
455     `endif
456 );
457 end
458 endgenerate
459
460
461 //////////////////////////////////////
462 // Formal Verification
463 //////////////////////////////////////
464 `ifdef FORMAL
465 wire all_halted = &w_halt;
466 //////////////////////////////////////
467 // Count number of clocks each core is spending on
468 // bus transactions
469 //////////////////////////////////////
470 reg [15:0] bus_core_times    [0:`CORES-1]; // bus work
471 reg [15:0] core_work_times   [0:`CORES-1]; // serial work
472 reg [15:0] instr_fetch_times [0:`CORES-1]; // instruction fetches
473 integer i2;
474 initial
475     for(i2 = 0; i2 < `CORES; i2 = i2 + 1) begin
476         bus_core_times[i2] = 0;
477         core_work_times[i2] = 0;
478     end
479
480 // total bus time

```

```

481 generate
482     genvar g2;
483     for (g2 = 0; g2 < `CORES; g2 = g2 + 1) begin : formal_for_times
484         always @(posedge clk) begin
485             if (w_PSELx[g2])
486                 bus_core_times[g2] <= bus_core_times[g2] + 1;
487
488             // Core working time
489             `ifndef DEF_CORE_HAS_INSTR_MEM
490                 if (!w_PSELx[g2] && !instr_w_PSELx[g2])
491                     `else
492                         if (!w_PSELx[g2])
493                             `endif
494                             if (!w_halt[g2])
495                                 core_work_times[g2] <= core_work_times[g2] + 1;
496
497         end
498     end
499 endgenerate
500
501 reg [15:0] bus_time_average = 0;
502 reg [15:0] bus_reqs_average = 0;
503 reg [15:0] fetch_time_average = 0;
504 reg [15:0] work_time_average = 0;
505 //
506 always @(all_halted) begin
507     for (i2 = 0; i2 < `CORES; i2 = i2 + 1) begin
508         bus_time_average = bus_time_average + bus_core_times[i2];
509         bus_reqs_average = bus_reqs_average + bus_core_reqs_count[i2];
510         work_time_average = work_time_average + core_work_times[i2];
511         fetch_time_average = fetch_time_average + instr_fetch_times[i2];
512     end
513
514     bus_time_average = bus_time_average / `CORES;
515     bus_reqs_average = bus_reqs_average / `CORES;
516     work_time_average = work_time_average / `CORES;
517     fetch_time_average = fetch_time_average / `CORES;
518 end
519
520 ////////////////////////////////////////////////////
521 // Count number of bus requests per core
522 ////////////////////////////////////////////////////
523 // 1 clock delay of w_PSELx
524 reg [`CORES-1:0] bus_core_reqs_last;
525 // rising edges of each
526 wire [`CORES-1:0] bus_core_reqs_real;
527 // storage for counters for each core
528 reg [15:0] bus_core_reqs_count [0:`CORES-1];
529 initial
530     for(i2 = 0; i2 < `CORES; i2 = i2 + 1)
531         bus_core_reqs_count[i2] = 0;
532
533 // 1 clk delay to detect rising edge
534 always @(posedge clk)
535     bus_core_reqs_last <= w_PSELx;
536
537 generate
538     genvar g3;
539     for (g3 = 0; g3 < `CORES; g3 = g3 + 1) begin : formal_for_reqs
540         // Detect new reqs for each core
541         assign bus_core_reqs_real[g3] = w_PSELx[g3] >
542             bus_core_reqs_last[g3];
543
544         always @(posedge clk)
545             if (bus_core_reqs_real[g3])
546                 bus_core_reqs_count[g3] <= bus_core_reqs_count[g3] + 1;
547     end
548 endgenerate
549
550
551 `ifndef DEF_CORE_HAS_INSTR_MEM
552     ////////////////////////////////////////////////////
553     // Time waiting for instruction fetches
554     // from global memory
555     ////////////////////////////////////////////////////
556     integer i3;
557     initial
558         for(i3 = 0; i3 < `CORES; i3 = i3 + 1)
559             instr_fetch_times[i3] = 0;
560
561     // total bus time
562     // Instruction fetches occur on the w2 master port
563     generate
564         genvar g4;
565         for (g4 = 0; g4 < `CORES; g4 = g4 + 1) begin : formal_for_fetch_times
566             always @(posedge clk)
567                 if (instr_w_PSELx[g4])
568                     instr_fetch_times[g4] <= instr_fetch_times[g4] + 1;
569         end
570     endgenerate

```

```

570         end
571     endgenerate
572 `endif
573
574
575 `endif // end FORMAL
576
577 endmodule

```

E.1.4 vmicro16.v

Vmicro16 CPU core module.

```

1  // This file contains multiple modules.
2  // Verilator likes 1 file for each module
3  /* verilator lint_off DECLFILENAME */
4  /* verilator lint_off UNUSED */
5  /* verilator lint_off BLKSEQ */
6  /* verilator lint_off WIDTH */
7
8  // Include Vmicro16 ISA containing definitions for the bits
9  `include "vmicro16_isa.v"
10
11  `include "clog2.v"
12  `include "formal.v"
13
14
15
16  // This module aims to be a SYNCHRONOUS, WRITE_FIRST BLOCK RAM
17  // https://www.xilinx.com/support/documentation/user_guides/ug473_7Series_Memory_Resources.pdf
18  // https://www.xilinx.com/support/documentation/user_guides/ug383.pdf
19  // https://www.xilinx.com/support/documentation/sw_manuals/xilinx2016_4/ug901-vivado-synthesis.pdf
20  module vmicro16_bram # (
21      parameter MEM_WIDTH      = 16,
22      parameter MEM_DEPTH      = 64,
23      parameter CORE_ID        = 0,
24      parameter USE_INITS      = 0,
25      parameter PARAM_DEFAULTS_R0 = 0,
26      parameter PARAM_DEFAULTS_R1 = 0,
27      parameter PARAM_DEFAULTS_R2 = 0,
28      parameter PARAM_DEFAULTS_R3 = 0,
29      parameter NAME            = "BRAM"
30  ) (
31      input clk,
32      input reset,
33
34      input      [`clog2(MEM_DEPTH)-1:0] mem_addr,
35      input      [MEM_WIDTH-1:0] mem_in,
36      input      mem_we,
37      output reg [MEM_WIDTH-1:0] mem_out
38  );
39  // memory vector
40  (* ram_style = "block" *)
41  reg [MEM_WIDTH-1:0] mem [0:MEM_DEPTH-1];
42
43  // not synthesizable
44  integer i;
45  initial begin
46      for (i = 0; i < MEM_DEPTH; i = i + 1) mem[i] = 0;
47      mem[0] = PARAM_DEFAULTS_R0;
48      mem[1] = PARAM_DEFAULTS_R1;
49      mem[2] = PARAM_DEFAULTS_R2;
50      mem[3] = PARAM_DEFAULTS_R3;
51
52      if (USE_INITS) begin
53          //`define TEST_SW
54          `ifdef TEST_SW
55              $readmemh("E:\\Projects\\uni\\vmicro16\\sw\\verilog_memh.txt", mem);
56          `endif
57
58          `define TEST_ASM
59          `ifdef TEST_ASM
60              $readmemh("E:\\Projects\\uni\\vmicro16\\sw\\asm.s.hex", mem);
61          `endif
62
63          //`define TEST_COND
64          `ifdef TEST_COND
65              mem[0] = {`VMICRO16_OP_MOVI, 3'h7, 8'hC0}; // lock
66              mem[0] = {`VMICRO16_OP_MOVI, 3'h7, 8'hC0}; // lock
67          `endif
68
69          //`define TEST_CMP
70          `ifdef TEST_CMP
71              mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'h0A};

```

```

72     mem[1] = {`VMICRO16_OP_MOVI,      3'h1, 8'h0B};
73     mem[2] = {`VMICRO16_OP_CMP,      3'h1, 3'h0, 5'h1};
74     `endif
75
76     //`define TEST_LWEX
77     `ifdef TEST_LWEX
78     mem[0] = {`VMICRO16_OP_MOVI,      3'h0, 8'hC5};
79     mem[1] = {`VMICRO16_OP_SW,        3'h0, 3'h0, 5'h1};
80     mem[2] = {`VMICRO16_OP_LW,        3'h2, 3'h0, 5'h1};
81     mem[3] = {`VMICRO16_OP_LWEX,      3'h2, 3'h0, 5'h1};
82     mem[4] = {`VMICRO16_OP_SWEX,      3'h3, 3'h0, 5'h1};
83     `endif
84
85     //`define TEST_MULTICORE
86     `ifdef TEST_MULTICORE
87     mem[0] = {`VMICRO16_OP_MOVI,      3'h0, 8'h90};
88     mem[1] = {`VMICRO16_OP_MOVI,      3'h1, 8'h33};
89     mem[2] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
90     mem[3] = {`VMICRO16_OP_MOVI,      3'h0, 8'h80};
91     mem[4] = {`VMICRO16_OP_LW,        3'h2, 3'h0, 5'h0};
92     mem[5] = {`VMICRO16_OP_MOVI,      3'h1, 8'h33};
93     mem[6] = {`VMICRO16_OP_MOVI,      3'h1, 8'h33};
94     mem[7] = {`VMICRO16_OP_MOVI,      3'h1, 8'h33};
95     mem[8] = {`VMICRO16_OP_MOVI,      3'h0, 8'h91};
96     mem[9] = {`VMICRO16_OP_SW,        3'h2, 3'h0, 5'h0};
97     `endif
98
99     //`define TEST_BR
100    `ifdef TEST_BR
101    mem[0] = {`VMICRO16_OP_MOVI,      3'h0, 8'h0};
102    mem[1] = {`VMICRO16_OP_MOVI,      3'h3, 8'h3};
103    mem[2] = {`VMICRO16_OP_MOVI,      3'h1, 8'h2};
104    mem[3] = {`VMICRO16_OP_ARITH_U,    3'h0, 3'h1, 5'b11111};
105    mem[4] = {`VMICRO16_OP_BR,        3'h3, `VMICRO16_OP_BR_U};
106    mem[5] = {`VMICRO16_OP_MOVI,      3'h0, 8'hFF};
107    `endif
108
109    //`define ALL_TEST
110    `ifdef ALL_TEST
111    // Standard all test
112    // REGSO
113    mem[0] = {`VMICRO16_OP_MOVI,      3'h0, 8'h81};
114    mem[1] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0}; // MMU[0x81] = 6
115    mem[2] = {`VMICRO16_OP_SW,        3'h2, 3'h0, 5'h1}; // MMU[0x82] = 6
116    // GPIO0
117    mem[3] = {`VMICRO16_OP_MOVI,      3'h0, 8'h90};
118    mem[4] = {`VMICRO16_OP_MOVI,      3'h1, 8'hD};
119    mem[5] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
120    mem[6] = {`VMICRO16_OP_LW,        3'h2, 3'h0, 5'h0};
121    // TIMO
122    mem[7] = {`VMICRO16_OP_MOVI,      3'h0, 8'h07};
123    mem[8] = {`VMICRO16_OP_LW,        3'h3, 3'h0, 5'h03};
124    // UART0
125    mem[9] = {`VMICRO16_OP_MOVI,      3'h0, 8'hA0}; // UART0
126    mem[10] = {`VMICRO16_OP_MOVI,     3'h1, 8'h41}; // ascii A
127    mem[11] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
128    mem[12] = {`VMICRO16_OP_MOVI,     3'h1, 8'h42}; // ascii B
129    mem[13] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
130    mem[14] = {`VMICRO16_OP_MOVI,     3'h1, 8'h43}; // ascii C
131    mem[15] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
132    mem[16] = {`VMICRO16_OP_MOVI,     3'h1, 8'h44}; // ascii D
133    mem[17] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
134    mem[18] = {`VMICRO16_OP_MOVI,     3'h1, 8'h45}; // ascii D
135    mem[19] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
136    mem[20] = {`VMICRO16_OP_MOVI,     3'h1, 8'h46}; // ascii E
137    mem[21] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
138    // BRAMO
139    mem[22] = {`VMICRO16_OP_MOVI,      3'h0, 8'hC0};
140    mem[23] = {`VMICRO16_OP_MOVI,      3'h1, 8'hA};
141    mem[24] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h5};
142    mem[25] = {`VMICRO16_OP_LW,        3'h2, 3'h0, 5'h5};
143    // GPIO1 (SSD 24-bit port)
144    mem[26] = {`VMICRO16_OP_MOVI,      3'h0, 8'h91};
145    mem[27] = {`VMICRO16_OP_MOVI,      3'h1, 8'h12};
146    mem[28] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
147    mem[29] = {`VMICRO16_OP_LW,        3'h2, 3'h0, 5'h0};
148    // GPIO2
149    mem[30] = {`VMICRO16_OP_MOVI,      3'h0, 8'h92};
150    mem[31] = {`VMICRO16_OP_MOVI,      3'h1, 8'h56};
151    mem[32] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h0};
152    `endif
153
154    //`define TEST_BRAM
155    `ifdef TEST_BRAM
156    // 2 core BRAMO test
157    mem[0] = {`VMICRO16_OP_MOVI,      3'h0, 8'hC0};
158    mem[1] = {`VMICRO16_OP_MOVI,      3'h1, 8'hA};
159    mem[2] = {`VMICRO16_OP_SW,        3'h1, 3'h0, 5'h5};
160    mem[3] = {`VMICRO16_OP_LW,        3'h2, 3'h0, 5'h5};

```

```

161         `endif
162     end
163 end
164
165 always @(posedge clk) begin
166     // synchronous WRITE_FIRST (page 13)
167     if (mem_we) begin
168         mem[mem_addr] <= mem_in;
169         $display($time, "\t\t%s[%h] <= %h",
170             NAME, mem_addr, mem_in);
171     end else
172         mem_out <= mem[mem_addr];
173 end
174
175 // TODO: Reset impl = every clock while reset is asserted, clear each cell
176 //       one at a time, mem[i++] <= 0
177 endmodule
178
179
180 module vmicro16_core_mmu # (
181     parameter MEM_WIDTH    = 16,
182     parameter MEM_DEPTH    = 64,
183
184     parameter CORE_ID      = 3'h0,
185     parameter CORE_ID_BITS = `clog2(`CORES)
186 ) (
187     input clk,
188     input reset,
189
190     input req,
191     output busy,
192
193     // From core
194     input [MEM_WIDTH-1:0] mmu_addr,
195     input [MEM_WIDTH-1:0] mmu_in,
196     input mmu_we,
197     input mmu_lwex,
198     input mmu_swex,
199     output reg [MEM_WIDTH-1:0] mmu_out,
200
201     // interrupts
202     output reg [`DATA_WIDTH*`DEF_NUM_INT-1:0] ints_vector,
203     output reg [`DEF_NUM_INT-1:0] ints_mask,
204
205     // TO APB interconnect
206     output reg [`APB_WIDTH-1:0] M_PADDR,
207     output reg M_PWRITE,
208     output reg M_PSELx,
209     output reg M_PENABLE,
210     output reg [MEM_WIDTH-1:0] M_PWDATA,
211     // from interconnect
212     input [MEM_WIDTH-1:0] M_PRDATA,
213     input M_PREADY
214 );
215 localparam MMU_STATE_T1 = 0;
216 localparam MMU_STATE_T2 = 1;
217 localparam MMU_STATE_T3 = 2;
218 reg [1:0] mmu_state = MMU_STATE_T1;
219
220 reg [MEM_WIDTH-1:0] per_out = 0;
221 wire [MEM_WIDTH-1:0] tim0_out;
222
223 assign busy = req || (mmu_state == MMU_STATE_T2);
224
225 // more luts than below but easier
226 //wire tim0_en = (mmu_addr >= `DEF_MMU_TIMO_S);
227 //    && (mmu_addr <= `DEF_MMU_TIMO_E);
228 //wire sreg_en = (mmu_addr >= `DEF_MMU_SREG_S);
229 //    && (mmu_addr <= `DEF_MMU_SREG_E);
230 //wire intv_en = (mmu_addr >= `DEF_MMU_INTSV_S);
231 //    && (mmu_addr <= `DEF_MMU_INTSV_E);
232 //wire intm_en = (mmu_addr >= `DEF_MMU_INTSM_S);
233 //    && (mmu_addr <= `DEF_MMU_INTSM_E);
234
235 wire tim0_en = ~mmu_addr[12] && ~mmu_addr[9] && ~mmu_addr[7];
236 wire sreg_en = mmu_addr[7] && ~mmu_addr[4] && ~mmu_addr[5];
237 wire intv_en = mmu_addr[8] && ~mmu_addr[3];
238 wire intm_en = mmu_addr[8] && mmu_addr[3];
239
240 wire apb_en = !(tim0_en, sreg_en, intv_en, intm_en);
241 wire tim0_we = (tim0_en && mmu_we);
242 wire intv_we = (intv_en && mmu_we);
243 wire intm_we = (intm_en && mmu_we);
244
245 // Special register selects
246 localparam SPECIAL_REGS = 8;
247 wire [MEM_WIDTH-1:0] sr_val;
248
249 // Interrupt vector and mask
250 initial ints_vector = 0;

```

```

251     initial ints_mask = 0;
252     wire [2:0] intv_addr = mmu_addr[`clog2(`DEF_NUM_INT)-1:0];
253     always @(posedge clk)
254         if (intv_we)
255             ints_vector[intv_addr*`DATA_WIDTH +: `DATA_WIDTH] <= mmu_in;
256
257     always @(posedge clk)
258         if (intm_we)
259             ints_mask <= mmu_in;
260
261
262     always @(ints_vector)
263         $display($time,
264             "\tC%d\t\tints_vector W: | %h %h %h %h | %h %h %h %h |",
265             CORE_ID,
266             ints_vector[0*`DATA_WIDTH +: `DATA_WIDTH],
267             ints_vector[1*`DATA_WIDTH +: `DATA_WIDTH],
268             ints_vector[2*`DATA_WIDTH +: `DATA_WIDTH],
269             ints_vector[3*`DATA_WIDTH +: `DATA_WIDTH],
270             ints_vector[4*`DATA_WIDTH +: `DATA_WIDTH],
271             ints_vector[5*`DATA_WIDTH +: `DATA_WIDTH],
272             ints_vector[6*`DATA_WIDTH +: `DATA_WIDTH],
273             ints_vector[7*`DATA_WIDTH +: `DATA_WIDTH]
274         );
275
276     always @(intm_we)
277         $display($time, "\tC%d\t\tintm_we W: %b", CORE_ID, ints_mask);
278
279     // Output port
280     always @(*)
281         if (tim0_en) mmu_out = tim0_out;
282         else if (sreg_en) mmu_out = sr_val;
283         else if (intv_en) mmu_out = ints_vector[mmu_addr[2:0]*`DATA_WIDTH
284             +: `DATA_WIDTH];
285         else if (intm_en) mmu_out = ints_mask;
286         else mmu_out = per_out;
287
288     // APB master to slave interface
289     always @(posedge clk)
290         if (reset) begin
291             mmu_state <= MMU_STATE_T1;
292             M_PENABLE <= 0;
293             M_PADDR <= 0;
294             M_PWDATA <= 0;
295             M_PSELx <= 0;
296             M_PWRITE <= 0;
297         end
298         else
299             casex (mmu_state)
300                 MMU_STATE_T1: begin
301                     if (req && apb_en) begin
302                         M_PADDR <= {mmu_lwex,
303                             mmu_swex,
304                             CORE_ID[CORE_ID_BITS-1:0],
305                             mmu_addr[MEM_WIDTH-1:0]};
306
307                         M_PWDATA <= mmu_in;
308                         M_PSELx <= 1;
309                         M_PWRITE <= mmu_we;
310
311                         mmu_state <= MMU_STATE_T2;
312                     end
313                 end
314
315                 `ifdef FIX_T3
316                 MMU_STATE_T2: begin
317                     M_PENABLE <= 1;
318
319                     if (M_PREADY == 1'b1) begin
320                         mmu_state <= MMU_STATE_T3;
321                     end
322                 end
323
324                 MMU_STATE_T3: begin
325                     // Slave has output a ready signal (finished)
326                     M_PENABLE <= 0;
327                     M_PADDR <= 0;
328                     M_PWDATA <= 0;
329                     M_PSELx <= 0;
330                     M_PWRITE <= 0;
331                     // Clock the peripheral output into a reg,
332                     // to output on the next clock cycle
333                     per_out <= M_PRDATA;
334
335                     mmu_state <= MMU_STATE_T1;
336                 end
337                 `else
338                 // No FIX_T3
339                 MMU_STATE_T2: begin
340                     if (M_PREADY == 1'b1) begin

```

```

341             M_PENABLE <= 0;
342             M_PADDR <= 0;
343             M_PWDATA <= 0;
344             M_PSELx <= 0;
345             M_PWRITE <= 0;
346             // Clock the peripheral output into a reg,
347             // to output on the next clock cycle
348             per_out <= M_PRDATA;
349
350             mmu_state <= MMU_STATE_T1;
351         end else begin
352             M_PENABLE <= 1;
353         end
354     end
355 `endif
356 endcase
357
358 (* ram_style = "block" *)
359 vmicro16_bram # (
360     .MEM_WIDTH (MEM_WIDTH),
361     .MEM_DEPTH (SPECIAL_REGS),
362     .USE_INITS (0),
363     .PARAM_DEFAULTS_R0 (CORE_ID),
364     .PARAM_DEFAULTS_R1 (`CORES),
365     .PARAM_DEFAULTS_R2 (`APB_BRAMO_CELLS),
366     .PARAM_DEFAULTS_R3 (`SLAVES),
367     .NAME ("ram_sr")
368 ) ram_sr (
369     .clk (clk),
370     .reset (reset),
371     .mem_addr (mmu_addr[`clog2(SPECIAL_REGS)-1:0]),
372     .mem_in (),
373     .mem_we (),
374     .mem_out (sr_val)
375 );
376
377 // Each M core has a TIMO scratch memory
378 (* ram_style = "block" *)
379 vmicro16_bram # (
380     .MEM_WIDTH (MEM_WIDTH),
381     .MEM_DEPTH (MEM_DEPTH),
382     .USE_INITS (0),
383     .NAME ("TIMO")
384 ) TIMO (
385     .clk (clk),
386     .reset (reset),
387     .mem_addr (mmu_addr[7:0]),
388     .mem_in (mmu_in),
389     .mem_we (tim0_we),
390     .mem_out (tim0_out)
391 );
392 endmodule
393
394
395
396 module vmicro16_regs # (
397     parameter CELL_WIDTH = 16,
398     parameter CELL_DEPTH = 8,
399     parameter CELL_SEL_BITS = `clog2(CELL_DEPTH),
400     parameter CELL_DEFAULTS = 0,
401     parameter DEBUG_NAME = "",
402     parameter CORE_ID = 0,
403     parameter PARAM_DEFAULTS_R0 = 16'h0000,
404     parameter PARAM_DEFAULTS_R1 = 16'h0000
405 ) (
406     input clk,
407     input reset,
408     // Dual port register reads
409     input [CELL_SEL_BITS-1:0] rs1, // port 1
410     output [CELL_WIDTH-1:0] rd1,
411     //input [CELL_SEL_BITS-1:0] rs2, // port 2
412     //output [CELL_WIDTH-1:0] rd2,
413     // EX/WB final stage write back
414     input we,
415     input [CELL_SEL_BITS-1:0] ws1,
416     input [CELL_WIDTH-1:0] wd
417 );
418 (* ram_style = "distributed" *)
419 reg [CELL_WIDTH-1:0] regs [0:CELL_DEPTH-1] /*verilator public_flat*/;
420
421 // Initialise registers with default values
422 // Really only used for special registers used by the soc
423 // TODO: How to do this on reset?
424 integer i;
425 initial
426     if (CELL_DEFAULTS)
427         $readmemh(CELL_DEFAULTS, regs);
428     else begin
429         for(i = 0; i < CELL_DEPTH; i = i + 1)

```



```

430         regs[i] = 0;
431         regs[0] = PARAM_DEFAULTS_R0;
432         regs[1] = PARAM_DEFAULTS_R1;
433     end
434
435     `ifdef ICARUS
436         always @(regs)
437             $display($time, "\tC%02h\t\t| %h %h %h %h | %h %h %h %h |",
438                     CORE_ID,
439                     regs[0], regs[1], regs[2], regs[3],
440                     regs[4], regs[5], regs[6], regs[7]);
441     `endif
442
443     always @(posedge clk)
444         if (reset) begin
445             for(i = 0; i < CELL_DEPTH; i = i + 1)
446                 regs[i] <= 0;
447             regs[0] <= PARAM_DEFAULTS_R0;
448             regs[1] <= PARAM_DEFAULTS_R1;
449         end
450         else if (we) begin
451             $display($time, "\tC%02h: REGS #s: Writing %h to reg[%d]",
452                     CORE_ID, DEBUG_NAME, wd, ws1);
453
454             // Perform the write
455             regs[ws1] <= wd;
456         end
457
458         // sync writes, async reads
459         assign rd1 = regs[rs1];
460         //assign rd2 = regs[rs2];
461     endmodule
462
463     module vmicro16_dec # (
464         parameter INSTR_WIDTH = 16,
465         parameter INSTR_OP_WIDTH = 5,
466         parameter INSTR_RS_WIDTH = 3,
467         parameter ALU_OP_WIDTH = 5
468     ) (
469         //input clk, // not used yet (all combinational)
470         //input reset, // not used yet (all combinational)
471
472         input [INSTR_WIDTH-1:0] instr,
473
474         output [INSTR_OP_WIDTH-1:0] opcode,
475         output [INSTR_RS_WIDTH-1:0] rd,
476         output [INSTR_RS_WIDTH-1:0] ra,
477         output [3:0] imm4,
478         output [7:0] imm8,
479         output [11:0] imm12,
480         output [4:0] simm5,
481
482         // This can be freely increased without affecting the isa
483         output reg [ALU_OP_WIDTH-1:0] alu_op,
484
485         output reg has_imm4,
486         output reg has_imm8,
487         output reg has_imm12,
488         output reg has_we,
489         output reg has_br,
490         output reg has_mem,
491         output reg has_mem_we,
492         output reg has_cmp,
493
494         output halt,
495         output intr,
496
497         output reg has_lwex,
498         output reg has_swex
499
500         // TODO: Use to identify bad instruction and
501         // raise exceptions
502         //, output is_bad
503     );
504
505     assign opcode = instr[15:11];
506     assign rd = instr[10:8];
507     assign ra = instr[7:5];
508     assign imm4 = instr[3:0];
509     assign imm8 = instr[7:0];
510     assign imm12 = instr[11:0];
511     assign simm5 = instr[4:0];
512
513     // exme_op
514     always @(*) case (opcode)
515         VMICRO16_OP_SPCL: casez(instr[11:0])
516             VMICRO16_OP_SPCL_NOP,
517             VMICRO16_OP_SPCL_HALT,
518             VMICRO16_OP_SPCL_INTR: alu_op = `VMICRO16_ALU_NOP;
519             default: alu_op = `VMICRO16_ALU_NOP; endcase

```

```

520     `VMICRO16_OP_LW:           alu_op = `VMICRO16_ALU_LW;
521     `VMICRO16_OP_SW:           alu_op = `VMICRO16_ALU_SW;
522     `VMICRO16_OP_LWEX:         alu_op = `VMICRO16_ALU_LW;
523     `VMICRO16_OP_SWEX:         alu_op = `VMICRO16_ALU_SW;
524
525     `VMICRO16_OP_MOV:          alu_op = `VMICRO16_ALU_MOV;
526     `VMICRO16_OP_MOVI:         alu_op = `VMICRO16_ALU_MOVI;
527
528     `VMICRO16_OP_BR:           alu_op = `VMICRO16_ALU_BR;
529     `VMICRO16_OP_MULT:         alu_op = `VMICRO16_ALU_MULT;
530
531     `VMICRO16_OP_CMP:          alu_op = `VMICRO16_ALU_CMP;
532     `VMICRO16_OP_SETC:         alu_op = `VMICRO16_ALU_SETC;
533
534     `VMICRO16_OP_BIT:          casez (simm5)
535     `VMICRO16_OP_BIT_OR:       alu_op = `VMICRO16_ALU_BIT_OR;
536     `VMICRO16_OP_BIT_XOR:      alu_op = `VMICRO16_ALU_BIT_XOR;
537     `VMICRO16_OP_BIT_AND:      alu_op = `VMICRO16_ALU_BIT_AND;
538     `VMICRO16_OP_BIT_NOT:      alu_op = `VMICRO16_ALU_BIT_NOT;
539     `VMICRO16_OP_BIT_LSHFT:    alu_op = `VMICRO16_ALU_BIT_LSHFT;
540     `VMICRO16_OP_BIT_RSHFT:    alu_op = `VMICRO16_ALU_BIT_RSHFT;
541     default:                    alu_op = `VMICRO16_ALU_BAD; endcase
542
543     `VMICRO16_OP_ARITH_U:       casez (simm5)
544     `VMICRO16_OP_ARITH_UADD:    alu_op = `VMICRO16_ALU_ARITH_UADD;
545     `VMICRO16_OP_ARITH_USUB:    alu_op = `VMICRO16_ALU_ARITH_USUB;
546     `VMICRO16_OP_ARITH_UADDI:   alu_op = `VMICRO16_ALU_ARITH_UADDI;
547     default:                    alu_op = `VMICRO16_ALU_BAD; endcase
548
549     `VMICRO16_OP_ARITH_S:       casez (simm5)
550     `VMICRO16_OP_ARITH_SADD:    alu_op = `VMICRO16_ALU_ARITH_SADD;
551     `VMICRO16_OP_ARITH_SSUB:    alu_op = `VMICRO16_ALU_ARITH_SSUB;
552     `VMICRO16_OP_ARITH_SSUBI:   alu_op = `VMICRO16_ALU_ARITH_SSUBI;
553     default:                    alu_op = `VMICRO16_ALU_BAD; endcase
554
555     default: begin
556         alu_op = `VMICRO16_ALU_NOP;
557         $display($time, "\tDEC: unknown opcode: %h ... NOPPING", opcode);
558     end
559 endcase
560
561 // Special opcodes
562 //assign nop == ((opcode == `VMICRO16_OP_SPCL) & (~instr[0]));
563 assign halt = ((opcode == `VMICRO16_OP_SPCL) & instr[0]);
564 assign intr = ((opcode == `VMICRO16_OP_SPCL) & instr[1]);
565
566 // Register writes
567 always @(*) case (opcode)
568     `VMICRO16_OP_LWEX,
569     `VMICRO16_OP_SWEX,
570     `VMICRO16_OP_LW,
571     `VMICRO16_OP_MOV,
572     `VMICRO16_OP_MOVI,
573     // `VMICRO16_OP_MOVI_L,
574     `VMICRO16_OP_ARITH_U,
575     `VMICRO16_OP_ARITH_S,
576     `VMICRO16_OP_SETC,
577     `VMICRO16_OP_BIT,
578     `VMICRO16_OP_MULT:         has_we = 1'b1;
579     default:                    has_we = 1'b0;
580 endcase
581
582 // Contains 4-bit immediate
583 always @(*)
584     if( ((opcode == `VMICRO16_OP_ARITH_U) && (simm5[4] == 0)) ||
585         ((opcode == `VMICRO16_OP_ARITH_S) && (simm5[4] == 0)) )
586         has_imm4 = 1'b1;
587     else
588         has_imm4 = 1'b0;
589
590 // Contains 8-bit immediate
591 always @(*) case (opcode)
592     `VMICRO16_OP_MOVI,
593     `VMICRO16_OP_BR:         has_imm8 = 1'b1;
594     default:                  has_imm8 = 1'b0;
595 endcase
596
597 // Contains 12-bit immediate
598 //always @(*) case (opcode)
599 //    `VMICRO16_OP_MOVI_L:     has_imm12 = 1'b1;
600 //    default:                  has_imm12 = 1'b0;
601 //endcase
602
603 // Will branch the pc
604 always @(*) case (opcode)
605     `VMICRO16_OP_BR:         has_br = 1'b1;
606     default:                  has_br = 1'b0;
607 endcase
608
609 // Requires external memory

```

```

610     always @(*) case (opcode)
611         `VMICRO16_OP_LW,
612         `VMICRO16_OP_SW,
613         `VMICRO16_OP_LWEX,
614         `VMICRO16_OP_SWEX: has_mem = 1'b1;
615         default:           has_mem = 1'b0;
616     endcase
617
618     // Requires external memory write
619     always @(*) case (opcode)
620         `VMICRO16_OP_SW,
621         `VMICRO16_OP_SWEX: has_mem_we = 1'b1;
622         default:           has_mem_we = 1'b0;
623     endcase
624
625     // Affects status registers (cmp instructions)
626     always @(*) case (opcode)
627         `VMICRO16_OP_CMP:   has_cmp = 1'b1;
628         default:           has_cmp = 1'b0;
629     endcase
630
631     // Performs exclusive checks
632     always @(*) case (opcode)
633         `VMICRO16_OP_LWEX:  has_lwex = 1'b1;
634         default:           has_lwex = 1'b0;
635     endcase
636
637     always @(*) case (opcode)
638         `VMICRO16_OP_SWEX:  has_swex = 1'b1;
639         default:           has_swex = 1'b0;
640     endcase
641 endmodule
642
643
644 module vmicro16_alu # (
645     parameter OP_WIDTH   = 5,
646     parameter DATA_WIDTH = 16,
647     parameter CORE_ID    = 0
648 ) (
649     // input clk, // TODO: make clocked
650
651     input      [OP_WIDTH-1:0] op,
652     input      [DATA_WIDTH-1:0] a, // rs1/dst
653     input      [DATA_WIDTH-1:0] b, // rs2
654     input      [3:0] flags,
655     output reg [DATA_WIDTH-1:0] c
656 );
657
658     localparam TOP_BIT = (DATA_WIDTH-1);
659     // 17-bit register
660     reg [DATA_WIDTH:0] cmp_tmp = 0; // = {carry, [15:0]}
661     wire r_setc;
662
663     always @(*) begin
664         cmp_tmp = 0;
665         case (op)
666             // branch/nop, output nothing
667             `VMICRO16_ALU_BR,
668             `VMICRO16_ALU_NOP: c = {DATA_WIDTH{1'b0}};
669             // load/store addresses (use value in rd2)
670             `VMICRO16_ALU_LW,
671             `VMICRO16_ALU_SW: c = b;
672             // bitwise operations
673             `VMICRO16_ALU_BIT_OR: c = a | b;
674             `VMICRO16_ALU_BIT_XOR: c = a ^ b;
675             `VMICRO16_ALU_BIT_AND: c = a & b;
676             `VMICRO16_ALU_BIT_NOT: c = ~(b);
677             `VMICRO16_ALU_BIT_LSHFT: c = a << b;
678             `VMICRO16_ALU_BIT_RSHFT: c = a >> b;
679
680             `VMICRO16_ALU_MOV: c = b;
681             `VMICRO16_ALU_MOVI: c = b;
682             `VMICRO16_ALU_MOVI_L: c = b;
683
684             `VMICRO16_ALU_ARITH_UADD: c = a + b;
685             `VMICRO16_ALU_ARITH_USUB: c = a - b;
686             // TODO: ALU should have simm5 as input
687             `VMICRO16_ALU_ARITH_UADDI: c = a + b;
688
689             `ifdef DEF_ALU_HW_MULT
690                 `VMICRO16_ALU_MULT: c = a * b;
691             `endif
692
693             `VMICRO16_ALU_ARITH_SADD: c = $signed(a) + $signed(b);
694             `VMICRO16_ALU_ARITH_SSUB: c = $signed(a) - $signed(b);
695             // TODO: ALU should have simm5 as input
696             `VMICRO16_ALU_ARITH_SSUBI: c = $signed(a) - $signed(b);
697
698             `VMICRO16_ALU_CMP: begin
699                 // TODO: Do a-b in 17-bit register
700                 // Set zero, overflow, carry, signed bits in result

```

```

700     cmp_tmp = a - b;
701     c = 0;
702
703     // N Negative condition code flag
704     // Z Zero condition code flag
705     // C Carry condition code flag
706     // V Overflow condition code flag
707     c[`VMICRO16_SFLAG_N] = cmp_tmp[TOP_BIT];
708     c[`VMICRO16_SFLAG_Z] = (cmp_tmp == 0);
709     c[`VMICRO16_SFLAG_C] = 0; //cmp_tmp[TOP_BIT+1]; // not used
710
711     // Overflow flag
712     // https://stackoverflow.com/questions/30957188/
713     // https://github.com/bendl/prco304/blob/master/prco_core/rtl/prco_alu.v#L50
714     case(cmp_tmp[TOP_BIT+1:TOP_BIT])
715         2'b01: c[`VMICRO16_SFLAG_V] = 1;
716         2'b10: c[`VMICRO16_SFLAG_V] = 1;
717         default: c[`VMICRO16_SFLAG_V] = 0;
718     endcase
719
720     $display($time, "\tC%02h: ALU CMP: %h %h = %h = %b", CORE_ID, a, b, cmp_tmp, c[3:0]);
721 end
722
723 `VMICRO16_ALU_SETC: c = { {15{1'b0}}, r_setc };
724
725 // TODO: Parameterise
726 default: begin
727     $display($time, "\tALU: unknown op: %h", op);
728     c = 0;
729     cmp_tmp = 0;
730 end
731     endcase
732 end
733
734 branch setc_check (
735     .flags      (flags),
736     .cond       (b[7:0]),
737     .en         (r_setc)
738 );
739 endmodule
740
741 // flags = 4 bit r_cmp_flags register
742 // cond = 8 bit VMICRO16_OP_BR_? value. See vmicro16_isa.v
743 module branch (
744     input [3:0] flags,
745     input [7:0] cond,
746     output reg en
747 );
748     always @(*)
749         case (cond)
750             `VMICRO16_OP_BR_U: en = 1;
751             `VMICRO16_OP_BR_E: en = (flags[`VMICRO16_SFLAG_Z] == 1);
752             `VMICRO16_OP_BR_NE: en = (flags[`VMICRO16_SFLAG_Z] == 0);
753             `VMICRO16_OP_BR_G: en = (flags[`VMICRO16_SFLAG_Z] == 0) &&
754                 (flags[`VMICRO16_SFLAG_N] == flags[`VMICRO16_SFLAG_V]);
755             `VMICRO16_OP_BR_L: en = (flags[`VMICRO16_SFLAG_Z] != flags[`VMICRO16_SFLAG_N]);
756             `VMICRO16_OP_BR_GE: en = (flags[`VMICRO16_SFLAG_Z] == flags[`VMICRO16_SFLAG_N]);
757             `VMICRO16_OP_BR_LE: en = (flags[`VMICRO16_SFLAG_Z] == 1) ||
758                 (flags[`VMICRO16_SFLAG_N] != flags[`VMICRO16_SFLAG_V]);
759             default: en = 0;
760         endcase
761 endmodule
762
763
764
765 module vmicro16_core # (
766     parameter DATA_WIDTH      = 16,
767     parameter MEM_INSTR_DEPTH = 64,
768     parameter MEM_SCRATCH_DEPTH = 64,
769     parameter MEM_WIDTH       = 16,
770
771     parameter CORE_ID          = 3'h0
772 ) (
773     input      clk,
774     input      reset,
775
776     output [7:0] debug,
777
778     output      halt,
779
780     // interrupt sources
781     input  [`DEF_NUM_INT-1:0] ints,
782     input  [`DEF_NUM_INT*DATA_WIDTH-1:0] ints_data,
783     output [`DEF_NUM_INT-1:0] ints_ack,
784
785     // APB master to slave interface (apb_intercon)
786     output  [`APB_WIDTH-1:0] w_PADDR,
787     output      w_PWRITE,
788     output      w_PSELx,
789     output      w_PENABLE,

```

```

790     output [DATA_WIDTH-1:0] w_PWDATA,
791     input  [DATA_WIDTH-1:0] w_PRDATA,
792     input  [DATA_WIDTH-1:0] w_PREADY
793
794     `ifndef DEF_CORE_HAS_INSTR_MEM
795     , // APB master interface to slave instruction memory
796     output reg [APB_WIDTH-1:0] w2_PADDR,
797     output reg [DATA_WIDTH-1:0] w2_PWRITE,
798     output reg [DATA_WIDTH-1:0] w2_PSELx,
799     output reg [DATA_WIDTH-1:0] w2_PENABLE,
800     output reg [DATA_WIDTH-1:0] w2_PWDATA,
801     input  [DATA_WIDTH-1:0] w2_PRDATA,
802     input  [DATA_WIDTH-1:0] w2_PREADY
803     `endif
804 );
805     localparam STATE_IF = 0;
806     localparam STATE_R1 = 1;
807     localparam STATE_R2 = 2;
808     localparam STATE_ME = 3;
809     localparam STATE_WB = 4;
810     localparam STATE_FE = 5;
811     localparam STATE_IDLE = 6;
812     localparam STATE_HALT = 7;
813     reg [2:0] r_state = STATE_IF;
814
815     reg [DATA_WIDTH-1:0] r_pc = 16'h0000;
816     reg [DATA_WIDTH-1:0] r_pc_saved = 16'h0000;
817     reg [DATA_WIDTH-1:0] r_instr = 16'h0000;
818     wire [DATA_WIDTH-1:0] w_mem_instr_out;
819     wire [DATA_WIDTH-1:0] w_halt;
820
821     assign debug = {7'h00, w_halt};
822     assign halt = w_halt;
823
824     wire [4:0] r_instr_opcode;
825     wire [4:0] r_instr_alu_op;
826     wire [2:0] r_instr_rsd;
827     wire [2:0] r_instr_rsa;
828     reg [DATA_WIDTH-1:0] r_instr_rdd = 0;
829     reg [DATA_WIDTH-1:0] r_instr_rda = 0;
830     wire [3:0] r_instr_imm4;
831     wire [7:0] r_instr_imm8;
832     wire [4:0] r_instr_simm5;
833     wire [4:0] r_instr_has_imm4;
834     wire [4:0] r_instr_has_imm8;
835     wire [4:0] r_instr_has_we;
836     wire [4:0] r_instr_has_br;
837     wire [4:0] r_instr_has_cmp;
838     wire [4:0] r_instr_has_mem;
839     wire [4:0] r_instr_has_mem_we;
840     wire [4:0] r_instr_halt;
841     wire [4:0] r_instr_has_lwex;
842     wire [4:0] r_instr_has_swex;
843
844     wire [DATA_WIDTH-1:0] r_alu_out;
845
846     wire [DATA_WIDTH-1:0] r_mem_scratch_addr = $signed(r_alu_out) + $signed(r_instr_simm5);
847     wire [DATA_WIDTH-1:0] r_mem_scratch_in = r_instr_rdd;
848     wire [DATA_WIDTH-1:0] r_mem_scratch_out;
849     wire [DATA_WIDTH-1:0] r_mem_scratch_we = r_instr_has_mem_we && (r_state == STATE_ME);
850     reg [DATA_WIDTH-1:0] r_mem_scratch_req = 0;
851     wire [DATA_WIDTH-1:0] r_mem_scratch_busy;
852
853     reg [2:0] r_reg_rs1 = 0;
854     wire [DATA_WIDTH-1:0] r_reg_rd1_s;
855     wire [DATA_WIDTH-1:0] r_reg_rd1_i;
856     wire [DATA_WIDTH-1:0] r_reg_rd1 = regs_use_int ? r_reg_rd1_i : r_reg_rd1_s;
857     //wire [15:0] r_reg_rd2;
858     wire [DATA_WIDTH-1:0] r_reg_wd = (r_instr_has_mem) ? r_mem_scratch_out : r_alu_out;
859     wire [DATA_WIDTH-1:0] r_reg_we = r_instr_has_we && (r_state == STATE_WB);
860
861     // branching
862     wire [4:0] w_intr;
863     wire [4:0] w_branch_en;
864     wire [4:0] w_branching = r_instr_has_br && w_branch_en;
865     reg [3:0] r_cmp_flags = 4'h00; // N, Z, C, V
866
867     always @(r_cmp_flags)
868     $display($time, "\tC%02h:\tALU CMP: %b", CORE_ID, r_cmp_flags);
869
870     // 2 cycle register fetch
871     always @(*) begin
872         r_reg_rs1 = 0;
873         if (r_state == STATE_R1)
874             r_reg_rs1 = r_instr_rsd;
875         else if (r_state == STATE_R2)
876             r_reg_rs1 = r_instr_rsa;
877         else
878             r_reg_rs1 = 3'h0;
879     end

```

```

880
881 reg regs_use_int = 0;
882 `ifndef DEF_ENABLE_INT
883 wire [`DEF_NUM_INT*`DATA_WIDTH-1:0] ints_vector;
884 wire [`DEF_NUM_INT-1:0] ints_mask;
885 wire has_int = ints & ints_mask;
886 reg int_pending = 0;
887 reg int_pending_ack = 0;
888 always @(posedge clk)
889     if (int_pending_ack)
890         // We've now branched to the isr
891         int_pending <= 0;
892     else if (has_int)
893         // Notify fsm to switch to the ints_vector at the last stage
894         int_pending <= 1;
895     else if (w_intr)
896         // Return to Interrupt instruction called,
897         // so we've finished with the interrupt
898         int_pending <= 0;
899 `endif
900
901 // Next program counter logic
902 reg [`DATA_WIDTH-1:0] next_pc = 0;
903 always @(posedge clk)
904     if (reset)
905         r_pc <= 0;
906     else if (r_state == STATE_WB) begin
907         `ifndef DEF_ENABLE_INT
908             if (int_pending) begin
909                 $display($time, "\tC%02h: Jumping to ISR: %h",
910                     CORE_ID,
911                     ints_vector[0 +: `DATA_WIDTH]);
912                 // TODO: check bounds
913                 // Save state
914                 r_pc_saved <= r_pc + 1;
915                 regs_use_int <= 1;
916                 int_pending_ack <= 1;
917                 // Jump to ISR
918                 r_pc <= ints_vector[0 +: `DATA_WIDTH];
919             end else if (w_intr) begin
920                 $display($time, "\tC%02h: Returning from ISR: %h",
921                     CORE_ID, r_pc_saved);
922             end
923             // Restore state
924             r_pc <= r_pc_saved;
925             regs_use_int <= 0;
926             int_pending_ack <= 0;
927         end else
928             `endif
929         if (w_branching) begin
930             $display($time, "\tC%02h: branching to %h", CORE_ID, r_instr_rdd);
931             r_pc <= r_instr_rdd;
932         end
933         `ifndef DEF_ENABLE_INT
934             int_pending_ack <= 0;
935         `endif
936     end else if (r_pc < (MEM_INSTR_DEPTH-1)) begin
937         // normal increment
938         // pc <= pc + 1
939         r_pc <= r_pc + 1;
940     end
941     `ifndef DEF_ENABLE_INT
942         int_pending_ack <= 0;
943     `endif
944 end
945 // end r_state == STATE_WB
946 else if (r_state == STATE_HALT) begin
947     `ifndef DEF_ENABLE_INT
948         // Only an interrupt can return from halt
949         // duplicate code form STATE_ME!
950         if (int_pending) begin
951             $display($time, "\tC%02h: Jumping to ISR: %h", CORE_ID, ints_vector[0 +: `DATA_WIDTH]);
952             // TODO: check bounds
953             // Save state
954             r_pc_saved <= r_pc; // + 1; HALT = stay with same PC
955             regs_use_int <= 1;
956             int_pending_ack <= 1;
957             // Jump to ISR
958             r_pc <= ints_vector[0 +: `DATA_WIDTH];
959         end else if (w_intr) begin
960             $display($time, "\tC%02h: Returning from ISR: %h", CORE_ID, r_pc_saved);
961             r_pc <= r_pc_saved;
962             regs_use_int <= 0;
963             int_pending_ack <= 0;
964         end
965     end
966     `endif
967 end
968 `ifndef DEF_CORE_HAS_INSTR_MEM
969     initial w2_PSELx = 0;

```

```

970     initial w2_PENABLE = 0;
971     initial w2_PADDR   = 0;
972 `endif
973
974     // cpu state machine
975     always @(posedge clk)
976         if (reset) begin
977             r_state      <= STATE_IF;
978             r_instr      <= 0;
979             r_mem_scratch_req <= 0;
980             r_instr_rdd   <= 0;
981             r_instr_rda   <= 0;
982         end
983         else begin
984
985 `ifdef DEF_CORE_HAS_INSTR_MEM
986             if (r_state == STATE_IF) begin
987                 r_instr <= w_mem_instr_out;
988
989                 $display("");
990                 $display($time, "\tC%02h: PC: %h", CORE_ID, r_pc);
991                 $display($time, "\tC%02h: INSTR: %h", CORE_ID, w_mem_instr_out);
992
993                 r_state <= STATE_R1;
994             end
995 `else
996             // wait for global instruction rom to give us our instruction
997             if (r_state == STATE_IF) begin
998                 // wait for ready signal
999                 if (!w2_PREADY) begin
1000                     w2_PSELx <= 1;
1001                     w2_PWRITE <= 0;
1002                     w2_PENABLE <= 1;
1003                     w2_PWDATA <= 0;
1004                     w2_PADDR <= r_pc;
1005                 end else begin
1006                     w2_PSELx <= 0;
1007                     w2_PWRITE <= 0;
1008                     w2_PENABLE <= 0;
1009                     w2_PWDATA <= 0;
1010
1011                     r_instr <= w2_PRDATA;
1012
1013                     $display("");
1014                     $display($time, "\tC%02h: PC: %h", CORE_ID, r_pc);
1015                     $display($time, "\tC%02h: INSTR: %h", CORE_ID, w2_PRDATA);
1016
1017                     r_state <= STATE_R1;
1018                 end
1019             end
1020 `endif
1021
1022         else if (r_state == STATE_R1) begin
1023             if (w_halt) begin
1024                 $display("");
1025                 $display($time, "\tC%02h: PC: %h HALT", CORE_ID, r_pc);
1026                 r_state <= STATE_HALT;
1027             end else begin
1028                 // primary operand
1029                 r_instr_rdd <= r_reg_rd1;
1030                 r_state <= STATE_R2;
1031             end
1032         end
1033     else if (r_state == STATE_R2) begin
1034         // Choose secondary operand (register or immediate)
1035         if (r_instr_has_imm8) r_instr_rda <= r_instr_imm8;
1036         else if (r_instr_has_imm4) r_instr_rda <= r_reg_rd1 + r_instr_imm4;
1037         else r_instr_rda <= r_reg_rd1;
1038
1039         if (r_instr_has_mem) begin
1040             r_state <= STATE_ME;
1041             // Pulse req
1042             r_mem_scratch_req <= 1;
1043         end else
1044             r_state <= STATE_WB;
1045     end
1046     else if (r_state == STATE_ME) begin
1047         // Pulse req
1048         r_mem_scratch_req <= 0;
1049         // Wait for MMU to finish
1050         if (!r_mem_scratch_busy)
1051             r_state <= STATE_WB;
1052     end
1053     else if (r_state == STATE_WB) begin
1054         if (r_instr_has_cmp) begin
1055             $display($time, "\tC%02h: CMP: %h", CORE_ID, r_alu_out[3:0]);
1056             r_cmp_flags <= r_alu_out[3:0];
1057         end
1058     end
1059

```



```

1060         r_state <= STATE_FE;
1061     end
1062     else if (r_state == STATE_FE)
1063         r_state <= STATE_IF;
1064     else if (r_state == STATE_HALT) begin
1065         `ifndef DEF_ENABLE_INT
1066             if (int_pending) begin
1067                 r_state <= STATE_FE;
1068             end
1069         `endif
1070     end
1071 end
1072
1073 `ifndef DEF_CORE_HAS_INSTR_MEM
1074 // Instruction ROM
1075 (* rom_style = "distributed" *)
1076 vmicro16_bram # (
1077     .MEM_WIDTH      (DATA_WIDTH),
1078     .MEM_DEPTH      (MEM_INSTR_DEPTH),
1079     .CORE_ID        (CORE_ID),
1080     .USE_INITS      (1),
1081     .NAME           ("INSTR_MEM")
1082 ) mem_instr (
1083     .clk            (clk),
1084     .reset          (reset),
1085     // port 1
1086     .mem_addr       (r_pc),
1087     .mem_in         (0),
1088     .mem_we         (1'b0), // ROM
1089     .mem_out        (w_mem_instr_out)
1090 );
1091 `endif
1092
1093 // MMU
1094 vmicro16_core_mmu # (
1095     .MEM_WIDTH      (DATA_WIDTH),
1096     .MEM_DEPTH      (MEM_SCRATCH_DEPTH),
1097     .CORE_ID        (CORE_ID)
1098 ) mmu (
1099     .clk            (clk),
1100     .reset          (reset),
1101     .req            (r_mem_scratch_req),
1102     .busy           (r_mem_scratch_busy),
1103     // interrupts
1104     .ints_vector    (ints_vector),
1105     .ints_mask      (ints_mask),
1106     // port 1
1107     .mmu_addr       (r_mem_scratch_addr),
1108     .mmu_in         (r_mem_scratch_in),
1109     .mmu_we         (r_mem_scratch_we),
1110     .mmu_lwex       (r_instr_has_lwex),
1111     .mmu_swex       (r_instr_has_swex),
1112     .mmu_out        (r_mem_scratch_out),
1113     // APB master to slave
1114     .M_PADDR        (w_PADDR),
1115     .M_PWRITE       (w_PWRITE),
1116     .M_PSELx        (w_PSELx),
1117     .M_PENABLE      (w_PENABLE),
1118     .M_PWDATA       (w_PWDATA),
1119     .M_PRDATA       (w_PRDATA),
1120     .M_PREADY       (w_PREADY)
1121 );
1122
1123 // Instruction decoder
1124 vmicro16_dec dec (
1125     // input
1126     .instr          (r_instr),
1127     // output async
1128     .opcode         (),
1129     .rd             (r_instr_rsd),
1130     .ra             (r_instr_rsa),
1131     .imm4           (r_instr_imm4),
1132     .imm8           (r_instr_imm8),
1133     .imm12          (),
1134     .simm5          (r_instr_simm5),
1135     .alu_op         (r_instr_alu_op),
1136     .has_imm4       (r_instr_has_imm4),
1137     .has_imm8       (r_instr_has_imm8),
1138     .has_we         (r_instr_has_we),
1139     .has_br         (r_instr_has_br),
1140     .has_cmp        (r_instr_has_cmp),
1141     .has_mem        (r_instr_has_mem),
1142     .has_mem_we     (r_instr_has_mem_we),
1143     .halt           (w_halt),
1144     .intr           (w_intr),
1145     .has_lwex       (r_instr_has_lwex),
1146     .has_swex       (r_instr_has_swex)
1147 );
1148

```



```

1149 // Software registers
1150 vmicro16_regs # (
1151     .CORE_ID (CORE_ID),
1152     .CELL_WIDTH (`DATA_WIDTH)
1153 ) regs (
1154     .clk (clk),
1155     .reset (reset),
1156     // async port 0
1157     .rs1 (r_reg_rs1),
1158     .rd1 (r_reg_rd1_s),
1159     // async port 1
1160     // .rs2 (),
1161     // .rd2 (),
1162     // write port
1163     .we (r_reg_we && ~regs_use_int),
1164     .ws1 (r_instr_rsd),
1165     .wd (r_reg_wd)
1166 );
1167
1168 // Interrupt replacement registers
1169 `ifdef DEF_ENABLE_INT
1170 vmicro16_regs # (
1171     .CORE_ID (CORE_ID),
1172     .CELL_WIDTH (`DATA_WIDTH),
1173     .DEBUG_NAME ("REGSINT")
1174 ) regs_intr (
1175     .clk (clk),
1176     .reset (reset),
1177     // async port 0
1178     .rs1 (r_reg_rs1),
1179     .rd1 (r_reg_rd1_i),
1180     // async port 1
1181     // .rs2 (),
1182     // .rd2 (),
1183     // write port
1184     .we (r_reg_we && regs_use_int),
1185     .ws1 (r_instr_rsd),
1186     .wd (r_reg_wd)
1187 );
1188 `endif
1189
1190 // ALU
1191 vmicro16_alu # (
1192     .CORE_ID(CORE_ID)
1193 ) alu (
1194     .op (r_instr_alu_op),
1195     .a (r_instr_rdd),
1196     .b (r_instr_rda),
1197     .flags (r_cmp_flags),
1198     // async output
1199     .c (r_alu_out)
1200 );
1201
1202 branch branch_check (
1203     .flags (r_cmp_flags),
1204     .cond (r_instr_imm8),
1205     .en (w_branch_en)
1206 );
1207
1208 endmodule

```

E.2 Peripheral Code Listing

Various memory-mapped APB peripherals, such as GPIO, UART, timers, and memory.

```

1 // Vmicro16 peripheral modules
2
3 `include "vmicro16_soc_config.v"
4 `include "formal.v"
5
6 // PSEL signal error detection peripheral
7 // No action is taken however.
8 module vmicro16_psel_err_apb (
9     input clk,
10     input reset,
11
12     // APB Slave to master interface
13     input [0:0] S_PADDR, // not used (optimised out)
14     input S_PWRITE,
15     input S_PSELx,
16     input S_PENABLE,

```

```

17     input  [0:0]                S_PWDATA,
18
19     // prdata not used
20     output [0:0]                S_PRDATA,
21     output                                S_PREADY,
22
23     // output an error interrupt signal
24     output err_i
25 );
26     assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
27     assign err_i    = S_PREADY;
28 endmodule
29
30
31 // Simple watchdog peripheral
32 module vmicro16_watchdog_apb # (
33     parameter BUS_WIDTH = 16,
34     parameter NAME      = "WD",
35     parameter CLK_HZ    = 50_000_000
36 ) (
37     input clk,
38     input reset,
39
40     // APB Slave to master interface
41     input [0:0]                S_PADDR, // not used (optimised out)
42     input                                S_PWRITE,
43     input                                S_PSELx,
44     input                                S_PENABLE,
45     input [0:0]                S_PWDATA,
46
47     // prdata not used
48     output [0:0]                S_PRDATA,
49     output                                S_PREADY,
50
51     // watchdog reset, active high
52     output reg                  wdreset
53 );
54 //assign S_PRDATA = (S_PSELx & S_PENABLE) ? gpio : 16'h0000;
55 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
56 wire we        = (S_PSELx & S_PENABLE & S_PWRITE);
57
58 // countdown timer
59 reg [`clog2(CLK_HZ)-1:0] timer = CLK_HZ;
60
61 wire w_wdreset = (timer == 0);
62
63 // infer a register to aid timing
64 initial wdreset = 0;
65 always @(posedge clk)
66     wdreset <= w_wdreset;
67
68 always @(posedge clk)
69     if (we) begin
70         $display($time, "\t\t%s <= RESET", NAME);
71         timer <= CLK_HZ;
72     end else begin
73         timer <= timer - 1;
74     end
75 endmodule
76
77 module timer_apb # (
78     parameter CLK_HZ = 50_000_000
79 ) (
80     input clk,
81     input reset,
82
83     input clk_en,
84
85     // 0 16-bit value R/W
86     // 1 16-bit control R    b0 = start, b1 = reset
87     // 2 16-bit prescaler
88     input [1:0]                S_PADDR,
89
90     input                                S_PWRITE,
91     input                                S_PSELx,
92     input                                S_PENABLE,
93     input [16:0]                S_PWDATA,
94
95     output reg [16:0]            S_PRDATA,
96     output                                S_PREADY,
97
98     output out,
99     output [16:0] int_data
100 );
101 //assign S_PRDATA = (S_PSELx & S_PENABLE) ? swex_success ? 16'hFOFO : 16'h0000;
102 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
103 wire en        = (S_PSELx & S_PENABLE);
104 wire we        = (en & S_PWRITE);
105
106 reg [16:0] r_counter = 0;

```

```

reg [`DATA_WIDTH-1:0] r_load = 0;
reg [`DATA_WIDTH-1:0] r_pres = 0;
reg [`DATA_WIDTH-1:0] r_ctrl = 0;

localparam CTRL_START = 0;
localparam CTRL_RESET = 1;
localparam CTRL_INT    = 2;

localparam ADDR_LOAD = 2'b00;
localparam ADDR_CTRL = 2'b01;
localparam ADDR_PRES = 2'b10;

always @(*) begin
    S_PRDATA = 0;
    if (en)
        case(S_PADDR)
            ADDR_LOAD: S_PRDATA = r_counter;
            ADDR_CTRL: S_PRDATA = r_ctrl;
            //ADDR_CTRL: S_PRDATA = r_pres;
            default:   S_PRDATA = 0;
        endcase
end

// prescaler counts from r_pres to 0, emitting a stb signal
// to enable the r_counter step
reg [`DATA_WIDTH-1:0] r_pres_counter = 0;
wire counter_en = (r_pres_counter == 0);
always @(posedge clk)
    if (r_pres_counter == 0)
        r_pres_counter <= r_pres;
    else
        r_pres_counter <= r_pres_counter - 1;

always @(posedge clk)
    if (we)
        case(S_PADDR)
            // Write to the load register:
            // Set load register
            // Set counter register
            ADDR_LOAD: begin
                r_load      <= S_PWDATA;
                r_counter    <= S_PWDATA;
                $display($time, "\t\ttimr0: WRITE LOAD: %h", S_PWDATA);
            end
            ADDR_CTRL: begin
                r_ctrl      <= S_PWDATA;
                $display($time, "\t\ttimr0: WRITE CTRL: %h", S_PWDATA);
            end
            ADDR_PRES: begin
                r_pres      <= S_PWDATA;
                $display($time, "\t\ttimr0: WRITE PRES: %h", S_PWDATA);
            end
        endcase
    else
        if (r_ctrl[CTRL_START]) begin
            if (r_counter == 0)
                r_counter <= r_load;
            else if(counter_en)
                r_counter <= r_counter - 1;
        end else if (r_ctrl[CTRL_RESET])
            r_counter <= r_load;

        // generate the output pulse when r_counter == 0
        // out = (counter reached zero && counter started)
        assign out     = (r_counter == 0) && r_ctrl[CTRL_START]; // && r_ctrl[CTRL_INT];
        assign int_data = `{DATA_WIDTH{1'b1}};
endmodule

// APB wrapped programmable vmicro16_bram
module vmicro16_bram_prog_apb # (
    parameter BUS_WIDTH    = 16,
    parameter MEM_WIDTH    = 16,
    parameter MEM_DEPTH    = 64,
    parameter APB_PADDR    = 0,
    parameter USE_INITS    = 0,
    parameter NAME         = "BRAMPROG",
    parameter CORE_ID      = 0
) (
    input clk,
    input reset,
    // APB Slave to master interface
    input  [`clog2(MEM_DEPTH)-1:0] S_PADDR,
    input                          S_PWRITE,
    input                          S_PSELx,
    input                          S_PENABLE,
    input [BUS_WIDTH-1:0]          S_PWDATA,

    output [BUS_WIDTH-1:0]          S_PRDATA,
    output                          S_PREADY,

```

```

197 // interface to program the instruction memory
198 input      [`clog2(`DEF_MEM_INSTR_DEPTH)-1:0] addr,
199 input      [`DATA_WIDTH-1:0] data,
200 input      we,
201 input      prog
202 );
203 wire [MEM_WIDTH-1:0] mem_out;
204
205 assign S_PRDATA = (S_PSELx & S_PENABLE) ? mem_out : 16'h0000;
206 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
207 wire s_we = (S_PSELx & S_PENABLE & S_PWRITE);
208
209 wire [`clog2(`DEF_MEM_INSTR_DEPTH)-1:0] mem_addr = we ? addr : S_PADDR;
210 wire [`DATA_WIDTH-1:0] mem_data = we ? data : S_PWDATA;
211 wire mem_we = we | s_we;
212
213 vmicro16_bram # (
214     .MEM_WIDTH (MEM_WIDTH),
215     .MEM_DEPTH (MEM_DEPTH),
216     .NAME ("BRAMPROG"),
217     .USE_INITS (0),
218     .CORE_ID (-1)
219 ) bram_apb (
220     .clk (clk),
221     .reset (reset),
222
223     .mem_addr (mem_addr),
224     .mem_in (mem_data),
225     .mem_we (mem_we),
226     .mem_out (mem_out)
227 );
228 endmodule
229
230 // APB wrapped vmicro16_bram
231 module vmicro16_bram_apb # (
232     parameter BUS_WIDTH = 16,
233     parameter MEM_WIDTH = 16,
234     parameter MEM_DEPTH = 64,
235     parameter APB_PADDR = 0,
236     parameter USE_INITS = 0,
237     parameter NAME = "BRAM",
238     parameter CORE_ID = 0
239 ) (
240     input clk,
241     input reset,
242     // APB Slave to master interface
243     input [`clog2(MEM_DEPTH)-1:0] S_PADDR,
244     input S_PWRITE,
245     input S_PSELx,
246     input S_PENABLE,
247     input [BUS_WIDTH-1:0] S_PWDATA,
248
249     output [BUS_WIDTH-1:0] S_PRDATA,
250     output S_PREADY
251 );
252 wire [MEM_WIDTH-1:0] mem_out;
253
254 assign S_PRDATA = (S_PSELx & S_PENABLE) ? mem_out : 16'h0000;
255 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
256 assign we = (S_PSELx & S_PENABLE & S_PWRITE);
257
258 always @(*)
259     if (S_PSELx && S_PENABLE)
260         $display($time, "\t\t%s => %h", NAME, mem_out);
261
262 always @(posedge clk)
263     if (we)
264         $display($time, "\t\t%s[%h] <= %h", NAME, S_PADDR, S_PWDATA);
265
266 vmicro16_bram # (
267     .MEM_WIDTH (MEM_WIDTH),
268     .MEM_DEPTH (MEM_DEPTH),
269     .NAME (NAME),
270     .USE_INITS (1),
271     .CORE_ID (-1)
272 ) bram_apb (
273     .clk (clk),
274     .reset (reset),
275
276     .mem_addr (S_PADDR),
277     .mem_in (S_PWDATA),
278     .mem_we (we),
279     .mem_out (mem_out)
280 );
281 endmodule
282
283 // Shared memory with hardware monitor (LWEX/SWEX)

```

```

286 module vmicro16_bram_ex_apb # (
287     parameter BUS_WIDTH    = 16,
288     parameter MEM_WIDTH    = 16,
289     parameter MEM_DEPTH    = 64,
290     parameter CORE_ID_BITS = 3,
291     parameter SWEX_SUCCESS = 16'h0000,
292     parameter SWEX_FAIL    = 16'h0001
293 ) (
294     input clk,
295     input reset,
296
297     // |19 |18 |16 |15 |0|
298     // | LWEX | SWEX | 3 bit CORE_ID | S_PADDR |
299     input  [`APB_WIDTH-1:0] S_PADDR,
300
301     input S_PWRITE,
302     input S_PSELx,
303     input S_PENABLE,
304     input [MEM_WIDTH-1:0] S_PWDATA,
305
306     output reg [MEM_WIDTH-1:0] S_PRDATA,
307     output S_PREADY
308 );
309 // exclusive flag checks
310 wire [MEM_WIDTH-1:0] mem_out;
311 reg swex_success = 0;
312
313 localparam ADDR_BITS = `clog2(MEM_DEPTH);
314
315 // hack to create a 1 clock delay to S_PREADY
316 // for bram to be ready
317 reg cdelay = 1;
318 always @(posedge clk)
319     if (S_PSELx)
320         cdelay <= 0;
321     else
322         cdelay <= 1;
323
324 //assign S_PRDATA = (S_PSELx & S_PENABLE) ? swex_success ? 16'hFOFO : 16'h0000;
325 assign S_PREADY = (S_PSELx & S_PENABLE & (!cdelay)) ? 1'b1 : 1'b0;
326 assign we       = (S_PSELx & S_PENABLE & S_PWRITE);
327 wire en         = (S_PSELx & S_PENABLE);
328
329 // Similar to:
330 // http://infocenter.arm.com/help/index.jsp?topic=/com.arm.doc.dui0204f/Cihbghef.html
331
332 // mem_wd is the CORE_ID sent in bits [18:16]
333 localparam TOP_BIT_INDEX = `APB_WIDTH - 1;
334 localparam PADDR_CORE_ID_MSB = TOP_BIT_INDEX - 2;
335 localparam PADDR_CORE_ID_LSB = PADDR_CORE_ID_MSB - (CORE_ID_BITS-1);
336
337 // [LWEX, CORE_ID, mem_addr] from S_PADDR
338 wire lwex = S_PADDR[TOP_BIT_INDEX];
339 wire swex = S_PADDR[TOP_BIT_INDEX-1];
340 wire [CORE_ID_BITS-1:0] core_id = S_PADDR[PADDR_CORE_ID_MSB:PADDR_CORE_ID_LSB];
341 // CORE_ID to write to ex_flags register
342 wire [ADDR_BITS-1:0] mem_addr = S_PADDR[ADDR_BITS-1:0];
343
344 wire [CORE_ID_BITS:0] ex_flags_read;
345 wire is_locked = |ex_flags_read;
346 wire is_locked_self = is_locked && (core_id == (ex_flags_read-1));
347
348 // Check exclusive access flags
349 always @(*) begin
350     swex_success = 0;
351     if (en)
352         // bug!
353         if (!swex && !lwex)
354             swex_success = 1;
355         else if (swex)
356             if (is_locked && !is_locked_self)
357                 // someone else has locked it
358                 swex_success = 0;
359             else if (is_locked && is_locked_self)
360                 swex_success = 1;
361 end
362
363 always @(*)
364     if (swex)
365         if (swex_success)
366             S_PRDATA = SWEX_SUCCESS;
367         else
368             S_PRDATA = SWEX_FAIL;
369     else
370         S_PRDATA = mem_out;
371
372 wire reg_we = en && ((lwex && !is_locked)
373     || (swex && swex_success));
374
375 reg [CORE_ID_BITS:0] reg_wd;

```

```

376     always @(*) begin
377         reg_wd = {{CORE_ID_BITS}{1'b0}};
378
379         if (en)
380             // if wanting to lock the addr
381             if (lwex)
382                 // and not already locked
383                 if (!is_locked) begin
384                     reg_wd = (core_id + 1);
385                 end
386             else if (swex)
387                 if (is_locked && is_locked_self)
388                     reg_wd = {{CORE_ID_BITS}{1'b0}};
389     end
390
391     // Exclusive flag for each memory cell
392     vmicro16_bram # (
393         .MEM_WIDTH  (CORE_ID_BITS + 1),
394         .MEM_DEPTH  (MEM_DEPTH),
395         .USE_INITS  (0),
396         .NAME        ("rexram")
397     ) ram_exflags (
398         .clk         (clk),
399         .reset        (reset),
400
401         .mem_addr     (mem_addr),
402         .mem_in        (reg_wd),
403         .mem_we        (reg_we),
404         .mem_out       (ex_flags_read)
405     );
406
407     always @(*)
408         if (S_PSELx && S_PENABLE)
409             $display($time, "\t\tBRAMex[%h] READ %h\tCORE: %h",
410                 mem_addr, mem_out, S_PADDR[16+: CORE_ID_BITS]);
411
412     always @(posedge clk)
413         if (we)
414             $display($time, "\t\tBRAMex[%h] WRITE %h\tCORE: %h",
415                 mem_addr, S_PWDATA, S_PADDR[16+: CORE_ID_BITS]);
416
417     vmicro16_bram # (
418         .MEM_WIDTH  (MEM_WIDTH),
419         .MEM_DEPTH  (MEM_DEPTH),
420         .USE_INITS  (0),
421         .NAME        ("BRAMexinst")
422     ) bram_apb (
423         .clk         (clk),
424         .reset        (reset),
425
426         .mem_addr     (mem_addr),
427         .mem_in        (S_PWDATA),
428         .mem_we        (we && swex_success),
429         .mem_out       (mem_out)
430     );
431 endmodule
432
433 // Simple APB memory-mapped register set
434 module vmicro16_regs_apb # (
435     parameter BUS_WIDTH      = 16,
436     parameter DATA_WIDTH    = 16,
437     parameter CELL_DEPTH     = 8,
438     parameter PARAM_DEFAULTS_R0 = 0,
439     parameter PARAM_DEFAULTS_R1 = 0
440 ) (
441     input clk,
442     input reset,
443     // APB Slave to master interface
444     input  [`clog2(CELL_DEPTH)-1:0] S_PADDR,
445     input                          S_PWRITE,
446     input                          S_PSELx,
447     input                          S_PENABLE,
448     input  [DATA_WIDTH-1:0] S_PWDATA,
449
450     output [DATA_WIDTH-1:0] S_PRDATA,
451     output                          S_PREADY
452 );
453 wire [DATA_WIDTH-1:0] rd1;
454
455 assign S_PRDATA = (S_PSELx & S_PENABLE) ? rd1 : 16'h0000;
456 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
457 assign reg_we   = (S_PSELx & S_PENABLE & S_PWRITE);
458
459 always @(*)
460     if (reg_we)
461         $display($time, "\t\tREGS-APB[%h] <= %h",
462             S_PADDR, S_PWDATA);
463
464 always @(*)

```

```

465     `rassert(reg_we == (S_PSELx & S_PENABLE & S_PWRITE))
466
467     vmicro16_regs # (
468         .CELL_DEPTH      (CELL_DEPTH),
469         .CELL_WIDTH      (DATA_WIDTH),
470         .PARAM_DEFAULTS_RO (PARAM_DEFAULTS_RO),
471         .PARAM_DEFAULTS_R1 (PARAM_DEFAULTS_R1)
472     ) regs_apb (
473         .clk      (clk),
474         .reset    (reset),
475         // port 1
476         .rs1      (S_PADDR),
477         .rd1      (rd1),
478         .we        (reg_we),
479         .ws1      (S_PADDR),
480         .wd        (S_PWDATA)
481         // port 2 unconnected
482         // .rs2    (),
483         // .rd2    ()
484     );
485 endmodule
486
487 // Simple GPIO write only peripheral
488 module vmicro16_gpio_apb # (
489     parameter BUS_WIDTH = 16,
490     parameter DATA_WIDTH = 16,
491     parameter PORTS = 8,
492     parameter NAME = "GPIO"
493 ) (
494     input clk,
495     input reset,
496     // APB Slave to master interface
497     input [0:0] S_PADDR, // not used (optimised out)
498     input S_PWRITE,
499     input S_PSELx,
500     input S_PENABLE,
501     input [DATA_WIDTH-1:0] S_PWDATA,
502
503     output [DATA_WIDTH-1:0] S_PRDATA,
504     output S_PREADY,
505     output reg [PORTS-1:0] gpio
506 );
507 assign S_PRDATA = (S_PSELx & S_PENABLE) ? gpio : 16'h0000;
508 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
509 assign ports_we = (S_PSELx & S_PENABLE & S_PWRITE);
510
511 always @(posedge clk)
512     if (reset)
513         gpio <= 0;
514     else if (ports_we) begin
515         $display($time, "\t\t%s <= %h", NAME, S_PWDATA[PORTS-1:0]);
516         gpio <= S_PWDATA[PORTS-1:0];
517     end
518 endmodule

```

E.3 Assembly Compiler Listing

The following python3 program is a text assembly to hex compiler for the Vmicro16's instruction. Users can include the compiled hex-stream in their SoC design by using \$readmemh on the file into the instruction memories (mem_instr or instr_rom_apb (shared)).

All assembly programs shown in this report can be compiled by this compiler.

Usage: python3 asm.py <filename.s>

Outputs: asm.s.hex file in the PWD.

```

1  import sys
2  import math
3
4  import argparse
5
6  parser = argparse.ArgumentParser(description='Parse assembly into vmicro16 instruction words.')
7  parser.add_argument('fname', metavar='fname', type=str, help="Filename containing assembly text")
8  args = parser.parse_args()
9  print(args.fname)
10

```

```

11  # Match lines using regex
12  import re
13  r_comment = re.compile("//.*")
14  r_label = re.compile("(\\w+):")
15  r_instr_rr = re.compile("\\s+(\\w+)\\s+r(\\d),\\s+r(\\d)")
16  r_instr_ri = re.compile("\\s+(\\w+)\\s+r(\\d),\\s+#0x([A-Fa-f0-9]+)")
17  r_instr_rif = re.compile("\\s+(\\w+)\\s+r(\\d),\\s+(\\w+)")
18  r_instr_br = re.compile("\\s+(br)\\s+(\\w+),\\s+#0x([A-Fa-f0-9]+)")
19  r_instr_lw = re.compile("\\s+(\\w+)\\s+r(\\d),\\s+r(\\d) \\s+ #0x([A-Fa-f0-9]+)")
20
21  all_instr = []
22  all_labels = []
23
24  num_errors = 0
25
26  class Comment:
27      pass
28
29  class Label:
30      name = ""
31      index = -1
32
33  class Instr:
34      op = "NOP"
35      rs1 = 0
36      rs2 = 0
37      imm8 = 0
38      imm5 = 0
39      index = -1
40      ref = ""
41      linestr = ""
42      label = None
43      def __str__(self):
44          return str(self.__class__) + ": " + str(self.__dict__)
45
46  def parse_line(l):
47      l = l.rstrip()
48      m = r_comment.match(l)
49      if m:
50          return None
51
52      m = r_instr_lw.match(l)
53      if m:
54          r = Instr()
55          r.op = m.group(1)
56          r.rs1 = int(m.group(2))
57          r.rs2 = int(m.group(3))
58          r.imm8 = int(m.group(4), 16)
59          r.linestr = l
60          return r
61
62      m = r_label.match(l)
63      if m:
64          r = Label()
65          r.addra = 0
66          r.name = m.group(1)
67          r.linestr = l
68          return r
69
70      m = r_instr_rr.match(l)
71      if m:
72          r = Instr()
73          r.op = m.group(1)
74          r.rs1 = int(m.group(2))
75          r.rs2 = int(m.group(3))
76          r.linestr = l
77          return r
78
79      m = r_instr_ri.match(l)
80      if m:
81          r = Instr()
82          r.op = m.group(1)
83          r.rs1 = int(m.group(2))
84          r.imm8 = int(m.group(3), 16)
85          r.linestr = l
86          return r
87
88      m = r_instr_br.match(l)
89      if m:
90          r = Instr()
91          r.op = m.group(1)
92          r.ref = m.group(2)
93          r.imm8 = int(m.group(3), 16)
94          r.linestr = l
95          return r
96
97      m = r_instr_rif.match(l)
98      if m:
99          r = Instr()
100         r.op = m.group(1)

```



```

101         r.rs1 = int(m.group(2))
102         r.ref = m.group(3)
103         r.linestr = 1
104         return r
105
106     print("Ignored!: {:s}".format(l))
107
108
109 def calc_offset(ls):
110     lsi = iter(ls)
111
112     index = 0
113     for l in lsi:
114         if isinstance(l, Instr):
115             l.index = index
116             index += 1
117
118     lsi = iter(ls)
119     for l in lsi:
120         if isinstance(l, Label):
121             # set label index = next instr index
122             n = next(lsi)
123             while(not isinstance(n, Instr)):
124                 n = next(lsi)
125             l.index = n.index
126             n.label = l
127
128 def find_str_label(s):
129     for l in all_labels:
130         if l.name == s:
131             return l
132     return None
133
134 def cg_replace_labels(xs):
135     global num_errors
136     # assert all items are of type Instr
137     assert(all(isinstance(x, Instr) for x in xs))
138
139     i = 0
140     x = xs[i]
141
142     while True:
143         if x.ref:
144             # it might be a label
145             label = find_str_label(x.ref)
146             if label:
147                 assert(label.index >= 0)
148                 #x.label = label
149                 x.imm8 = label.index
150             else:
151                 label = cg_str_to_imm(x.ref)
152                 if label != None:
153                     x.imm8 = label
154                 else:
155                     sys.stderr.write("Unknown label '{:s}'".format(x.ref))
156                     num_errors += 1
157         try:
158             i += 1
159             x = xs[i]
160         except:
161             break
162
163 def cg_str_to_imm(str):
164     global num_errors
165     if str == "BR_U":
166         return 0
167     elif str == "BR_E":
168         return 1
169     elif str == "BR_NE":
170         return 2
171     elif str == "BR_G":
172         return 3
173     elif str == "BR_GE":
174         return 4
175     elif str == "BR_L":
176         return 5
177     elif str == "BR_LE":
178         return 6
179     elif str == "BR_S":
180         return 7
181     elif str == "BR_SS":
182         return 8
183     else:
184         sys.stderr.write("cg_str_to_imm for {:s} not implemented!".format(str))
185         num_errors += 1
186         return None
187
188 def cg(xs):
189     global num_errors
190     # assert all items are of type Instr
191     assert(all(isinstance(x, Instr) for x in xs))

```

```

192     binstr = []
193     for x in xs:
194         #print("Cg for {}:{}".format(x.op))
195         op = 0
196         if x.op == "movi":
197             op |= 0b00101 << 11
198             op |= x.rs1 << 8
199             op |= x.imm8 << 0
200             binstr.append(op)
201         elif x.op == "mov":
202             op |= 0b00100 << 11
203             op |= x.rs1 << 8
204             op |= x.rs2 << 5
205             binstr.append(op)
206         elif x.op == "mult":
207             op |= 0b01011 << 11
208             op |= x.rs1 << 8
209             op |= x.rs2 << 5
210             binstr.append(op)
211         elif x.op == "lshft":
212             op |= 0b00011 << 11
213             op |= x.rs1 << 8
214             op |= x.rs2 << 5
215             op |= 0b00100
216             binstr.append(op)
217         elif x.op == "rshft":
218             op |= 0b00011 << 11
219             op |= x.rs1 << 8
220             op |= x.rs2 << 5
221             op |= 0b00101
222             binstr.append(op)
223         elif x.op == "xor":
224             op |= 0b00011 << 11
225             op |= x.rs1 << 8
226             op |= x.rs2 << 5
227             op |= 0b00001
228             binstr.append(op)
229         elif x.op == "nop":
230             op = 0x0000
231             binstr.append(op)
232         elif x.op == "add":
233             op |= 0b00110 << 11
234             op |= x.rs1 << 8
235             op |= x.rs2 << 5
236             op |= 0b11111 << 0;
237             binstr.append(op)
238         elif x.op == "addi":
239             op |= 0b00110 << 11
240             op |= x.rs1 << 8
241             op |= x.rs2 << 5
242             op |= x.imm8 << 0
243             binstr.append(op)
244         elif x.op == "subi":
245             op |= 0b00111 << 11
246             op |= x.rs1 << 8
247             op |= x.rs2 << 5
248             op |= x.imm8 << 0
249             binstr.append(op)
250         elif x.op == "sub":
251             op |= 0b00110 << 11
252             op |= x.rs1 << 8
253             op |= x.rs2 << 5
254             op |= 0b10000 << 0;
255             binstr.append(op)
256         elif x.op == "setc":
257             op |= 0b01010 << 11
258             op |= x.rs1 << 8
259             op |= x.imm8
260             binstr.append(op)
261         elif x.op == "br":
262             op |= 0b01000 << 11
263             op |= x.rs1 << 8
264             op |= x.imm8 << 0
265             binstr.append(op)
266         elif x.op == "cmp":
267             op |= 0b01001 << 11
268             op |= x.rs1 << 8
269             op |= x.rs2 << 5
270             binstr.append(op)
271         elif x.op == "lw":
272             op |= 0b00001 << 11
273             op |= x.rs1 << 8
274             op |= x.rs2 << 5
275             assert(x.imm8 >= -16 and x.imm8 <= 15)
276             op |= x.imm8 << 0
277             binstr.append(op)
278         elif x.op == "sw":
279             op |= 0b00010 << 11
280

```

```

281         op |= x.rs1 << 8
282         op |= x.rs2 << 5
283         assert(x.imm8 >= -16 and x.imm8 <= 15)
284         op |= x.imm8 << 0
285         binstr.append(op)
286     elif x.op == "halt":
287         op = 0x0001
288         binstr.append(op)
289     elif x.op == "intr":
290         op = 0x0002
291         binstr.append(op)
292     elif x.op == "lwex":
293         op |= 0b01101 << 11
294         op |= x.rs1 << 8
295         op |= x.rs2 << 5
296         assert(x.imm8 >= -16 and x.imm8 <= 15)
297         op |= x.imm8 << 0
298         binstr.append(op)
299     elif x.op == "swex":
300         op |= 0b01110 << 11
301         op |= x.rs1 << 8
302         op |= x.rs2 << 5
303         assert(x.imm8 >= -16 and x.imm8 <= 15)
304         op |= x.imm8 << 0
305         binstr.append(op)
306     else:
307         sys.stderr.write("Cg for '{:s}' not implemented!".format(x.op))
308         num_errors += 1
309
310     # check op fits within 16-bits
311     assert((op >= 0x0000) and (op <= 0xFFFF))
312
313     return binstr
314
315 with open(args.fname, "r") as f:
316     # Apply a structure to each line
317     lines = list(map(parse_line, f.readlines()))
318     # Removes empty information
319     lines = list(filter(lambda x: x != None, lines))
320     # Calculates instruction offsets
321     calc_offset(list(lines))
322
323     all_instr = list(filter(lambda x: isinstance(x, Instr), lines))
324     all_labels = list(filter(lambda x: isinstance(x, Label), lines))
325
326     print("\nFound {:d} LABELS".format(len(all_labels)))
327     print("Found {:d} INSTR".format(len(all_instr)))
328
329     print("\nReplacing labels...")
330     cg_replace_labels(all_instr)
331
332     # Write hex words to verilog memh file
333     binstr = cg(all_instr)
334     # Ensure instructions fit within the instruction memory size 64/4096
335     assert(len(binstr) <= 64)
336     assert(len(binstr) <= 4096)
337
338     print("\n{:s} produces:".format(args.fname))
339     print("=====\n")
340     with open("asm.s.hex", "w") as out:
341         for i, b in enumerate(binstr):
342             if all_instr[i].label:
343                 print("{:s}:".format(all_instr[i].label.name))
344                 print("\t{:x}\t{:s}\t\t{:04x}".format(i, all_instr[i].linestr, b), end = '')
345
346                 print("")
347                 out.write("{:04x}\n".format(b))
348     print("\nWritten asm.s.hex file!")
349
350 if num_errors:
351     print("\nERRORS {:d}".format(num_errors))

```

E.4 Text Compiler Listing

A text-based programming language compiler was also used to write high-level software code for the Vmicro16 processor. The PRCO304 [12] compiler was extended to support the Vmicro16 instruction set architecture and some extra language features (arrays, inline assembly, pointer writing, etc.). However, the compiler ended up not being not used in favour of the assembly compiler which could be easily used to write multi-threaded code.

The code changes to extend the compiler are available as a .patch for users to patch the compiler themselves. The patch is available from: <https://github.com/bendl/vmicro16/tree/master/sw/patch>.

Code files are found in `sw/demos/prco/*.prco`.