

Multi-core RISC Processor Design and Implementation

(Rev. 2.02)

ELEC5881M - Final Report

Ben David Lancaster
Student ID: 201280376

Submitted in accordance with the requirements for the degree of
Master of Science (MSc)
in Embedded Systems Engineering

Supervisor: Dr. David Cowell
Assessor: Mr David Moore

University of Leeds
School of Electrical and Electronic Engineering

August 27, 2019

Word count: 4689

Abstract

This interim report details the 4-month progress on a project to design, implement, and verify, a multi-core FPGA RISC processor. The project has been split into two stages: firstly to build a functional single-core RISC processor, and then secondly to add multiprocessor principles and functionality to it.

Current multiprocessor and network-on-chip communication methods have been discussed and how they could be included in this multi-core RISC design. To-date, a 16-bit instruction set architecture has been designed featuring common load/store instructions, comparison, and bitwise operations. A single-core processor has been implemented in Verilog and verified using simulations/test benches running various simple software programs.

Future tasks have been planned and will focus on the second stage of the project. Work will start on designing a loosely coupled multiprocessor communication interface and bringing them to the single-core processor.

Declaration of Academic Integrity

The candidate confirms that the work submitted is his/her own, except where work which has formed part of jointly-authored publications has been included. The contribution of the candidate and the other authors to this work has been explicitly indicated in the report. The candidate confirms that appropriate credit has been given within the report where reference has been made to the work of others.

This copy has been supplied on the understanding that no quotation from the report may be published without proper acknowledgement. The candidate, however, confirms his/her consent to the University of Leeds copying and distributing all or part of this work in any forms and using third parties, who might be outside the University, to monitor breaches of regulations, to verify whether this work contains plagiarised material, and for quality assurance purposes.

The candidate confirms that the details of any mitigating circumstances have been submitted to the Student Support Office at the School of Electronic and Electrical Engineering, at the University of Leeds.

Name: Ben David Lancaster

Date: August 27, 2019

Acknowledgements

I would like to thank my supervisor, David Cowell, and assessor, David Moore, for giving me the opportunity to explore a project of my choosing.

Revision History

Date	Version	Changes
10/04/2019	2.02	Update future stages.
05/04/2019	2.01	Fix processor RTL diagram.
04/04/2019	2.00	Initial processor RTL diagram.
01/04/2019	1.00	Initial section outline.

Document revisions.

Table of Contents

1	Introduction	11
1.1	Why Multi-core?	11
1.2	Why RISC?	12
1.3	Why FPGA?	12
2	Background	13
2.1	Amdahl's Law and Parallelism	13
2.2	Loosely and Tightly Coupled Processors	13
2.3	Network-on-chip Architectures	14
3	Project Overview	16
3.1	Project Deliverables	16
3.1.1	Core Deliverables (CD)	16
3.1.2	Extended Deliverables (ED)	17
3.2	Project Timeline	18
3.2.1	Project Stages	18
3.2.2	Project Stage Detail	18
3.2.3	Timeline	20
3.3	Resources	20
3.3.1	Hardware Resources	21
3.3.2	Software Resources	21
3.4	Legal and Ethical Considerations	22
4	Single-core Design	23
4.1	Introduction	23
4.2	Design and Implementation	23
4.2.1	Instruction Set Architecture	24
4.2.2	Memory Management Unit	25
4.2.3	Instruction and Data Memory	25
4.2.4	ALU Design	25
4.2.5	Decoder Design	27
4.2.6	Pipelining	27
4.2.7	Design Optimisations	28
4.3	Interrupts	28

4.3.1	Overview	29
4.3.2	Hardware Implementation	29
4.3.3	Software Interface	30
4.3.4	Design Improvements	31
4.4	Verification	32
5	Interconnect	33
5.1	Introduction	33
5.1.1	Comparison of On-chip Buses	33
5.2	Overview	34
5.2.1	Design Considerations	35
5.3	Interfaces	36
5.3.1	Master to Slave Interface	36
5.3.2	Multi-master Support	37
5.4	Further Work	38
6	Memory Mapping	40
6.1	Introduction	40
6.2	Address Decoding	40
6.2.1	Decoder Optimisations	41
6.3	Memory Map	43
7	Multi-core Communication	45
7.1	Introduction	45
7.1.1	Design Goals	45
7.1.2	Context Identification	45
7.1.3	Thread Synchronisation	46
8	Analysis & Results	49
8.1	Introduction	49
8.2	Implementation Analysis	49
8.2.1	Design Size	49
8.2.2	Maximum Frequency	50
8.3	Scenario Performance	50
8.3.1	Scenario Overview	50
8.3.2	Performance Measurements	52
8.3.3	Performance Results	52
	References	55
	Appendices	56
A	Peripheral Information	56
A.1	Special Registers	56
A.2	Watchdog Timer	57

A.3	GPIO Interface	57
A.4	Timer with Interrupt	58
A.5	UART Interface	58
B	Additional Figures	59
B.1	Register Set Multiplex	59
B.2	Instruction Set Architecture	60
C	Configuration Options	61
C.1	System-on-chip Configuration Options	61
C.2	Core Options	62
C.3	Peripheral Options	63
D	Viva Demonstration Examples	64
D.1	2-core Timer Interrupt and ISR	64
D.2	1-160 Core Parallel Summation	66
E	Code Listing	68
E.1	SoC Code Listing	68
E.1.1	vmicro16_soc_config.v	68
E.1.2	top_ms.v	70
E.1.3	vmicro16_soc.v	71
E.1.4	vmicro16.v	77
E.2	Peripheral Code Listing	91

List of Figures

2.1	A loosely coupled multiprocessor system. Each node features it's own memory and IO modules and uses a Message Transfer System to perform inter-node communication. Image source: [1].	14
2.2	A tightly coupled multiprocessor system. Nodes are directly connected to memory and IO modules. Image source: [1].	14
2.3	A multiprocessor network-on-chip architecture with 16 processing nodes. Nodes are connected in a grid formation with routers and links. Image source: [2]. . .	15
3.1	Project stages in a Gantt chart.	20
3.2	Terasic DE1-SoC development board featuring the Altera Cyclone V FPGA and many peripherals. Image source: [3].	21
3.3	Minispartan-6+ development board featuring the Xilinx Spartan 6 XC6SLX9. Note that the XC6SLX9 and XC6SLX25 FPGAs share the same board. Image source: [4].	22
4.1	Vmicro16 RISC 5-stage RTL diagram showing: instruction pipelining (data passed forward through clocked register banks at each stage); branch address calculation; ALU operand calculation (rd2 or imm); and program counter incrementing.	24
4.2	Vmicro16 ALU diagram showing clocked inputs from the previous IDEX stage being	26
4.3	Time diagram showing the TIMR0 peripheral emitting a 1us periodic interrupt signal (out) to the processor. The processor acknowledges the interrupt (int_pending_ack) and enters the interrupt mode (regs_use_int) for a period of time. When the interrupt handler reaches the Interrupt Return instruction (indicated by w_intr) the processor returns to normal mode and restores the normal state.	30
4.4	The interrupt vector (0x0100 - 0x0107) consists of eight 16-bit values that point to memory addresses of the instruction memory to jump to.	30
4.5	Interrupt Mask register (0x0108). Each bit corresponds to an interrupt source. 1 signifies the interrupt is enabled for/visible to the core. Bits [7:2] are left to the designer to assign. Bit 0 is assigned to TIMR0's interval timer. Bit 1 is assigned to the UART0's receiver (unassigned if DEF_USE_REPROG is enabled).	31
5.1	Waveform showing an APB read transaction.	34

5.2	Block diagram of the Vmicro16 system-on-chip.	35
5.3	Foo	38
6.1	Schematic showing the address decoder (addr_dec) accepting the active PADDR signal and outputting PSEL chip enable signals to each peripheral.	41
6.2	Example 4-bit binary comparator which compares the bits (a, b, c, d) to the constant value 1010. The 0s of the constant are inverted and then all are passed to a wide-AND.	41
6.3	Bits [7:3] of an 8-bit PADDR signal are used as inputs to 5-bit LUTs to generate a PSEL signal. In addition, a default error case is shown allowing the address decoder to detect incorrect PADDR values (e.g. if no PSEL signals are generated).	42
6.4	Partial address decoding used by the Vmicro16 SoC design. Each peripheral shown only needs to decode a signal bit to determine if it is enabled.	42
6.5	Memory map showing addresses of various memory sections.	44
7.1	Block diagram showing the main multi-processing components: the CPU array and a peripheral interconnect used for core synchronisation.	46
7.2	Vmicro16 Special Registers layout (0x0080 - 0x008F).	46
7.3	Assembly code for locking a mutex. r1 is the address to lock. r3 is zero. r4 is the branch address.	47
8.1	A theoretical FPGA device with 3 BRAM blocks running a 4-core design. Each core can map onto a BRAM block, however as there are more cores than BRAM blocks available, some core memories will be implemented as distributed RAM, or in the worse case using ALMs.	50
8.2	Cyclone V maximum design frequency for various core count configurations.	51
8.3	Chart showing how the communication times (Tbus) and serial times (Tsum) changes with core count.	52
8.4	Similar to Figure 8.3 but using shared instruction memory to reduce block memory requirements per core.	53
A.1	Vmicro16 Special Registers layout (0x0080 - 0x008F).	57
B.1	Normal mode (bottom) and interrupt mode (top) register sets are multiplexed to switch between contexts.	59
B.2	Vmicro16 instruction set architecture.	60

List of Tables

3.1	Project stages throughout the life cycle of the project.	19
C.1	SoC Configuration Options	61
C.2	Core Options	62
C.3	Peripheral Options	63

List of Listings

1	ALU branch detection using flags: zero (Z), overflow (V), and negative (N). . . .	26
2	Vmicro16's ALU implementation named vmicro16_alu. vmicro16.v	27
3	Vmicro16's ALU implementation named vmicro16_alu. vmicro16.v	27
4	Vmicro16's decoder module code showing nested bit switches to determine the intended opcode. vmicro16.v	27
5	RAM and lock memories instantiated by the shared memory peripheral.	47
6	Assembly code for a memory barrier. Threads will wait in the barrier_wait function until all other threads have reached that code point.	48
7	Variable size inputs and outputs to the interconnect.	59

Chapter 1

Introduction

1.1 Why Multi-core?	11
1.2 Why RISC?	12
1.3 Why FPGA?	12

This project will detail the design, implementation, and verification, of a new multi-core RISC processor aimed at FPGA devices. This project was chosen due to my interest in processor design, in which I have only previously designed single-core RISC processors, and wish to extend this knowledge to gain a basic understanding of multi-core communication, design considerations, and the challenges of software and hardware parallelism first hand.

I will use this opportunity to further develop my knowledge of FPGA and processor design by implementing, designing, and verifying, a multi-core RISC processor from scratch, including the design of a communication interface between multiple cores.

1.1 Why Multi-core?

Moore's Law states that the number of transistors in a chip will double every 2 years []. CPU designers would utilize the additional transistors to add more pipeline stages in the processor to reduce the propagation delay [] which would allow for higher clock frequencies.

The size of transistors have been decreasing [] and today can be manufactured in sub-10 nanometer range. However, the extremely small transistor size increases electrical leakage and other negative effects resulting in unreliability and potential damage to the transistor []. The high transistor count produces large amounts of heat and requires increasing power to supply the chip. These trade-offs are currently managed by reducing the input voltage, utilising complex cooling techniques, and reducing clock frequency. These factors limit the performance of the chip significantly. These are contributing factors to Moore's Law *slowing* down. The capacity limit of the current-generation planar transistors is approaching and so in order for performance increases to continue, other approaches such as alternate transistor technologies like Multigate transistors [5], software and hardware optimisations, and multi-processor architectures are employed.

This report will focus on the latter: to produce a small multi-core processor that can utilise software-based parallelism to gain performance benefits, compared to a larger single-core

design.

1.2 Why RISC?

RISC architectures feature simpler and fewer instructions compared to CISC, which emphasises instructions that perform larger tasks. A single CISC instruction might be performed with multiple RISC instructions. Because of the fewer and simpler instructions, RISC machines rely heavily on software optimisations for performance. RISC instruction sets are based on load/store architectures, where most instructions are either register-to-register or memory reading and writing [6]. This constraint greatly reduces complexity.

RISC architectures are easier to design implement, especially for beginners, due to their simpler instructions that share the same pipeline, compared to CISC where there may be different pipeline for each instruction, which would greatly consume FPGA resources.

1.3 Why FPGA?

Field programmable gate arrays (FPGA) are a great choice for prototyping digital logic designs due to their programmable nature and quick development times.

My previous experience with FPGAs in previous projects will reduce risk and learning times and allow for more time to be spent on adding and extending features (discusses further in section 3.1).

FPGAs, however, may not be suitable for prototyping all register-transistor logic (RTL) projects. Larger RTL projects, such as large commercial processors, may greatly exceed the logic cell resources available in today's high-end FPGA devices and may only be prototyped through silicon fabrication, which can be expensive. This resource limitation will not be problem as the project aims to produce a small and minimal design specifically for learning about multi-core architectures.

Chapter 2

Background

2.1 Amdahl's Law and Parallelism	13
2.2 Loosely and Tightly Coupled Processors	13
2.3 Network-on-chip Architectures	14

2.1 Amdahl's Law and Parallelism

In many applications, not restricted to software, there may exist many opportunities for processes or algorithms to be performed in parallel. These algorithms can be split into two parts: a serial part that cannot be parallelised, and a part that can be parallelised. Amdahl's Law defines a formula for calculating the maximum *speedup* of a process with potential parallelism opportunities when ran in parallel with n many processors. Speedup is a term used to describe the potential performance improvements of an algorithm using an enhanced resource (in this case, adding parallel processors) compared to the original algorithm. Amdahl's Law is defined below, where the potential speedup S_p is dependant on the portion of program that can be parallelised p and the number of processing cores n :

$$S_p = \frac{1}{(1 - p) + \frac{p}{n}} \quad (2.1)$$

This formula will be used throughout the project to gauge the performance of the multi-core design running various software algorithms.

2.2 Loosely and Tightly Coupled Processors

Multiprocessor systems can be generalised into two architectures: loosely and tightly coupled, and each architecture has advantages and disadvantages. In loosely coupled systems, each processing node is self-contained – each node has its own dedicated memory and IO modules. Communication between nodes is performed over a *Message Transfer System (MTS)* [1] in a master-slave control architecture.

Scalability in loosely coupled systems is generally easier to implement as each node can simply be appended to the shared MTS interface without large modifications to the rest of

the system. Scalability is an important concern in this project as I wish to test the developed solution with a range of processing nodes.

As loosely coupled system's nodes feature their own memory and IO modules, they generally perform better in cases where interaction between nodes is not prominent – each node can store a separate part of the software program in its memory module allowing simultaneous executing of the program.

In scenarios where inter-node communication is prominent however, access to the MTS interface must be scheduled to avoid access conflicts which introduces delays and idle times in the software programs execution, resulting in lower throughput. Figure 2.1 shows a general layout of a loosely coupled multiprocessor system.

Tightly coupled systems feature processing nodes that do not have their own dedicated memory or IO modules – each node is directly connected to a shared memory module using a dedicated port. In scenarios where inter-node communication is prominent, tightly coupled systems are generally better suited as nodes are directly connected to a shared memory and do not need to wait to use a shared bus.

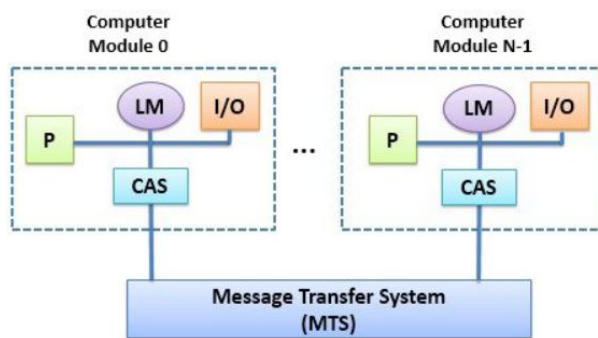


Figure 2.1: A loosely coupled multiprocessor system. Each node features its own memory and IO modules and uses a Message Transfer System to perform inter-node communication. Image source: [1].

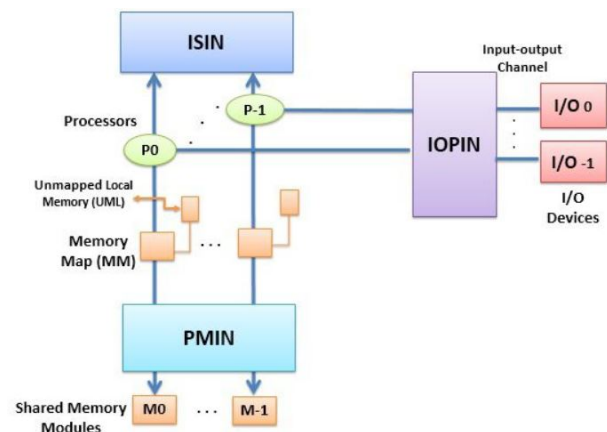


Figure 2.2: A tightly coupled multiprocessor system. Nodes are directly connected to memory and IO modules. Image source: [1].

This project will utilise a loosely coupled architecture due to its easier scalability implementation and my previous experience with the design of single-core processors. Although it will require a scheduler to access the MTS, the experience and knowledge gained from this task will be greatly beneficial for future projects.

2.3 Network-on-chip Architectures

Network-on-chip (NoC) architectures implement on-chip communication mechanisms that are based on network communication principles, such as routing, switching, and massive scalability [7]. NoC's can generally support hundreds to millions of processing cores. Figure 2.3 shows an example 16-core network-on-chip architecture. NoC's can scale to very large sizes while not sacrificing performance because each processor core is able to drive the network rather than needing to wait for a shared bus to become free before doing so.

The greater the number of cores in a network-on-chip design, the greater quality of service

(QoS) problems arise. As such, network-on-chip architectures suffer the same problems as networks, such as fairness and throughput [8].

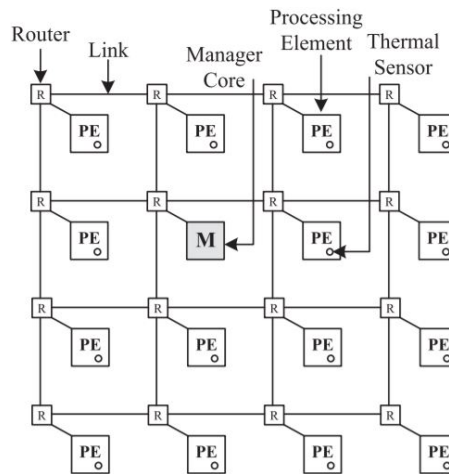


Figure 2.3: A multiprocessor network-on-chip architecture with 16 processing nodes. Nodes are connected in a grid formation with routers and links. Image source: [2].

Chapter 3

Project Overview

3.1	Project Deliverables	16
3.1.1	Core Deliverables (CD)	16
3.1.2	Extended Deliverables (ED)	17
3.2	Project Timeline	18
3.2.1	Project Stages	18
3.2.2	Project Stage Detail	18
3.2.3	Timeline	20
3.3	Resources	20
3.3.1	Hardware Resources	21
3.3.2	Software Resources	21
3.4	Legal and Ethical Considerations	22

This chapter discusses the the project’s requirements, goals, and structure.

3.1 Project Deliverables

The project’s deliverables are split into two sections: core deliverables (CD) – each deliverable must be satisfied for the project to be a minimum viable product (MVP), and extended deliverables (ED) – deliverables that are not required for a MVP – features that only improve upon an existing feature.

3.1.1 Core Deliverables (CD)

The project’s core deliverables are described below.

CD1 Design a compact 16-bit RISC instruction set architecture.

The instruction set will be the primary interface to control the processor from software. An instruction set will be required to implement the custom multi-core communication interface.

It was decided to design a new instruction set rather than to extend an existing architecture as this will increase my knowledge of the constraints to consider when designing instruction sets and processors.

CD2 Design and implement a Verilog RISC core that implements the ISA in CD1.

The Verilog RISC core will be able to run software program written for the instruction set architecture.

CD3 Design and implement an on-chip interconnect for multi-core processing (2 to 32 cores) using the RISC core from CD2.

The interconnect will be a chief requirement to enable multi-core communication. The interconnect should support up to 32 cores, however FPGA implementation constraints may limit this due to limited resources.

The interconnect will control communication between the cores to enable software parallelism.

CD4 Analyse performance of serial and parallel software algorithms, such as parallel DFT, on the processor.

To evaluate the effectiveness of the developed solution, a serial and parallel implementation of a simple computing algorithm (parallel reduction, sorting) will be ran on the processor and it's performance analysed. Effectiveness will be rated on total algorithm run-time and the speed-up gained by adding more cores.

CD5 Allow the RISC core to be easily compiled to multiple FPGA vendors (Xilinx, Altera).

The developed solution should be generic and portable to allow it to be used across a wide-range of FPGA vendors and devices.

Verilog is a generic implementation-independent hardware-description language and so designing implementation specific modules is recommended.

A key consideration for this requirement is to consider the varying hard IP provided by the FPGA vendors (such as BRAM, ethernet, and PCIe [9, 10]). To overcome this problem, the developed Verilog code will conditionally compile where vendor specific requirements are present.

3.1.2 Extended Deliverables (ED)

The project's extended deliverables are described below.

ED1 Design a RISC core with an instructions-per-clock (IPC) rating of at least 1.0 (a single-cycle CPU).

ED2 Design a RISC core with a pipe-lined data path to increase the design's clock speed.

ED3 Design a scalable multi-core interconnect supporting arbitrary (more than 32) RISC core instances (manycore) using Network-on-Chip (NoC) architecture.

ED4 Design a compiler-backend for the PRCO304 [11] compiler to support the ISA from CD1. This will make it easier to build complex multi-core software for the processor.

ED5 The RISC core can communicate to peripherals via a memory-mapped addresses using the Wishbone bus.

ED6 Implement various memory-mapped peripherals such as UART, GPIO, LCD, to aid visual representation of the processor during the demonstration viva.

ED7 Store instruction memory in SPI flash.

ED8 Reprogram instruction memory at runtime from host computer.

ED9 Processor external debugger using host-processor link.

3.2 Project Timeline

3.2.1 Project Stages

The project is split up into many stages to aid planning and management of the project. There are 8 unique stage areas: 1. Initial project conception; 2 Basic RISC core development; 3. Extended RISC core development; 4. Multi-core development; 5. Processor quality-of-life (QoL) improvements; 6. Compiler development; 7. Demo preparation, and 8. Final report.

The project stages are shown in Table [3.1](#).

3.2.2 Project Stage Detail

Stages 1.0 through 1.2 – Research and Project Conception

These stages cover initial research of existing problems and solutions in the multiprocessor area. The instruction set architecture is also proposed that later stages will implement.

Stages 2.1 through 2.3 – Processor module Design, Implementation, and Integration

These stages cover the design, implementation, and integration of key processor core modules such as the instruction decoder, register sets and local memory. Integration of all the modules is a challenging task because some modules have both asynchronous and synchronous signals that need to be timed correctly in order for other modules to receive valid data. An example of this is the register set which has asynchronous read ports that are later clocked in the instruction decode stage.

Stages 3.1 through 3.4 – Advanced Processor Implementation

These stages add advanced features to the processor to provide a more functional product. Although these stages are classified as extended, their technical requirement to design and implement is not great and so are have time allocations in the project schedule. The extended features that these stages introduce are: pipelined processor stages – to drastically increase processor performance; provide a memory-mapped peripheral interface through the MMU; provide a Wishbone master interface to the MMU – allowing external peripherals such as GPIO and LCD displays to be utilised in a modular fashion; and to implement a cache memory for each processor core.

Stage	Title	Start Date	Days	Core	Applicable Deliverables
1.0	Research	Feb 04	7	x	
1.1	Requirement gathering/review	Feb 11	14	x	
1.1	Processor specification, architecture, ISA	Feb 18	100	x	CD1
1.2	Stage/Time Allocation Planning	Feb 25	7	x	
2.1	Decoder, Register Set, impl & integration	Feb 25	14	x	CD2
2.2	Register set impl & integration	Mar 04	14	x	CD2
2.3	Local memory impl & integration	Mar 11	14	x	CD2
3.1	Memory mapped register layout & impl	Apr 01	21		ED5
3.2	Wishbone peripheral bus connected to MMU	Apr 08	21		ED5
3.3	Pipelined implementation and verification	Apr 15	21		ED2
3.4	Cache memory design & impl	Apr 22	28		ED2
4.1	Multi-core communication interface	TBD	TBD	x	CD3
4.2	Shared-memory controller	TBD	TBD	x	CD3
4.3	Scalable multi-core interface (10s of cores)	TBD	TBD	x	CD3
4.4	Multi-core example program (reduction)	TBD	TBD	x	CD4
5.1	SPI-FPGA interface for OTG programming	TBD	TBD		ED7
5.2	FPGA-PC interfacing	TBD	TBD		ED9
5.3	FPGA-PC debugging (instruction breakpoints)	TBD	TBD		ED9
6.1	Compiler backend for vmicro16	TBD	TBD		ED4
6.2	Compiler support for multi-core codegen	TBD	TBD		ED4
7.1	Wishbone peripherals for demo	TBD	TBD	x	CD4
8.1	Final Report	TBD	TBD	x	

Table 3.1: Project stages throughout the life cycle of the project.

Stages 4.1 through 4.4 – Multiprocessor Functionality

These stages are dedicated to adding multiprocessor functionality using a loosely coupled architecture to the processor.

Stages 5.1 through 5.3 – Debugging Features

These stages cover debugging features and are classified as extended due to the large development time required to implement them as well as not being related to multiprocessor systems.

Stages 6.1 through 6.2 – Compiler Backends

These stages cover the implementation of a compiler backend to ease software writing and programming of the processor.

Stage 7.1 – Wishbone Peripherals

Additional Wishbone peripherals, such as SPI and timers will be added to produce a more useful multiprocessor system.

Stage 8.1 – Final Report

This stage is dedicated to the final report write-up. It is expected to be an iterative task that is active throughout the lifespan of the project.

3.2.3 Timeline

The project stages from Table 3.1 are displayed below in a Gantt chart.

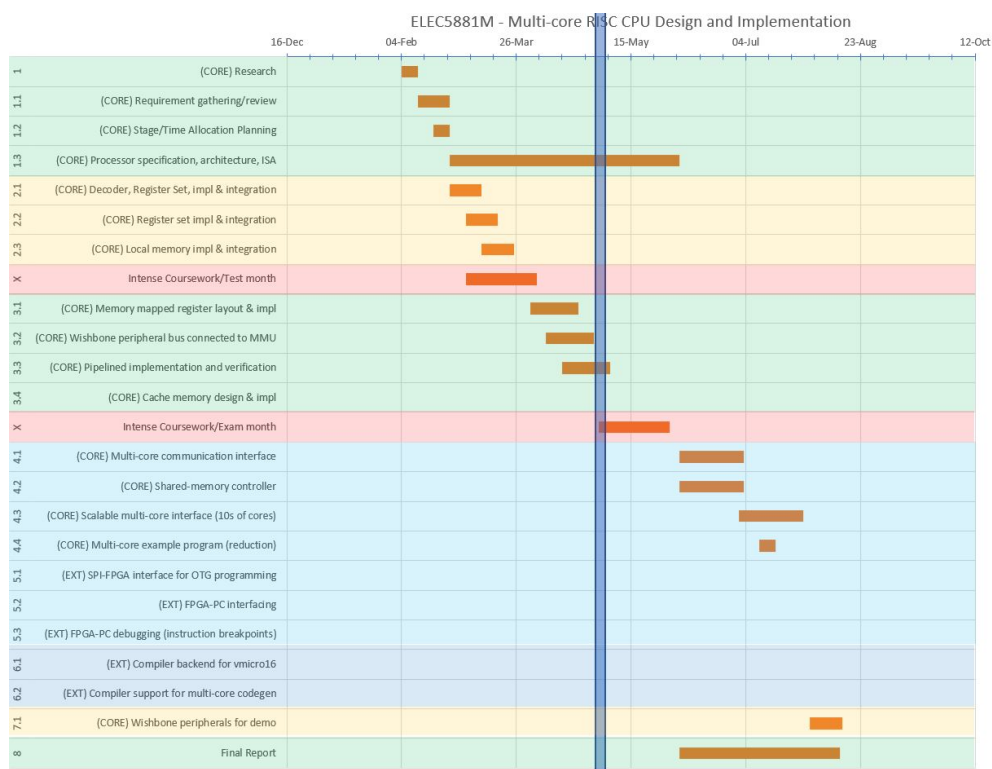


Figure 3.1: Project stages in a Gantt chart.

3.3 Resources

This section describes the hardware and software resources required to fulfil the project.

3.3.1 Hardware Resources

Core deliverable [CD5](#) requires the designed RISC core to be implemented and demonstrated on multiple FPGA devices. Although my design should synthesise for physical IC implementation, due to high costs and lengthy production times, it is not a primary development target. Due to having past experience with Xilinx FPGAs from my placement work and experience with Altera from university modules it was decided to target the Xilinx Spartan 6 XC6SLX9 and the Altera Cyclone V.

Terasic DE1-SoC Development Board

The Terasic DE1-SoC development board features a large Cyclone V FPGA and many peripherals, such as seven-segment displays, 64 MB SDRAM, ADCs, and buttons and switches, which will aid demonstration of the project. The development board is available through the university so the cost is negligible. [Figure 3.2](#) shows the peripherals (green) available to the FPGA.

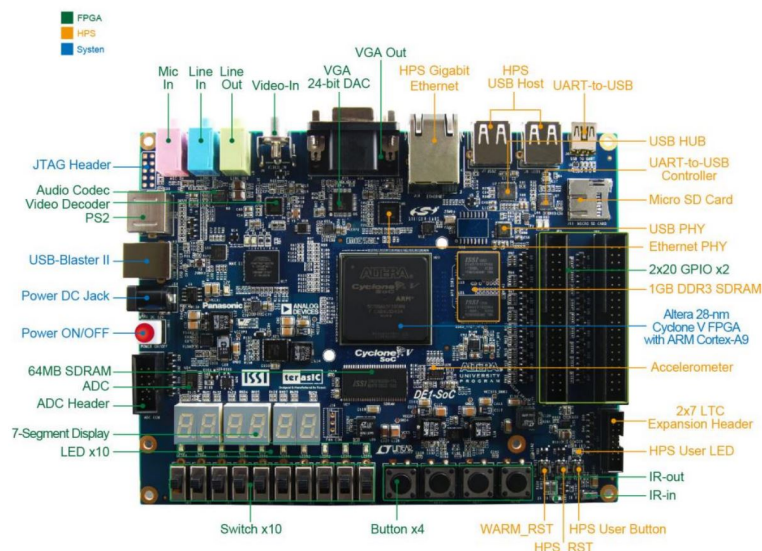


Figure 3.2: Terasic DE1-SoC development board featuring the Altera Cyclone V FPGA and many peripherals. Image source: [\[3\]](#).

Minispartan 6+ FPGA Development Board

The Minispartan 6+ is a hobbyist FPGA development board with fewer peripherals than the DE1-SoC. The board features a Xilinx Spartan 6 XC6SLX9 which has far fewer resources than the DE1-SoC's Cyclone V however it's simplicity and my familiarity with Xilinx's software suite will speed up development. The development board is shown in [Figure 3.3](#).

3.3.2 Software Resources

Intel Quartus

Intel Quartus Prime is a paid-for SoC, CPLD, and FPGA software suite targeting Intel's Stratix, Arria, and Cyclone based FPGAs. The university provides student licences which will be used

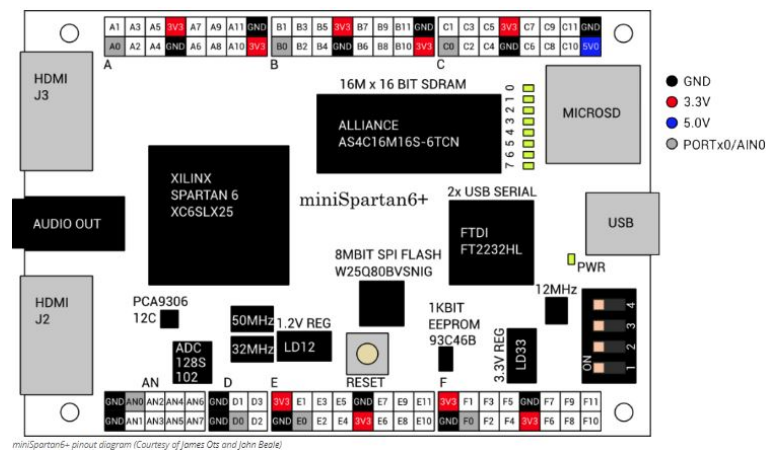


Figure 3.3: Minispartan-6+ development board featuring the Xilinx Spartan 6 XC6SLX9. Note that the XC6SLX9 and XC6SLX25 FPGAs share the same board. Image source: [4].

via VPN.

Xilinx ISE Webpack

Xilinx ISE Webpack is Xilinx's free software suite for FPGA development for Spartan 6 based FPGAs. Due to ISE's intuitive and fast work flow, most of the initial simulation and verification processes will be performed using ISE. This will greatly improve development times.

Verilator

Verilator is an open-source Verilog to C++ transpiler which provides a C++ interface to simulate Verilog modules and read/write values similar to a test bench. Verilator will be used for specific modules within the RISC core such as the ALU and decoder as Verilator is useful when performing exhaustive verification.

3.4 Legal and Ethical Considerations

The RISC core is designed to be used as an academic research and educational tool to aid learning and understanding of RISC and multi-core machines. It should not be use for roles where mission critical or safety is a factor.

The processor does not provide any memory protection features and any software running on the processor has full access to all memory.

The processor does not store/track/predict software instructions. The processor uses pipelining techniques to improve performance which results in future instructions entering the pipeline even if the software's logical sequence does not include these instructions. This could result in security vulnerabilities similar to Intel's Spectre vulnerability [12].

Chapter 4

Single-core Design

4.1	Introduction	23
4.2	Design and Implementation	23
4.2.1	Instruction Set Architecture	24
4.2.2	Memory Management Unit	25
4.2.3	Instruction and Data Memory	25
4.2.4	ALU Design	25
4.2.5	Decoder Design	27
4.2.6	Pipelining	27
4.2.7	Design Optimisations	28
4.3	Interrupts	28
4.3.1	Overview	29
4.3.2	Hardware Implementation	29
4.3.3	Software Interface	30
4.3.4	Design Improvements	31
4.4	Verification	32

4.1 Introduction

While the majority of this report will focus on the multi-processing functionality of this project, it is important to understand the design decisions of the single core to understand the features and limitations of the multi-core system-on-chip as a whole.

4.2 Design and Implementation

The single-core design is a traditional 5-stage RISC processor (fetch, decode, execute, memory, write-back). The core uses separate instruction and data memories in the style of a Harvard architecture [?].

To satisfy **CD5**, the Verilog code will be self-contained in a single file. This reduces the hierarchical complexity and eases cross-vendor project set-up as only a single file is required to be included.

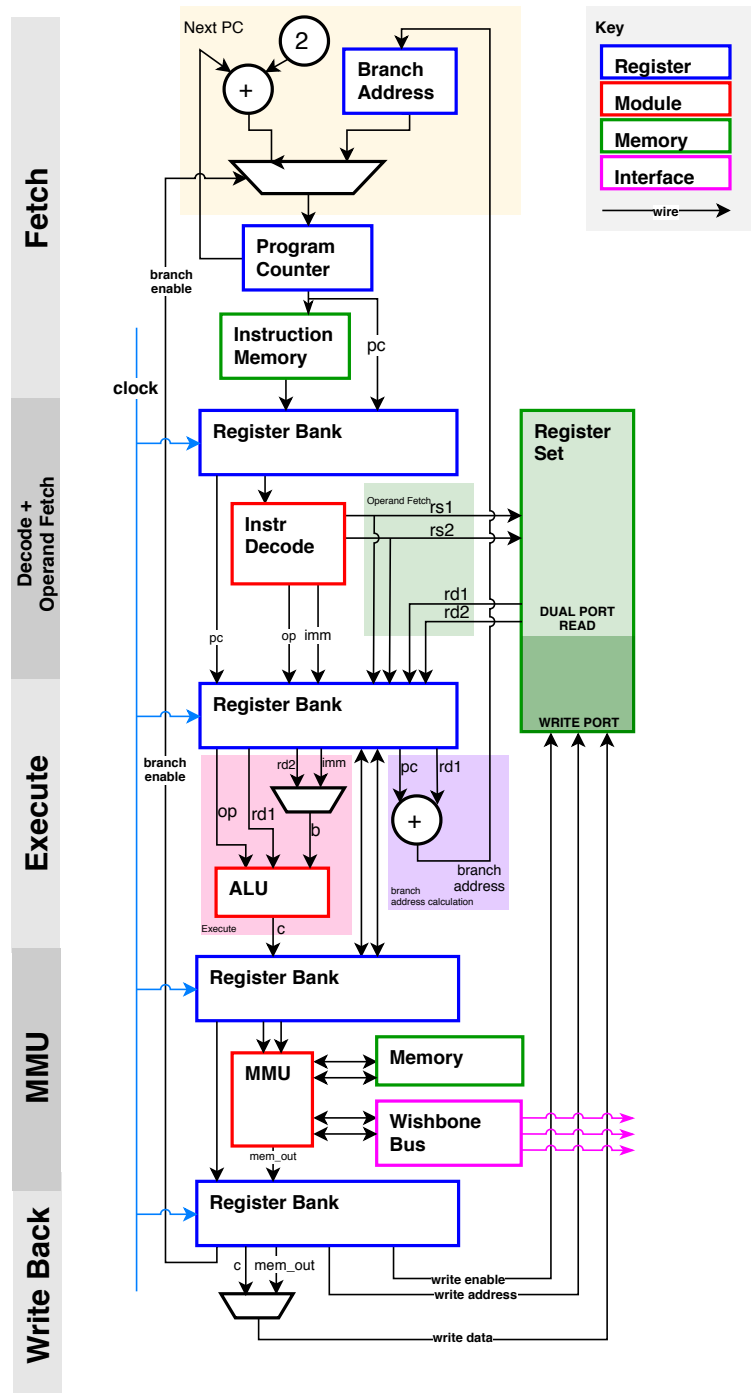


Figure 4.1: Vmicro16 RISC 5-stage RTL diagram showing: instruction pipelining (data passed forward through clocked register banks at each stage); branch address calculation; ALU operand calculation (rd2 or imm); and program counter incrementing.

A small reduction in size within the single-core will result in substantial size reductions in

4.2.1 Instruction Set Architecture

Core deliverable **CD1** details the background for the requirement of a custom instruction set architecture. The 16-bit instruction set listing is shown in [Figure B.2](#).

In this proposed architecture, most instructions are *destructive* meaning that source operands also act as the destination, hence effectively *destroying* the original source data. This design

decision reduces the complexity of the ISA as traditional three operand instructions, for example `add r0, r0, r1`, can be encoded using only two operands `add r0, r1`. However, this does increase the complexity of compilers as they may need to make temporary copies of registers as the instructions will *destroy* the original source data.

The instruction set is split into 7 categories (highlighted by colours in [Figure B.2](#)):

- Special instructions, such as halting and interrupt returns;
- Bitwise operations, such as XOR and AND;
- Signed arithmetic;
- Unsigned arithmetic;
- Conditional branches and compare instructions;
- and Load/store instructions, with their atomic equivalents.

4.2.2 Memory Management Unit

It was decided to use a memory management unit (MMU) to make it easier and extensible to communicate with external peripherals or additional registers. This method transparently uses the existing `LW[EX]` / `SW[EX]` to easily provide an arbitrary number of peripherals/special purpose addresses to the software running on the processor.

4.2.3 Instruction and Data Memory

The design uses separate instruction and data memories similar to a Harvard architecture computer. This architecture was chosen due because it is generally easier to implement, however later resulted in design challenges in large multi-core designs. This is discussed later in the report.

Each single-core has it's own *scratch* memory – a small RAM-like memory which can be used for stack-space and arrays too large to fit into the 8 registers. These memories are provided as is – meaning it's up to the software to implement and provide any stack-frame, function, and calling, functionality. Each core also features it's own read-only instruction memory that is programmed at compile time of the design, or via the UART0 reciever interface (discussed later). Both of these memories map onto synchronous, read-first, single-port, FPGA block RAMs to minimise LUT requirements.

Users can customise the size of these memories by tweaking the following parameters in the `vmicro16_soc_config.v` file: `DEF_MEM_INSTR_DEPTH` for the instruction memory, and `DEF_MEM_SCRATCH_DEPTH` for the scratch memory.

4.2.4 ALU Design

The Vmicro16's ALU is an asynchronous module that has 3 inputs: data a; data b; and opcode `op`; and outputs data c. The ALU is able to operate on both register data (`rd1` and `rd2`) and

immediate values. A switch is used to set the *b* input to either the *rd2* or *imm* value from the previous stage.

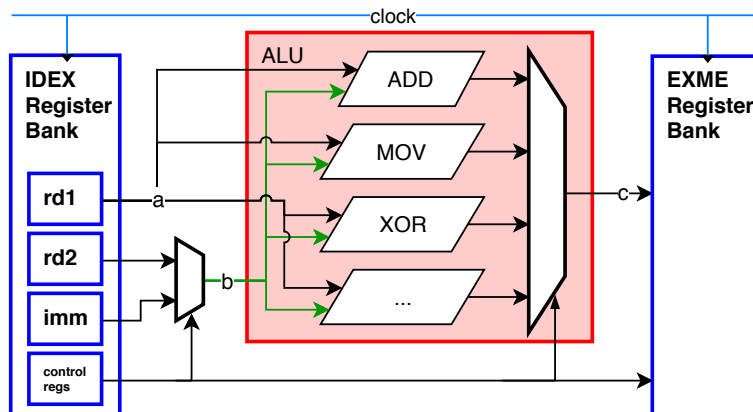


Figure 4.2: Vmicro16 ALU diagram showing clocked inputs from the previous IDEX stage being

The ALU also performs comparison (CMP) operations in which it returns flags similar to X86's overflow, signed, and zero, flags. The combination of these flags can be used to easily compute relationships between the two input operands. For example, if the zero flag is not equal to the signed flag, then the relationship between inputs *a* and *b* is that $a < b$.

```

1  module branch (
2      input [3:0] flags,
3      input [7:0] cond,
4      output reg en
5  );
6      always @(*)
7          case (cond)
8              `VMICRO16_OP_BR_U: en = 1;
9              `VMICRO16_OP_BR_E: en = (flags[`VMICRO16_SFLAG_Z] == 1);
10             `VMICRO16_OP_BR_NE: en = (flags[`VMICRO16_SFLAG_Z] == 0);
11             `VMICRO16_OP_BR_G: en = (flags[`VMICRO16_SFLAG_Z] == 0) &&
12                 (flags[`VMICRO16_SFLAG_N] == flags[`VMICRO16_SFLAG_V]);
13             `VMICRO16_OP_BR_L: en = (flags[`VMICRO16_SFLAG_Z] != flags[`VMICRO16_SFLAG_N]);
14             `VMICRO16_OP_BR_GE: en = (flags[`VMICRO16_SFLAG_Z] == flags[`VMICRO16_SFLAG_N]);
15             `VMICRO16_OP_BR_LE: en = (flags[`VMICRO16_SFLAG_Z] == 1) ||
16                 (flags[`VMICRO16_SFLAG_N] != flags[`VMICRO16_SFLAG_V]);
17             default: en = 0;
18         endcase
19     endmodule

```

Listing 1: ALU branch detection using flags: zero (Z), overflow (V), and negative (N).

The Verilog implementation of the ALU is shown in Listing 2. The ALU's asynchronous output is clocked with other registers, such as destination register *rs1* and other control signals, in the EXME register bank.

```

1  always @(*) case (op)
2      // branch/nop, output nothing
3      `VMICRO16_ALU_BR,
4      `VMICRO16_ALU_NOP:      c = {DATA_WIDTH{1'b0}};
5      // load/store addresses (use value in rd2)
6      `VMICRO16_ALU_LW,
7      `VMICRO16_ALU_SW:      c = b;
8      // bitwise operations
9      `VMICRO16_ALU_BIT_OR:      c = a | b;
10     `VMICRO16_ALU_BIT_XOR:      c = a ^ b;
11     `VMICRO16_ALU_BIT_AND:      c = a & b;
12     `VMICRO16_ALU_BIT_NOT:      c = ~(b);
13     `VMICRO16_ALU_BIT_LSHFT:    c = a << b;
14     `VMICRO16_ALU_BIT_RSHFT:    c = a >> b;

```

Listing 2: Vmicro16's ALU implementation named vmicro16.alu. vmicro16.v

4.2.5 Decoder Design

Instruction decoding occurs in the between the IFID and IDEX stages. The decoder extracts register selects and operands from the input instruction. The decoder outputs are asynchronous which allows the register selects to be passed to the register set and register data to be read asynchronously. The register selects and register read data is then clocked into the IDEX register bank.

```

1  always @(*) case (instr[15:11])
2      `VMICRO16_OP_BR:      alu_op = `VMICRO16_ALU_BR;
3      `VMICRO16_OP_MULT:    alu_op = `VMICRO16_ALU_MULT;
4
5      `VMICRO16_OP_CMP:      alu_op = `VMICRO16_ALU_CMP;
6      `VMICRO16_OP_SETC:    alu_op = `VMICRO16_ALU_SETC;
7
8      `VMICRO16_OP_BIT:      casez (instr[4:0])
9          `VMICRO16_OP_BIT_OR:      alu_op = `VMICRO16_ALU_BIT_OR;
10         `VMICRO16_OP_BIT_XOR:      alu_op = `VMICRO16_ALU_BIT_XOR;
11         `VMICRO16_OP_BIT_AND:      alu_op = `VMICRO16_ALU_BIT_AND;
12         `VMICRO16_OP_BIT_NOT:      alu_op = `VMICRO16_ALU_BIT_NOT;
13         `VMICRO16_OP_BIT_LSHFT:    alu_op = `VMICRO16_ALU_BIT_LSHFT;
14         `VMICRO16_OP_BIT_RSHFT:    alu_op = `VMICRO16_ALU_BIT_RSHFT;
15     default:                alu_op = `VMICRO16_ALU_BAD; endcase

```

Listing 4: Vmicro16's decoder module code showing nested bit switches to determine the intended opcode. vmicro16.v

In Listing 4, it can be seen that the first 4 opcode cases (BR, MULT, CMP, SETC) are represented using the same 15-11 (opcode) bits, however the BIT instructions share the same opcode and so require another bit range to be compared to determine the output function.

4.2.6 Pipelining

In the interim progress update, the processor design featured *instruction pipelining* to meet requirement ED1. Instruction pipelining allows instructions executions to be overlapped in the pipeline, resulting in higher throughput (up to one instruction per clock) at the expense of 5-6 clocks of latency and *significant* code complexity. As the development of the project shifted from single-core to multi-core, it became obvious that the complexity of the pipelined processor would inhibit the integration of multi-core functionality. It was decided to remove the instruction pipelining functionality and use a simpler state-machine based pipeline that is much simpler to extend and would cause fewer challenges later in the project.

4.2.7 Design Optimisations

In a design that has many instantiations of the same component, a small resource saving improvement within the component can have a significant overall savings improvement if it is instantiated many times. Project requirement [CD5](#) requires the design to be compiled for a range of FPGA sizes, and so space saving optimisations are considered.

Register Set Size Improvements

A register set in a CPU is a fast, temporary, and small memory that software instructions directly manipulate to perform computation. In the Vmicro16 instruction set, eight registers named r0 to r7 are available to software. The instruction set allows up to two registers to be references in most instructions, for example the instruction `add r0, r1` tells the processor to perform the following actions:

- Clock 1.** Fetch r0 and r1 from the register set
- Clock 2.** Add the two values together in the ALU
- Clock 3.** Store the result back the register set in r0

For Clock 1, it was originally decided to use a dual port register set (meaning that two data reads can be performed in a single clock, in this case r0 and r1), however due to the asynchronous design of the register set (for speed) the RTL produced consumed a significant amount of FPGA resources, approximately 256 flip-flops ($16 \text{ (data width)} * 8 \text{ (registers)} * 2 \text{ (ports)}$). To reduce this, it was decided to split task 1 into two steps over two clock cycles using a single-port register set. This required the processor pipe-line to use another clock cycle resulting in slightly lower performance, however the size improvements will allow for more cores to be instantiated in the design. This optimisation is also applied to the interrupt register set, resulting in a saving of approximately 256 flip-flops per core (128 in the normal mode register set, and 128 in the interrupt register set). As shown, adding a single clock delay saves a significant amount of LUTs. This saving will be amplified in designs with many cores.

4.3 Interrupts

Interrupts are a technique used by processors to run software functions when an event occurs within the processor, such as exceptions, or signalled from an external source, such as a UART receiver signalling it has received new data. Today, it is common for micro-controllers, soft-processors, and desktop processors, to all feature interrupts. Modern implementations support an *interrupt vector* which is a memory array that contains addresses to different *interrupt handlers* (a software function called when a particular interrupt is received).

Although interrupts are not a requirement for a multi-core system, it was decided to implement this functionality to boost my understanding of such systems. In addition, example demos provided with this project are better visualised with a interrupt functionality.

4.3.1 Overview

The interrupt functionality in this project supports the following:

- Per-core 8 cell interrupt vector accessible to software.
Software programs running on the Vmicro16 processor can edit the interrupt vector to add their own interrupt handlers at runtime.
- Fast context switching.
A dedicated interrupt register set is multiplexed with the normal mode register set to provide faster context switching. It should be noted that only the registers are saved during a context switch. This means that the stack is not saved. A schematic of the register multiplex is shown in [Figure B.1](#).
- Parametrised interrupt sources and widths.
Users can configure the width of the interrupt in signals and the data width per interrupt source via the `vmicro16_soc_config.v`. By default, 8 interrupt sources are available and each can provide 8-bits of data.

4.3.2 Hardware Implementation

Context Switching

When acting upon an incoming interrupt the current state the processor must be saved so that changes from the interrupt handler, such as register writes and branches, do not affect the current state. After the interrupt handler function signals it has finished (by using the *Interrupt Return* INTR instruction) the saved state is restored. In the case of the Vmicro16 processor, the program counter `r_pc[15:0]` and register set `regs` instance are the only states that are saved. Going forth, the terms *normal mode* and *interrupt mode* are used to describe what registers the processor should use when executing instructions.

When saving the state, to avoid clocking 128 bits (8 registers of 16 bits) into another register (which would increase timing delays and logic elements), a dedicated register set for the interrupt mode (`regs_isr`) is multiplexed with the normal mode register set (`regs`). Then depending on the mode (identified by the register `regs_use_int`) the processor can easily switch between the two large states without significantly affecting timing.

The timing diagram in [Figure 4.3](#) shows the behavioural logic for the TIMR0 interrupt source.



Figure 4.3: Time diagram showing the TIMR0 peripheral emitting a 1us periodic interrupt signal (out) to the processor. The processor acknowledges the interrupt (int_pending_ack) and enters the interrupt mode (regs.use_int) for a period of time. When the interrupt handler reaches the Interrupt Return instruction (indicated by w_intr) the processor returns to normal mode and restores the normal state.

4.3.3 Software Interface

A memory-mapped software interface is provided through the MMU to allow easy software control of the interrupt behaviour. The interface is provided at the address range 0x0100 to 0x0108. This interface is per-core allowing each core to individually control what interrupts it receives and what functions to call upon an interrupt. This enables complex functionality, such as allowing each core to execute different functions upon the same interrupt.

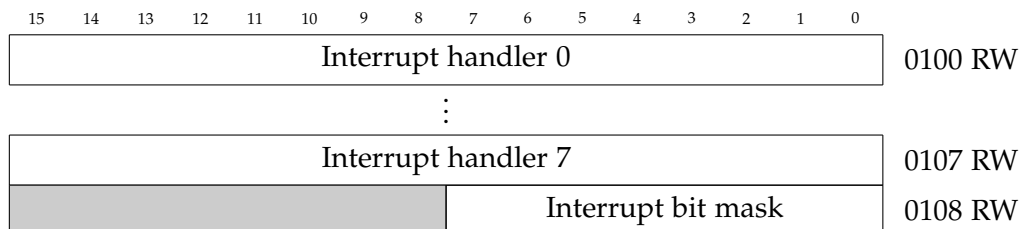


Figure 4.4: The interrupt vector (0x0100 - 0x0107) consists of eight 16-bit values that point to memory addresses of the instruction memory to jump to.

Interrupt Vector (0x0100-0x0107)

The interrupt vector is a per-core register that is used to store the addresses of interrupt handlers. An interrupt handler is simply a software function residing in instruction memory that is branched to when a particular interrupt is received.

Interrupt Mask (0x0108)

The interrupt mask is a per-core register that is used to mask/listen specific interrupt sources. This enables processing cores to individually select which interrupts they respond to. This allows for multi-processor designs where each core can be used for a particular interrupt

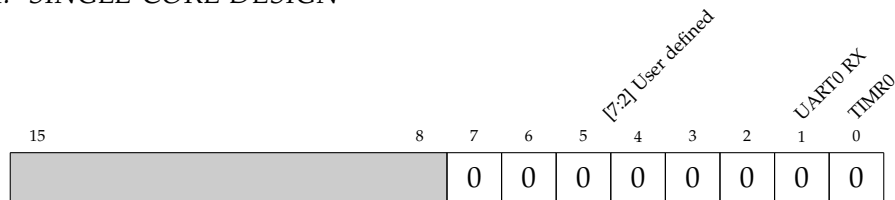


Figure 4.5: Interrupt Mask register (0x0108). Each bit corresponds to an interrupt source. 1 signifies the interrupt is enabled for/visible to the core. Bits [7:2] are left to the designer to assign. Bit 0 is assigned to TIMR0's interval timer. Bit 1 is assigned to the UART0's receiver (unassigned if DEF.USE_REPROG is enabled).

source, improving the time response to the interrupt for time critical programs. The Interrupt Mask register is an 8-bit read/write register where each bit corresponds to a particular interrupt source and each bit corresponds with the interrupt handler in the interrupt vector. The interrupt mask register is shown in Figure 4.5.

Software Example

To better understand the usage of the described interrupt registers, a simple software program is described below. The following software program produces a simple and power efficient routine to initialise the interrupt vector and interrupt mask.

```

1  setup_interruptions:
2      // Set interrupt vector at 0x100
3      // Move address of isr0 function to vector[0]
4      movi    r0, isr0
5      // create 0x100 value by left shifting 1 8 bits
6      movi    r1, #0x1
7      movi    r2, #0x8
8      lshft   r1, r2
9      // write isr0 address to vector[0]
10     sw      r0, r1
11
12  enable_interruptions:
13     // enable all interrupts by writing 0x0f to 0x108
14     movi    r0, #0x0f
15     sw      r0, r1 + #0x8 // (0x100 + 0x8 = 0x108)
16     halt
17     // enter low power idle state
18
19  isr0:
20     // arbitrary name
21     movi    r0, #0xff // do something
22     intr    // return from interrupt

```

A more complex example software program utilising interrupts and the TIMR0 interrupt is described in section D.1.

4.3.4 Design Improvements

The hardware and software interrupt design have changed throughout the projects cycle. In initial versions of the interrupt implementation, the software program, while waiting for an interrupt, would be in a tight infinite loop (branching to the same instruction). This resulted in the processor using all pipeline stages during this time. The pipeline stages produce many logic transitions and memory fetches which raise power consumption and temperatures. This is quite noticeable especially when running on the Spartan-6 LX9 FPGA.

To improve this, it was decided to implement a new state within the processor's state machine that, when entered, did not produce high frequency logic transitions or memory

fetches. The HALT instruction was modified to enter this state and the only way to leave is from an interrupt or top-level reset. This removes the need for a software infinite loop that produces high frequency logic transitions (decoding, ALU, register reads, etc.) and memory fetches.

4.4 Verification

Various verification techniques are employed to ensure correct operation of the processor.

The first technique involves using static assertions to identify incorrect configuration parameters at compile time, such as having zero instruction memory and scratch memory depth. These assertions use the `static_assert` for top level checks and `static_assert_ng` for checks inside generate blocks.

The second verification technique is to use assertions in always blocks to identify incorrect behavioural states. This is done using the `rassert` (run-time assert) macro.

The third verification technique is to use automatic verifying test benches. These test benches drive components of the processor, such as the ALU and decoder, and check the output against the correct value. This uses the `rassert` macro.

The final method of verification is to verify the complete design via a behavioural test bench. The design is passed a compiled software program with a known expected output, and is ran until the `r_halt` signal is raised. The test bench then checks the value on the `debug0`, `debug1`, and `debug2` signals against the expected value. If this matches, then it is assumed that sub-components of the design also operate correctly. This technique does not monitor the states of sub-components and statistics (such as time taken to execute an instruction), there leaves the possibility that some components could have entered an illegal state.

Chapter 5

Interconnect

5.1	Introduction	33
5.1.1	Comparison of On-chip Buses	33
5.2	Overview	34
5.2.1	Design Considerations	35
5.3	Interfaces	36
5.3.1	Master to Slave Interface	36
5.3.2	Multi-master Support	37
5.4	Further Work	38

5.1 Introduction

The Vmicro16 processor needs to communicate with multiple peripheral modules (such as UART, timers, GPIO, and more) to provide useful functionality for the end user.

Previous peripheral interface designs of mine have been directly connected to a main driver with unique inputs and outputs that the peripheral required. For example, a timer peripheral would have dedicated wires for it's load and prescaler values, wires for enabling and resetting, and wires for reading. A memory peripheral would have wires for it's address, read and write data, and a write enable signal. This resulted in each peripheral having a unique interface and unique logic for driving the peripheral, which consumed significant amounts of limited FPGA resources.

It can be seen that many of the peripherals need similar inputs and outputs (for example read and write data signals, write enables, and addresses), and because of this, a standard interface can be used to interface with each peripheral. Using a standard interface can reduce logic requirements as each peripheral can be driven by a single driver.

5.1.1 Comparison of On-chip Buses

The choice of on-chip interconnect has changed multiple times over the life-cycle of this project, primary due to ease of implementation and resource requirements.

Originally, it was planned to use the Wishbone bus [?] due to it's popularity within open-source FPGA modules and good quality documentation.

Late in the project, it was decided to use the AMBA APB protocol [?] as it is more commonly used in large commercial designs and understanding how the interface worked would better benefit myself. APB describes an intuitive and easy to implement 2-state interface aimed at communicating with low-throughput devices, such as UARTs, timers, and watchdogs.

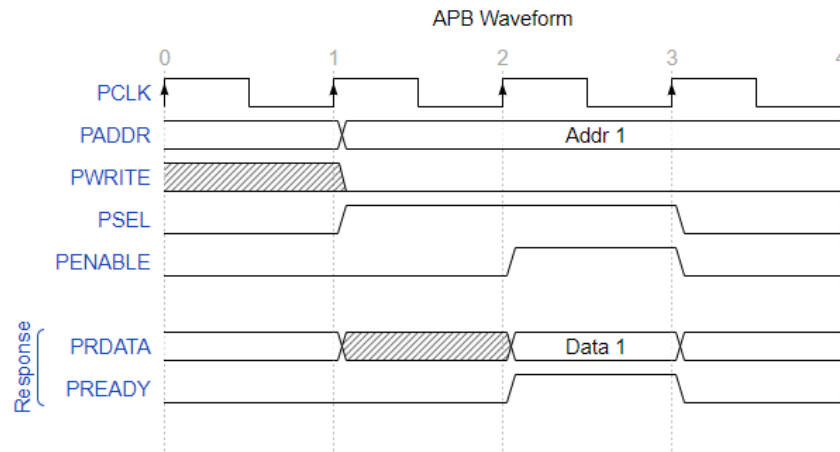


Figure 5.1: Waveform showing an APB read transaction.

5.2 Overview

The system-on-chip design is split into 3 main parts: peripheral interconnect (red), CPU array (gray), and the instruction memory interconnect (green).

A block diagram of this project is shown in [Figure 5.2](#)

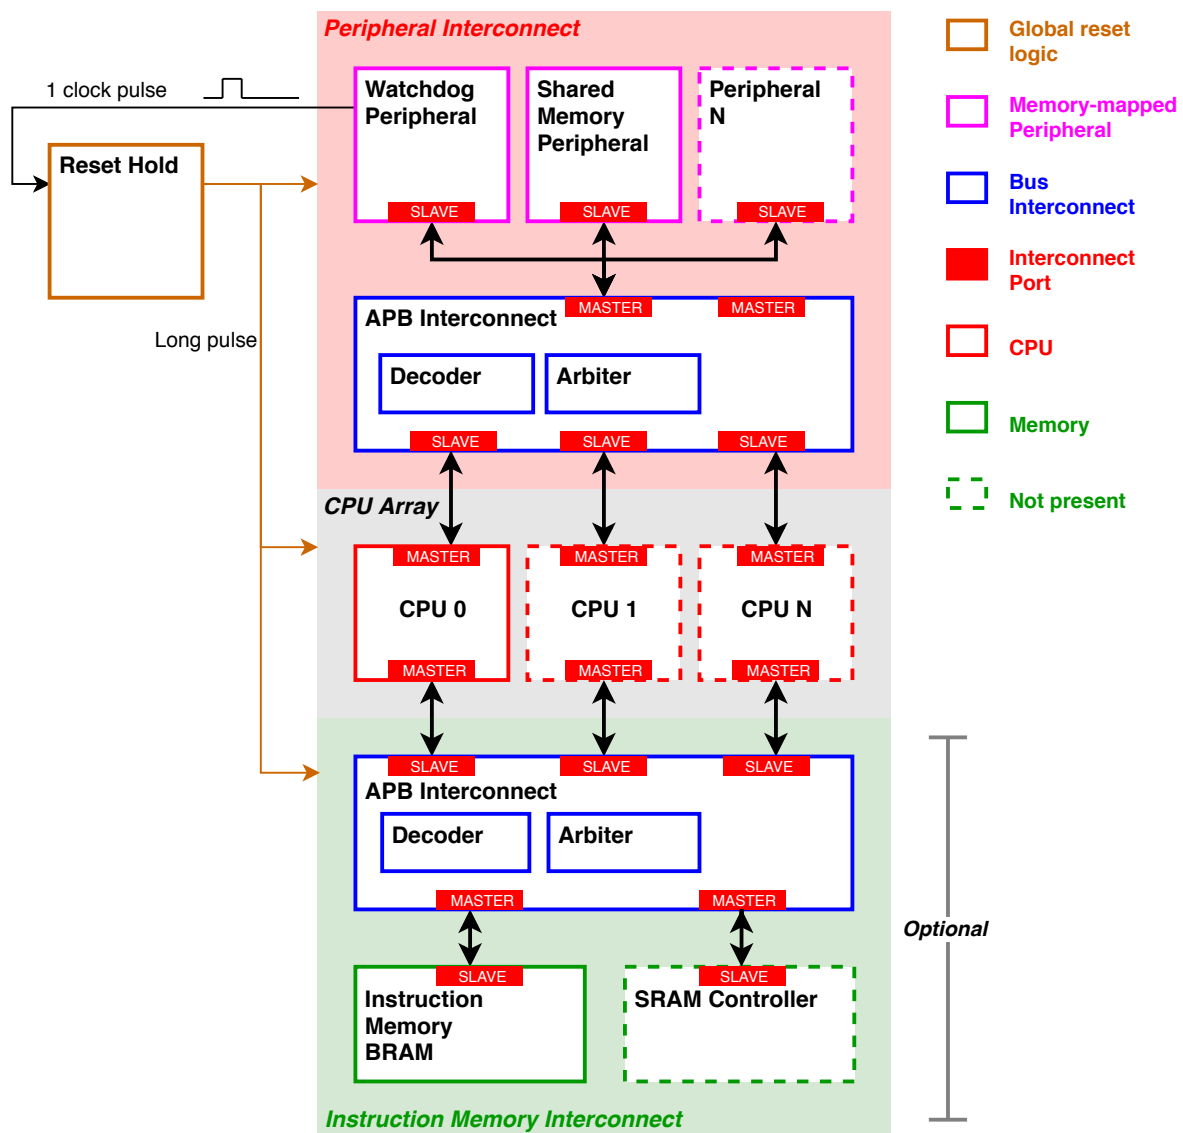


Figure 5.2: Block diagram of the Vmicro16 system-on-chip.

5.2.1 Design Considerations

There are several design issues to consider for this project. These are listed below:

- **Design size limitations**

The target devices for this project are small to medium sized FPGAs (featuring approximately 10,000 to 30,000 logic cells). Because of this, it is important to use a bus interconnect that has a small logic footprint yet is able to scale reasonably well.

- **Ease of implementation**

The interconnect and any peripherals should be easy to implement within the time allocations specified in Figure 3.1.

- **Scalable**

The interconnect should allow for easy scalability of master and slave interfaces with minimal code changes.

5.3.2 Multi-master Support

In this design, each processor can act as an APB master to communicate with peripherals, for example to write a value to UART or to the shared memory peripheral. Because each core runs independently from other cores, it is likely, especially in many-core systems, that two or more processors will want to use the peripheral bus at the same time.

As the peripheral and instruction interconnects use a shared one-to-many (one master to many slaves) bus architecture, only one master can use the bus at any-time. To enable multiple masters to use the bus, a device called an *arbiter* must be used to control which master gets access to drive the shared interconnect.

Arbiters can vary in complexity, mostly relative to throughput requirements.

An ideal arbiter for this interconnect, which ideally features many, possibly tens of, high-throughput masters, would likely feature a priority-based and pipelined arbiter with various devices to improve performance such as cache-coherencies.

Overview

Due to this project's limited time, and my personal knowledge in this area, a simple rotating arbiter is used. This arbitration scheme is likely the simplest that can be thought of. A schematic of arbiter interconnect is shown in [Figure 5.3](#).

In this scheme, access to the bus is given incrementally to each master port, even if the master port has not requested to use the bus. The active master port can use the bus for as long as it requires, and signals it has finished by lowering the PSEL signal. When the PSEL signal is lowered, the arbiter grants access to the next master port. If this next master port has not raised its PSEL signal (i.e. it has not requested access to the bus) then the arbiter grants access to the next master port, and so on. In Verilog, this is simply an incremental counter which is used to index the master ports array. To support a variable number of master ports, the width of each APB signal is multiplied by the number of cores, as shown in [Listing 7](#).

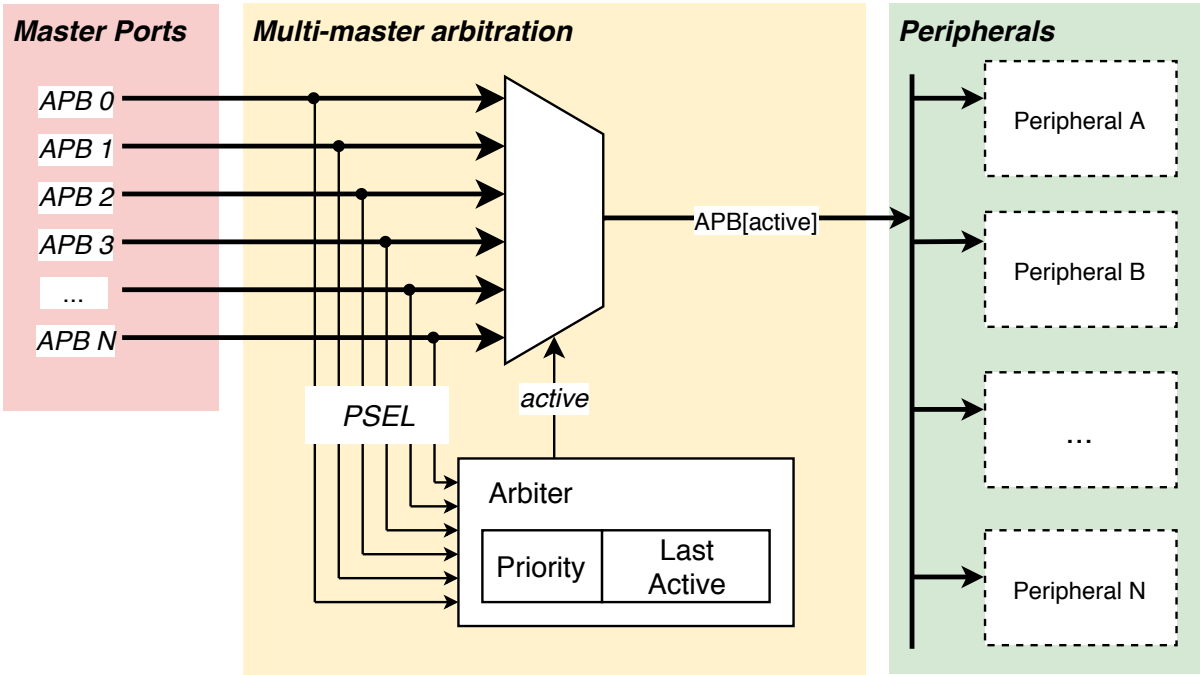


Figure 5.3: Foo

83	62	41	20	0
Core N-1	...	Core 1	Core 0	

5.4 Further Work

The submitted design is acceptable for a multi-core system as it fulfils the following requirements:

- Support an arbitrary number of peripherals.
- Supports memory-mapped address decoding.
- Supports multiple master interfaces.

Arbiter Performance Improvements

However, it fails in the performance aspect. A one clock penalty occurs if the next master port has not requested the bus. This may seem a small price to pay for such a simple arbiter design, however it can add up significantly in many-core designs. For example, if core #0 performs some action on the bus, but core #10 is the next master that wants to use the bus, then the arbiter will waste time incremental granting access to cores #1 to #9 which do not need the bus. This is also made worse when one of the cores is blocking access to a peripheral resource, such as through a mutex or semaphore.

To overcome this penalty, a scheme could use an algorithm to find the next master port requesting access, and grant access directly to it when the current master has finished. Another scheme could be to use a priority encoder. Here, a hard-coded lookup table (LUT) could be used, where the inputs are each master port’s PSEL signal (acting as a bus request line) and

the output being which master to grant access to. As this is targetting FPGA devices, this implemented would require few LUT resources for the arbiter, due to the hard-coded LUT approach. An example of this is given in M. Weber's *Arbiter: Design Ideas and Coding Styles* [13, p. 2].

APB Bus Errors and Recovery

This project's implementation of a multi-master APB interconnect does not provide a method of detecting errors and stalls. This is mainly due to time constraints.

An easy error that could be detected is PADDR addresses that do not fall into a memory-mapped address range. This can easily and cheaply be detected in the address decoding module. This will be discussed in detail in the next chapter.

As previously stated, the active bus master can take control of the bus for as long as it wants to. This is useful for high-throughput transactions, such as memory operations to global memory, but detecting a stalled or glitched operation is not immediately identifiable. If an active master stalls or glitches, it may not be able to lower the PSEL line which appears to the arbiter that the transaction is still happening normally. To overcome this, a timer could be used to detect stalled operations and reset the affected peripheral (essential a watchdog but for an interconnect).

Chapter 6

Memory Mapping

6.1	Introduction	40
6.2	Address Decoding	40
6.2.1	Decoder Optimisations	41
6.3	Memory Map	43

The Vmicro16 processor uses a memory-mapping scheme to communicate with peripherals and other cores. This chapter describes the design decisions and implementation of the memory-map used in this project.

6.1 Introduction

Memory mapping is a common technique used by CPUs, micro-controllers, and other system-on-chip devices, that enables peripherals and other devices to be accessed via a memory address on a common bus. In a processor use-case, this allows for the reuse of existing instructions (commonly memory load/store instructions) to communicate with external peripherals with little additional logic.

6.2 Address Decoding

An address decoder is used to determine the peripheral that the address is requesting. The address decoder module, `addr_dec` in `apb_intercon.v`, takes the 16-bit `PADDR` from the active APB interface and checks for set bits to determine which peripheral to select. The decoder outputs a chip enable signal `PSEL` for the selected peripheral. For example, if bit 12 is set in `PADDR` then the shared memory peripheral's `PSEL` is set high and others to low. A schematic for the decoder is shown in [Figure 6.1](#).

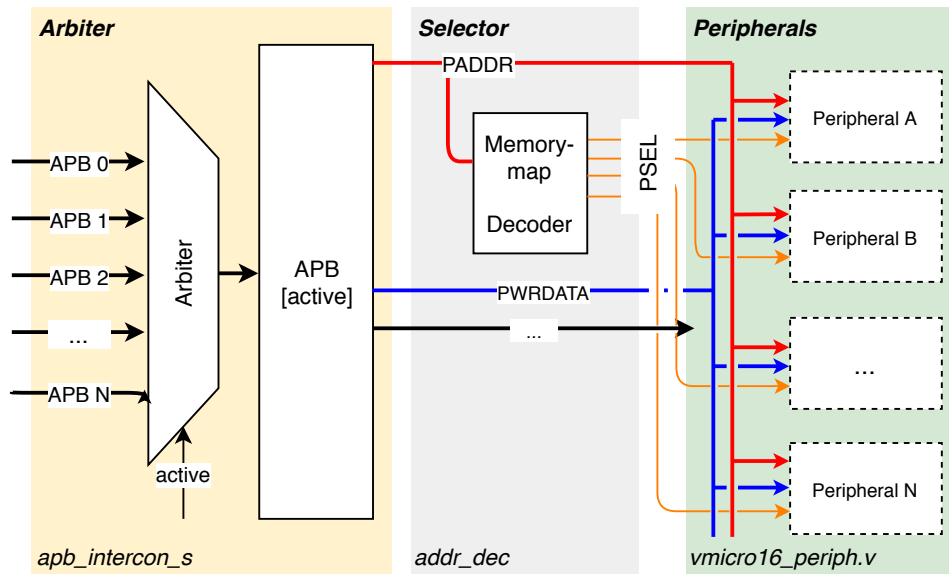


Figure 6.1: Schematic showing the address decoder (*addr_dec*) accepting the active PADDR signal and outputting PSEL chip enable signals to each peripheral.

6.2.1 Decoder Optimisations

Performing a 16-bit equality comparison of the PADDR signal against each peripheral memory address consumes a significant amount of logic. Depending on the synthesis tools and FPGA features, a 16-bit comparator might require a fixed 16-bit value input to compare against (where the 0s are inverted) and a wide-AND to reduce and compare [14, 15]. An example 4-bit comparator is shown below in Figure 6.2.

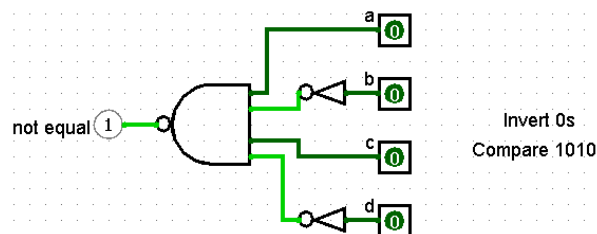


Figure 6.2: Example 4-bit binary comparator which compares the bits (*a*, *b*, *c*, *d*) to the constant value 1010. The 0s of the constant are inverted and then all are passed to a wide-AND.

As we are targeting FPGAs, which use LUTs to implement combinatorial logic, we can conveniently utilise Verilog's `==` operator on fairly large operands without worrying about consuming too many resources. The targeted FPGA devices in this project, the Cyclone V and Spartan 6, feature 6-input LUTs which allow 64 different configurations [16, 17]. Knowing this, we can design the address decoder to utilise the FPGA's LUTs more effectively and reduce its footprint significantly.

We can use part of the PADDR signal as a chip select and the other bits as sub-addresses to interface with the peripheral. The addressing bits are passed into the FPGA's 6-input LUTs which are programmed (via the bitstream) to output 1 or 0 depending on the address. Figure 6.3 below shows a LUT based approach to address decoding which will utilise approximately one ALM/CLB module per peripheral chip select (PSEL) and one for error detection. This method

of comparison (LUT based) is utilised in the `addr_dec` module in `apb_intercon.v`.

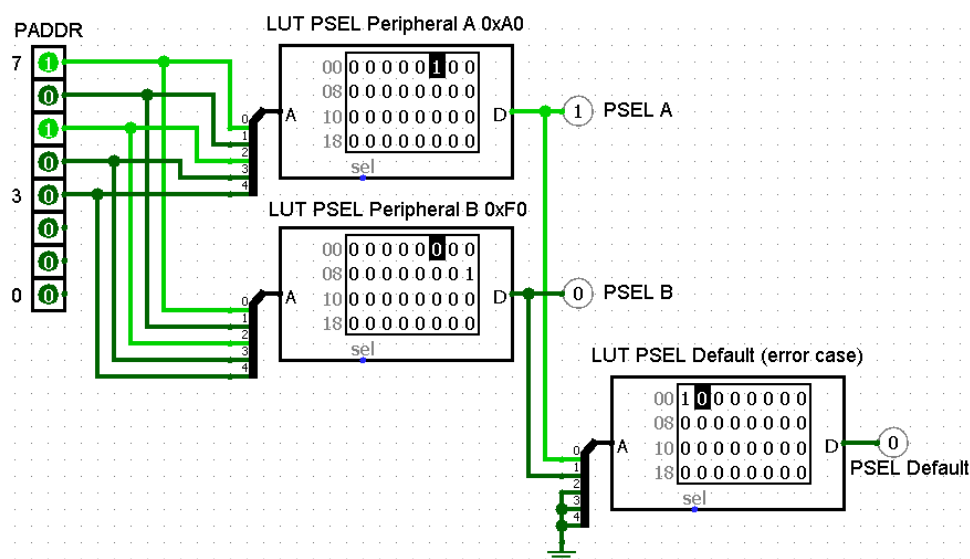


Figure 6.3: Bits [7:3] of an 8-bit `PADDR` signal are used as inputs to 5-bit LUTs to generate a `PSEL` signal. In addition, a default error case is shown allowing the address decoder to detect incorrect `PADDR` values (e.g. if no `PSEL` signals are generated).

The address decoding methods discussed above are examples of *full-address* decoding, where each bit (whether required or not) is compared. It is possible to further reduce the required logic by utilising *partial-address* decoding [18]. Partial-address decoding can reduce logic requirements by not using all bits. For example, if bits in address `0x0100` do not conflict with bits in other addresses (i.e. bit 8 is high in more than 1 address), then the address decoder needs only concern bit 8, not the other bits. This is visualised in Figure 6.4 below. This method is utilised in the MMU's address decoder (module `vmicro16_mmu` in `vmicro16.v`:181). As this is an optimisation per core, significant resources can be saved when a large number of cores are used.

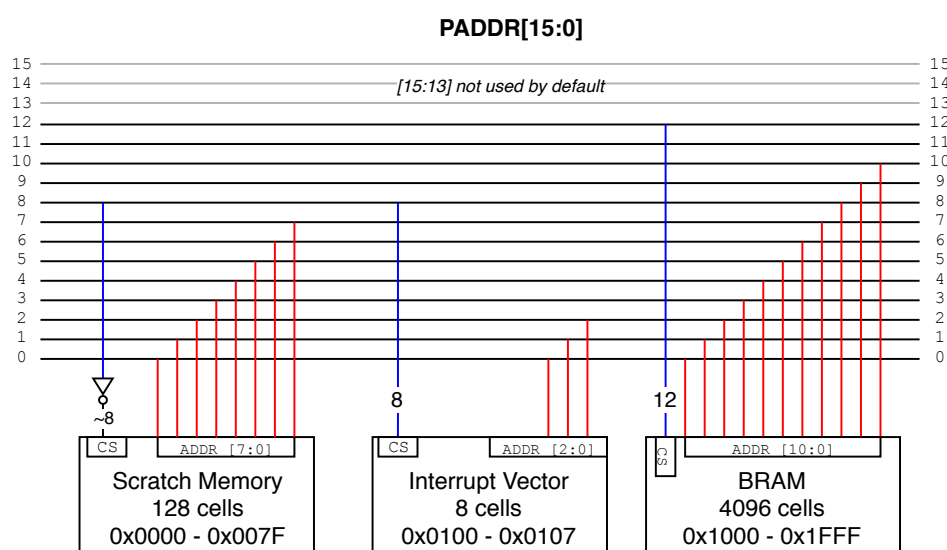


Figure 6.4: Partial address decoding used by the Vmicro16 SoC design. Each peripheral shown only needs to decode a signal bit to determine if it is enabled.

6.3 Memory Map

The system-on-chip's memory map is shown below in [Figure 6.5](#). The addresses for each peripheral have been carefully chosen for both:

- Easy software access – creating addresses via software requires few instructions (normally one to four MOVI and LSHIFT instructions to address 0x0000 to 0xffff), which increases software performance.
- and Reducing address decoding logic – most addresses can be decoded using partial decoding techniques.

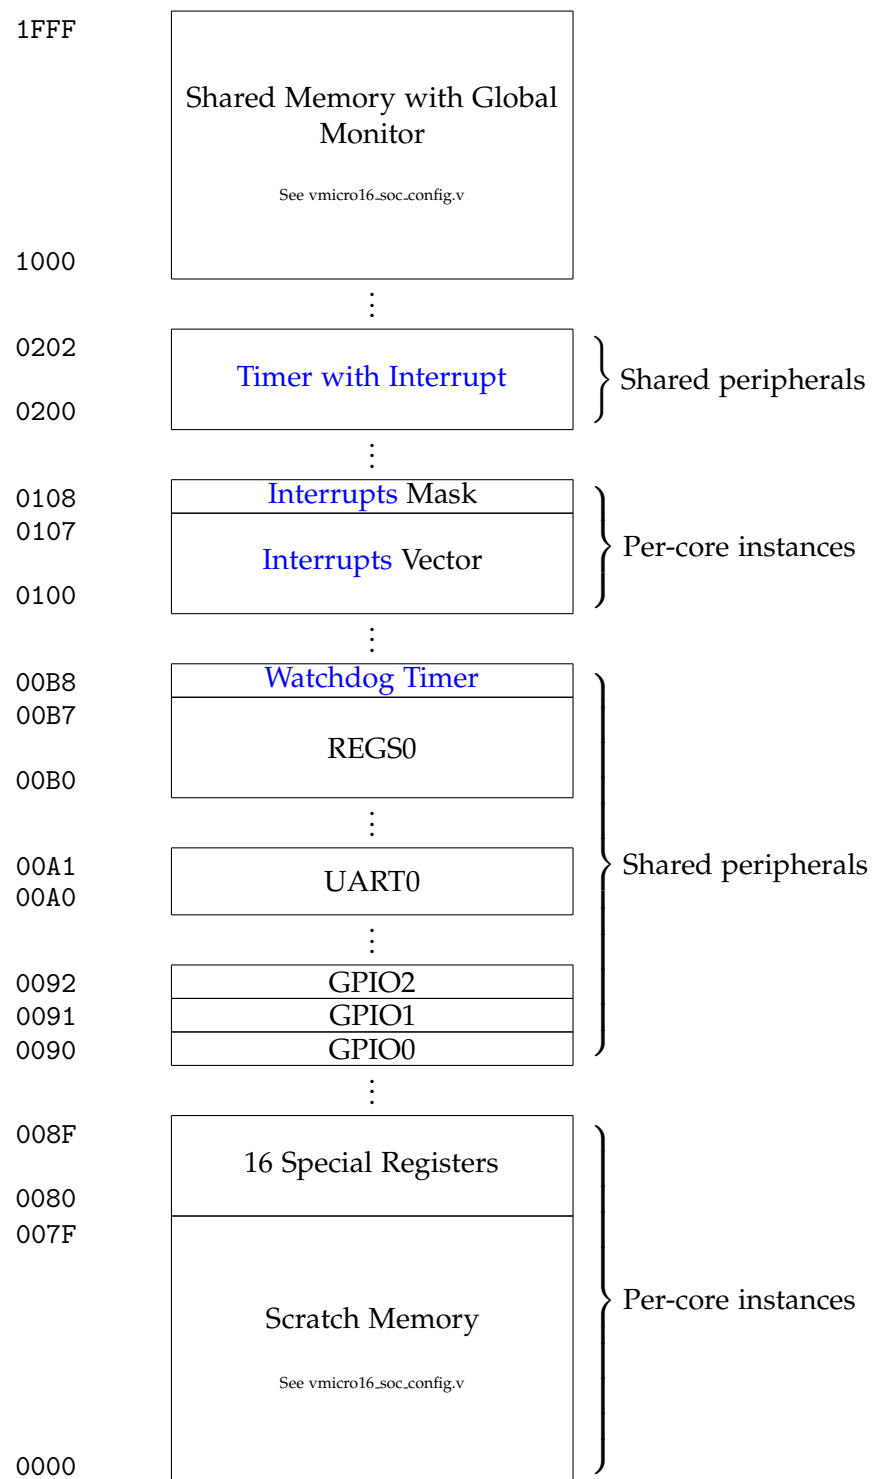


Figure 6.5: Memory map showing addresses of various memory sections.

Chapter 7

Multi-core Communication

7.1	Introduction	45
7.1.1	Design Goals	45
7.1.2	Context Identification	45
7.1.3	Thread Synchronisation	46

So far we have discussed the features and design of the Vmicro16 system-on-chip. This section will discuss the multi-processing functionality and how to use it.

7.1 Introduction

Multi-processing functionality is the primary deliverable of this project.

7.1.1 Design Goals

- **Support common synchronisation primitives.**

Software should be able to implement common synchronisation primitives, such as mutexes, semaphores, and memory barriers, to perform atomic operations and avoid race conditions, which are critical in parallel and concurrent software applications.

- **Context identification.**

The SoC should expose configuration information such as: the number of processing cores, amount of shared and scratch memory, and the `CORE_ID`, to each thread.

7.1.2 Context Identification

A goal of the multi-processing functionality of this project is allow software written for it to be run on any number of cores. This means that a software program will scale to use all cores in the SoC without needing to rewrite the software. To enable this functionality, the software must be able to read contextual information about the SoC, such as the number of cores, how much global and scratch memory is available, and what the `CORE_ID` of the current core is.

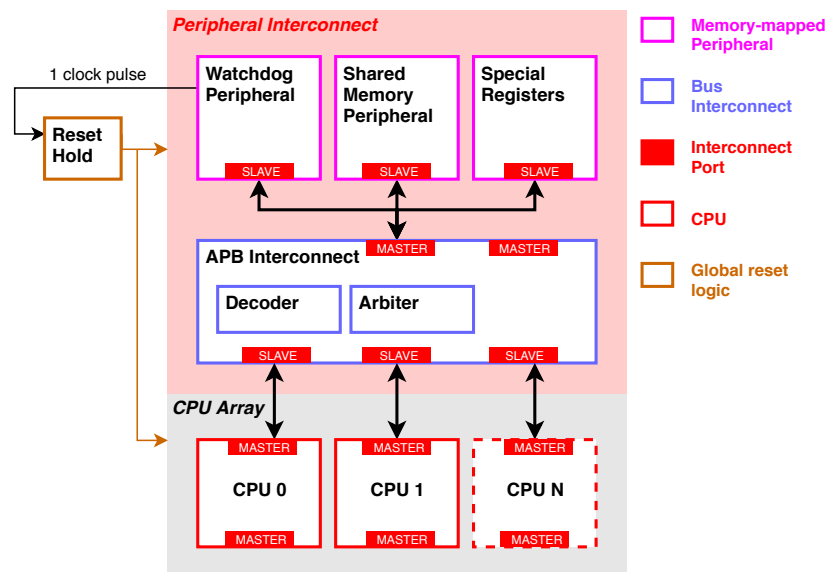


Figure 7.1: Block diagram showing the main multi-processing components: the CPU array and a peripheral interconnect used for core synchronisation.

This information is provided through the Special Registers peripheral (0x0080 - 0x008F), shown in Figure 7.1. This register set provides relevant information for writing software that can dynamically scale for various SoC configurations.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
								CORE_ID								0080	R
								NUM_CORES								0081	R
SHARED_MEMORY cells (default 4096)																0082	R
								NUM_PERIPHERALS								0083	R
SCRATCH_MEMORY cells (default 64)																0084	RW
User defined																0085	RW
⋮																	
User defined																008F	RW

Figure 7.2: Vmicro16 Special Registers layout (0x0080 - 0x008F).

7.1.3 Thread Synchronisation

In multi-threaded software it is important

The mutex functionality is implemented using a similar scheme to that of ARM's *Global Monitor* [?].

Mutexes

In software, a mutex is an object used to control access to a shared resource. The term *object* is used as it's implementation is normally platform dependant, meaning that the processor may provide a hardware mechanism or is left for the operating system to provide.

In this project, mutexes are provided by the processor through the Shared Memory Peripheral (0x1000 to 0x1FFF) which provides a large RAM-style memory accessible by all cores through the peripheral interconnect bus. This large memory is explicitly defined to use the FPGA's BRAM blocks using Xilinx's Verilog `ram_style="block"` attribute to avoid wasting LUTs when using high core counts. The peripheral allows each memory cell to be *locked*, meaning that only the cell owner can modify it's contents. This is implemented by using another large memory, `locks`, to store the `CORE_ID + 1` of the owner, as shown in Listing 5. In this system, a lock containing the value 0 indicates an unlocked cell. As `CORE_ID`s are indexed from zero, 1 is arithmetically added to each cell. For example, if core #2 wants to lock a memory cell, the value 3 is written to the lock.

```

1  reg [15:0] ram [0:8191]; // 16KB large RAM memory
2  reg [clog2(CORES):0] locks [0:8181]; // memory cell owner

```

Listing 5: RAM and lock memories instantiated by the shared memory peripheral.

To lock and unlock cells, the instructions `LWEX` and `SWEX` instructions are used. These instructions are similar to the `LW`/`SW` instructions but provide locking functionality. The `EX` in the instruction names indicate *exclusive access*. `LWEX` is used to read memory contents (like `LW`) and also lock the cell if not already locked. If a core attempts to lock an already locked cell, the lock does not change. Unlocking is done by the `SWEX` instruction, which conditionally writes to the memory cell if it is locked by the same core. Unlike `SW`, `SWEX` returns a zero for success and one for failure if it is locked by another core.

Figure 7.3 shows a simple assembly function to lock a memory cell.

```

1  lock_mutex:
2      // attempt lock
3      lwex r0, r1
4      // check success
5      swex r0, r1
6      cmp r0, r3
7      // if not equal (NE), retry
8      movi r4, lock_mutex
9      br r4, BR_NE
10 critical:
11     // core has the mutex

```

Figure 7.3: Assembly code for locking a mutex. `r1` is the address to lock. `r3` is zero. `r4` is the branch address.

Barriers

Barriers are a useful software sequence used to block execution until all other threads (or a subset) have reached the same point. Barriers are often used for broadcast and gather actions (sending values to each core or receiving them). They are also used to synchronise program execution if some threads have more work to do than others.

The Vmicro16 processor provides barrier synchronisation through the Shared Memory Peripheral. Like the mutex code, the barrier code uses the `LWEX` and `SWEX` instructions to lock a memory cell. Instead of immediately checking the lock as an abstract object, the barrier code treats the cell as a normal memory cell containing a numeric value. Listing 6 shows a software example of this. When the `barrier_reached` code is reached, the code will increment the shared memory value by 1, indicating that the number of threads that have reached this

```
1  barrier_reached:
2      // load latest count
3      lwex    r0, r5
4      // try increment count
5      // increment by 1
6      addi    r0, r3 + #0x01
7      // attempt store
8      swex    r0, r5
9
10     // check success (== 0)
11     cmp     r0, r3
12     // branch if failed
13     movi    r4, barrier_reached
14     br      r4, BR_NE
15
16 barrier_wait:
17     // load the count
18     lw      r0, r5
19     // compare with number of threads
20     cmp     r0, r7
21     // jump back to barrier if not equal
22     movi    r4, barrier_wait
23     br      r4, BR_NE
```

Listing 6: Assembly code for a memory barrier. Threads will wait in the `barrier_wait` function until all other threads have reached that code point.

point has increased by one (r5). The `barrier_wait` function is then entered which waits until this numeric value (r5) is equal to the number of threads (r7) in the system. If this is true, then all threads have reached the `barrier_wait` function and can continue with normal program execution.

Chapter 8

Analysis & Results

8.1	Introduction	49
8.2	Implementation Analysis	49
8.2.1	Design Size	49
8.2.2	Maximum Frequency	50
8.3	Scenario Performance	50
8.3.1	Scenario Overview	50
8.3.2	Performance Measurements	52
8.3.3	Performance Results	52

So far the system's design, implementation, and example usage, has been presented and discussed.

8.1 Introduction

This chapter presents analytic information

8.2 Implementation Analysis

This section analysis the synthesised and implemented system-on-chip design to see the effect of increasing core counts.

8.2.1 Design Size

Implementation

On a minimal system-on-chip configuration, with one core and minimal peripherals and features (no reprogramming, no interrupts, no UART), the design requires as few as 500 LUTs with the processor core requiring approximately 300-400 LUTs.

Constraints

As discussed in [Chapter 4 Single-core Design](#), each processor core features two memories: instruction and scratch memory, which can both map onto synchronous, single-port, FPGA BRAM blocks. While this will reduce LUT requirements in designs with few cores, it becomes a non-trivial problem as the core counts increase. FPGAs have a fixed number of hard-BRAM blocks available for inference by the HDL compiler, for example the low-end Xilinx Spartan-6 XC6SLX9 FGPA features 32 18 Kb BRAM blocks [19, p. 2], and the Cyclone V 5CSEMA5F31C6N (used in the DE1-SoC) has 397 10 Kb blocks [20, p. 22].

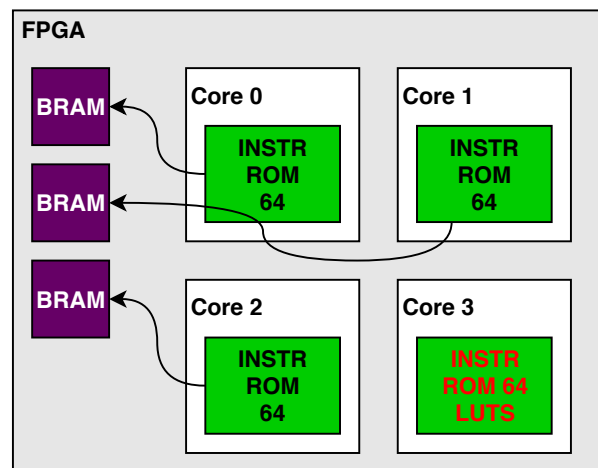


Figure 8.1: A theoretical FPGA device with 3 BRAM blocks running a 4-core design. Each core can map onto a BRAM block, however as there are more cores than BRAM blocks available, some core memories will be implemented as distributed RAM, or in the worse case using ALMs.

As shown in [Figure 8.1](#), as the number of processor cores increasing they eventually outnumber the available BRAM blocks, resulting in their memories being implemented in either distributed RAMs or ALMs, both of which can consume significant logic resources of the FPGA which reduces the maximum possible core count.

8.2.2 Maximum Frequency

8.3 Scenario Performance

To evaluate the performance of the system-on-chip, scenarios encompassing computational problems that are reflective of real-world applications are compiled and ran on the design.

8.3.1 Scenario Overview

The scenario is a software program that runs a parallel implementation of the summation function, i.e. $\text{sum } [1..10]$ which returns 55. While this may seem too simple at first to measure performance of a multi-core system-on-chip, the function is actually quite appropriate as it encompasses various parallel problems, such as: a fixed time/size serial part; broadcasting of the data set (in this case the range of the summation); thread synchronisation (to know when the data is ready and to schedule gathering of intermediary results); and is highly scalable.




Figure 8.2: Cyclone V maximum design frequency for various core count configurations.

The summation task flow is as follows:

1. Root (core #0) broadcasts the range of the summation (i.e. sum 1 to 10) to all cores via the global shared memory.
2. Non-root cores wait for this broadcast to finish (memory barrier), then calculate their own subset of the range to sum. For example, if Root broadcasts that there are 240 samples and 10 cores in the system, each core calculates the subset size:

$$240/10 = 24 \quad (8.1)$$

calculations starting from:

$$ID_{CORE} * 24 \quad (8.2)$$

For example, Core #5 will start its 24 sample subset summation from

$$5 * 24 = 120 \quad (8.3)$$

effectively performing sum [120..123].

3. All cores perform an intermediary summation over their subset of the range (serial part).
4. All cores attempt to add their intermediary result to a global sum value in global shared memory (mutex).
5. All cores halt, signalling that their work has been committed to the global shared memory and have finished the program.

This program is written in assembly in the file `sw/demos/asm/sum64.s` and can be compiled using the assembly compiler (developed for deliverable [ED4](#)) using the command below. The assembly compiler outputs the file `asm.s.hex` containing hex instruction words for use in Verilog's `$readmemh` function. This data is used for each core's instruction memory. The assembly program is also shown in [Section D.2](#).

```
python sw/asm.py sw/demos/asm/sum64.s
```

8.3.2 Performance Measurements

Behavioural simulation will be used to measure the following metrics to estimate general performance of the system-on-chip:

- Total program run-time.
This is the time from when the reset signal is de-asserted to when all cores have halted. Each core has an output halt signal which the SoC can use to determine if all cores have halted using `wire all_halted = &core_halts;`.
- Time spent on the serial part.
The serial part of this scenario consists of the intermediary summation of it's subset range. As each core is performing this task, the average will be used.
- Time spent on communication.
This includes time spent on thread synchronisation, i.e. waiting for the global memory to become available and waiting on the root to finish broadcast. Again, the average time will be used.
- Time spent fetching instructions.
Instruction fetches occur during stage STAGE_IF of the pipeline. The behavioural test bench will record the number of clock cycles each core spends in this state, then calculate the average time spent fetching instructions.

8.3.3 Performance Results

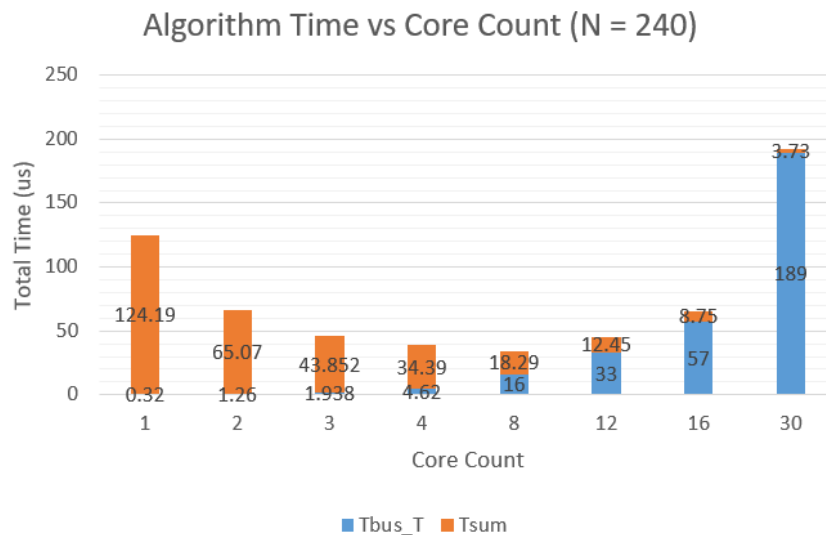


Figure 8.3: Chart showing how the communication times (Tbus) and serial times (Tsum) changes with core count.

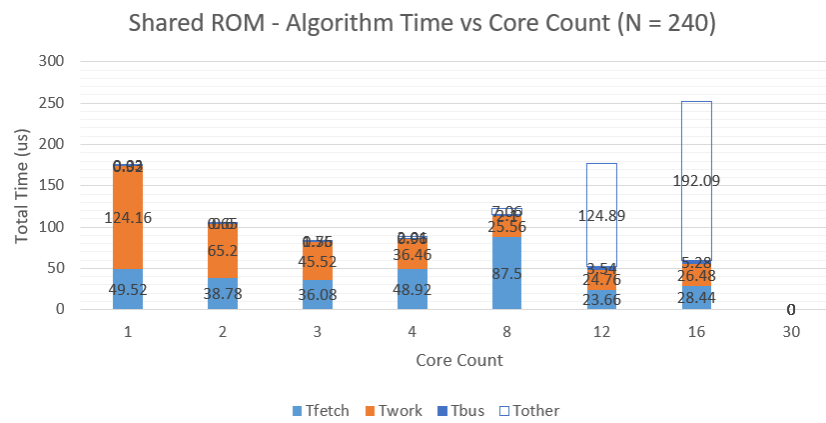


Figure 8.4: Similar to Figure 8.3 but using shared instruction memory to reduce block memory requirements per core.

References

- [1] Tech Differences, "Difference between loosely coupled and tightly coupled multiprocessor system (with comparison chart)," Jul 2017. [Online]. Available: <https://techdifferences.com/difference-between-loosely-coupled-and-tightly-coupled-multiprocessor-system.html> (Accessed 2019-04-20).
- [2] N. Chatterjee, S. Paul, and S. Chattopadhyay, "Fault-tolerant dynamic task mapping and scheduling for network-on-chip-based multicore platform," *ACM Transactions on Embedded Computing Systems*, vol. 16, pp. 1–24, 05 2017.
- [3] Terasic Technologies, "SoC Platform - Cyclone - DE1-SoC Board." [Online]. Available: <https://www.terasic.com.tw/cgi-bin/page/archive.pl?Language=English&No=836> (Accessed 2019-04-20).
- [4] *MiniSpartan6+*, Scarab Hardware, 2014. [Online]. Available: <https://www.scarabhardware.com/minispartan6/> (Accessed 2019-04-20).
- [5] V. Subramanian, "Multiple gate field-effect transistors for future CMOS technologies," *IETE Technical review*, vol. 27, no. 6, pp. 446–454, 2010.
- [6] M. J. Flynn, *Computer architecture: Pipelined and parallel processor design*. Jones & Bartlett Learning, 1995.
- [7] L. Benini and G. De Micheli, "Networks on Chips: A new SoC paradigm," *Computer*, vol. 35, pp. 70–78, 02 2002.
- [8] D. Zhu, L. Chen, S. Yue, T. M. Pinkston, and M. Pedram, "Balancing On-Chip Network Latency in Multi-application Mapping for Chip-Multiprocessors," in *2014 IEEE 28th International Parallel and Distributed Processing Symposium*, May 2014, pp. 872–881.
- [9] Xilinx, *Spartan-6 FPGA Block RAM Resources*, Xilinx.
- [10] Altera, *Recommended HDL Coding Styles - QII51007-9.0.0*, Altera.
- [11] B. Lancaster, "FPGA-based RISC Microprocessor and Compiler," vol. 3.14, pp. 37–50. [Online]. Available: <https://github.com/bendl/prco304> (Accessed March 2018).
- [12] P. Kocher, D. Genkin, D. Gruss, W. Haas, M. Hamburg, M. Lipp, S. Mangard, T. Prescher, M. Schwarz, and Y. Yarom, "Spectre attacks: Exploiting speculative execution," *arXiv preprint arXiv:1801.01203*, 2018.
- [13] M. Weber, "Arbiters: Design Ideas and Coding Styles," p. 2.
- [14] A. Palchoudhuri and R. S. Chakraborty, *High Performance Integer Arithmetic Circuit Design on FPGA: Architecture, Implementation and Design Automation*. Springer, 2015, vol. 51.
- [15] V. Salauyou and M. Gruszewski, "Designing of hierarchical structures for binary comparators on fpga/soc," in *IFIP International Conference on Computer Information Systems and Industrial Management*. Springer, 2015, pp. 386–396.

- [16] Xilinx, *Spartan-6 FPGA Configurable Logic Block - User Guide - UG384*, Xilinx.
- [17] Altera, *Cyclone V Device Handbook - Device Interfaces and Integration - CV-5V2*, Altera.
- [18] A. S. Tanenbaum, *Structured Computer Organization*. Pearson Education India, 2016.
- [19] Xilinx, *Spartan-6 Family Overview - DS160*, Xilinx.
- [20] Intel, *Cyclone V Device Overview - CV-51001*, Intel.

Appendix A

Peripheral Information

A.1 Special Registers	56
A.2 Watchdog Timer	57
A.3 GPIO Interface	57
A.4 Timer with Interrupt	58
A.5 UART Interface	58
B.1 Register Set Multiplex	59
B.2 Instruction Set Architecture	60

To provide user's with useful functionality, common system-on-chip peripherals were created. This section describes each peripheral and it's design decisions. The full memory-map is shown in [Figure 6.5](#).

A.1 Special Registers

From the software perspective, it is important for both the developer and software algorithms to know the target system's architecture to better utilise the resources available to them. Software written for one architecture with N cores must also run on an architecture with M cores. To enable such portability, the software must query the system for information such as: number of processor cores and the current core identifier. Without this information, the developer would be required to produce software for each individual architecture (e.g. an Intel i5 with 4 cores or an Intel i7 with 8 cores, or an NVIDIA GTX 970 with 1664 CUDA cores.

The special register peripheral is shown below.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
								CORE_ID								0080	R
								NUM_CORES								0081	R
SHARED_MEMORY cells (default 4096)																0082	R
								NUM_PERIPHERALS								0083	R
SCRATCH_MEMORY cells (default 64)																0084	RW
User defined																0085	RW
⋮																	
User defined																008F	RW

Figure A.1: Vmicro16 Special Registers layout (0x0080 - 0x008F).

A.2 Watchdog Timer

In any multi-threaded system there exists the possibility for a deadlock – a state where all threads are in a waiting state – and algorithm execution is forever blocked. This can occur either by poor software programming or incorrect thread arbitration by the processor. A common method of detecting a deadlock is to make each thread signal that it is not blocked by resetting a countdown timer. If the countdown timer is not reset, it will eventually reach zero and it is assumed that all threads are blocked as none have reset the countdown.

In this system-on-chip design, software can reset the watchdog timer by writing any 16-bit value to the address 0x00B8.

This peripheral is optional and can be enabled using the configuration parameters described in [Configuration Options](#).

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Reset Watchdog																00B8 W

A.3 GPIO Interface

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
								GPIO0 Output								0090 RW
GPIO1 Output																0091 RW
								GPIO2 Output								0092 RW
								GPIO3 Input								0093 R

On the DE1-SoC board, GPIO0 is assigned to the LEDs, and GPIO1 and GPIO2 to the 6 seven-segment displays.

A.4 Timer with Interrupt

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Load Value																0200 RW
													I	R	S	0201 W
Prescaler																0202 W

Clock Frequency Uses top level FPGA clock (normally 50 MHz).

Load Value Value to count down from each clock.

I Interrupt enable bit. Default 0.

R Reset Load Value and Prescaler values to their last written value.

S Start the timer countdown. 1 = start. 0 = stop.

Prescaler Number of clocks per FPGA clock to wait between each decrement.

A.5 UART Interface

15	8	7	1	0			
				Transmit Data		00A0 W	
				Receive Data		00A1 R	
					E	I	00A2 RW

E Enable the UART component.

I Enable an interrupt upon receiving new data. Default 1.

Note: If DEF_USE_REPROG is enabled in `vmicro16_soc_config.v` then the receiver port will be reserved for programming the instruction memory, resulting in reads and writes to addresses 0x00A1 and 0x00A2 to return 0.

Appendix B

Additional Figures

```
1  input      [MASTER_PORTS*BUS_WIDTH-1:0] S_PADDR,  
2  input      [MASTER_PORTS-1:0]          S_PWRITE,  
3  input      [MASTER_PORTS-1:0]          S_PSELx,  
4  input      [MASTER_PORTS-1:0]          S_PENABLE,  
5  input      [MASTER_PORTS*DATA_WIDTH-1:0] S_PWDATA,  
6  output reg [MASTER_PORTS*DATA_WIDTH-1:0] S_PRDATA,  
7  output reg [MASTER_PORTS-1:0]          S_PREADY,
```

Listing 7: Variable size inputs and outputs to the interconnect.

B.1 Register Set Multiplex

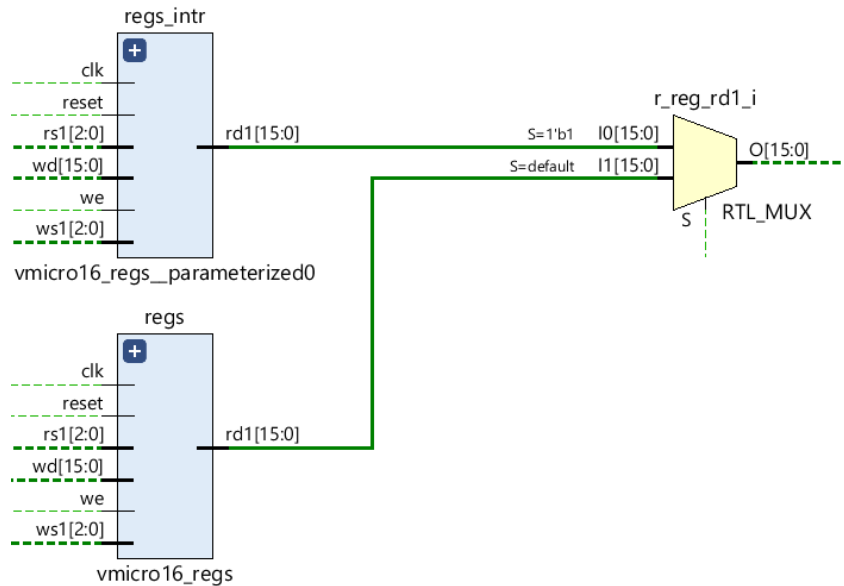


Figure B.1: Normal mode (bottom) and interrupt mode (top) register sets are multiplexed to switch between contexts.

B.2 Instruction Set Architecture

	15-11	10-8	7-5	4-0	rd ra simm5
	15-11	10-8	7-0		rd imm8
	15-11	10-0			nop
	15	14:12	11:0		extended immediate
SPCL	00000	11 bits			NOP
SPCL	00000	11h'000			NOP
SPCL	00000	11h'001			HALT
SPCL	00000	11h'002			Return from interrupt
LW	00001	Rd	Ra	s5	Rd <= RAM[Ra+s5]
SW	00010	Rd	Ra	s5	RAM[Ra+s5] <= Rd
BIT	00011	Rd	Ra	s5	bitwise operations
BIT_OR	00011	Rd	Ra	00000	Rd <= Rd Ra
BIT_XOR	00011	Rd	Ra	00001	Rd <= Rd ^ Ra
BIT_AND	00011	Rd	Ra	00010	Rd <= Rd & Ra
BIT_NOT	00011	Rd	Ra	00011	Rd <= ~Ra
BIT_LSHFT	00011	Rd	Ra	00100	Rd <= Rd << Ra
BIT_RSHFT	00011	Rd	Ra	00101	Rd <= Rd >> Ra
MOV	00100	Rd	Ra	X	Rd <= Ra
MOVI	00101	Rd		i8	Rd <= i8
ARITH_U	00110	Rd	Ra	s5	unsigned arithmetic
ARITH_UADD	00110	Rd	Ra	11111	Rd <= uRd + uRa
ARITH_USUB	00110	Rd	Ra	10000	Rd <= uRd - uRa
ARITH_UADDI	00110	Rd	Ra	0AAAA	Rd <= uRd + Ra + AAAA
ARITH_S	00111	Rd	Ra	s5	signed arithmetic
ARITH_SADD	00111	Rd	Ra	11111	Rd <= sRd + sRa
ARITH_SSUB	00111	Rd	Ra	10000	Rd <= sRd - sRa
ARITH_SSUBI	00111	Rd	Ra	0AAAA	Rd <= sRd - sRa + AAAA
BR	01000	Rd		i8	conditional branch
BR_U	01000	Rd		0000 0000	Any
BR_E	01000	Rd		0000 0001	Z=1
BR_NE	01000	Rd		0000 0010	Z=0
BR_G	01000	Rd		0000 0011	Z=0 and S=0
BR_GE	01000	Rd		0000 0100	S=0
BR_L	01000	Rd		0000 0101	S != 0
BR_LE	01000	Rd		0000 0110	Z=1 or (S != 0)
BR_S	01000	Rd		0000 0111	S=1
BR_NS	01000	Rd		0000 1000	S=0
CMP	01001	Rd	Ra	X	SZO <= CMP(Rd, Ra)
SETC	01010	Rd		Imm8	Rd <= (Imm8_f_SZO) ? 1 : 0
MULT	01011	Rd	Ra	X	Rd <= uRd * uRa
HALT	01100			X	
LWEX	01101	Rd	Ra	s5	Rd <= RAM[Ra+s5]
SWEX	01110	Rd	Ra	s5	RAM[Ra+s5] <= Rd Rd <= 0 1 if success

Figure B.2: Vmicro16 instruction set architecture.

Appendix C

Configuration Options

C.1 System-on-chip Configuration Options	61
C.2 Core Options	62
C.3 Peripheral Options	63

The following configuration options are defined in `vmicro16_soc_config.v`.

Defaults with empty/blank values signifies that the preprocessor define is commented out/not defined/disabled by default/computed by other parameters.

C.1 System-on-chip Configuration Options

Macro	Default	Purpose
CORES	4	Number of CPU cores in the SoC
SLAVES	8	Number of peripherals
DEF_USE_WATCHDOG	//	Enable watchdog module to recover from deadlocks and infinite loops
DEF_GLOBAL_RESET	//	Enable synchronous reset logic. Will consume more LUT resources. Does not reset BRAM blocks.

Table C.1: SoC Configuration Options

C.2 Core Options

Macro	Default	Purpose
DATA_WIDTH	16	Width of CPU registers in bits
DEF_CORE_HAS_INSTR_MEM	//	Enable a per core instruction memory cache
DEF_MEM_INSTR_DEPTH	64	Instruction memory cache per core
DEF_MEM_SCRATCH_DEPTH	64	RW RAM per core
DEF_ALU_HW_MULT	1	Enable/disable HW multiply (1 clock)
FIX_T3	//	Enable a T3 state for the APB transaction
DEF_USE_REPROG	//	Programme instruction memory via UART0. Requires DEF_GLOBAL_RESET. Enabling this will reserve the UART0 RX port for exclusive use for programming the instruction memory. Software reads of UART0 RX will return 0.

Table C.2: Core Options

C.3 Peripheral Options

Macro	Default	Purpose
APB.WIDTH		AMBA APB PADDR signal width
APB.PSELX_GPIO0	0	GPIO0 index
APB.PSELX_UART0	1	UART0 index
APB.PSELX_REGS0	2	REGS0 index
APB.PSELX_BRAM0	3	BRAM0 index
APB.PSELX_GPIO1	4	GPIO1 index
APB.PSELX_GPIO2	5	GPIO2 index
APB.PSELX_TIMR0	6	TIMR0 index
APB.BRAM0_CELLS	4096	Shared memory words
DEF_MMU_TIM0_S	16'h0000	Per core scratch memory start/end address
DEF_MMU_TIM0_E	16'h007F	"
DEF_MMU_SREG_S	16'h0080	Per core special registers start/end address
DEF_MMU_SREG_E	16'h008F	"
DEF_MMU_GPIO0_S	16'h0090	Shared GPIO0 start/end address
DEF_MMU_GPIO0_E	16'h0090	"
DEF_MMU_GPIO1_S	16'h0091	"
DEF_MMU_GPIO1_E	16'h0091	"
DEF_MMU_GPIO2_S	16'h0092	"
DEF_MMU_GPIO2_E	16'h0092	"
DEF_MMU_UART0_S	16'h00A0	Shared UART start/end address
DEF_MMU_UART0_E	16'h00A1	"
DEF_MMU_REGS0_S	16'h00B0	Shared registers start/end address
DEF_MMU_REGS0_E	16'h00B7	"
DEF_MMU_BRAM0_S	16'h1000	Shared memory with global monitor start/end address
DEF_MMU_BRAM0_E	16'h1FFF	"
DEF_MMU_TIMR0_S	16'h0200	Shared timer peripheral start/end address
DEF_MMU_TIMR0_E	16'h0202	"

Table C.3: Peripheral Options

Appendix D

Viva Demonstration Examples

D.1 2-core Timer Interrupt and ISR	64
D.2 1-160 Core Parallel Summation	66

D.1 2-core Timer Interrupt and ISR

This example demo, shown during the viva, blinks an LED every 0.5 seconds via a timer interrupt. Core 0 sets up the interrupt vector (by writing the isr0 function address to the interrupt vector) and enables all interrupt sources. Core 1 sets up the timer interval peripheral to produce an interrupt every 0.5 seconds. Core 1 also performs the interrupt handler (isr0): toggle an LED, write the state to UART0, and resets the watchdog.

```
1 // interrupts.s
2 // Toggle LED in ISR
3
4 entry:
5 // get core idx 0x80 in r7
6 movi r7, #0x80
7 lw r7, r7
8
9 // core1 sets up the timer
10 // Core0 enables interrupts and performs the isr
11 cmp r7, r0
12 movi r0, timer
13 br r0, BR_NE
14
15 // Set interrupt vector (0)
16 movi r0, isr0
17 movi r1, #0x1
18 movi r2, #0x08
19 lshft r1, r2
20 sw r0, r1
21
22 // enable all interrupts
23 movi r0, #0x0f
24 sw r0, r1 + #0x8
25
26 // enter idle state
27 halt r0, r0
28
29 timer:
30 // set timr0 address 0x200 into r0
31 // shift left 8 places
32 movi r0, #0x01
33 movi r1, #0x09
34 lshft r0, r1
35
36 // Set load value
37 //movi r1, #0x31
38 //sw r1, r0
39 // test we the expected value back
40 //lw r2, r0
41
42 // set load = 0x3000
```

```

43     movi    r1, #0x3
44     movi    r2, #0x0C
45     //movi   r2, #0x04
46     lshft   r1, r2
47     sw      r1, r0
48
49     // Set prescale value to 0x1000
50     // 20ns * load * prescaler = nanosecond delay
51     // 20ns * 10000 * 5000     = 1.0s
52     // 20.0 * 0x3000 * 0x1000  = ~1.0s
53     movi    r1, #0x1
54     // 1.0 second
55     //movi   r2, #0x0C
56     // 0.5 second
57     movi    r2, #0x0B
58     // 0.25 second
59     //movi   r2, #0x0a
60     // 0.0625 second
61     //movi   r2, #0x04
62     lshft   r1, r2
63     sw      r1, r0 + #0x02
64
65     // Start the timer (write 0x0001 to 0x0101)
66     movi    r1, #0x01
67     sw      r1, r0 + #0x01
68
69     exit:
70     // enter idle state
71     halt    r0, r0
72
73     isr0:
74     movi    r0, #0x90
75     lw      r1, r0
76     // xor with 1
77     movi    r2, #0x1
78     xor     r1, r2
79     // write back
80     sw      r1, r0
81
82     // write ascii value to uart0
83     movi    r0, #0xa0
84     movi    r2, #0x30
85     add     r1, r2
86     sw      r1, r0
87
88     // reset watchdog
89     movi    r0, #0xb8
90     sw      r1, r0
91
92     // return from interrupt
93     intr    r0, r0

```

D.2 1-160 Core Parallel Summation

This example demo performs a parallel summation of numbers 1 to 320. The algorithm *assigns* each core a subset of the summation space. It does this using the core's ID and the number of cores in the system. The following formulas determine where the subset begins and ends for each core. Core 0 broadcasts the number to sum to then each core calculates its subset start and end positions. Each core then performs a summation over it's subset then adds the result to a global shared value. After pushes it's results, the global shared value will contain the final summation result.

$$N_{samples} = 320 \quad (D.1)$$

$$N_{threads} = 64 \quad (D.2)$$

$$subset = N_{samples} / N_{threads} \quad (D.3)$$

$$start = ID * subset \quad (D.4)$$

$$end = start + subset \quad (D.5)$$

```

1  // sum64.s
2  // Simple 1-160 core summation program
3
4  // Set up common values, such as: Core id (r6),
5  // number of threads (cores) (r7), shared memory addresses (r5)
6  entry:
7  // Core id in r6
8  movi r0, #0x80
9  lw r0, r0
10 // store in r6
11 mov r6, r0
12
13 // get number of threads
14 movi r0, #0x81
15 lw r0, r0
16 // store in r7
17 mov r7, r0
18
19 // BRAMO shared memory 0x1000
20 movi r5, #0x01
21 movi r2, #0x0C
22 lshft r5, r2
23
24 jmp_to_barrier:
25 // NOT_ROOT
26 // wait at barrier
27 cmp r6, r3
28 movi r4, barrier_arrive
29 br r4, BR_NE
30
31 // ROOT
32 // calculates nsamples_per_thread
33 // ns = 100
34 // nst = ns / (num_threads)
35 // nst = ns >> (num_threads - 1)
36 // r0 = (num_threads - 1) WRONG!!!
37
38 root_broadcast:
39 // The root (core idx 0) broadcasts the number of samples
40 // 16 cores
41 //movi r4, #0x14
42 // 32 cores
43 //movi r4, #0x0a
44 // 64 cores
45 movi r4, #0x05
46 // 80 cores
47 //movi r4, #0x04
48 // 160 cores
49 //movi r4, #0x02
50
51 // ROOT
52 // Do the broadcast
53 // write nsamples_per_thread to shared bram (broadcast)
54 // 0x1001
55 sw r4, r5 + #0x01
56

```

```

57 // Reach the barrier to tell everyone
58 // that we have arrived
59 barrier_arrive:
60 // load latest count
61 lwex    r0, r5
62 // try increment count
63 // increment by 1
64 addi    r0, r3 + #0x01
65 // attempt store
66 swex    r0, r5
67 // check success (== 0)
68 cmp     r0, r3
69 // branch if failed
70 movi    r4, barrier_arrive
71 br      r4, BR_NE
72
73 // Wait in an infinite loop
74 // for all cores to 'arrive'
75 barrier:
76 // load the count
77 lw      r0, r5
78 // compare with number of threads
79 cmp     r0, r7
80 // jump back to barrier if not equal
81 movi    r4, barrier
82 br      r4, BR_NE
83
84 // EACH CORE
85 // All cores have arrived and in sync
86 synced1:
87 // Retrieve load the nsamples_per_thread
88 lw      r4, r5 + #0x01
89 // Calculate nstart = idx * nsamples_per_thread
90 // in r2
91 mov     r2, r6
92 mult    r2, r4
93
94 // Loop limit in r4
95 // samples_per_thread -> samples_per_thread + nstart
96 add     r4, r2
97
98 // Perform the summation in a tight for loop
99 // Sum numbers from nstart to limit
100 sum_loop:
101 // sum += i
102 add     r1, r2
103 // increment i
104 addi    r2, r3 + #0x01
105 // check end
106 cmp     r2, r4
107 movi    r0, sum_loop
108 br      r0, BR_NE
109
110 // Summation of the subset finished, result is in r1
111 // Now use a mutex to add it to the global sum value in shared mem
112 sum_mutex:
113 // load latest count
114 lwex    r0, r5 + #0x2
115 // try increment count
116 // increment by 1
117 add     r0, r1
118 // make copy as swex has a return value
119 mov     r2, r0
120 // attempt store
121 swex    r0, r5 + #0x02
122 // check success (== 0)
123 cmp     r0, r3
124 // branch if failed
125 movi    r4, sum_mutex
126 br      r4, BR_NE
127
128 // Write the latest global sum value to gpio1
129 write_gpio:
130 movi    r3, #0x91
131 sw      r2, r3
132
133 // Write the latest global sum value to uart0 tx
134 write_uart_done:
135 movi    r3, #0xa0
136 movi    r2, #0x30
137 add     r2, r6
138 sw      r2, r3
139
140 // This core has finished
141 // Enter a low power state
142 exit:
143 halt    r0, r0

```

Appendix E

Code Listing

E.1	SoC Code Listing	68
E.1.1	vmicro16_soc_config.v	68
E.1.2	top_ms.v	70
E.1.3	vmicro16_soc.v	71
E.1.4	vmicro16.v	77
E.2	Peripheral Code Listing	91

E.1 SoC Code Listing

E.1.1 vmicro16_soc_config.v

Configuration file for configuring the vmicro16_soc.v and vmicro16.v features.

```
1 // Configuration defines for the vmicro16_soc and vmicro16 cpu.
2
3 `ifndef VMICRO16_SOC_CONFIG_H
4 `define VMICRO16_SOC_CONFIG_H
5
6 `include "clog2.v"
7
8 `define FORMAL
9
10 `define CORES 32
11 `define SLAVES 8
12
13 ///////////////////////////////////////////////////////////////////
14 // Core parameters
15 ///////////////////////////////////////////////////////////////////
16 // Per core instruction memory
17 // Set this to give each core its own instruction memory cache
18 //`define DEF_CORE_HAS_INSTR_MEM
19
20 // Top level data width for registers, memory cells, bus widths
21 `define DATA_WIDTH 16
22
23 // Set this to use a workaround for the MMU's APB T2 clock
24 //`define FIX_T3
25
26 // Instruction memory (read only)
27 // Must be large enough to support software program.
28 `ifndef DEF_CORE_HAS_INSTR_MEM
29 // 64 16-bit words per core
30 `define DEF_MEM_INSTR_DEPTH 64
31 `else
32 // 4096 16-bit words global
33 `define DEF_MEM_INSTR_DEPTH 4096
34 `endif
35
36 // Scratch memory (read/write) on each core.
37 // See `DEF_MMU_TIMO_* defines for info.
38 `define DEF_MEM_SCRATCH_DEPTH 64
39
```

```

40 // Enables hardware multiplier and mult rr instruction
41 `define DEF_ALU_HW_MULT 1
42
43 // Enables global reset (requires more luts)
44 `define DEF_GLOBAL_RESET
45
46 // Enable a watch dog timer to reset the soc if threadlocked
47 `define DEF_USE_WATCHDOG
48
49 // Enables instruction memory programming via UART0
50 `define DEF_USE_REPROG
51
52 `ifdef DEF_USE_REPROG
53     `ifndef DEF_GLOBAL_RESET
54         `error_DEF_USE_REPROG_requires_DEF_GLOBAL_RESET
55     `endif
56 `endif
57
58 //////////////////////////////////////
59 // Memory mapping
60 //////////////////////////////////////
61 `define APB_WIDTH      (2 + `clog2(`CORES) + `DATA_WIDTH)
62
63 `define APB_PSELX_GPIO0 0
64 `define APB_PSELX_UART0 1
65 `define APB_PSELX_REGSO 2
66 `define APB_PSELX_BRAMO 3
67 `define APB_PSELX_GPIO1 4
68 `define APB_PSELX_GPIO2 5
69 `define APB_PSELX_TIMRO 6
70 `define APB_PSELX_WDOGO 7
71
72 `define APB_GPIO0_PINS 8
73 `define APB_GPIO1_PINS 16
74 `define APB_GPIO2_PINS 8
75
76 // Shared memory words
77 `define APB_BRAMO_CELLS 4096
78
79 //////////////////////////////////////
80 // Memory mapping
81 //////////////////////////////////////
82 // TIMO
83 // Number of scratch memory cells per core
84 `define DEF_MMU_TIMO_CELLS 64
85 `define DEF_MMU_TIMO_S 16'h0000
86 `define DEF_MMU_TIMO_E 16'h007F
87 // SREG
88 `define DEF_MMU_SREG_S 16'h0080
89 `define DEF_MMU_SREG_E 16'h008F
90 // GPIO0
91 `define DEF_MMU_GPIO0_S 16'h0090
92 `define DEF_MMU_GPIO0_E 16'h0090
93 // GPIO1
94 `define DEF_MMU_GPIO1_S 16'h0091
95 `define DEF_MMU_GPIO1_E 16'h0091
96 // GPIO2
97 `define DEF_MMU_GPIO2_S 16'h0092
98 `define DEF_MMU_GPIO2_E 16'h0092
99 // UART0
100 `define DEF_MMU_UART0_S 16'h00A0
101 `define DEF_MMU_UART0_E 16'h00A1
102 // REGSO
103 `define DEF_MMU_REGSO_S 16'h00B0
104 `define DEF_MMU_REGSO_E 16'h00B7
105 // WDOGO
106 `define DEF_MMU_WDOGO_S 16'h00B8
107 `define DEF_MMU_WDOGO_E 16'h00B8
108 // BRAMO
109 `define DEF_MMU_BRAMO_S 16'h1000
110 `define DEF_MMU_BRAMO_E 16'h1fff
111 // TIMRO
112 `define DEF_MMU_TIMRO_S 16'h0200
113 `define DEF_MMU_TIMRO_E 16'h0202
114
115 //////////////////////////////////////
116 // Interrupts
117 //////////////////////////////////////
118 // Enable/disable interrupts
119 // Disabling will free up resources for other features
120 `define DEF_ENABLE_INT
121 // Number of interrupt in signals
122 `define DEF_NUM_INT 8
123 // Default interrupt bitmask (0 = hidden, 1 = enabled)
124 `define DEF_INT_MASK 0
125 // Bit position of the TIMRO interrupt signal
126 `define DEF_INT_TIMRO 0
127 // Interrupt vector memory location
128 `define DEF_MMU_INTSV_S 16'h0100
129 `define DEF_MMU_INTSV_E 16'h0107

```

```

130 // Interrupt vector memory location
131 `define DEF_MMU_INTSM_S 16'h0108
132 `define DEF_MMU_INTSM_E 16'h0108
133
134 `endif

```

E.1.2 top_ms.v

Top level module that connects the SoC design to hardware pins on the FPGA.

```

1  module seven_display # (
2      parameter INVERT = 1
3  ) (
4      input  [3:0] n,
5      output [6:0] segments
6  );
7      reg [6:0] bits;
8      assign segments = (INVERT ? ~bits : bits);
9
10     always @(n)
11     case (n)
12         4'h0: bits = 7'b0111111; // 0
13         4'h1: bits = 7'b0000110; // 1
14         4'h2: bits = 7'b1011011; // 2
15         4'h3: bits = 7'b1001111; // 3
16         4'h4: bits = 7'b1100110; // 4
17         4'h5: bits = 7'b1101101; // 5
18         4'h6: bits = 7'b1111101; // 6
19         4'h7: bits = 7'b0000111; // 7
20         4'h8: bits = 7'b1111111; // 8
21         4'h9: bits = 7'b1100111; // 9
22         4'hA: bits = 7'b1110111; // A
23         4'hB: bits = 7'b1111100; // B
24         4'hC: bits = 7'b0111100; // C
25         4'hD: bits = 7'b1011110; // D
26         4'hE: bits = 7'b1111001; // E
27         4'hF: bits = 7'b1110001; // F
28     endcase
29 endmodule
30
31
32 // minispartan6+ XC6SLX9
33 module top_ms # (
34     parameter GPIO_PINS = 8
35 ) (
36     input          CLK50,
37     input  [3:0]   SW,
38     // UART
39     input          RXD,
40     output         TXD,
41     // Peripherals
42     output [7:0]   LEDS,
43
44     // 3v3 input from the s6 on the deisoc
45     input          S6_3v3,
46
47     // SSDs
48     output [6:0]   ssd0,
49     output [6:0]   ssd1,
50     output [6:0]   ssd2,
51     output [6:0]   ssd3,
52     output [6:0]   ssd4,
53     output [6:0]   ssd5
54 );
55     //wire [15:0] M_PADDR;
56     //wire M_PWRITE;
57     //wire [5-1:0] M_PSELx; // not shared
58     //wire M_PENABLE;
59     //wire [15:0] M_PWDATA;
60     //wire [15:0] M_PRDATA; // input to intercon
61     //wire M_PREADY; // input to intercon
62
63     wire [7:0] gpio0;
64     wire [15:0] gpio1;
65     wire [7:0] gpio2;
66
67     vmicro16_soc soc (
68         .clk      (CLK50),
69         .reset    (~SW[0]),
70
71         // .M_PADDR (M_PADDR),
72         // .M_PWRITE (M_PWRITE),
73         // .M_PSELx (M_PSELx),

```



```

74      //M_PENABLE (M_PENABLE),
75      //M_PWDATA (M_PWDATA),
76      //M_PRDATA (M_PRDATA),
77      //M_PREADY (M_PREADY),
78
79      // UART
80      .uart_tx (TXD),
81      .uart_rx (RXD),
82
83      // GPIO
84      .gpio0 (LEDS[3:0]),
85      .gpio1 (gpio1),
86      .gpio2 (gpio2),
87
88      // DEBUG
89      .debug0 (LEDS[4])
90      //debug1 (LEDS[7:4])
91 );
92
93 assign LEDS[7:5] = {TXD, RXD, S6_3v3};
94
95 // SSD displays (split across 2 gpio ports 1 and 2)
96 wire [3:0] ssd_chars [0:5];
97 assign ssd_chars[0] = gpio1[3:0];
98 assign ssd_chars[1] = gpio1[7:4];
99 assign ssd_chars[2] = gpio1[11:8];
100 assign ssd_chars[3] = gpio1[15:12];
101 assign ssd_chars[4] = gpio2[3:0];
102 assign ssd_chars[5] = gpio2[7:4];
103 seven_display ssd_0 (.n(ssd_chars[0]), .segments (ssd0));
104 seven_display ssd_1 (.n(ssd_chars[1]), .segments (ssd1));
105 seven_display ssd_2 (.n(ssd_chars[2]), .segments (ssd2));
106 seven_display ssd_3 (.n(ssd_chars[3]), .segments (ssd3));
107 seven_display ssd_4 (.n(ssd_chars[4]), .segments (ssd4));
108 seven_display ssd_5 (.n(ssd_chars[5]), .segments (ssd5));
109
110 endmodule

```

E.1.3 vmicro16_soc.v

```

1  //
2  //
3
4  `include "vmicro16_soc_config.v"
5  `include "clog2.v"
6  `include "formal.v"
7
8  module pow_reset # (
9      parameter INIT = 1,
10     parameter N = 8
11 ) (
12     input clk,
13     input reset,
14     output reg resethold
15 );
16     initial resethold = INIT ? (N-1) : 0;
17
18     always @(*)
19         resethold = |hold;
20
21     reg [`clog2(N)-1:0] hold = (N-1);
22     always @(posedge clk)
23         if (reset)
24             hold <= N-1;
25         else
26             if (hold)
27                 hold <= hold - 1;
28 endmodule
29
30 // Vmicro16 multi-core SoC with various peripherals
31 // and interrupts
32 module vmicro16_soc (
33     input clk,
34     input reset,
35
36     // UART0
37     input uart_rx,
38     output uart_tx,
39     //
40     output [`APB_GPIO0_PINS-1:0] gpio0,
41     output [`APB_GPIO1_PINS-1:0] gpio1,
42     output [`APB_GPIO2_PINS-1:0] gpio2,
43     //
44     output halt,
45     //

```

```

46     output    [`CORES-1:0]      debug0,
47     output    [`CORES*8-1:0]    debug1
48 );
49 wire [`CORES-1:0] w_halt;
50 assign halt = &w_halt;
51
52 assign debug0 = w_halt;
53
54 // Watchdog reset pulse signal.
55 // Passed to pow_reset to generate a longer reset pulse
56 wire wreset;
57 wire prog_prog;
58
59 // soft register reset hold for brams and registers
60 wire soft_reset;
61 `ifdef DEF_GLOBAL_RESET
62     pow_reset # (
63         .INIT      (1),
64         .N          (8)
65     ) por_inst (
66         .clk        (clk),
67         `ifdef DEF_USE_WATCHDOG
68         .reset      (reset | wreset | prog_prog),
69         `else
70         .reset      (reset),
71         `endif
72         .resethold  (soft_reset)
73     );
74 `else
75     assign soft_reset = 0;
76 `endif
77
78 // Peripherals (master to slave)
79 wire [`APB_WIDTH-1:0] M_PADDR;
80 wire M_PWRITE;
81 wire [`SLAVES-1:0] M_PSELx; // not shared
82 wire M_PENABLE;
83 wire [`DATA_WIDTH-1:0] M_PWDATA;
84 wire [`SLAVES*`DATA_WIDTH-1:0] M_PRDATA; // input to intercon
85 wire [`SLAVES-1:0] M_PREADY; // input
86
87 // Master apb interfaces
88 wire [`CORES*`APB_WIDTH-1:0] w_PADDR;
89 wire [`CORES-1:0] w_PWRITE;
90 wire [`CORES-1:0] w_PSELx;
91 wire [`CORES-1:0] w_PENABLE;
92 wire [`CORES*`DATA_WIDTH-1:0] w_PWDATA;
93 wire [`CORES*`DATA_WIDTH-1:0] w_PRDATA;
94 wire [`CORES-1:0] w_PREADY;
95
96 // Interrupts
97 `ifdef DEF_ENABLE_INT
98     wire [`DEF_NUM_INT-1:0] ints;
99     wire [`DEF_NUM_INT*`DATA_WIDTH-1:0] ints_data;
100     assign ints[7:1] = 0;
101     assign ints_data[`DEF_NUM_INT*`DATA_WIDTH-1:`DATA_WIDTH] =
102         {`DEF_NUM_INT*(`DATA_WIDTH-1){1'b0}};
103 `endif
104
105 apb_intercon_s # (
106     .MASTER_PORTS  (`CORES),
107     .SLAVE_PORTS    (`SLAVES),
108     .BUS_WIDTH      (`APB_WIDTH),
109     .DATA_WIDTH     (`DATA_WIDTH),
110     .HAS_PSELX_ADDR (1)
111 ) apb (
112     .clk            (clk),
113     .reset          (soft_reset),
114     // APB master to slave
115     .S_PADDR        (w_PADDR),
116     .S_PWRITE        (w_PWRITE),
117     .S_PSELx         (w_PSELx),
118     .S_PENABLE       (w_PENABLE),
119     .S_PWDATA         (w_PWDATA),
120     .S_PRDATA        (w_PRDATA),
121     .S_PREADY        (w_PREADY),
122     // shared bus
123     .M_PADDR         (M_PADDR),
124     .M_PWRITE         (M_PWRITE),
125     .M_PSELx         (M_PSELx),
126     .M_PENABLE       (M_PENABLE),
127     .M_PWDATA         (M_PWDATA),
128     .M_PRDATA        (M_PRDATA),
129     .M_PREADY        (M_PREADY)
130 );
131
132 `ifdef DEF_USE_WATCHDOG
133     vmicro16_watchdog_apb # (
134         .BUS_WIDTH  (`APB_WIDTH),

```

```

135     .NAME          ("WDOG0")
136 ) wdog0_apb (
137     .clk           (clk),
138     .reset         (),
139     // apb slave to master interface
140     .S_PADDR       (),
141     .S_PWRITE      (M_PWRITE),
142     .S_PSELx       (M_PSELx[`APB_PSELX_WDOG0]),
143     .S_PENABLE     (M_PENABLE),
144     .S_PWDATA      (),
145     .S_PRDATA      (),
146     .S_PREADY      (M_PREADY[`APB_PSELX_WDOG0]),
147
148     .wdreset       (wdreset)
149 );
150 `endif
151
152 vmicro16_gpio_apb # (
153     .BUS_WIDTH     (`APB_WIDTH),
154     .DATA_WIDTH    (`DATA_WIDTH),
155     .PORTS         (`APB_GPIO0_PINS),
156     .NAME          ("GPIO0")
157 ) gpio0_apb (
158     .clk           (clk),
159     .reset         (soft_reset),
160     // apb slave to master interface
161     .S_PADDR       (M_PADDR),
162     .S_PWRITE      (M_PWRITE),
163     .S_PSELx       (M_PSELx[`APB_PSELX_GPIO0]),
164     .S_PENABLE     (M_PENABLE),
165     .S_PWDATA      (M_PWDATA),
166     .S_PRDATA      (M_PRDATA[`APB_PSELX_GPIO0*`DATA_WIDTH +: `DATA_WIDTH]),
167     .S_PREADY      (M_PREADY[`APB_PSELX_GPIO0]),
168     .gpio          (gpio0)
169 );
170
171 // GPIO1 for Seven segment displays (16 pin)
172 vmicro16_gpio_apb # (
173     .BUS_WIDTH     (`APB_WIDTH),
174     .DATA_WIDTH    (`DATA_WIDTH),
175     .PORTS         (`APB_GPIO1_PINS),
176     .NAME          ("GPIO1")
177 ) gpio1_apb (
178     .clk           (clk),
179     .reset         (soft_reset),
180     // apb slave to master interface
181     .S_PADDR       (M_PADDR),
182     .S_PWRITE      (M_PWRITE),
183     .S_PSELx       (M_PSELx[`APB_PSELX_GPIO1]),
184     .S_PENABLE     (M_PENABLE),
185     .S_PWDATA      (M_PWDATA),
186     .S_PRDATA      (M_PRDATA[`APB_PSELX_GPIO1*`DATA_WIDTH +: `DATA_WIDTH]),
187     .S_PREADY      (M_PREADY[`APB_PSELX_GPIO1]),
188     .gpio          (gpio1)
189 );
190
191 // GPIO2 for Seven segment displays (8 pin)
192 vmicro16_gpio_apb # (
193     .BUS_WIDTH     (`APB_WIDTH),
194     .DATA_WIDTH    (`DATA_WIDTH),
195     .PORTS         (`APB_GPIO2_PINS),
196     .NAME          ("GPIO2")
197 ) gpio2_apb (
198     .clk           (clk),
199     .reset         (soft_reset),
200     // apb slave to master interface
201     .S_PADDR       (M_PADDR),
202     .S_PWRITE      (M_PWRITE),
203     .S_PSELx       (M_PSELx[`APB_PSELX_GPIO2]),
204     .S_PENABLE     (M_PENABLE),
205     .S_PWDATA      (M_PWDATA),
206     .S_PRDATA      (M_PRDATA[`APB_PSELX_GPIO2*`DATA_WIDTH +: `DATA_WIDTH]),
207     .S_PREADY      (M_PREADY[`APB_PSELX_GPIO2]),
208     .gpio          (gpio2)
209 );
210
211 apb_uart_tx # (
212     .DATA_WIDTH    (8),
213     .ADDR_EXP      (4) //2^4 = 16 FIFO words
214 ) uart0_apb (
215     .clk           (clk),
216     .reset         (soft_reset),
217     // apb slave to master interface
218     .S_PADDR       (M_PADDR),
219     .S_PWRITE      (M_PWRITE),
220     .S_PSELx       (M_PSELx[`APB_PSELX_UART0]),
221     .S_PENABLE     (M_PENABLE),
222     .S_PWDATA      (M_PWDATA),
223     .S_PRDATA      (M_PRDATA[`APB_PSELX_UART0*`DATA_WIDTH +: `DATA_WIDTH]),

```

```

224     .S_PREADY    (M_PREADY[`APB_PSELX_UART0]),
225     // uart wires
226     .tx_wire     (uart_tx),
227     .rx_wire     ()
228 );
229
230 timer_apb timr0 (
231     .clk          (clk),
232     .reset        (soft_reset),
233     // apb slave to master interface
234     .S_PADDR      (M_PADDR),
235     .S_PWRITE     (M_PWRITE),
236     .S_PSELx      (M_PSELx[`APB_PSELX_TIMRO]),
237     .S_PENABLE    (M_PENABLE),
238     .S_PWDATA     (M_PWDATA),
239     .S_PRDATA     (M_PRDATA[`APB_PSELX_TIMRO*`DATA_WIDTH +: `DATA_WIDTH]),
240     .S_PREADY     (M_PREADY[`APB_PSELX_TIMRO])
241     //
242     `ifdef DEF_ENABLE_INT
243     ,.out          (ints    [`DEF_INT_TIMRO]),
244     .int_data     (ints_data[`DEF_INT_TIMRO*`DATA_WIDTH +: `DATA_WIDTH])
245     `endif
246 );
247
248 // Shared register set for system-on-chip info
249 // R0 = number of cores
250 vmicro16_regs_apb # (
251     .BUS_WIDTH      (`APB_WIDTH),
252     .DATA_WIDTH     (`DATA_WIDTH),
253     .CELL_DEPTH     (8),
254     .PARAM_DEFAULTS_R0 (`CORES),
255     .PARAM_DEFAULTS_R1 (`SLAVES)
256 ) regs0_apb (
257     .clk          (clk),
258     .reset        (soft_reset),
259     // apb slave to master interface
260     .S_PADDR      (M_PADDR),
261     .S_PWRITE     (M_PWRITE),
262     .S_PSELx      (M_PSELx[`APB_PSELX_REGSO]),
263     .S_PENABLE    (M_PENABLE),
264     .S_PWDATA     (M_PWDATA),
265     .S_PRDATA     (M_PRDATA[`APB_PSELX_REGSO*`DATA_WIDTH +: `DATA_WIDTH]),
266     .S_PREADY     (M_PREADY[`APB_PSELX_REGSO])
267 );
268
269 vmicro16_bram_ex_apb # (
270     .BUS_WIDTH      (`APB_WIDTH),
271     .MEM_WIDTH      (`DATA_WIDTH),
272     .MEM_DEPTH      (`APB_BRAMO_CELLS),
273     .CORE_ID_BITS   (`clog2(`CORES))
274 ) bram_apb (
275     .clk          (clk),
276     .reset        (soft_reset),
277     // apb slave to master interface
278     .S_PADDR      (M_PADDR),
279     .S_PWRITE     (M_PWRITE),
280     .S_PSELx      (M_PSELx[`APB_PSELX_BRAMO]),
281     .S_PENABLE    (M_PENABLE),
282     .S_PWDATA     (M_PWDATA),
283     .S_PRDATA     (M_PRDATA[`APB_PSELX_BRAMO*`DATA_WIDTH +: `DATA_WIDTH]),
284     .S_PREADY     (M_PREADY[`APB_PSELX_BRAMO])
285 );
286
287 // There must be atleast 1 core
288 `static_assert(`CORES > 0)
289 `static_assert(`DEF_MEM_INSTR_DEPTH > 0)
290 `static_assert(`DEF_MMU_TIMO_CELLS > 0)
291
292
293 // Single instruction memory
294 `ifndef DEF_CORE_HAS_INSTR_MEM
295 // slave input/outputs from interconnect
296 wire [ `APB_WIDTH-1:0]      instr_M_PADDR;
297 wire                        instr_M_PWRITE;
298 wire [1-1:0]                instr_M_PSELx; // not shared
299 wire                        instr_M_PENABLE;
300 wire [ `DATA_WIDTH-1:0]     instr_M_PWDATA;
301 wire [1*`DATA_WIDTH-1:0]    instr_M_PRDATA; // slave response
302 wire [1-1:0]                instr_M_PREADY; // slave response
303
304 // Master apb interfaces
305 wire [ `CORES*`APB_WIDTH-1:0] instr_w_PADDR;
306 wire [ `CORES-1:0]           instr_w_PWRITE;
307 wire [ `CORES-1:0]           instr_w_PSELx;
308 wire [ `CORES-1:0]           instr_w_PENABLE;
309 wire [ `CORES*`DATA_WIDTH-1:0] instr_w_PWDATA;
310 wire [ `CORES*`DATA_WIDTH-1:0] instr_w_PRDATA;
311 wire [ `CORES-1:0]           instr_w_PREADY;
312

```

```

313 `ifdef DEF_USE_REPROG
314     wire [`clog2(`DEF_MEM_INSTR_DEPTH)-1:0] prog_addr;
315     wire [`DATA_WIDTH-1:0] prog_data;
316     wire prog_we;
317     uart_prog_rom_prog (
318         .clk          (clk),
319         .reset         (reset | wdreset),
320         // input stream
321         .uart_rx       (uart_rx),
322         // programmer
323         .addr          (prog_addr),
324         .data          (prog_data),
325         .we            (prog_we),
326         .prog          (prog_prog)
327     );
328 `endif
329
330 `ifdef DEF_USE_REPROG
331     vmicro16_bram_prog_apb
332 `else
333     vmicro16_bram_apb
334 `endif
335 # (
336     .BUS_WIDTH        (`APB_WIDTH),
337     .MEM_WIDTH        (`DATA_WIDTH),
338     .MEM_DEPTH        (`DEF_MEM_INSTR_DEPTH),
339     .USE_INITS        (1),
340     .NAME             ("INSTR_ROM_G")
341 ) instr_rom_apb (
342     .clk              (clk),
343     .reset            (reset),
344     .S_PADDR          (instr_M_PADDR),
345     .S_PWRITE         (0),
346     .S_PSELx          (instr_M_PSELx),
347     .S_PENABLE        (instr_M_PENABLE),
348     .S_PWDATA         (0),
349     .S_PRDATA         (instr_M_PRDATA),
350     .S_PREADY         (instr_M_PREADY)
351
352     `ifdef DEF_USE_REPROG
353     ,
354     .addr            (prog_addr),
355     .data            (prog_data),
356     .we              (prog_we),
357     .prog            (prog_prog)
358     `endif
359 );
360
361 apb_intercon_s # (
362     .MASTER_PORTS    (`CORES),
363     .SLAVE_PORTS      (1),
364     .BUS_WIDTH        (`APB_WIDTH),
365     .DATA_WIDTH       (`DATA_WIDTH),
366     .HAS_PSELX_ADDR   (0)
367 ) apb_instr_intercon (
368     .clk              (clk),
369     .reset            (soft_reset),
370     // APB master from cores
371     // master
372     .S_PADDR          (instr_w_PADDR),
373     .S_PWRITE         (instr_w_PWRITE),
374     .S_PSELx          (instr_w_PSELx),
375     .S_PENABLE        (instr_w_PENABLE),
376     .S_PWDATA         (instr_w_PWDATA),
377     .S_PRDATA         (instr_w_PRDATA),
378     .S_PREADY         (instr_w_PREADY),
379     // shared bus slaves
380     // slave outputs
381     .M_PADDR          (instr_M_PADDR),
382     .M_PWRITE         (instr_M_PWRITE),
383     .M_PSELx          (instr_M_PSELx),
384     .M_PENABLE        (instr_M_PENABLE),
385     .M_PWDATA         (instr_M_PWDATA),
386     .M_PRDATA         (instr_M_PRDATA),
387     .M_PREADY         (instr_M_PREADY)
388 );
389 `endif
390
391 genvar i;
392 generate for(i = 0; i < `CORES; i = i + 1) begin : cores
393
394     vmicro16_core # (
395         .CORE_ID        (i),
396         .DATA_WIDTH     (`DATA_WIDTH),
397
398         .MEM_INSTR_DEPTH (`DEF_MEM_INSTR_DEPTH),
399         .MEM_SCRATCH_DEPTH (`DEF_MMU_TIMO_CELLS)
400     ) c1 (
401         .clk            (clk),

```

```

402         .reset      (soft_reset),
403
404         // debug
405         .halt        (w_halt[i]),
406
407         // interrupts
408         .ints         (ints),
409         .ints_data    (ints_data),
410
411         // Output master port 1
412         .w_PADDR      (w_PADDR  [`APB_WIDTH*i +: `APB_WIDTH] ),
413         .w_PWRITE     (w_PWRITE  [i] ),
414         .w_PSELx      (w_PSELx   [i] ),
415         .w_PENABLE    (w_PENABLE [i] ),
416         .w_PWDATA     (w_PWDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
417         .w_PRDATA     (w_PRDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
418         .w_PREADY     (w_PREADY  [i] )
419
420 `ifndef DEF_CORE_HAS_INSTR_MEM
421     // APB instruction rom
422     , // Output master port 2
423     .w2_PADDR      (instr_w_PADDR  [`APB_WIDTH*i +: `APB_WIDTH] ),
424     // .w2_PWRITE    (instr_w_PWRITE [i] ),
425     .w2_PSELx      (instr_w_PSELx   [i] ),
426     .w2_PENABLE    (instr_w_PENABLE [i] ),
427     // .w2_PWDATA    (instr_w_PWDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
428     .w2_PRDATA     (instr_w_PRDATA  [`DATA_WIDTH*i +: `DATA_WIDTH]),
429     .w2_PREADY     (instr_w_PREADY  [i] )
430 `endif
431 );
432 end
433 endgenerate
434
435
436 //////////////////////////////////////
437 // Formal Verification
438 //////////////////////////////////////
439 `ifndef FORMAL
440 wire all_halted = &w_halt;
441 //////////////////////////////////////
442 // Count number of clocks each core is spending on
443 // bus transactions
444 //////////////////////////////////////
445 reg [15:0] bus_core_times    [0:`CORES-1]; // bus work
446 reg [15:0] core_work_times  [0:`CORES-1]; // serial work
447 reg [15:0] instr_fetch_times [0:`CORES-1]; // instruction fetches
448 integer i2;
449 initial
450     for(i2 = 0; i2 < `CORES; i2 = i2 + 1) begin
451         bus_core_times[i2] = 0;
452         core_work_times[i2] = 0;
453     end
454
455 // total bus time
456 generate
457     genvar g2;
458     for (g2 = 0; g2 < `CORES; g2 = g2 + 1) begin : formal_for_times
459         always @(posedge clk) begin
460             if (w_PSELx[g2])
461                 bus_core_times[g2] <= bus_core_times[g2] + 1;
462
463             // Core working time
464             `ifndef DEF_CORE_HAS_INSTR_MEM
465                 if (!w_PSELx[g2] && !instr_w_PSELx[g2])
466                     `else
467                         if (!w_PSELx[g2])
468                             `endif
469                             if (!w_halt[g2])
470                                 core_work_times[g2] <= core_work_times[g2] + 1;
471
472             end
473         end
474     endgenerate
475
476 reg [15:0] bus_time_average = 0;
477 reg [15:0] bus_reqs_average = 0;
478 reg [15:0] fetch_time_average = 0;
479 reg [15:0] work_time_average = 0;
480 //
481 always @(all_halted) begin
482     for (i2 = 0; i2 < `CORES; i2 = i2 + 1) begin
483         bus_time_average = bus_time_average + bus_core_times[i2];
484         bus_reqs_average = bus_reqs_average + bus_core_reqs_count[i2];
485         work_time_average = work_time_average + core_work_times[i2];
486         fetch_time_average = fetch_time_average + instr_fetch_times[i2];
487     end
488
489     bus_time_average = bus_time_average / `CORES;
490     bus_reqs_average = bus_reqs_average / `CORES;

```

```

491         work_time_average = work_time_average / `CORES;
492         fetch_time_average = fetch_time_average / `CORES;
493     end
494
495     //////////////////////////////////////
496     // Count number of bus requests per core
497     //////////////////////////////////////
498     // 1 clock delay of w_PSELx
499     reg [`CORES-1:0] bus_core_reqs_last;
500     // rising edges of each
501     wire [`CORES-1:0] bus_core_reqs_real;
502     // storage for counters for each core
503     reg [15:0] bus_core_reqs_count [0:`CORES-1];
504     initial
505         for(i2 = 0; i2 < `CORES; i2 = i2 + 1)
506             bus_core_reqs_count[i2] = 0;
507
508     // 1 clk delay to detect rising edge
509     always @(posedge clk)
510         bus_core_reqs_last <= w_PSELx;
511
512     generate
513         genvar g3;
514         for (g3 = 0; g3 < `CORES; g3 = g3 + 1) begin : formal_for_reqs
515             // Detect new reqs for each core
516             assign bus_core_reqs_real[g3] = w_PSELx[g3] >
517                                     bus_core_reqs_last[g3];
518
519             always @(posedge clk)
520                 if (bus_core_reqs_real[g3])
521                     bus_core_reqs_count[g3] <= bus_core_reqs_count[g3] + 1;
522
523         end
524     endgenerate
525
526     `ifndef DEF_CORE_HAS_INSTR_MEM
527         //////////////////////////////////////
528         // Time waiting for instruction fetches
529         // from global memory
530         //////////////////////////////////////
531         integer i3;
532         initial
533             for(i3 = 0; i3 < `CORES; i3 = i3 + 1)
534                 instr_fetch_times[i3] = 0;
535
536         // total bus time
537         // Instruction fetches occur on the w2 master port
538         generate
539             genvar g4;
540             for (g4 = 0; g4 < `CORES; g4 = g4 + 1) begin : formal_for_fetch_times
541                 always @(posedge clk)
542                     if (instr_w_PSELx[g4])
543                         instr_fetch_times[g4] <= instr_fetch_times[g4] + 1;
544
545             end
546         endgenerate
547     `endif
548
549     `endif // end FORMAL
550
551 endmodule
552

```

E.1.4 vmicro16.v

Vmicro16 CPU core module.

```

1  // This file contains multiple modules.
2  // Verilator likes 1 file for each module
3  /* verilator lint_off DECLFILENAME */
4  /* verilator lint_off UNUSED */
5  /* verilator lint_off BLKSEQ */
6  /* verilator lint_off WIDTH */
7
8  // Include Vmicro16 ISA containing definitions for the bits
9  `include "vmicro16_isa.v"
10
11  `include "clog2.v"
12  `include "formal.v"
13
14
15
16  // This module aims to be a SYNCHRONOUS, WRITE_FIRST BLOCK RAM

```

```

17 // https://www.xilinx.com/support/documentation/user_guides/ug473_7Series_Memory_Resources.pdf
18 // https://www.xilinx.com/support/documentation/user_guides/ug383.pdf
19 // https://www.xilinx.com/support/documentation/sw_manuals/xilinx2016_4/ug901-vivado-synthesis.pdf
20 module vmicro16_bram # (
21     parameter MEM_WIDTH      = 16,
22     parameter MEM_DEPTH      = 64,
23     parameter CORE_ID        = 0,
24     parameter USE_INITS      = 0,
25     parameter PARAM_DEFAULTS_R0 = 0,
26     parameter PARAM_DEFAULTS_R1 = 0,
27     parameter PARAM_DEFAULTS_R2 = 0,
28     parameter PARAM_DEFAULTS_R3 = 0,
29     parameter NAME           = "BRAM"
30 ) (
31     input clk,
32     input reset,
33
34     input      [`clog2(MEM_DEPTH)-1:0] mem_addr,
35     input      [MEM_WIDTH-1:0] mem_in,
36     input      mem_we,
37     output reg [MEM_WIDTH-1:0] mem_out
38 );
39 // memory vector
40 (* ram_style = "block" *)
41 reg [MEM_WIDTH-1:0] mem [0:MEM_DEPTH-1];
42
43 // not synthesizable
44 integer i;
45 initial begin
46     for (i = 0; i < MEM_DEPTH; i = i + 1) mem[i] = 0;
47     mem[0] = PARAM_DEFAULTS_R0;
48     mem[1] = PARAM_DEFAULTS_R1;
49     mem[2] = PARAM_DEFAULTS_R2;
50     mem[3] = PARAM_DEFAULTS_R3;
51
52     if (USE_INITS) begin
53         // define TEST_SW
54         `ifdef TEST_SW
55             $readmemh("E:\\Projects\\uni\\vmicro16\\sw\\verilog_memh.txt", mem);
56         `endif
57
58         `define TEST_ASM
59         `ifdef TEST_ASM
60             $readmemh("E:\\Projects\\uni\\vmicro16\\sw\\asm.s.hex", mem);
61         `endif
62
63         //`define TEST_COND
64         `ifdef TEST_COND
65             mem[0] = {`VMICRO16_OP_MOVI, 3'h7, 8'hC0}; // lock
66             mem[0] = {`VMICRO16_OP_MOVI, 3'h7, 8'hC0}; // lock
67         `endif
68
69         //`define TEST_CMP
70         `ifdef TEST_CMP
71             mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'h0A};
72             mem[1] = {`VMICRO16_OP_MOVI, 3'h1, 8'h0B};
73             mem[2] = {`VMICRO16_OP_CMP, 3'h1, 3'h0, 5'h1};
74         `endif
75
76         //`define TEST_LWEX
77         `ifdef TEST_LWEX
78             mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'hC5};
79             mem[1] = {`VMICRO16_OP_SW, 3'h0, 3'h0, 5'h1};
80             mem[2] = {`VMICRO16_OP_LW, 3'h2, 3'h0, 5'h1};
81             mem[3] = {`VMICRO16_OP_LWEX, 3'h2, 3'h0, 5'h1};
82             mem[4] = {`VMICRO16_OP_SWEX, 3'h3, 3'h0, 5'h1};
83         `endif
84
85         //`define TEST_MULTICORE
86         `ifdef TEST_MULTICORE
87             mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'h90};
88             mem[1] = {`VMICRO16_OP_MOVI, 3'h1, 8'h33};
89             mem[2] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
90             mem[3] = {`VMICRO16_OP_MOVI, 3'h0, 8'h80};
91             mem[4] = {`VMICRO16_OP_LW, 3'h2, 3'h0, 5'h0};
92             mem[5] = {`VMICRO16_OP_MOVI, 3'h1, 8'h33};
93             mem[6] = {`VMICRO16_OP_MOVI, 3'h1, 8'h33};
94             mem[7] = {`VMICRO16_OP_MOVI, 3'h1, 8'h33};
95             mem[8] = {`VMICRO16_OP_MOVI, 3'h0, 8'h91};
96             mem[9] = {`VMICRO16_OP_SW, 3'h2, 3'h0, 5'h0};
97         `endif
98
99         //`define TEST_BR
100        `ifdef TEST_BR
101            mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'h0};
102            mem[1] = {`VMICRO16_OP_MOVI, 3'h3, 8'h3};
103            mem[2] = {`VMICRO16_OP_MOVI, 3'h1, 8'h2};
104            mem[3] = {`VMICRO16_OP_ARITH_U, 3'h0, 3'h1, 5'b1111};
105            mem[4] = {`VMICRO16_OP_BR, 3'h3, `VMICRO16_OP_BR_U};
106            mem[5] = {`VMICRO16_OP_MOVI, 3'h0, 8'hFF};

```



```

107         `endif
108
109         //`define ALL_TEST
110         `ifdef ALL_TEST
111         // Standard all test
112         // REGS0
113         mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'h81};
114         mem[1] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0}; // MMU[0x81] = 6
115         mem[2] = {`VMICRO16_OP_SW, 3'h2, 3'h0, 5'h1}; // MMU[0x82] = 6
116         // GPIO0
117         mem[3] = {`VMICRO16_OP_MOVI, 3'h0, 8'h90};
118         mem[4] = {`VMICRO16_OP_MOVI, 3'h1, 8'hD};
119         mem[5] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
120         mem[6] = {`VMICRO16_OP_LW, 3'h2, 3'h0, 5'h0};
121         // TIMO
122         mem[7] = {`VMICRO16_OP_MOVI, 3'h0, 8'h07};
123         mem[8] = {`VMICRO16_OP_LW, 3'h3, 3'h0, 5'h03};
124         // UART0
125         mem[9] = {`VMICRO16_OP_MOVI, 3'h0, 8'hA0}; // UART0
126         mem[10] = {`VMICRO16_OP_MOVI, 3'h1, 8'h41}; // ascii A
127         mem[11] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
128         mem[12] = {`VMICRO16_OP_MOVI, 3'h1, 8'h42}; // ascii B
129         mem[13] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
130         mem[14] = {`VMICRO16_OP_MOVI, 3'h1, 8'h43}; // ascii C
131         mem[15] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
132         mem[16] = {`VMICRO16_OP_MOVI, 3'h1, 8'h44}; // ascii D
133         mem[17] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
134         mem[18] = {`VMICRO16_OP_MOVI, 3'h1, 8'h45}; // ascii D
135         mem[19] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
136         mem[20] = {`VMICRO16_OP_MOVI, 3'h1, 8'h46}; // ascii E
137         mem[21] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
138         // BRAMO
139         mem[22] = {`VMICRO16_OP_MOVI, 3'h0, 8'hC0};
140         mem[23] = {`VMICRO16_OP_MOVI, 3'h1, 8'hA};
141         mem[24] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h5};
142         mem[25] = {`VMICRO16_OP_LW, 3'h2, 3'h0, 5'h5};
143         // GPIO1 (SSD 24-bit port)
144         mem[26] = {`VMICRO16_OP_MOVI, 3'h0, 8'h91};
145         mem[27] = {`VMICRO16_OP_MOVI, 3'h1, 8'h12};
146         mem[28] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
147         mem[29] = {`VMICRO16_OP_LW, 3'h2, 3'h0, 5'h0};
148         // GPIO2
149         mem[30] = {`VMICRO16_OP_MOVI, 3'h0, 8'h92};
150         mem[31] = {`VMICRO16_OP_MOVI, 3'h1, 8'h56};
151         mem[32] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h0};
152         `endif
153
154         //`define TEST_BRAM
155         `ifdef TEST_BRAM
156         // 2 core BRAMO test
157         mem[0] = {`VMICRO16_OP_MOVI, 3'h0, 8'hC0};
158         mem[1] = {`VMICRO16_OP_MOVI, 3'h1, 8'hA};
159         mem[2] = {`VMICRO16_OP_SW, 3'h1, 3'h0, 5'h5};
160         mem[3] = {`VMICRO16_OP_LW, 3'h2, 3'h0, 5'h5};
161         `endif
162     end
163 end
164
165 always @(posedge clk) begin
166     // synchronous WRITE_FIRST (page 13)
167     if (mem_we) begin
168         mem[mem_addr] <= mem_in;
169         $display($time, "\t\t%s[%h] <= %h",
170             NAME, mem_addr, mem_in);
171     end else
172         mem_out <= mem[mem_addr];
173 end
174
175 // TODO: Reset impl = every clock while reset is asserted, clear each cell
176 // one at a time, mem[i++] <= 0
177 endmodule
178
179
180 module vmicro16_core_mmu # (
181     parameter MEM_WIDTH = 16,
182     parameter MEM_DEPTH = 64,
183
184     parameter CORE_ID = 3'h0,
185     parameter CORE_ID_BITS = `clog2(`CORES)
186 ) (
187     input clk,
188     input reset,
189
190     input req,
191     output busy,
192
193     // From core
194     input [MEM_WIDTH-1:0] mmu_addr,
195     input [MEM_WIDTH-1:0] mmu_in,

```

```

196     input                                mmu_we,
197     input                                mmu_lwex,
198     input                                mmu_swex,
199     output reg [MEM_WIDTH-1:0] mmu_out,
200
201     // interrupts
202     output reg [`DATA_WIDTH*`DEF_NUM_INT-1:0] ints_vector,
203     output reg [`DEF_NUM_INT-1:0] ints_mask,
204
205     // TO APB interconnect
206     output reg [`APB_WIDTH-1:0] M_PADDR,
207     output reg M_PWRITE,
208     output reg M_PSELx,
209     output reg M_PENABLE,
210     output reg [MEM_WIDTH-1:0] M_PWDATA,
211     // from interconnect
212     input [MEM_WIDTH-1:0] M_PRDATA,
213     input M_PREADY
214 );
215     localparam MMU_STATE_T1 = 0;
216     localparam MMU_STATE_T2 = 1;
217     localparam MMU_STATE_T3 = 2;
218     reg [1:0] mmu_state = MMU_STATE_T1;
219
220     reg [MEM_WIDTH-1:0] per_out = 0;
221     wire [MEM_WIDTH-1:0] tim0_out;
222
223     assign busy = req || (mmu_state == MMU_STATE_T2);
224
225     // more luts than below but easier
226     //wire tim0_en = (mmu_addr >= `DEF_MMU_TIMO_S)
227     //              && (mmu_addr <= `DEF_MMU_TIMO_E);
228     //wire sreg_en = (mmu_addr >= `DEF_MMU_SREG_S)
229     //              && (mmu_addr <= `DEF_MMU_SREG_E);
230     //wire intv_en = (mmu_addr >= `DEF_MMU_INTSV_S)
231     //              && (mmu_addr <= `DEF_MMU_INTSV_E);
232     //wire intm_en = (mmu_addr >= `DEF_MMU_INTSM_S)
233     //              && (mmu_addr <= `DEF_MMU_INTSM_E);
234
235     wire tim0_en = ~mmu_addr[12] && ~mmu_addr[9] && ~mmu_addr[7];
236     wire sreg_en = mmu_addr[7] && ~mmu_addr[4] && ~mmu_addr[5];
237     wire intv_en = mmu_addr[8] && ~mmu_addr[3];
238     wire intm_en = mmu_addr[8] && mmu_addr[3];
239
240     wire apb_en = !(tim0_en, sreg_en, intv_en, intm_en);
241     wire tim0_we = (tim0_en && mmu_we);
242     wire intv_we = (intv_en && mmu_we);
243     wire intm_we = (intm_en && mmu_we);
244
245     // Special register selects
246     localparam SPECIAL_REGS = 8;
247     wire [MEM_WIDTH-1:0] sr_val;
248
249     // Interrupt vector and mask
250     initial ints_vector = 0;
251     initial ints_mask = 0;
252     wire [2:0] intv_addr = mmu_addr[`clog2(`DEF_NUM_INT)-1:0];
253     always @(posedge clk)
254         if (intv_we)
255             ints_vector[intv_addr*`DATA_WIDTH +: `DATA_WIDTH] <= mmu_in;
256
257     always @(posedge clk)
258         if (intm_we)
259             ints_mask <= mmu_in;
260
261
262     always @(ints_vector)
263         $display($time,
264             "\tC%d\t\tints_vector W: | %h %h %h %h | %h %h %h %h |",
265             CORE_ID,
266             ints_vector[0*`DATA_WIDTH +: `DATA_WIDTH],
267             ints_vector[1*`DATA_WIDTH +: `DATA_WIDTH],
268             ints_vector[2*`DATA_WIDTH +: `DATA_WIDTH],
269             ints_vector[3*`DATA_WIDTH +: `DATA_WIDTH],
270             ints_vector[4*`DATA_WIDTH +: `DATA_WIDTH],
271             ints_vector[5*`DATA_WIDTH +: `DATA_WIDTH],
272             ints_vector[6*`DATA_WIDTH +: `DATA_WIDTH],
273             ints_vector[7*`DATA_WIDTH +: `DATA_WIDTH]
274         );
275
276     always @(intm_we)
277         $display($time, "\tC%d\t\tintm_we W: %b", CORE_ID, ints_mask);
278
279     // Output port
280     always @(*)
281         if (tim0_en) mmu_out = tim0_out;
282         else if (sreg_en) mmu_out = sr_val;
283         else if (intv_en) mmu_out = ints_vector[mmu_addr[2:0]*`DATA_WIDTH
284             +: `DATA_WIDTH];

```

```

285         else if (intm_en) mmu_out = ints_mask;
286         else
287             mmu_out = per_out;
288
289         // APB master to slave interface
290         always @(posedge clk)
291             if (reset) begin
292                 mmu_state <= MMU_STATE_T1;
293                 M_PENABLE <= 0;
294                 M_PADDR <= 0;
295                 M_PWDATA <= 0;
296                 M_PSELx <= 0;
297                 M_PWRITE <= 0;
298             end
299             else
300                 casex (mmu_state)
301                     MMU_STATE_T1: begin
302                         if (req && apb_en) begin
303                             M_PADDR <= {mmu_lwex,
304                                     mmu_swex,
305                                     CORE_ID[CORE_ID_BITS-1:0],
306                                     mmu_addr[MEM_WIDTH-1:0]};
307
308                             M_PWDATA <= mmu_in;
309                             M_PSELx <= 1;
310                             M_PWRITE <= mmu_we;
311
312                             mmu_state <= MMU_STATE_T2;
313                         end
314                     end
315                     `ifdef FIX_T3
316                     MMU_STATE_T2: begin
317                         M_PENABLE <= 1;
318
319                         if (M_PREADY == 1'b1) begin
320                             mmu_state <= MMU_STATE_T3;
321                         end
322                     end
323                     MMU_STATE_T3: begin
324                         // Slave has output a ready signal (finished)
325                         M_PENABLE <= 0;
326                         M_PADDR <= 0;
327                         M_PWDATA <= 0;
328                         M_PSELx <= 0;
329                         M_PWRITE <= 0;
330                         // Clock the peripheral output into a reg,
331                         // to output on the next clock cycle
332                         per_out <= M_PRDATA;
333
334                         mmu_state <= MMU_STATE_T1;
335                     end
336                     `else
337                     // No FIX_T3
338                     MMU_STATE_T2: begin
339                         if (M_PREADY == 1'b1) begin
340                             M_PENABLE <= 0;
341                             M_PADDR <= 0;
342                             M_PWDATA <= 0;
343                             M_PSELx <= 0;
344                             M_PWRITE <= 0;
345                             // Clock the peripheral output into a reg,
346                             // to output on the next clock cycle
347                             per_out <= M_PRDATA;
348
349                             mmu_state <= MMU_STATE_T1;
350                         end
351                     end
352                     else begin
353                         M_PENABLE <= 1;
354                     end
355                 end
356             endcase
357
358         (* ram_style = "block" *)
359         vmicro16_bram # (
360             .MEM_WIDTH (MEM_WIDTH),
361             .MEM_DEPTH (SPECIAL_REGS),
362             .USE_INITS (0),
363             .PARAM_DEFAULTS_RO (CORE_ID),
364             .PARAM_DEFAULTS_R1 (`CORES),
365             .PARAM_DEFAULTS_R2 (`APB_BRAMO_CELLS),
366             .PARAM_DEFAULTS_R3 (`SLAVES),
367             .NAME ("ram_sr")
368         ) ram_sr (
369             .clk (clk),
370             .reset (reset),
371             .mem_addr (mmu_addr[`clog2(SPECIAL_REGS)-1:0]),
372             .mem_in (),
373             .mem_we (),
374             .mem_out (sr_val)

```

```

375 );
376
377 // Each M core has a TIMO scratch memory
378 (* ram_style = "block" *)
379 vmicro16_bram # (
380     .MEM_WIDTH    (MEM_WIDTH),
381     .MEM_DEPTH    (MEM_DEPTH),
382     .USE_INITS    (0),
383     .NAME         ("TIMO")
384 ) TIMO (
385     .clk          (clk),
386     .reset        (reset),
387     .mem_addr     (mmu_addr[7:0]),
388     .mem_in       (mmu_in),
389     .mem_we       (tim0_we),
390     .mem_out      (tim0_out)
391 );
392 endmodule
393
394
395
396 module vmicro16_regs # (
397     parameter CELL_WIDTH      = 16,
398     parameter CELL_DEPTH     = 8,
399     parameter CELL_SEL_BITS   = `clog2(CELL_DEPTH),
400     parameter CELL_DEFAULTS  = 0,
401     parameter DEBUG_NAME     = "",
402     parameter CORE_ID        = 0,
403     parameter PARAM_DEFAULTS_R0 = 16'h0000,
404     parameter PARAM_DEFAULTS_R1 = 16'h0000
405 ) (
406     input clk,
407     input reset,
408     // Dual port register reads
409     input [CELL_SEL_BITS-1:0] rs1, // port 1
410     output [CELL_WIDTH-1:0] rd1,
411     //input [CELL_SEL_BITS-1:0] rs2, // port 2
412     //output [CELL_WIDTH-1:0] rd2,
413     // EX/WB final stage write back
414     input we,
415     input [CELL_SEL_BITS-1:0] ws1,
416     input [CELL_WIDTH-1:0] wd
417 );
418 (* ram_style = "distributed" *)
419 reg [CELL_WIDTH-1:0] regs [0:CELL_DEPTH-1] /*verilator public_flat*/;
420
421 // Initialise registers with default values
422 // Really only used for special registers used by the soc
423 // TODO: How to do this on reset?
424 integer i;
425 initial
426     if (CELL_DEFAULTS)
427         $readmemh(CELL_DEFAULTS, regs);
428     else begin
429         for(i = 0; i < CELL_DEPTH; i = i + 1)
430             regs[i] = 0;
431         regs[0] = PARAM_DEFAULTS_R0;
432         regs[1] = PARAM_DEFAULTS_R1;
433     end
434
435 `ifndef ICARUS
436     always @(regs)
437         $display($time, "\tC%02h\t\t| %h %h %h %h | %h %h %h %h |",
438             CORE_ID,
439             regs[0], regs[1], regs[2], regs[3],
440             regs[4], regs[5], regs[6], regs[7]);
441 `endif
442
443 always @(posedge clk)
444     if (reset) begin
445         for(i = 0; i < CELL_DEPTH; i = i + 1)
446             regs[i] <= 0;
447         regs[0] <= PARAM_DEFAULTS_R0;
448         regs[1] <= PARAM_DEFAULTS_R1;
449     end
450     else if (we) begin
451         $display($time, "\tC%02h: REGS #s: Writing %h to reg[%d]",
452             CORE_ID, DEBUG_NAME, wd, ws1);
453
454         // Perform the write
455         regs[ws1] <= wd;
456     end
457
458 // sync writes, async reads
459 assign rd1 = regs[rs1];
460 //assign rd2 = regs[rs2];
461 endmodule
462
463 module vmicro16_dec # (

```

```

464     parameter INSTR_WIDTH    = 16,
465     parameter INSTR_OP_WIDTH = 5,
466     parameter INSTR_RS_WIDTH = 3,
467     parameter ALU_OP_WIDTH   = 5
468 ) (
469     //input clk,    // not used yet (all combinational)
470     //input reset, // not used yet (all combinational)
471
472     input  [INSTR_WIDTH-1:0] instr,
473
474     output [INSTR_OP_WIDTH-1:0] opcode,
475     output [INSTR_RS_WIDTH-1:0] rd,
476     output [INSTR_RS_WIDTH-1:0] ra,
477     output [3:0] imm4,
478     output [7:0] imm8,
479     output [11:0] imm12,
480     output [4:0] simm5,
481
482     // This can be freely increased without affecting the isa
483     output reg [ALU_OP_WIDTH-1:0] alu_op,
484
485     output reg has_imm4,
486     output reg has_imm8,
487     output reg has_imm12,
488     output reg has_we,
489     output reg has_br,
490     output reg has_mem,
491     output reg has_mem_we,
492     output reg has_cmp,
493
494     output halt,
495     output intr,
496
497     output reg has_lwex,
498     output reg has_swex
499
500     // TODO: Use to identify bad instruction and
501     //        raise exceptions
502     //, output is_bad
503 );
504 assign opcode = instr[15:11];
505 assign rd     = instr[10:8];
506 assign ra     = instr[7:5];
507 assign imm4   = instr[3:0];
508 assign imm8   = instr[7:0];
509 assign imm12  = instr[11:0];
510 assign simm5  = instr[4:0];
511
512 // exme_op
513 always @(*) case (opcode)
514     `VMICRO16_OP_SPCL: casez(instr[11:0])
515         `VMICRO16_OP_SPCL_NOP,
516         `VMICRO16_OP_SPCL_HALT,
517         `VMICRO16_OP_SPCL_INTR: alu_op = `VMICRO16_ALU_NOP;
518         default: alu_op = `VMICRO16_ALU_NOP; endcase
519
520     `VMICRO16_OP_LW: alu_op = `VMICRO16_ALU_LW;
521     `VMICRO16_OP_SW: alu_op = `VMICRO16_ALU_SW;
522     `VMICRO16_OP_LWEX: alu_op = `VMICRO16_ALU_LW;
523     `VMICRO16_OP_SWEX: alu_op = `VMICRO16_ALU_SW;
524
525     `VMICRO16_OP_MOV: alu_op = `VMICRO16_ALU_MOV;
526     `VMICRO16_OP_MOVI: alu_op = `VMICRO16_ALU_MOVI;
527
528     `VMICRO16_OP_BR: alu_op = `VMICRO16_ALU_BR;
529     `VMICRO16_OP_MULT: alu_op = `VMICRO16_ALU_MULT;
530
531     `VMICRO16_OP_CMP: alu_op = `VMICRO16_ALU_CMP;
532     `VMICRO16_OP_SETC: alu_op = `VMICRO16_ALU_SETC;
533
534     `VMICRO16_OP_BIT: casez (simm5)
535         `VMICRO16_OP_BIT_OR: alu_op = `VMICRO16_ALU_BIT_OR;
536         `VMICRO16_OP_BIT_XOR: alu_op = `VMICRO16_ALU_BIT_XOR;
537         `VMICRO16_OP_BIT_AND: alu_op = `VMICRO16_ALU_BIT_AND;
538         `VMICRO16_OP_BIT_NOT: alu_op = `VMICRO16_ALU_BIT_NOT;
539         `VMICRO16_OP_BIT_LSHFT: alu_op = `VMICRO16_ALU_BIT_LSHFT;
540         `VMICRO16_OP_BIT_RSHFT: alu_op = `VMICRO16_ALU_BIT_RSHFT;
541         default: alu_op = `VMICRO16_ALU_BAD; endcase
542
543     `VMICRO16_OP_ARITH_U: casez (simm5)
544         `VMICRO16_OP_ARITH_UADD: alu_op = `VMICRO16_ALU_ARITH_UADD;
545         `VMICRO16_OP_ARITH_USUB: alu_op = `VMICRO16_ALU_ARITH_USUB;
546         `VMICRO16_OP_ARITH_UADDI: alu_op = `VMICRO16_ALU_ARITH_UADDI;
547         default: alu_op = `VMICRO16_ALU_BAD; endcase
548
549     `VMICRO16_OP_ARITH_S: casez (simm5)
550         `VMICRO16_OP_ARITH_SADD: alu_op = `VMICRO16_ALU_ARITH_SADD;
551         `VMICRO16_OP_ARITH_SSUB: alu_op = `VMICRO16_ALU_ARITH_SSUB;
552         `VMICRO16_OP_ARITH_SSUBI: alu_op = `VMICRO16_ALU_ARITH_SSUBI;
553         default: alu_op = `VMICRO16_ALU_BAD; endcase

```

```

554
555         default: begin
556             alu_op = `VMICRO16_ALU_NOP;
557             $display($time, "\tDEC: unknown opcode: %h ... NOPPING", opcode);
558         end
559     endcase
560
561     // Special opcodes
562     //assign nop == ((opcode == `VMICRO16_OP_SPCL) & (~instr[0]));
563     assign halt = ((opcode == `VMICRO16_OP_SPCL) & instr[0]);
564     assign intr = ((opcode == `VMICRO16_OP_SPCL) & instr[1]);
565
566     // Register writes
567     always @(*) case (opcode)
568         `VMICRO16_OP_LWEX,
569         `VMICRO16_OP_SWEX,
570         `VMICRO16_OP_LW,
571         `VMICRO16_OP_MOV,
572         `VMICRO16_OP_MOVI,
573         // `VMICRO16_OP_MOVI_L,
574         `VMICRO16_OP_ARITH_U,
575         `VMICRO16_OP_ARITH_S,
576         `VMICRO16_OP_SETC,
577         `VMICRO16_OP_BIT,
578         `VMICRO16_OP_MULT:      has_we = 1'b1;
579         default:                has_we = 1'b0;
580     endcase
581
582     // Contains 4-bit immediate
583     always @(*)
584         if( ((opcode == `VMICRO16_OP_ARITH_U) && (simm5[4] == 0)) ||
585             ((opcode == `VMICRO16_OP_ARITH_S) && (simm5[4] == 0)) )
586             has_imm4 = 1'b1;
587         else
588             has_imm4 = 1'b0;
589
590     // Contains 8-bit immediate
591     always @(*) case (opcode)
592         `VMICRO16_OP_MOVI,
593         `VMICRO16_OP_BR:      has_imm8 = 1'b1;
594         default:              has_imm8 = 1'b0;
595     endcase
596
597     /// Contains 12-bit immediate
598     //always @(*) case (opcode)
599     //     `VMICRO16_OP_MOVI_L:      has_imm12 = 1'b1;
600     //     default:                  has_imm12 = 1'b0;
601     //endcase
602
603     // Will branch the pc
604     always @(*) case (opcode)
605         `VMICRO16_OP_BR:      has_br = 1'b1;
606         default:              has_br = 1'b0;
607     endcase
608
609     // Requires external memory
610     always @(*) case (opcode)
611         `VMICRO16_OP_LW,
612         `VMICRO16_OP_SW,
613         `VMICRO16_OP_LWEX,
614         `VMICRO16_OP_SWEX:    has_mem = 1'b1;
615         default:              has_mem = 1'b0;
616     endcase
617
618     // Requires external memory write
619     always @(*) case (opcode)
620         `VMICRO16_OP_SW,
621         `VMICRO16_OP_SWEX:    has_mem_we = 1'b1;
622         default:              has_mem_we = 1'b0;
623     endcase
624
625     // Affects status registers (cmp instructions)
626     always @(*) case (opcode)
627         `VMICRO16_OP_CMP:      has_cmp = 1'b1;
628         default:              has_cmp = 1'b0;
629     endcase
630
631     // Performs exclusive checks
632     always @(*) case (opcode)
633         `VMICRO16_OP_LWEX:      has_lwex = 1'b1;
634         default:                has_lwex = 1'b0;
635     endcase
636
637     always @(*) case (opcode)
638         `VMICRO16_OP_SWEX:      has_swex = 1'b1;
639         default:                has_swex = 1'b0;
640     endcase
641 endmodule
642
643

```

```

644 module vmicro16_alu # (
645     parameter OP_WIDTH = 5,
646     parameter DATA_WIDTH = 16,
647     parameter CORE_ID = 0
648 ) (
649     // input clk, // TODO: make clocked
650
651     input [OP_WIDTH-1:0] op,
652     input [DATA_WIDTH-1:0] a, // rs1/dst
653     input [DATA_WIDTH-1:0] b, // rs2
654     input [3:0] flags,
655     output reg [DATA_WIDTH-1:0] c
656 );
657 localparam TOP_BIT = (DATA_WIDTH-1);
658 // 17-bit register
659 reg [DATA_WIDTH:0] cmp_tmp = 0; // = {carry, [15:0]}
660 wire r_setc;
661
662 always @(*) begin
663     cmp_tmp = 0;
664     case (op)
665         // branch/nop, output nothing
666         `VMICRO16_ALU_BR,
667         `VMICRO16_ALU_NOP: c = {DATA_WIDTH{1'b0}};
668         // load/store addresses (use value in rd2)
669         `VMICRO16_ALU_LW,
670         `VMICRO16_ALU_SW: c = b;
671         // bitwise operations
672         `VMICRO16_ALU_BIT_OR: c = a | b;
673         `VMICRO16_ALU_BIT_XOR: c = a ^ b;
674         `VMICRO16_ALU_BIT_AND: c = a & b;
675         `VMICRO16_ALU_BIT_NOT: c = ~(b);
676         `VMICRO16_ALU_BIT_LSHFT: c = a << b;
677         `VMICRO16_ALU_BIT_RSHFT: c = a >> b;
678
679         `VMICRO16_ALU_MOV: c = b;
680         `VMICRO16_ALU_MOVI: c = b;
681         `VMICRO16_ALU_MOVI_L: c = b;
682
683         `VMICRO16_ALU_ARITH_UADD: c = a + b;
684         `VMICRO16_ALU_ARITH_USUB: c = a - b;
685         // TODO: ALU should have simm5 as input
686         `VMICRO16_ALU_ARITH_UADDI: c = a + b;
687
688         `ifdef DEF_ALU_HW_MULT
689             `VMICRO16_ALU_MULT: c = a * b;
690         `endif
691
692         `VMICRO16_ALU_ARITH_SADD: c = $signed(a) + $signed(b);
693         `VMICRO16_ALU_ARITH_SSUB: c = $signed(a) - $signed(b);
694         // TODO: ALU should have simm5 as input
695         `VMICRO16_ALU_ARITH_SSUBI: c = $signed(a) - $signed(b);
696
697         `VMICRO16_ALU_CMP: begin
698             // TODO: Do a-b in 17-bit register
699             // Set zero, overflow, carry, signed bits in result
700             cmp_tmp = a - b;
701             c = 0;
702
703             // N Negative condition code flag
704             // Z Zero condition code flag
705             // C Carry condition code flag
706             // V Overflow condition code flag
707             c[`VMICRO16_SFLAG_N] = cmp_tmp[TOP_BIT];
708             c[`VMICRO16_SFLAG_Z] = (cmp_tmp == 0);
709             c[`VMICRO16_SFLAG_C] = 0; //cmp_tmp[TOP_BIT+1]; // not used
710
711             // Overflow flag
712             // https://stackoverflow.com/questions/30957188/
713             // https://github.com/bendl/prco304/blob/master/prco_core/rtl/prco_alu.v#L50
714             case(cmp_tmp[TOP_BIT+1:TOP_BIT])
715                 2'b01: c[`VMICRO16_SFLAG_V] = 1;
716                 2'b10: c[`VMICRO16_SFLAG_V] = 1;
717                 default: c[`VMICRO16_SFLAG_V] = 0;
718             endcase
719
720             $display($time, "\tC%02h: ALU CMP: %h %h = %h = %b", CORE_ID, a, b, cmp_tmp, c[3:0]);
721         end
722
723         `VMICRO16_ALU_SETC: c = { {15{1'b0}}, r_setc };
724
725         // TODO: Parameterise
726         default: begin
727             $display($time, "\tALU: unknown op: %h", op);
728             c = 0;
729             cmp_tmp = 0;
730         end
731     endcase
732 end
733

```



```

734     branch setc_check (
735         .flags      (flags),
736         .cond       (b[7:0]),
737         .en         (r_setc)
738     );
739 endmodule
740
741 // flags = 4 bit r_cmp_flags register
742 // cond = 8 bit VMICRO16_OP_BR_? value. See vmicro16_isa.v
743 module branch (
744     input [3:0] flags,
745     input [7:0] cond,
746     output reg en
747 );
748     always @(*)
749         case (cond)
750             `VMICRO16_OP_BR_U: en = 1;
751             `VMICRO16_OP_BR_E: en = (flags[`VMICRO16_SFLAG_Z] == 1);
752             `VMICRO16_OP_BR_NE: en = (flags[`VMICRO16_SFLAG_Z] == 0);
753             `VMICRO16_OP_BR_G: en = (flags[`VMICRO16_SFLAG_Z] == 0) &&
754                 (flags[`VMICRO16_SFLAG_N] == flags[`VMICRO16_SFLAG_V]);
755             `VMICRO16_OP_BR_L: en = (flags[`VMICRO16_SFLAG_Z] != flags[`VMICRO16_SFLAG_N]);
756             `VMICRO16_OP_BR_GE: en = (flags[`VMICRO16_SFLAG_Z] == flags[`VMICRO16_SFLAG_N]);
757             `VMICRO16_OP_BR_LE: en = (flags[`VMICRO16_SFLAG_Z] == 1) ||
758                 (flags[`VMICRO16_SFLAG_N] != flags[`VMICRO16_SFLAG_V]);
759             default: en = 0;
760         endcase
761 endmodule
762
763
764 module vmicro16_core # (
765     parameter DATA_WIDTH      = 16,
766     parameter MEM_INSTR_DEPTH = 64,
767     parameter MEM_SCRATCH_DEPTH = 64,
768     parameter MEM_WIDTH       = 16,
769     parameter CORE_ID         = 3'h0
770 ) (
771     input      clk,
772     input      reset,
773
774     output [7:0] dbug,
775
776     output      halt,
777
778     // interrupt sources
779     input  [`DEF_NUM_INT-1:0] ints,
780     input  [`DEF_NUM_INT*DATA_WIDTH-1:0] ints_data,
781     output [`DEF_NUM_INT-1:0] ints_ack,
782
783     // APB master to slave interface (apb_intercon)
784     output  [`APB_WIDTH-1:0] w_PADDR,
785     output  w_PWRITE,
786     output  w_PSELx,
787     output  w_PENABLE,
788     output [DATA_WIDTH-1:0] w_PWDATA,
789     input  [DATA_WIDTH-1:0] w_PRDATA,
790     input  w_PREADY
791
792     `ifndef DEF_CORE_HAS_INSTR_MEM
793     , // APB master interface to slave instruction memory
794     output reg  [`APB_WIDTH-1:0] w2_PADDR,
795     output reg  w2_PWRITE,
796     output reg  w2_PSELx,
797     output reg  w2_PENABLE,
798     output reg  [DATA_WIDTH-1:0] w2_PWDATA,
799     input  [DATA_WIDTH-1:0] w2_PRDATA,
800     input  w2_PREADY
801     `endif
802 );
803
804     localparam STATE_IF = 0;
805     localparam STATE_R1 = 1;
806     localparam STATE_R2 = 2;
807     localparam STATE_ME = 3;
808     localparam STATE_WB = 4;
809     localparam STATE_FE = 5;
810     localparam STATE_IDLE = 6;
811     localparam STATE_HALT = 7;
812     reg [2:0] r_state = STATE_IF;
813
814     reg [DATA_WIDTH-1:0] r_pc      = 16'h0000;
815     reg [DATA_WIDTH-1:0] r_pc_saved = 16'h0000;
816     reg [DATA_WIDTH-1:0] r_instr  = 16'h0000;
817     wire [DATA_WIDTH-1:0] w_mem_instr_out;
818     wire w_halt;
819
820     assign dbug = {7'h00, w_halt};
821     assign halt = w_halt;
822
823

```



```

824 wire [4:0] r_instr_opcode;
825 wire [4:0] r_instr_alu_op;
826 wire [2:0] r_instr_rsd;
827 wire [2:0] r_instr_rsa;
828 reg [DATA_WIDTH-1:0] r_instr_rdd = 0;
829 reg [DATA_WIDTH-1:0] r_instr_rda = 0;
830 wire [3:0] r_instr_imm4;
831 wire [7:0] r_instr_imm8;
832 wire [4:0] r_instr_simm5;
833 wire r_instr_has_imm4;
834 wire r_instr_has_imm8;
835 wire r_instr_has_we;
836 wire r_instr_has_br;
837 wire r_instr_has_cmp;
838 wire r_instr_has_mem;
839 wire r_instr_has_mem_we;
840 wire r_instr_halt;
841 wire r_instr_has_lwex;
842 wire r_instr_has_swex;
843
844 wire [DATA_WIDTH-1:0] r_alu_out;
845
846 wire [DATA_WIDTH-1:0] r_mem_scratch_addr = $signed(r_alu_out) + $signed(r_instr_simm5);
847 wire [DATA_WIDTH-1:0] r_mem_scratch_in = r_instr_rdd;
848 wire [DATA_WIDTH-1:0] r_mem_scratch_out;
849 wire r_mem_scratch_we = r_instr_has_mem_we && (r_state == STATE_ME);
850 reg r_mem_scratch_req = 0;
851 wire r_mem_scratch_busy;
852
853 reg [2:0] r_reg_rs1 = 0;
854 wire [DATA_WIDTH-1:0] r_reg_rd1_s;
855 wire [DATA_WIDTH-1:0] r_reg_rd1_i;
856 wire [DATA_WIDTH-1:0] r_reg_rd1 = regs_use_int ? r_reg_rd1_i : r_reg_rd1_s;
857 //wire [15:0] r_reg_rd2;
858 wire [DATA_WIDTH-1:0] r_reg_wd = (r_instr_has_mem) ? r_mem_scratch_out : r_alu_out;
859 wire r_reg_we = r_instr_has_we && (r_state == STATE_WB);
860
861 // branching
862 wire w_intr;
863 wire w_branch_en;
864 wire w_branching = r_instr_has_br && w_branch_en;
865 reg [3:0] r_cmp_flags = 4'h00; // N, Z, C, V
866
867 always @(r_cmp_flags)
868     $display($time, "\tC%02h: \tALU CMP: %b", CORE_ID, r_cmp_flags);
869
870 // 2 cycle register fetch
871 always @(*) begin
872     r_reg_rs1 = 0;
873     if (r_state == STATE_R1)
874         r_reg_rs1 = r_instr_rsd;
875     else if (r_state == STATE_R2)
876         r_reg_rs1 = r_instr_rsa;
877     else
878         r_reg_rs1 = 3'h0;
879 end
880
881 reg regs_use_int = 0;
882 `ifdef DEF_ENABLE_INT
883 wire [`DEF_NUM_INT*`DATA_WIDTH-1:0] ints_vector;
884 wire [`DEF_NUM_INT-1:0] ints_mask;
885 wire has_int = ints & ints_mask;
886
887 reg int_pending = 0;
888 reg int_pending_ack = 0;
889 always @(posedge clk)
890     if (int_pending_ack)
891         // We've now branched to the isr
892         int_pending <= 0;
893     else if (has_int)
894         // Notify fsm to switch to the ints_vector at the last stage
895         int_pending <= 1;
896     else if (w_intr)
897         // Return to Interrupt instruction called,
898         // so we've finished with the interrupt
899         int_pending <= 0;
900 `endif
901
902 // Next program counter logic
903 reg [`DATA_WIDTH-1:0] next_pc = 0;
904 always @(posedge clk)
905     if (reset)
906         r_pc <= 0;
907     else if (r_state == STATE_WB) begin
908         `ifdef DEF_ENABLE_INT
909         if (int_pending) begin
910             $display($time, "\tC%02h: Jumping to ISR: %h",
911                     CORE_ID,
912                     ints_vector[0 +: `DATA_WIDTH]);
913             // TODO: check bounds
914             // Save state

```

```

914         r_pc_saved      <= r_pc + 1;
915         regs_use_int     <= 1;
916         int_pending_ack  <= 1;
917         // Jump to ISR
918         r_pc             <= ints_vector[0 +: `DATA_WIDTH];
919     end else if (w_intr) begin
920         $display($time, "\tC%02h: Returning from ISR: %h",
921             CORE_ID, r_pc_saved);
922
923         // Restore state
924         r_pc             <= r_pc_saved;
925         regs_use_int     <= 0;
926         int_pending_ack  <= 0;
927     end else
928     `endif
929     if (w_branching) begin
930         $display($time, "\tC%02h: branching to %h", CORE_ID, r_instr_rdd);
931         r_pc             <= r_instr_rdd;
932
933         `ifdef DEF_ENABLE_INT
934             int_pending_ack <= 0;
935         `endif
936     end else if (r_pc < (MEM_INSTR_DEPTH-1)) begin
937         // normal increment
938         r_pc <= pc + 1
939         r_pc             <= r_pc + 1;
940
941         `ifdef DEF_ENABLE_INT
942             int_pending_ack <= 0;
943         `endif
944     end
945 end // end r_state == STATE_WB
946 else if (r_state == STATE_HALT) begin
947     `ifdef DEF_ENABLE_INT
948         // Only an interrupt can return from halt
949         // duplicate code form STATE_ME!
950     if (int_pending) begin
951         $display($time, "\tC%02h: Jumping to ISR: %h", CORE_ID, ints_vector[0 +: `DATA_WIDTH]);
952         // TODO: check bounds
953         // Save state
954         r_pc_saved      <= r_pc; // + 1; HALT = stay with same PC
955         regs_use_int     <= 1;
956         int_pending_ack  <= 1;
957         // Jump to ISR
958         r_pc             <= ints_vector[0 +: `DATA_WIDTH];
959     end else if (w_intr) begin
960         $display($time, "\tC%02h: Returning from ISR: %h", CORE_ID, r_pc_saved);
961         r_pc             <= r_pc_saved;
962         regs_use_int     <= 0;
963         int_pending_ack  <= 0;
964     end
965     `endif
966 end
967
968 `ifndef DEF_CORE_HAS_INSTR_MEM
969     initial w2_PSELx     = 0;
970     initial w2_PENABLE   = 0;
971     initial w2_PADDR     = 0;
972 `endif
973
974 // cpu state machine
975 always @(posedge clk)
976     if (reset) begin
977         r_state          <= STATE_IF;
978         r_instr           <= 0;
979         r_mem_scratch_req <= 0;
980         r_instr_rdd       <= 0;
981         r_instr_rda       <= 0;
982     end
983     else begin
984
985     `ifdef DEF_CORE_HAS_INSTR_MEM
986         if (r_state == STATE_IF) begin
987             r_instr <= w_mem_instr_out;
988
989             $display("");
990             $display($time, "\tC%02h: PC: %h", CORE_ID, r_pc);
991             $display($time, "\tC%02h: INSTR: %h", CORE_ID, w_mem_instr_out);
992
993             r_state <= STATE_R1;
994         end
995     `else
996         // wait for global instruction rom to give us our instruction
997         if (r_state == STATE_IF) begin
998             // wait for ready signal
999             if (!w2_PREADY) begin
1000                 w2_PSELx <= 1;
1001                 w2_PWRITE <= 0;
1002                 w2_PENABLE <= 1;
1003                 w2_PWDATA <= 0;

```

```

1004         w2_PADDR    <= r_pc;
1005     end else begin
1006         w2_PSELx    <= 0;
1007         w2_PWRITE    <= 0;
1008         w2_PENABLE   <= 0;
1009         w2_PWDATA    <= 0;
1010
1011         r_instr <= w2_PRDATA;
1012
1013         $display("");
1014         $display($time, "\tC%02h: PC: %h", CORE_ID, r_pc);
1015         $display($time, "\tC%02h: INSTR: %h", CORE_ID, w2_PRDATA);
1016
1017         r_state <= STATE_R1;
1018     end
1019 end
1020 `endif
1021
1022 else if (r_state == STATE_R1) begin
1023     if (w_halt) begin
1024         $display("");
1025         $display($time, "\tC%02h: PC: %h HALT", CORE_ID, r_pc);
1026         r_state <= STATE_HALT;
1027     end else begin
1028         // primary operand
1029         r_instr_rdd <= r_reg_rd1;
1030         r_state <= STATE_R2;
1031     end
1032 end
1033
1034 else if (r_state == STATE_R2) begin
1035     // Choose secondary operand (register or immediate)
1036     if (r_instr_has_imm8) r_instr_rda <= r_instr_imm8;
1037     else if (r_instr_has_imm4) r_instr_rda <= r_reg_rd1 + r_instr_imm4;
1038     else r_instr_rda <= r_reg_rd1;
1039
1040     if (r_instr_has_mem) begin
1041         r_state <= STATE_ME;
1042         // Pulse req
1043         r_mem_scratch_req <= 1;
1044     end else
1045         r_state <= STATE_WB;
1046 end
1047
1048 else if (r_state == STATE_ME) begin
1049     // Pulse req
1050     r_mem_scratch_req <= 0;
1051     // Wait for MMU to finish
1052     if (!r_mem_scratch_busy)
1053         r_state <= STATE_WB;
1054 end
1055
1056 else if (r_state == STATE_WB) begin
1057     if (r_instr_has_cmp) begin
1058         $display($time, "\tC%02h: CMP: %h", CORE_ID, r_alu_out[3:0]);
1059         r_cmp_flags <= r_alu_out[3:0];
1060     end
1061
1062     r_state <= STATE_FE;
1063 end
1064
1065 else if (r_state == STATE_FE)
1066     r_state <= STATE_IF;
1067
1068 else if (r_state == STATE_HALT) begin
1069     `ifndef DEF_ENABLE_INT
1070         if (int_pending) begin
1071             r_state <= STATE_FE;
1072         end
1073     end
1074 end
1075 `endif
1076
1077 end
1078
1079 `ifndef DEF_CORE_HAS_INSTR_MEM
1080 // Instruction ROM
1081 (* rom_style = "distributed" *)
1082 vmicro16_bram # (
1083     .MEM_WIDTH      (DATA_WIDTH),
1084     .MEM_DEPTH      (MEM_INSTR_DEPTH),
1085     .CORE_ID        (CORE_ID),
1086     .USE_INITS      (1),
1087     .NAME            ("INSTR_MEM")
1088 ) mem_instr (
1089     .clk             (clk),
1090     .reset           (reset),
1091     // port 1
1092     .mem_addr        (r_pc),
1093     .mem_in          (0),
1094     .mem_we          (1'b0), // ROM
1095     .mem_out          (w_mem_instr_out)
1096 );
1097 `endif
1098
1099 // MMU

```

```

1094 vmicro16_core_mmu # (
1095     .MEM_WIDTH      (DATA_WIDTH),
1096     .MEM_DEPTH      (MEM_SCRATCH_DEPTH),
1097     .CORE_ID        (CORE_ID)
1098 ) mmu (
1099     .clk             (clk),
1100     .reset           (reset),
1101     .req             (r_mem_scratch_req),
1102     .busy            (r_mem_scratch_busy),
1103     // interrupts
1104     .ints_vector     (ints_vector),
1105     .ints_mask       (ints_mask),
1106     // port 1
1107     .mmu_addr        (r_mem_scratch_addr),
1108     .mmu_in          (r_mem_scratch_in),
1109     .mmu_we          (r_mem_scratch_we),
1110     .mmu_lwex        (r_instr_has_lwex),
1111     .mmu_swex        (r_instr_has_swex),
1112     .mmu_out         (r_mem_scratch_out),
1113     // APB master to slave
1114     .M_PADDR         (w_PADDR),
1115     .M_PWRITE        (w_PWRITE),
1116     .M_PSELx         (w_PSELx),
1117     .M_PENABLE       (w_PENABLE),
1118     .M_PWDATA        (w_PWDATA),
1119     .M_PRDATA        (w_PRDATA),
1120     .M_PREADY        (w_PREADY)
1121 );
1122
1123 // Instruction decoder
1124 vmicro16_dec dec (
1125     // input
1126     .instr           (r_instr),
1127     // output async
1128     .opcode          (),
1129     .rd              (r_instr_rsd),
1130     .ra              (r_instr_rsa),
1131     .imm4            (r_instr_imm4),
1132     .imm8            (r_instr_imm8),
1133     .imm12           (),
1134     .simm5           (r_instr_simm5),
1135     .alu_op          (r_instr_alu_op),
1136     .has_imm4        (r_instr_has_imm4),
1137     .has_imm8        (r_instr_has_imm8),
1138     .has_we          (r_instr_has_we),
1139     .has_br          (r_instr_has_br),
1140     .has_cmp         (r_instr_has_cmp),
1141     .has_mem         (r_instr_has_mem),
1142     .has_mem_we      (r_instr_has_mem_we),
1143     .halt            (w_halt),
1144     .intr            (w_intr),
1145     .has_lwex        (r_instr_has_lwex),
1146     .has_swex        (r_instr_has_swex)
1147 );
1148
1149 // Software registers
1150 vmicro16_regs # (
1151     .CORE_ID         (CORE_ID),
1152     .CELL_WIDTH      (`DATA_WIDTH)
1153 ) regs (
1154     .clk             (clk),
1155     .reset           (reset),
1156     // async port 0
1157     .rs1             (r_reg_rs1),
1158     .rd1             (r_reg_rd1_s),
1159     // async port 1
1160     //.rs2            (),
1161     //.rd2            (),
1162     // write port
1163     .we              (r_reg_we && ~regs_use_int),
1164     .ws1            (r_instr_rsd),
1165     .wd              (r_reg_wd)
1166 );
1167
1168 // Interrupt replacement registers
1169 `ifndef DEF_ENABLE_INT
1170 vmicro16_regs # (
1171     .CORE_ID         (CORE_ID),
1172     .CELL_WIDTH      (`DATA_WIDTH),
1173     .DEBUG_NAME      ("REGSINT")
1174 ) regs_intr (
1175     .clk             (clk),
1176     .reset           (reset),
1177     // async port 0
1178     .rs1             (r_reg_rs1),
1179     .rd1             (r_reg_rd1_i),
1180     // async port 1
1181     //.rs2            (),
1182     //.rd2            (),

```

```

1183         // write port
1184         .we      (r_reg_we && regs_use_int),
1185         .ws1     (r_instr_rsd),
1186         .wd      (r_reg_wd)
1187     );
1188 `endif
1189
1190     // ALU
1191     vmicro16_alu # (
1192         .CORE_ID(CORE_ID)
1193     ) alu (
1194         .op      (r_instr_alu_op),
1195         .a       (r_instr_rdd),
1196         .b       (r_instr_rda),
1197         .flags   (r_cmp_flags),
1198         // async output
1199         .c       (r_alu_out)
1200     );
1201
1202     branch branch_check (
1203         .flags   (r_cmp_flags),
1204         .cond    (r_instr_imm8),
1205         .en      (w_branch_en)
1206     );
1207
1208 endmodule

```

E.2 Peripheral Code Listing

Various memory-mapped APB peripherals, such as GPIO, UART, timers, and memory.

```

1  // Vmicro16 peripheral modules
2
3  `include "vmicro16_soc_config.v"
4  `include "formal.v"
5
6  // Simple watchdog peripheral
7  module vmicro16_watchdog_apb # (
8      parameter BUS_WIDTH = 16,
9      parameter NAME      = "WD",
10     parameter CLK_HZ     = 50_000_000
11 ) (
12     input clk,
13     input reset,
14
15     // APB Slave to master interface
16     input [0:0]      S_PADDR, // not used (optimised out)
17     input            S_PWRITE,
18     input            S_PSELx,
19     input            S_PENABLE,
20     input [0:0]      S_PWDATA,
21
22     // prdata not used
23     output [0:0]     S_PRDATA,
24     output           S_PREADY,
25
26     // watchdog reset, active high
27     output reg       wdreset
28 );
29 //assign S_PRDATA = (S_PSELx & S_PENABLE) ? gpio : 16'h0000;
30 assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
31 wire we        = (S_PSELx & S_PENABLE & S_PWRITE);
32
33 // countdown timer
34 reg [`clog2(CLK_HZ)-1:0] timer = CLK_HZ;
35
36 wire w_wdreset = (timer == 0);
37
38 // infer a register to aid timing
39 initial wdreset = 0;
40 always @(posedge clk)
41     wdreset <= w_wdreset;
42
43 always @(posedge clk)
44     if (we) begin
45         $display($time, "\t\t%s <= RESET", NAME);
46         timer <= CLK_HZ;
47     end else begin
48         timer <= timer - 1;
49     end
50 endmodule
51

```

```

module timer_apb # (
    parameter CLK_HZ = 50_000_000
) (
    input clk,
    input reset,

    input clk_en,

    // 0 16-bit value R/W
    // 1 16-bit control R b0 = start, b1 = reset
    // 2 16-bit prescaler
    input [1:0] S_PADDR,

    input S_PWRITE,
    input S_PSELx,
    input S_PENABLE,
    input [`DATA_WIDTH-1:0] S_PWDATA,

    output reg [`DATA_WIDTH-1:0] S_PRDATA,
    output S_PREADY,

    output out,
    output [`DATA_WIDTH-1:0] int_data
);
//assign S_PRDATA = (S_PSELx & S_PENABLE) ? swex_success ? 16'hFOFO : 16'h0000;
assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
wire en = (S_PSELx & S_PENABLE);
wire we = (en & S_PWRITE);

reg [`DATA_WIDTH-1:0] r_counter = 0;
reg [`DATA_WIDTH-1:0] r_load = 0;
reg [`DATA_WIDTH-1:0] r_pres = 0;
reg [`DATA_WIDTH-1:0] r_ctrl = 0;

localparam CTRL_START = 0;
localparam CTRL_RESET = 1;
localparam CTRL_INT = 2;

localparam ADDR_LOAD = 2'b00;
localparam ADDR_CTRL = 2'b01;
localparam ADDRPres = 2'b10;

always @(*) begin
    S_PRDATA = 0;
    if (en)
        case(S_PADDR)
            ADDR_LOAD: S_PRDATA = r_counter;
            ADDR_CTRL: S_PRDATA = r_ctrl;
            //ADDR_CTRL: S_PRDATA = r_pres;
            default: S_PRDATA = 0;
        endcase
end

// prescaler counts from r_pres to 0, emitting a stb signal
// to enable the r_counter step
reg [`DATA_WIDTH-1:0] r_pres_counter = 0;
wire counter_en = (r_pres_counter == 0);
always @(posedge clk)
    if (r_pres_counter == 0)
        r_pres_counter <= r_pres;
    else
        r_pres_counter <= r_pres_counter - 1;

always @(posedge clk)
    if (we)
        case(S_PADDR)
            // Write to the load register:
            // Set load register
            // Set counter register
            ADDR_LOAD: begin
                r_load <= S_PWDATA;
                r_counter <= S_PWDATA;
                $display($time, "\t\ttimrO: WRITE LOAD: %h", S_PWDATA);
            end
            ADDR_CTRL: begin
                r_ctrl <= S_PWDATA;
                $display($time, "\t\ttimrO: WRITE CTRL: %h", S_PWDATA);
            end
            ADDRPres: begin
                r_pres <= S_PWDATA;
                $display($time, "\t\ttimrO: WRITE PRES: %h", S_PWDATA);
            end
        endcase
    else
        if (r_ctrl[CTRL_START]) begin
            if (r_counter == 0)
                r_counter <= r_load;
            else if(counter_en)
                r_counter <= r_counter -1;
            end else if (r_ctrl[CTRL_RESET])

```

```

142         r_counter <= r_load;
143
144         // generate the output pulse when r_counter == 0
145         // out = (counter reached zero && counter started)
146         assign out = (r_counter == 0) && r_ctrl[CTRL_START]; // && r_ctrl[CTRL_INT];
147         assign int_data = {`DATA_WIDTH{1'b1}};
148     endmodule
149
150
151     // APB wrapped programmable vmicro16_bram
152     module vmicro16_bram_prog_apb # (
153         parameter BUS_WIDTH    = 16,
154         parameter MEM_WIDTH    = 16,
155         parameter MEM_DEPTH    = 64,
156         parameter APB_PADDR    = 0,
157         parameter USE_INITS    = 0,
158         parameter NAME         = "BRAMPROG",
159         parameter CORE_ID      = 0
160     ) (
161         input clk,
162         input reset,
163         // APB Slave to master interface
164         input  [`clog2(MEM_DEPTH)-1:0] S_PADDR,
165         input                               S_PWRITE,
166         input                               S_PSELx,
167         input                               S_PENABLE,
168         input  [BUS_WIDTH-1:0]          S_PWDATA,
169
170         output [BUS_WIDTH-1:0]          S_PRDATA,
171         output                               S_PREADY,
172
173         // interface to program the instruction memory
174         input  [`clog2(`DEF_MEM_INSTR_DEPTH)-1:0] addr,
175         input  [ `DATA_WIDTH-1:0] data,
176         input                               we,
177         input                               prog
178     );
179     wire [MEM_WIDTH-1:0] mem_out;
180
181     assign S_PRDATA = (S_PSELx & S_PENABLE) ? mem_out : 16'h0000;
182     assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
183     wire s_we = (S_PSELx & S_PENABLE & S_PWRITE);
184
185     wire [`clog2(`DEF_MEM_INSTR_DEPTH)-1:0] mem_addr = we ? addr : S_PADDR;
186     wire [ `DATA_WIDTH-1:0] mem_data = we ? data : S_PWDATA;
187     wire mem_we = we | s_we;
188
189     vmicro16_bram # (
190         .MEM_WIDTH (MEM_WIDTH),
191         .MEM_DEPTH (MEM_DEPTH),
192         .NAME      ("BRAMPROG"),
193         .USE_INITS (0),
194         .CORE_ID   (-1)
195     ) bram_apb (
196         .clk      (clk),
197         .reset     (reset),
198
199         .mem_addr  (mem_addr),
200         .mem_in    (mem_data),
201         .mem_we    (mem_we),
202         .mem_out   (mem_out)
203     );
204 endmodule
205
206     // APB wrapped vmicro16_bram
207     module vmicro16_bram_apb # (
208         parameter BUS_WIDTH    = 16,
209         parameter MEM_WIDTH    = 16,
210         parameter MEM_DEPTH    = 64,
211         parameter APB_PADDR    = 0,
212         parameter USE_INITS    = 0,
213         parameter NAME         = "BRAM",
214         parameter CORE_ID      = 0
215     ) (
216         input clk,
217         input reset,
218         // APB Slave to master interface
219         input  [`clog2(MEM_DEPTH)-1:0] S_PADDR,
220         input                               S_PWRITE,
221         input                               S_PSELx,
222         input                               S_PENABLE,
223         input  [BUS_WIDTH-1:0]          S_PWDATA,
224
225         output [BUS_WIDTH-1:0]          S_PRDATA,
226         output                               S_PREADY
227     );
228     wire [MEM_WIDTH-1:0] mem_out;
229
230     assign S_PRDATA = (S_PSELx & S_PENABLE) ? mem_out : 16'h0000;
231     assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;

```

```

232     assign we          = (S_PSELx & S_PENABLE & S_PWRITE);
233
234     always @(*)
235         if (S_PSELx && S_PENABLE)
236             $display($time, "\t\t%s => %h", NAME, mem_out);
237
238     always @(posedge clk)
239         if (we)
240             $display($time, "\t\t%s[%h] <= %h", NAME,
241                 S_PADDR, S_PWDATA);
242
243     vmicro16_bram # (
244         .MEM_WIDTH    (MEM_WIDTH),
245         .MEM_DEPTH    (MEM_DEPTH),
246         .NAME          (NAME),
247         .USE_INITS     (1),
248         .CORE_ID       (-1)
249     ) bram_apb (
250         .clk           (clk),
251         .reset         (reset),
252
253         .mem_addr      (S_PADDR),
254         .mem_in        (S_PWDATA),
255         .mem_we        (we),
256         .mem_out       (mem_out)
257     );
258 endmodule
259
260 // Shared memory with hardware monitor (LWEX/SWEX)
261 module vmicro16_bram_ex_apb # (
262     parameter BUS_WIDTH    = 16,
263     parameter MEM_WIDTH    = 16,
264     parameter MEM_DEPTH    = 64,
265     parameter CORE_ID_BITS = 3,
266     parameter SWEX_SUCCESS = 16'h0000,
267     parameter SWEX_FAIL   = 16'h0001
268 ) (
269     input clk,
270     input reset,
271
272     // |19 |18 |16 |15 |0|
273     // | LWEX | SWEX | 3 bit CORE_ID | S_PADDR |
274     input  [APB_WIDTH-1:0] S_PADDR,
275
276     input S_PWRITE,
277     input S_PSELx,
278     input S_PENABLE,
279     input [MEM_WIDTH-1:0] S_PWDATA,
280
281     output reg [MEM_WIDTH-1:0] S_PRDATA,
282     output S_PREADY
283 );
284 // exclusive flag checks
285 wire [MEM_WIDTH-1:0] mem_out;
286 reg swex_success = 0;
287
288 localparam ADDR_BITS = `clog2(MEM_DEPTH);
289
290 // hack to create a 1 clock delay to S_PREADY
291 // for bram to be ready
292 reg cdelay = 1;
293 always @(posedge clk)
294     if (S_PSELx)
295         cdelay <= 0;
296     else
297         cdelay <= 1;
298
299 //assign S_PRDATA = (S_PSELx & S_PENABLE) ? swex_success ? 16'hFOFO : 16'h0000;
300 assign S_PREADY = (S_PSELx & S_PENABLE & (!cdelay)) ? 1'b1 : 1'b0;
301 assign we       = (S_PSELx & S_PENABLE & S_PWRITE);
302 wire en         = (S_PSELx & S_PENABLE);
303
304 // Similar to:
305 // http://infocenter.arm.com/help/index.jsp?topic=/com.arm.doc.dui0204f/Cihbghef.html
306
307 // mem_wd is the CORE_ID sent in bits [18:16]
308 localparam TOP_BIT_INDEX = `APB_WIDTH - 1;
309 localparam PADDR_CORE_ID_MSB = TOP_BIT_INDEX - 2;
310 localparam PADDR_CORE_ID_LSB = PADDR_CORE_ID_MSB - (CORE_ID_BITS-1);
311
312 // [LWEX, CORE_ID, mem_addr] from S_PADDR
313 wire lwex = S_PADDR[TOP_BIT_INDEX];
314 wire swex = S_PADDR[TOP_BIT_INDEX-1];
315 wire [CORE_ID_BITS-1:0] core_id = S_PADDR[PADDR_CORE_ID_MSB:PADDR_CORE_ID_LSB];
316 // CORE_ID to write to ex_flags register
317 wire [ADDR_BITS-1:0] mem_addr = S_PADDR[ADDR_BITS-1:0];
318
319 wire [CORE_ID_BITS:0] ex_flags_read;
320 wire is_locked = ex_flags_read;
321 wire is_locked_self = is_locked && (core_id == (ex_flags_read-1));

```



```

322
323 // Check exclusive access flags
324 always @(*) begin
325     swex_success = 0;
326     if (en)
327         // bug!
328         if (!swex && !lwex)
329             swex_success = 1;
330         else if (swex)
331             if (is_locked && !is_locked_self)
332                 // someone else has locked it
333                 swex_success = 0;
334             else if (is_locked && is_locked_self)
335                 swex_success = 1;
336     end
337
338 always @(*)
339     if (swex)
340         if (swex_success)
341             S_PRDATA = SWEX_SUCCESS;
342         else
343             S_PRDATA = SWEX_FAIL;
344     else
345         S_PRDATA = mem_out;
346
347 wire reg_we = en && ((lwex && !is_locked)
348                     || (swex && swex_success));
349
350 reg [CORE_ID_BITS:0] reg_wd;
351 always @(*) begin
352     reg_wd = {{CORE_ID_BITS}{1'b0}};
353
354     if (en)
355         // if wanting to lock the addr
356         if (lwex)
357             // and not already locked
358             if (!is_locked) begin
359                 reg_wd = (core_id + 1);
360             end
361         else if (swex)
362             if (is_locked && is_locked_self)
363                 reg_wd = {{CORE_ID_BITS}{1'b0}};
364     end
365
366 // Exclusive flag for each memory cell
367 vmicro16_bram # (
368     .MEM_WIDTH (CORE_ID_BITS + 1),
369     .MEM_DEPTH (MEM_DEPTH),
370     .USE_INITS (0),
371     .NAME ("rexram")
372 ) ram_exflags (
373     .clk (clk),
374     .reset (reset),
375
376     .mem_addr (mem_addr),
377     .mem_in (reg_wd),
378     .mem_we (reg_we),
379     .mem_out (ex_flags_read)
380 );
381
382 always @(*)
383     if (S_PSELx && S_PENABLE)
384         $display($time, "\t\tBRAMex[%h] READ %h\tCORE: %h",
385                 mem_addr, mem_out, S_PADDR[16 +: CORE_ID_BITS]);
386
387 always @(posedge clk)
388     if (we)
389         $display($time, "\t\tBRAMex[%h] WRITE %h\tCORE: %h",
390                 mem_addr, S_PWDATA, S_PADDR[16 +: CORE_ID_BITS]);
391
392 vmicro16_bram # (
393     .MEM_WIDTH (MEM_WIDTH),
394     .MEM_DEPTH (MEM_DEPTH),
395     .USE_INITS (0),
396     .NAME ("BRAMexinst")
397 ) bram_apb (
398     .clk (clk),
399     .reset (reset),
400
401     .mem_addr (mem_addr),
402     .mem_in (S_PWDATA),
403     .mem_we (we && swex_success),
404     .mem_out (mem_out)
405 );
406 endmodule
407
408 // Simple APB memory-mapped register set
409 module vmicro16_regs_apb # (
410     parameter BUS_WIDTH = 16,

```

```

411     parameter DATA_WIDTH      = 16,
412     parameter CELL_DEPTH      = 8,
413     parameter PARAM_DEFAULTS_RO = 0,
414     parameter PARAM_DEFAULTS_R1 = 0
415 ) (
416     input clk,
417     input reset,
418     // APB Slave to master interface
419     input  [`clog2(CELL_DEPTH)-1:0] S_PADDR,
420     input                               S_PWRITE,
421     input                               S_PSELx,
422     input                               S_PENABLE,
423     input  [DATA_WIDTH-1:0]         S_PWDATA,
424
425     output [DATA_WIDTH-1:0]         S_PRDATA,
426     output                               S_PREADY
427 );
428     wire [DATA_WIDTH-1:0] rd1;
429
430     assign S_PRDATA = (S_PSELx & S_PENABLE) ? rd1 : 16'h0000;
431     assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
432     assign reg_we   = (S_PSELx & S_PENABLE & S_PWRITE);
433
434     always @(*)
435         if (reg_we)
436             $display($time, "\t\tREGS-APB[%h] <= %h",
437                     S_PADDR, S_PWDATA);
438
439     always @(*)
440         `rassert(reg_we == (S_PSELx & S_PENABLE & S_PWRITE))
441
442     vmicro16_regs # (
443         .CELL_DEPTH      (CELL_DEPTH),
444         .CELL_WIDTH      (DATA_WIDTH),
445         .PARAM_DEFAULTS_RO (PARAM_DEFAULTS_RO),
446         .PARAM_DEFAULTS_R1 (PARAM_DEFAULTS_R1)
447     ) regs_apb (
448         .clk      (clk),
449         .reset    (reset),
450         // port 1
451         .rs1      (S_PADDR),
452         .rd1      (rd1),
453         .we        (reg_we),
454         .ws1      (S_PADDR),
455         .wd        (S_PWDATA)
456         // port 2 unconnected
457         //.rs2      (),
458         //.rd2      ()
459     );
460     endmodule
461
462     // Simple GPIO write only peripheral
463     module vmicro16_gpio_apb # (
464         parameter BUS_WIDTH      = 16,
465         parameter DATA_WIDTH    = 16,
466         parameter PORTS          = 8,
467         parameter NAME           = "GPIO"
468     ) (
469         input clk,
470         input reset,
471         // APB Slave to master interface
472         input  [0:0]             S_PADDR, // not used (optimised out)
473         input                               S_PWRITE,
474         input                               S_PSELx,
475         input                               S_PENABLE,
476         input  [DATA_WIDTH-1:0]         S_PWDATA,
477
478         output [DATA_WIDTH-1:0]         S_PRDATA,
479         output                               S_PREADY,
480         output reg [PORTS-1:0]         gpio
481     );
482     assign S_PRDATA = (S_PSELx & S_PENABLE) ? gpio : 16'h0000;
483     assign S_PREADY = (S_PSELx & S_PENABLE) ? 1'b1 : 1'b0;
484     assign ports_we = (S_PSELx & S_PENABLE & S_PWRITE);
485
486     always @(posedge clk)
487         if (reset)
488             gpio <= 0;
489         else if (ports_we) begin
490             $display($time, "\t\t%s <= %h", NAME, S_PWDATA[PORTS-1:0]);
491             gpio <= S_PWDATA[PORTS-1:0];
492         end
493     endmodule

```