

Multi-core RISC SoC Design & Implementation

Demonstration Viva

Ben Lancaster

201280376
ELEC5881M - Main Project

July 22, 2019

Quick Links

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

- GitHub repository: <https://github.com/bendl/vmicro16>
- Full Report: https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M_Ben_Lancaster_201280376_Final.pdf
- This presentation: https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M_Ben_Lancaster_201280376_viva.pdf
- About me: <https://bendl.me/>

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

- 1 Introduction
 - Why a project on CPUs?
 - Why Multi-core?
 - Why RISC?
- 2 Top Level Design
 - Overview
 - Memory Map
 - Interconnect
 - ISA
 - Interrupts
- 3 Multi-core Functionality
 - HW/SW Requirements
 - Context Identification
 - Atomics
 - Design Challenges

- 4 Results
 - Summation
 - Per-core Memory VS Global
 - Memory for instructions
- 5 Conclusion
 - Accomplishments
 - Future Improvements
 - Q&A

Main Project

B. Lancaster

Introduction

Why a project on CPUs?

Why Multi-core?

Why RISC?

Top Level Design

Multi-core Functionality

Results

Conclusion

- 1 Introduction
 - Why a project on CPUs?
 - Why Multi-core?
 - Why RISC?

- 2 Top Level Design

- 3 Multi-core Functionality

- 4 Results

- 5 Conclusion

Why a project on CPUs?

Main Project

B. Lancaster

Introduction

Why a project on
CPUs?

Why Multi-core?

Why RISC?

Top Level
Design

Multi-core
Functionality

Results

Conclusion

- **CPUs will be used for the rest of humanity**
1000s of years
- **Understand design constraints and considerations**
Be a better engineer
- **Prepare myself for future employment/work**

Why Multi-core?

Main Project

B. Lancaster

Introduction

Why a project on
CPUs?

Why Multi-core?

Why RISC?

Top Level
Design

Multi-core
Functionality

Results

Conclusion

- **Rate of single-core speed improvements slowing**
Pipelining, register renaming, branch predictions, OoOE, clock speeds
- **Future of computing = parallel**
 - Identifying parallel opportunities
 - Massively parallel (NoC's)
 - Higher throughput

Why RISC?

Main Project

B. Lancaster

Introduction

Why a project on
CPUs?

Why Multi-core?
Why RISC?

Top Level
Design

Multi-core
Functionality

Results

Conclusion

- Easier design & impl
- Smaller = fit more cores on a chip
- FPGA size limitations
- Previous experience + future work
- I'm a RISC purist

1 Introduction

2 Top Level Design

Overview

Memory Map

Interconnect

ISA

Interrupts

3 Multi-core Functionality

4 Results

5 Conclusion

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core
Functionality

Results

Conclusion

Overview

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core
Functionality

Results

Conclusion

What this project produces:

- **System-on-Chip with multi-processor functionality**
Tested on FPGA hardware with 1-96 CPU cores.

Overview

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core
Functionality

Results

Conclusion

What this project produces:

- **System-on-Chip with multi-processor functionality**
Tested on FPGA hardware with 1-96 CPU cores.
- **Custom 16-bit RISC CPU**
With interrupts and its own Instruction Set Architecture (ISA).

Overview

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core
Functionality

Results

Conclusion

What this project produces:

- **System-on-Chip with multi-processor functionality**
Tested on FPGA hardware with 1-96 CPU cores.
- **Custom 16-bit RISC CPU**
With interrupts and its own Instruction Set Architecture (ISA).
- **Software/Assembly compiler**
PRCO304 programming language/Intel assembly syntax.

Overview

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core
Functionality

Results

Conclusion

What this project produces:

- **System-on-Chip with multi-processor functionality**
Tested on FPGA hardware with 1-96 CPU cores.
- **Custom 16-bit RISC CPU**
With interrupts and its own Instruction Set Architecture (ISA).
- **Software/Assembly compiler**
PRCO304 programming language/Intel assembly syntax.
- **Aimed at Design Engineers, not end users**
Project is provided as source code/design files for Design Engineers to customise and implement in hardware themselves.

Conclusion



Memory Map

Main Project

B. Lancaster

Introduction

Top Level Design

Overview

Memory Map

Interconnect


ISA

Interrupts

Multi-core Functionality

Results

Conclusion



- **Shared Memory with Global Monitor**
- Timer with Interrupt
- Per-core Interrupt Vector and Mask
- Shared Register Set
- UART Transceiver
- Multiple GPIO ports
- Per-core scratch memory
- **Per-core Special Registers**
- Customisable by designers

Interconnect

Main Project

B. Lancaster

Introduction

Top Level Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core Functionality

Results

Conclusion

- AMBA APB Bus
- Tristate & Non-tristate (mux) impl
- Originally Wishbone, now APB
- AHB too complex (limited time)
- Various schedulers

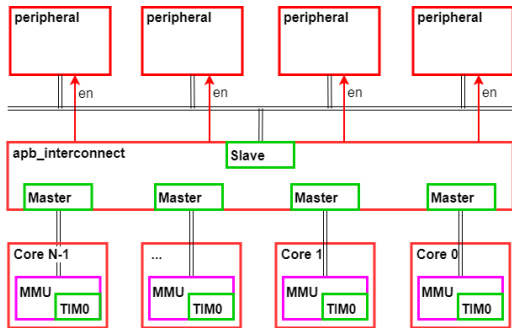


Figure: Vmicro16 interconnect

Interconnect Schematic

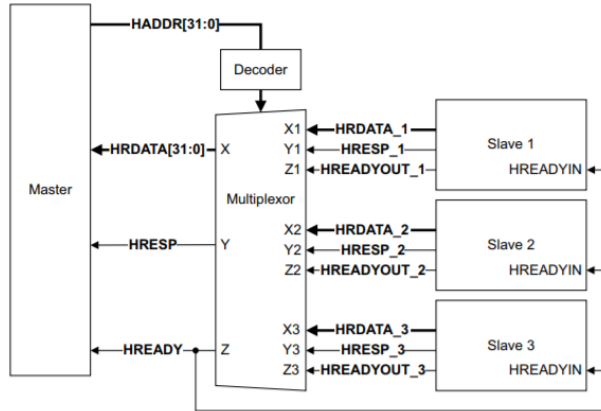


Figure: Source: ARM AHB-Lite Protocol Specification Figure 4-2.

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview
Memory Map
Interconnect

ISA
Interrupts

Multi-core
Functionality

Results

Conclusion

Instruction Set Architecture

Main Project

B. Lancaster

Introduction

Top Level

Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core

Functionality

Results

Conclusion

- + 16-bits, 38 instructions
- + Simple load/store arch
- + (Un)signed instructions
- + Compact
- + Optional hardware multiply instruction
- Only 8 registers
- No good compiler support

	15-11	10-8	7-5	4-0	rd ra simm5
	15-11	10-8	7-0		rd imm8
	15-11	10-0			nop
	15	14:12	11:0		extended immediate
SPCL	00000	11 bits			NOP
SPCL	00000	11h'000			NOP
SPCL	00000	11h'001			HALT
SPCL	00000	11h'002			Return from interrupt
LW	00001	Rd	Ra	s5	Rd <= RAM[Ra+s5]
SW	00010	Rd	Ra	s5	RAM[Ra+s5] <= Rd
BIT	00011	Rd	Ra	s5	bitwise operations
BIT_OR	00011	Rd	Ra	00000	Rd <= Rd Ra
BIT_XOR	00011	Rd	Ra	00001	Rd <= Rd ^ Ra
BIT_AND	00011	Rd	Ra	00010	Rd <= Rd & Ra
BIT_NOT	00011	Rd	Ra	00011	Rd <= ~Ra
BIT_LSHFT	00011	Rd	Ra	00100	Rd <= Rd << Ra
BIT_RSHFT	00011	Rd	Ra	00101	Rd <= Rd >> Ra
MOV	00100	Rd	Ra	X	Rd <= Ra
MOVI	00101	Rd		i8	Rd <= i8
ARITH_U	00110	Rd	Ra	s5	unsigned arithmetic
ARITH_UADD	00110	Rd	Ra	11111	Rd <= uRd + uRa
ARITH_USUB	00110	Rd	Ra	10000	Rd <= uRd - uRa
ARITH_UADDI	00110	Rd	Ra	0AAAA	Rd <= uRd + Ra + AAAA
ARITH_S	00111	Rd	Ra	s5	signed arithmetic
ARITH_SADD	00111	Rd	Ra	11111	Rd <= sRd + sRa
ARITH_SSUB	00111	Rd	Ra	10000	Rd <= sRd - sRa
ARITH_SSUBI	00111	Rd	Ra	0AAAA	Rd <= sRd - sRa + AAAA
BR	01000	Rd		i8	conditional branch
BR_U	01000	Rd		0000 0000	Any
BR_E	01000	Rd		0000 0001	Z=1
BR_NE	01000	Rd		0000 0010	Z=0
BR_G	01000	Rd		0000 0011	Z=0 and S=0
BR_GE	01000	Rd		0000 0100	S=0
BR_L	01000	Rd		0000 0101	S != 0
BR_LE	01000	Rd		0000 0110	Z=1 or (S != 0)
BR_S	01000	Rd		0000 0111	S=1
BR_NS	01000	Rd		0000 1000	S=0
CMP	01001	Rd	Ra	X	SZO <= CMP(Rd, Ra)
SETC	01010	Rd		Imm8	Rd <= (Imm8_f_ SZO) ? 1 : 0
MULT	01011	Rd	Ra	X	Rd <= uRd * uRa
HALT	01100			X	
LWEX	01101	Rd	Ra	s5	Rd <= RAM[Ra+s5]
SWEX	01110	Rd	Ra	s5	RAM[Ra+s5] <= Rd Rd <= 0 1 if success

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview
Memory Map
Interconnect
ISA
Interrupts

Multi-core
Functionality

Results

Conclusion

SW Example

Interrupts

Main Project

B. Lancaster

Introduction

Top Level Design

Overview

Memory Map

Interconnect



ISA

Interrupts

Multi-core Functionality

Results

Conclusion



```
1 entry:
2     // Set interrupt vector at 0x100
3     // Move address of isr0 function to vector[0]
4     movi    r0, isr0
5     // create 0x100 value by left shifting 1 8 bits
6     movi    r1, #0x1
7     movi    r2, #0x8
8     lshft   r1, r2
9     // write isr0 address to vector[0]
10    sw      r0, r1
11
12    // enable all interrupts by writing 0x0f to 0x108
13    movi    r0, #0x0f
14    sw      r0, r1 + #0x8
15    halt    // enter low power idle state
16
17 isr0:
18     movi    r0, #0xff // arbitrary name
19     intr    // do something
20     intr    // return from interrupt
```

Demo: 2 Core LED toggle (GPIO0) with TIMR0 1s interrupt (interrupts_2.s)

Timer Interrupt Example

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview

Memory Map

Interconnect

ISA

Interrupts

Multi-core
Functionality

Results

Conclusion



Figure: TIMR0 1us interrupt with context switching

Timer Peripheral Registers

Main Project

B. Lancaster

Introduction

Top Level
Design

Overview
Memory Map
Interconnect
ISA
Interrupts

Multi-core
Functionality

Results

Conclusion




Figure: $t = 20ns * load * prescaler$

Resolution (32-bit timer): 20ns to 85s.

Examples:

- For 1us: Load = 0x32, Prescaler = 0 ($20ns * 0x32 = 1000ns$)
- For 1s: Load = 0x1000, Prescaler = 0x3000 (demo)
($20ns * 0x1000 * 0x3000 = \text{approx. } 1s$)

1 Introduction

2 Top Level Design

3 Multi-core Functionality
HW/SW Requirements
Context Identification
Atomics
Design Challenges

4 Results

5 Conclusion

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

HW/SW Requirements

Main Project

B. Lancaster

Hardware:

Software:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements

Context Identification

Atomics

Design Challenges

Results

Conclusion

HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Hardware:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)
- **Atomic functions**
(atomic versions of load/store to prevent race conditions)

Software:

HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Hardware:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)
- **Atomic functions**
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory

Software:

HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Hardware:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)
- **Atomic functions**
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

Software:

HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Hardware:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)
- **Atomic functions**
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

Software:

- **Semaphores/Mutexes**
(exclusive memory access)

HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Hardware:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)
- **Atomic functions**
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

Software:

- **Semaphores/Mutexes**
(exclusive memory access)
- **Thread synchronisation**
(memory barriers)

HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Hardware:

- **Bus Arbitration**
(scheduling: priority, rotating, etc.)
- **Atomic functions**
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

Software:

- **Semaphores/Mutexes**
(exclusive memory access)
- **Thread synchronisation**
(memory barriers)
- **Context identification**
What core am I?
How many cores?
How much memory?

Context Identification

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements

Context Identification

Atomics

Design Challenges

Results

Conclusion




Figure: Special Registers 0x0080 to 0x008F

```
entry:
    // get core idx 0x80 in r7
    movi    r7, #0x80
    lw      r7, r7

    // Branch away if not core 0
    cmp     r7, r0
    movi    r0, exit
    br      r0, BR_NE

    // Core 0 only instructions
    nop
    nop
    nop

exit:
    halt
```

Atomic Instructions

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification

Atomics

Design Challenges

Results

Conclusion

- Enables semaphores, mutexes, memory barriers
- Prevent race conditions between threads/cores
- LW[EX] and SW[EX]
- Implementation in next slide

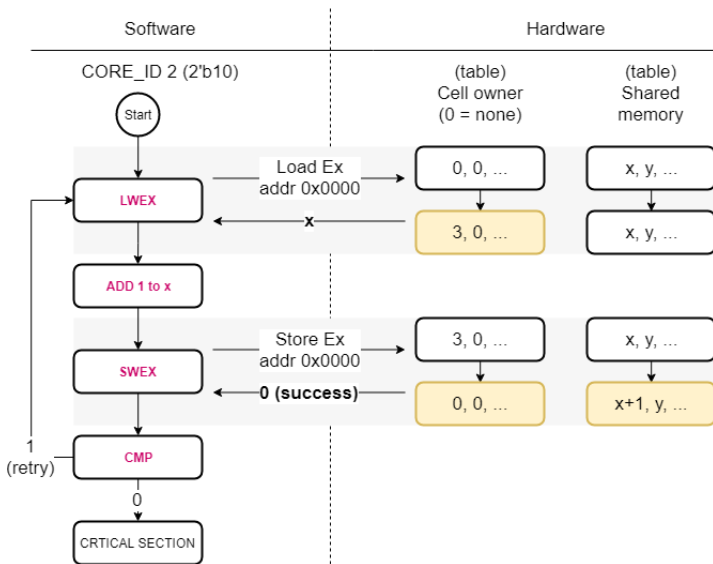
```
try_inc:
    // load and lock
    // (if not already locked)
    lwex    r0, r1
    // do something
    // (i.e. add 1 (semaphore))
    addi    r0, #0x01
    // attempt store
    swex    r0, r1

    // check success (== 0)
    cmp     r0, r3

    // if not equal (NE), retry
    movi    r4, try_inc
    br      r4, BR_NE

critical:
    // r0 is latest value
```

Exclusive Access Flow Chart



Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification

Atomics
Design Challenges

Results

Conclusion

HW - How do I know which core this lwex/swex is from?

Main Project

B. Lancaster

Introduction

Top Level Design

Multi-core Functionality

HW/SW Requirements


Context Identification

Atomics


Design Challenges

Results

Conclusion



The Core Idx is sent with each MMU request to the shared bus.



PADDR*NUMCORES-1:0 interconnect input.

Exclusive Access

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

```
mutex_claim:
    // load and lock
    // (if not already locked)
    lwex    r0, r1
    // do something
    // (i.e. add 1 (semaphore))
    addi    r0, #0x01
    // attempt store
    swex    r0, r1

    // check success (== 0)
    cmp     r0, r3

    // if not equal (NE), retry
    movi    r4, mutex_claim
    br      r4, BR_NE

critical:
    nop
```

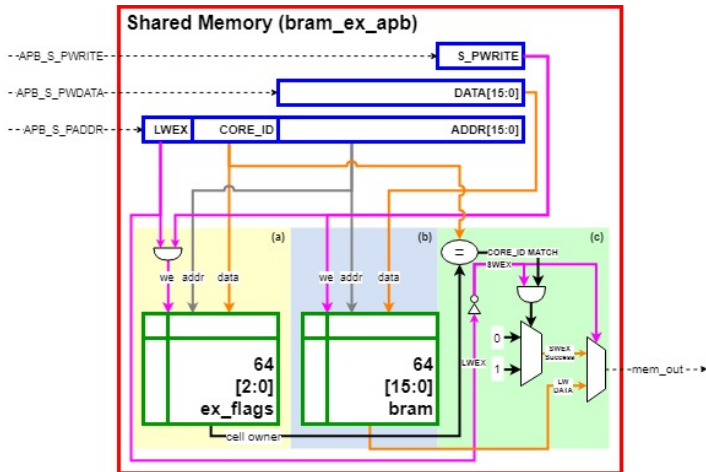


Figure: HW impl

Demo: 8 core number summation (sum.s)

Design Challenges

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

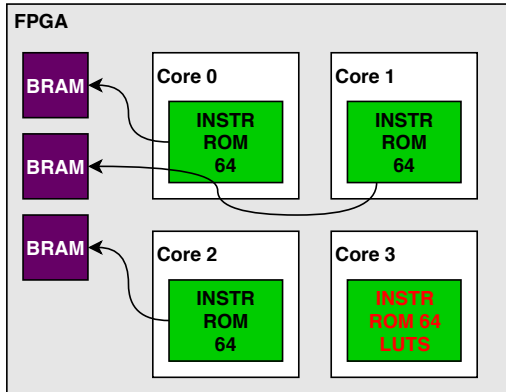
Results

Conclusion

Memory Limitations

Each core has it's own instruction memory

- + Fast fetching and branching
- Requires a dedicated BRAM (FPGA) per core
- Limited BRAM blocks available
 - Low consumption (DEF_MEM_INSTR_DEPTH)
 - More cores = some implemented in LUT/regs (distributed)
- **Reduces maximum core count**




Design Challenges

Memory Limitations - Solution

Global instruction ROM

- + Reduce duplicate ROMs
- + Single BRAM requirement (expandable)
- + Smaller core size
- **Slow access times** (exponential)
 - Requires another interconnect/arbiter/scheduler = logic
- + **Increases maximum core count**



Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

HW/SW
Requirements
Context Identification
Atomics
Design Challenges

Results

Conclusion

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Summation
Per-core Memory VS
Global Memory for
instructions

Conclusion

- 1 Introduction
- 2 Top Level Design
- 3 Multi-core Functionality
- 4 Results**
 - Summation
 - Per-core Memory VS Global Memory for instructions
- 5 Conclusion

Summation - Multi-core vs Single-Core

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Summation

Per-core Memory VS
Global Memory for
instructions

Conclusion

- Each core has low work load
 - Sum subset of numbers in for loop
- Ideal scenario for parallelism
 - Highly parallelisable
 - Few inter-thread dependencies

Summation - Multi-core vs Single-Core

240 samples (@30 cores = 8 samples per core)

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality


Results

Summation

Per-core Memory VS
Global Memory for
instructions

Conclusion


Algorithm Time vs Core Count (N = 240)



Multi-core vs Single-Core for Summation

240 samples (@30 cores = 8 samples per core)

Algorithm Time vs Core Count (N = 240)



Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Summation
Per-core Memory VS
Global Memory for
instructions

Conclusion

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

Accomplishments
Future Improvements
Q&A

- 1 Introduction
- 2 Top Level Design
- 3 Multi-core Functionality
- 4 Results
- 5 Conclusion
 - Accomplishments
 - Future Improvements
 - Q&A

Accomplishments

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

Accomplishments

Future Improvements
Q&A

- **Near complete System-on-Chip design with various peripherals**
Timers, GPIO, UART, Registers, Memory
- **Common multi-thread/core synchronisation primitives**
Semaphores, Mutexes, Memory Barriers, Atomic Instructions
- **AMBA APB bus interface with Global Monitor**
Timers, GPIO, UART, Registers, Memory
- **Working shared bus arbitration**
Schedules access to shared resources
- **Working FPGA implementation for a 96 core design**
Nearly fills Cyclone V FPGA on the DE1-SoC
- **Interrupts with hardware context-switching**
Low latency to react to interrupt
- **Acknowledges design limitations and attempts to overcome**
LUT resources, block memories, power and temperature requirements

Future Improvements

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

Accomplishments
Future Improvements
Q&A

- **Working Global Reset**

Global resets are expensive (LUT resources)

Resetting block memories is not trivial

- **On-chip Programming**

Use the UART0 receiver to program each cores flash memory

- **Per-core gating/enabling**

Improve power efficiency for ASIC implementation by disabling cores at run-time via software.

- **Improve memory bottleneck**

Each core requires it's own memory - reduce by multiplexing access to a single large memory.

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

Accomplishments
Future Improvements
Q&A

Q&A

Info

Main Project

B. Lancaster

Introduction

Top Level
Design

Multi-core
Functionality

Results

Conclusion

Accomplishments
Future Improvements
Q&A

- GitHub repository: <https://github.com/bendl/vmicro16>
- Full Report: https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M_Ben_Lancaster_201280376_Final.pdf
- Presentation tools:
 - Latex Beamer
 - `\usecolortheme{orchid}`
 - `\useoutertheme[hideothersubsections]{sidebar}`