

# **Multi-core RISC SoC Design & Implementation**

Demonstration Viva

Ben Lancaster

201280376  
ELEC5881M - Main Project

July 22, 2019

# Quick Links

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

- GitHub repository: <https://github.com/bendl/vmicro16>
- Full Report: [https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M\\_Ben\\_Lancaster\\_201280376\\_Final.pdf](https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M_Ben_Lancaster_201280376_Final.pdf)
- This presentation: [https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M\\_Ben\\_Lancaster\\_201280376\\_viva.pdf](https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M_Ben_Lancaster_201280376_viva.pdf)
- About me: <https://bendl.me/>

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

## 1 Introduction

Why Multi-core?

Why RISC?

## 2 Top Level Design

Overview

Memory Map

Interconnect

Interrupts

## 3 Multi-core Functionality

HW/SW Requirements

Context Identification

Atomics

## 4 Results

Results 1

## 5 Conclusion

Accomplishments

Future Improvements

Q&A

Main Project

B. Lancaster

Introduction

Why Multi-core?

Why RISC?

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

# 1 Introduction

## Why Multi-core?

## Why RISC?

# 2 Top Level Design

# 3 Multi-core Functionality

# 4 Results

# 5 Conclusion

# Why a project on CPUs?

Main Project

B. Lancaster

Introduction

Why Multi-core?

Why RISC?

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

- **CPUs will be used for the rest of humanity**  
1000s of years
- **Understand constrains and considerations**  
1000s of years
- **Prepare myself for future employment/work**

# Why Multi-core?

Main Project

B. Lancaster

Introduction

Why Multi-core?

Why RISC?

Top Level

Design

Multi-core

Functionality

Results

Conclusion

- **Rate of single-core speed improvements slowing**  
Pipelining, register renaming, branch predictions, clock speeds
- **Future of computing = parallel**
  - Identifying parallel opportunities
  - Massively parallel (NoC's)
  - Higher throughput

# Why RISC?

Main Project

B. Lancaster

Introduction

Why Multi-core?

Why RISC?

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

- FPGA size limitations
- Easier design & impl
- Smaller = fit more cores on a chip
- Previous experience + future work
- RISC purist

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview  
Memory Map  
Interconnect  
Interrupts

Multi-core  
Functionality

Results

Conclusion

1 Introduction

2 Top Level Design  
Overview  
Memory Map  
Interconnect  
Interrupts

3 Multi-core Functionality

4 Results

5 Conclusion



# Overview

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

Interconnect

Interrupts

Multi-core  
Functionality

Results

Conclusion

## What this project produces:

- **System-on-Chip with multi-processor functionality**  
Tested on FPGA hardware with 1-96 CPU cores.

# Overview

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

Interconnect

Interrupts

Multi-core  
Functionality

Results

Conclusion

## What this project produces:

- **System-on-Chip with multi-processor functionality**  
Tested on FPGA hardware with 1-96 CPU cores.
- **Custom 16-bit RISC CPU**  
With interrupts and its own Instruction Set Architecture (ISA).

# Overview

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

Interconnect

Interrupts

Multi-core  
Functionality

Results

Conclusion

## What this project produces:

- **System-on-Chip with multi-processor functionality**  
Tested on FPGA hardware with 1-96 CPU cores.
- **Custom 16-bit RISC CPU**  
With interrupts and its own Instruction Set Architecture (ISA).
- **Software/Assembly compiler**  
PRCO304 programming language/Intel assembly syntax.

# Overview

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

Interconnect

Interrupts

Multi-core  
Functionality

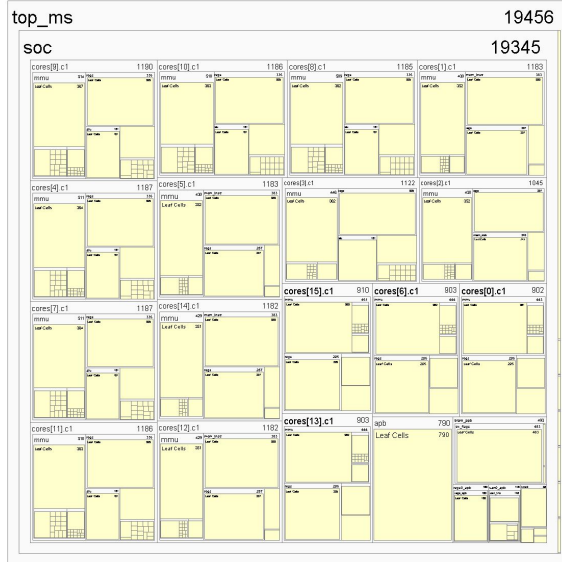
Results

Conclusion

## What this project produces:

- **System-on-Chip with multi-processor functionality**  
Tested on FPGA hardware with 1-96 CPU cores.
- **Custom 16-bit RISC CPU**  
With interrupts and its own Instruction Set Architecture (ISA).
- **Software/Assembly compiler**  
PRCO304 programming language/Intel assembly syntax.
- **Aimed at Design Engineers, not end users**  
Project is provided as source code/design files for Design Engineers to customise and implement in hardware themselves.

# Top Level Hierarchy



Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

Interconnect

Interrupts

Multi-core  
Functionality

Results

Conclusion

# Memory Map

## Main Project

## B. Lancaster

## Introduction

## Top Level Design

### Overview

### Memory Map

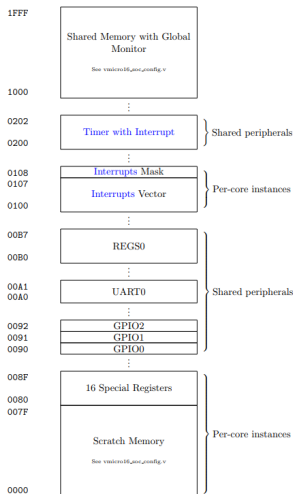
### Interconnect

### Interrupts

## Multi-core Functionality

## Results

## Conclusion



- **Shared Memory with Global Monitor**
- Timer with Interrupt
- Per-core Interrupt Vector and Mask
- Shared Register Set
- UART Transceiver
- Multiple GPIO ports
- Per-core scratch memory
- **Per-core Special Registers**
- Customisable by designers

# Interconnect

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

**Interconnect**

Interrupts

Multi-core  
Functionality

Results

Conclusion

# Interrupts

## Main Project

### B. Lancaster

#### Introduction

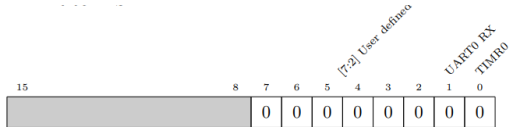
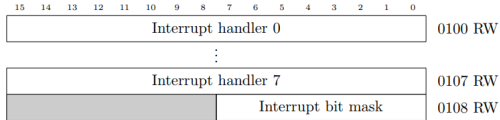
#### Top Level Design

Overview  
Memory Map  
Interconnect  
Interrupts

#### Multi-core Functionality

#### Results

#### Conclusion



```
1 entry:
2     // Set interrupt vector at 0x100
3     // Move address of isr0 function to vector[0]
4     movi    r0, isr0
5     // create 0x100 value by left shifting 1 8 bits
6     movi    r1, #0x1
7     movi    r2, #0x8
8     lshft   r1, r2
9     // write isr0 address to vector[0]
10    sw      r0, r1
11
12    // enable all interrupts by writing 0x0f to 0x108
13    movi    r0, #0x0f
14    sw      r0, r1 + #0x8
15    halt    // enter low power idle state
16
17 isr0:
18     movi    r0, #0xff // arbitrary name
19     intr    // do something
20     // return from interrupt
```

Demo: 2 Core LED toggle (GPIO0) with TIMR0 1s interrupt (interrupts\_2.s)



# Timer Interrupt Example

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview

Memory Map

Interconnect

Interrupts

Multi-core  
Functionality

Results

Conclusion



Figure: TIMR0 1us interrupt with context switching

# Timer Peripheral Registers

Main Project

B. Lancaster

Introduction

Top Level  
Design

Overview  
Memory Map  
Interconnect

Interrupts

Multi-core  
Functionality

Results

Conclusion

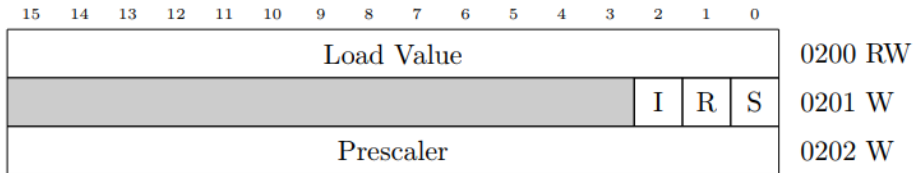


Figure:  $t = 20ns * load * prescaler$

Resolution (32-bit timer): 20ns to 85s.

Examples:

- For 1us: Load = 0x32, Prescaler = 0 ( $20ns * 0x32 = 1000ns$ )
- For 1s: Load = 0x1000, Prescaler = 0x3000 (demo)  
( $20ns * 0x1000 * 0x3000 = \text{approx. } 1s$ )

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

1 Introduction

2 Top Level Design

3 Multi-core Functionality  
HW/SW Requirements  
Context Identification  
Atomics

4 Results

5 Conclusion

# HW/SW Requirements

Main Project

B. Lancaster

Hardware:

Software:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

# HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

Hardware:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)
- **Atomic functions**  
(atomic versions of load/store to prevent race conditions)

Software:

# HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

Hardware:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)
- **Atomic functions**  
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory

Software:

# HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

Hardware:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)
- **Atomic functions**  
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

Software:

# HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

## Hardware:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)
- **Atomic functions**  
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

## Software:

- **Semaphores/Mutexes**  
(exclusive memory access)



# HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

## Hardware:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)
- **Atomic functions**  
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

## Software:

- **Semaphores/Mutexes**  
(exclusive memory access)
- **Thread synchronisation**  
(memory barriers)

# HW/SW Requirements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification  
Atomics

Results

Conclusion

## Hardware:

- **Bus Arbitration**  
(scheduling: priority, rotating, etc.)
- **Atomic functions**  
(atomic versions of load/store to prevent race conditions)
- Per-core instruction memory
- Per-core context-switching for interrupt handling

## Software:

- **Semaphores/Mutexes**  
(exclusive memory access)
- **Thread synchronisation**  
(memory barriers)
- **Context identification**  
What core am I?  
How many cores?  
How much memory?

# Context Identification

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

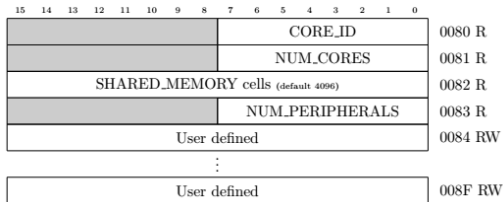
HW/SW  
Requirements

Context Identification

Atomics

Results

Conclusion



**Figure:** Special Registers 0x0080 to 0x008F

```
entry:
    // get core idx 0x80 in r7
    movi    r7, #0x80
    lw      r7, r7

    // Branch away if not core 0
    cmp     r7, r0
    movi    r0, exit
    br      r0, BR_NE

    // Core 0 only instructions
    nop
    nop
    nop

exit:
    halt
```

# Atomic Instructions

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification

Atomics

Results

Conclusion

- Enables semaphores, mutexes, memory barriers
- Prevent race conditions between threads/cores
- LW[EX] and SW[EX]
- Implementation in next slide

Example:

```
try_inc:
    // load and lock (if not already locked)
    lwex    r0, r1
    // do something (i.e. add 1 (semaphore))
    addi    r0, #0x01
    // attempt store
    swex    r0, r1

    // check success (== 0)
    cmp     r0, r3

    // if not equal (NE), retry
    movi    r4, try_inc
    br      r4, BR_NE

critical:
    // r0 is latest value
```

# Exclusive Access Flow Chart

Main Project

B. Lancaster

Introduction

Top Level  
Design

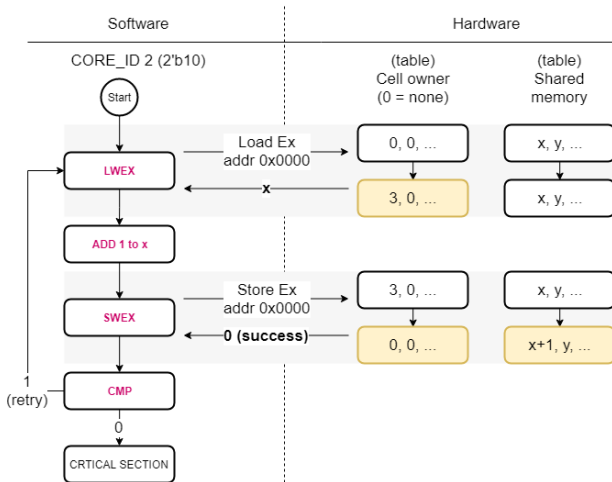
Multi-core  
Functionality

HW/SW  
Requirements  
Context Identification

Atomics

Results

Conclusion



## Main Project

## B. Lancaster

20	19	18	17	16	15												0	
LE	SE	CORE_ID				Address											PADDR[20:0]	
						Write data											PWDATA[15:0]	
						Read Data											PRDATA[15:0]	
																	WE	PWRITE[0:0]
																	EN	PENABLE[0:0]

The Core Idx is sent with each MMU request to the shared bus.

83	62	41	20	0
Core N-1	$\dots$	Core 1	Core 0	

PADDR\*NUMCORES-1:0 interconnect input.

## Introduction

## Top Level Design

## Multi-core Functionality

HW/SW  
Requirements  
Context Identification  
**Atomics**

## Results

## Conclusion

# Exclusive Access

## Main Project

### B. Lancaster

```
mutex_claim:
    // load and lock (if not already locked)
    lwex    r0, r1
    // do something (i.e. add 1 (semaphore))
    addi    r0, #0x01
    // attempt store
    swex    r0, r1

    // check success (== 0)
    cmp     r0, r3

    // if not equal (NE), retry
    movi    r4, mutex_claim
    br      r4, BR_NE
```

critical:

1.

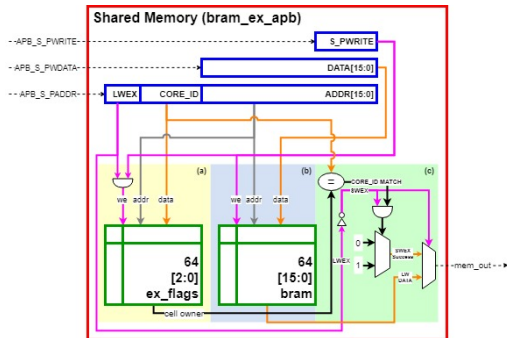


Figure: HW impl

Demo: 8 core number summation (sum.s)

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Results 1

Conclusion

- 1 Introduction
- 2 Top Level Design
- 3 Multi-core Functionality
- 4 Results**  
Results 1
- 5 Conclusion



# Multi-core vs Single-Core for Summation

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Results 1

Conclusion

- Each core has low work load
  - Sum subset of numbers in for loop
- Ideal scenario for parallelism
  - Highly parallelisable
  - Few inter-thread dependencies

Insert graph showing core count vs total time

# Results 1

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Results 1

Conclusion

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

Accomplishments  
Future Improvements  
Q&A

1 Introduction

2 Top Level Design

3 Multi-core Functionality

4 Results

5 Conclusion  
Accomplishments  
Future Improvements  
Q&A

# Accomplishments

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

Accomplishments

Future Improvements  
Q&A

- **Near complete System-on-Chip design with various peripherals**  
Timers, GPIO, UART, Registers, Memory
- **Common multi-thread/core synchronisation primitives**  
Semaphores, Mutexes, Memory Barriers, Atomic Instructions
- **AMBA APB bus interface with Global Monitor**  
Timers, GPIO, UART, Registers, Memory
- **Working shared bus arbitration**  
Schedules access to shared resources
- **Working FPGA implementation for a 96 core design**  
Nearly fills Cyclone V FPGA on the DE1-SoC
- **Interrupts with hardware context-switching**  
Low latency to react to interrupt
- **Acknowledges design limitations and attempts to overcome**  
LUT resources, block memories, power and temperature requirements

# Future Improvements

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

Accomplishments  
Future Improvements  
Q&A

- **Working Global Reset**

Global resets are expensive (LUT resources)

Resetting block memories is not trivial

- **On-chip Programming**

Use the UART0 receiver to program each cores flash memory

- **Per-core gating/enabling**

Improve power efficiency for ASIC implementation by disabling cores at run-time via software.

- **Improve memory bottleneck**

Each core requires it's own memory - reduce by multiplexing access to a single large memory.

Main Project

B. Lancaster

Introduction

Top Level  
Design

Multi-core  
Functionality

Results

Conclusion

Accomplishments  
Future Improvements  
**Q&A**

# Q&A

# Info

## Main Project

B. Lancaster

## Introduction

Top Level  
Design

Multi-core  
Functionality

## Results

## Conclusion

Accomplishments  
Future Improvements  
Q&A

- GitHub repository: <https://github.com/bendl/vmicro16>
- Full Report: [https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M\\_Ben\\_Lancaster\\_201280376\\_Final.pdf](https://github.com/bendl/vmicro16/blob/master/docs/reports/build/ELEC5881M_Ben_Lancaster_201280376_Final.pdf)
- Presentation tools:
  - Latex Beamer
  - `\usecolortheme{orchid}`
  - `\useoutertheme[hideothersubsections]{sidebar}`